

CS2200

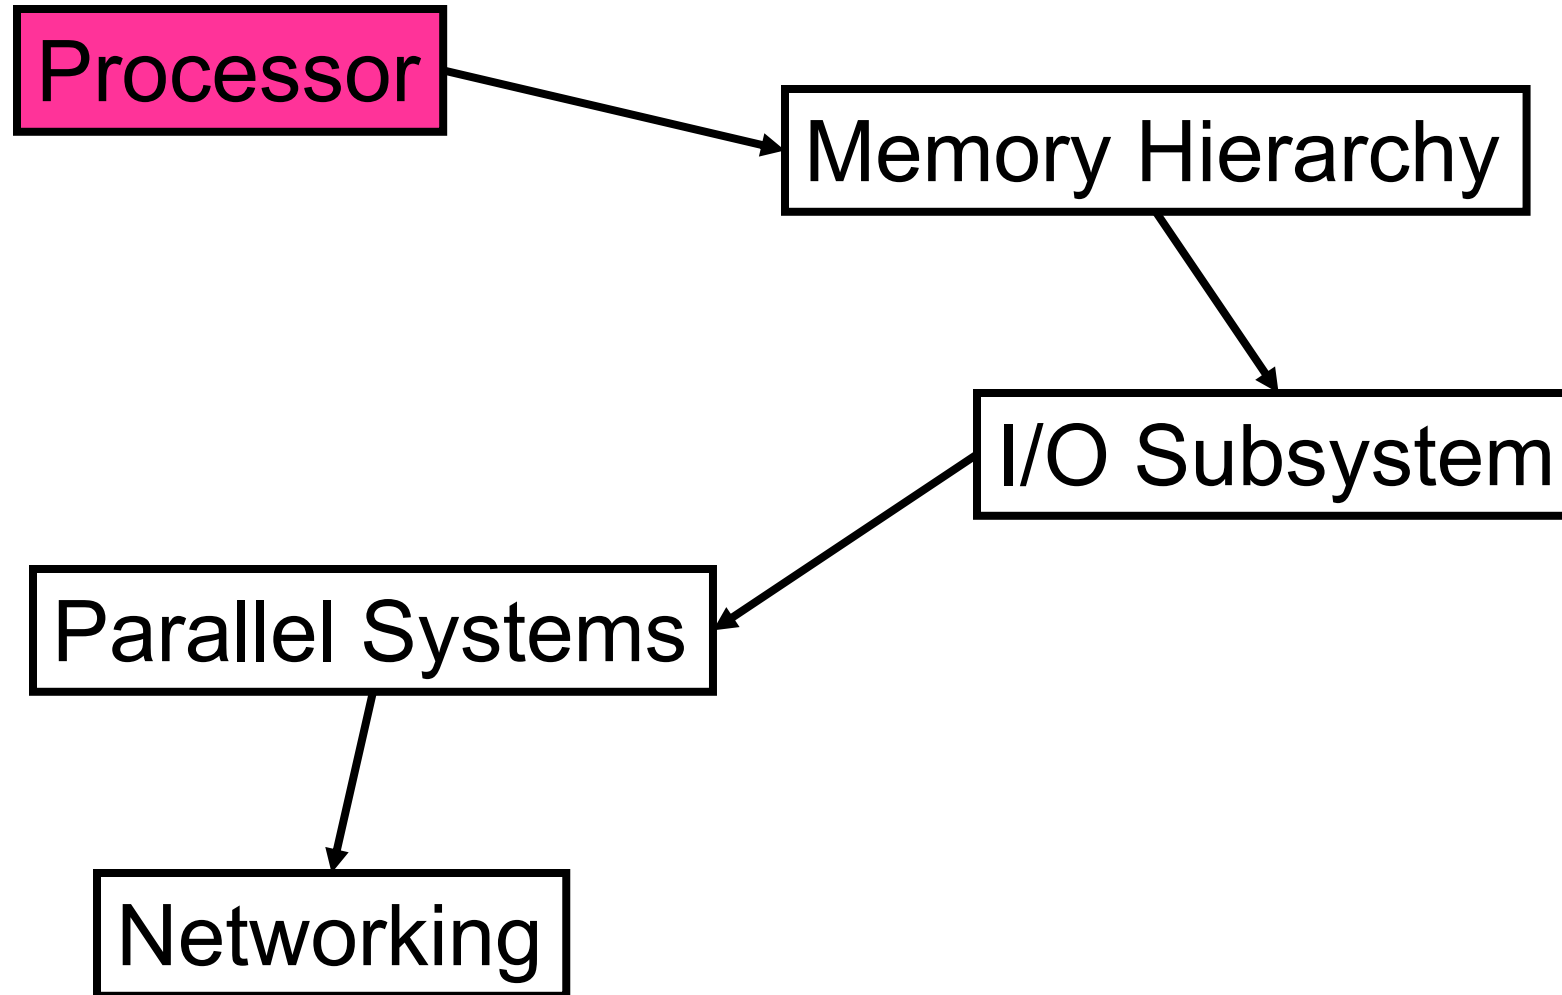
Systems and Networks

Spring 2022


Lecture 1: Processors

Alexandros (Alex) Daglis
School of Computer Science
Georgia Institute of Technology
adaglis@gatech.edu

Our Road Map




Five Classic Components



What does the processor do?

- Knows where it is in program
 - Can get and put data into memory
 - Can do some arithmetic
 - Can make tests and take different paths depending on the results
-
- Do you need a language to make a computer run?

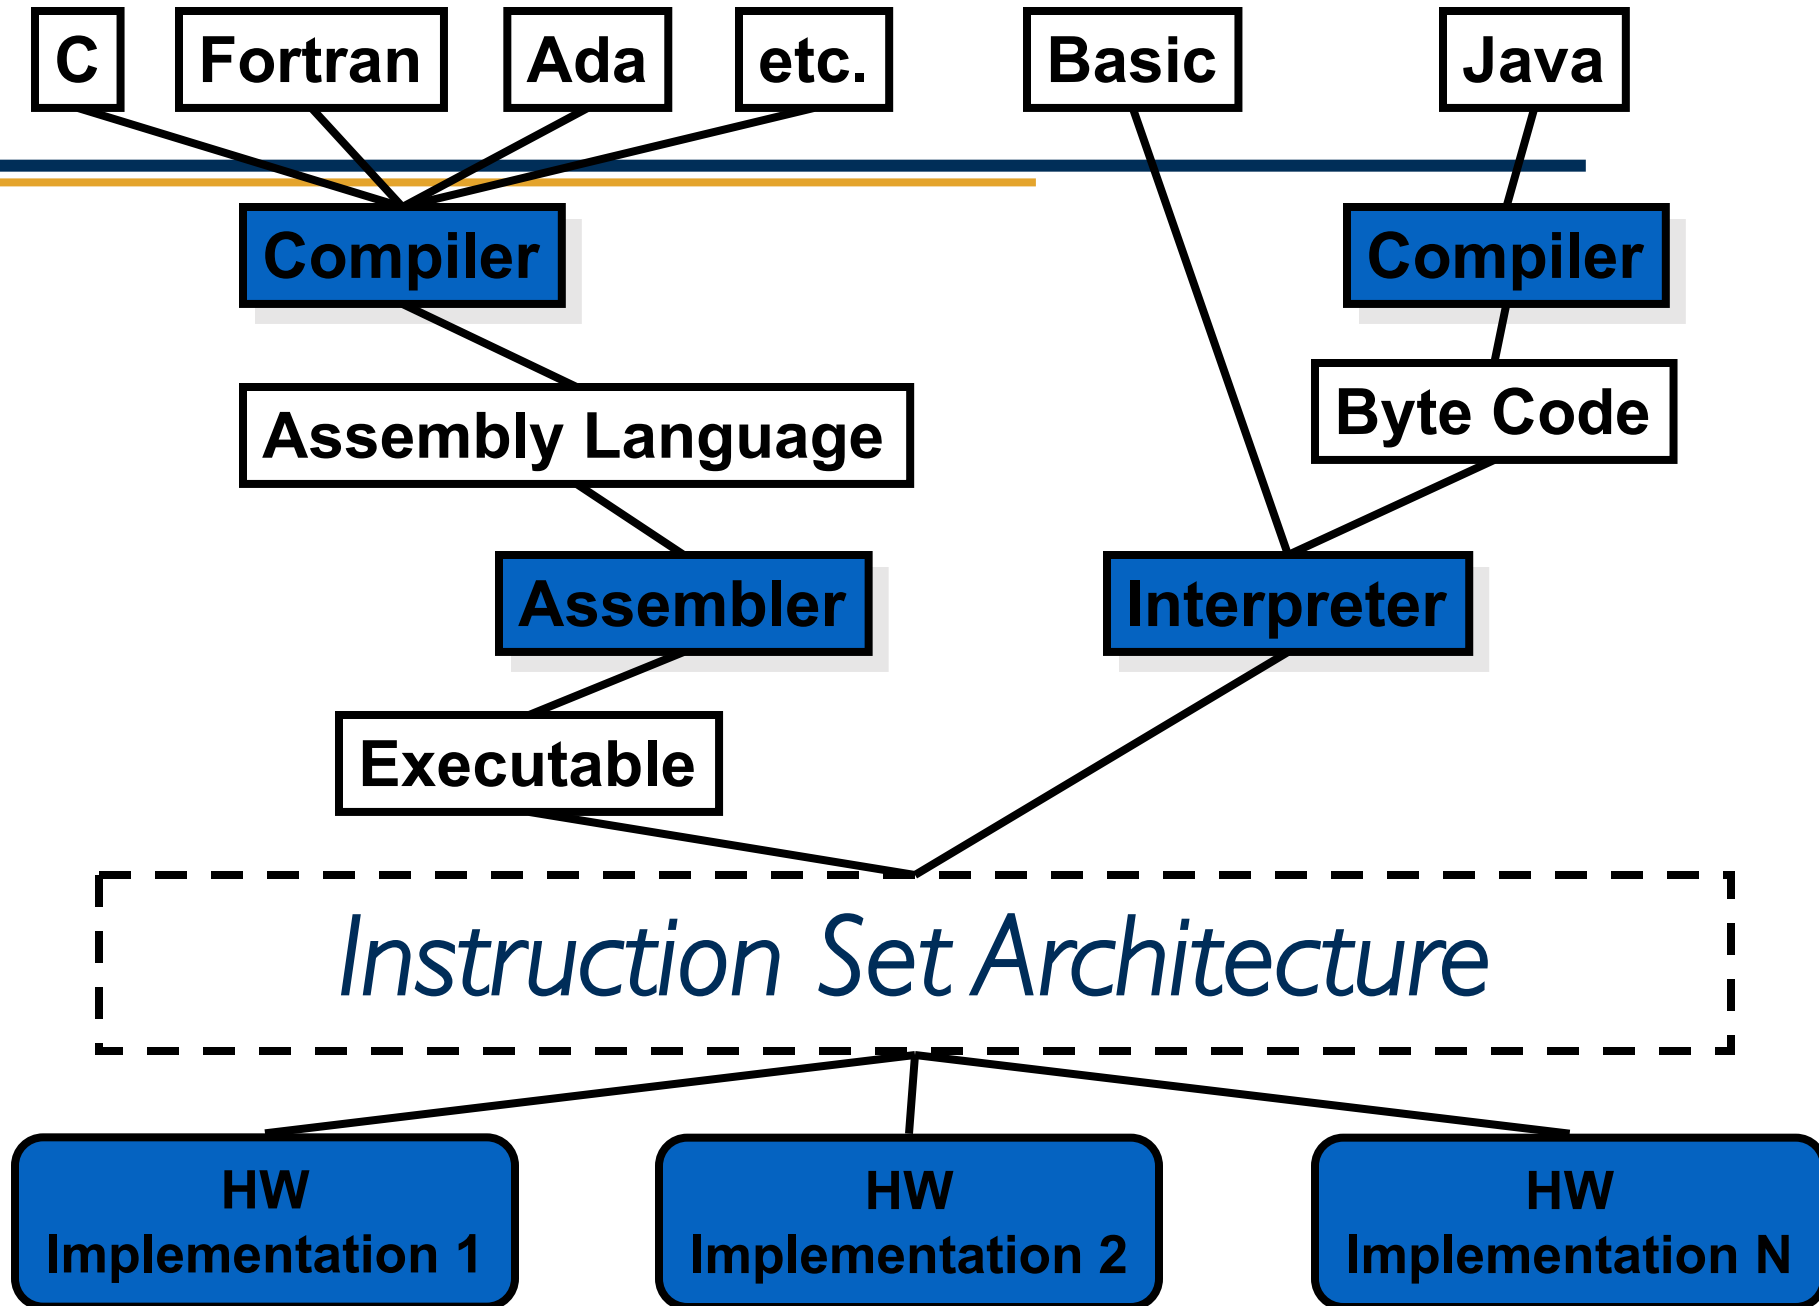
A Little History



Replacing a bad tube meant checking among ENIAC's 19,000 possibilities.

A Little History

- First computers programmed by hand
1000110010100000
- Somewhat tedious, so invented:
- Assembler
add A,B
- If we can convert from Assembly Language to machine code why not from some higher level language to Assembler?
 $A + B$




Instructions


- Language of the machine
- Vocabulary is the instruction set (ISA)
- Two levels
 - Human readable (assembly)
 - Machine readable (machine code)

Computing Evolution in 70 Years

5k ops/s in 1800sqft



5T ops/s in 16sq in




1 billion ENIACs in your palm

16 trillion times higher compute density

Moving forward


- Instruction set design from HLL constructs
 - Expressions, assignments => ALU instructions
 - Data abstraction => Addressing modes
 - Conditional & loop statements => Branch instructions
 - Procedure calls/returns => stack management
- Please note the reading assignments in the schedule: Start reading chapter 2

Simple Machine Model




- Remember the LC-3? It used a greatly simplified ARM instruction set
- We'll be introducing the LC-2200, a greatly simplified MIPS instruction set
 - Architecture that's similar, but not the same as the LC-3.

Simple Machine Model



- Let's consider the execution of a HLL
- $a = a + 1$
 $c = a + b$
if ($c == d$) {
 ...
}

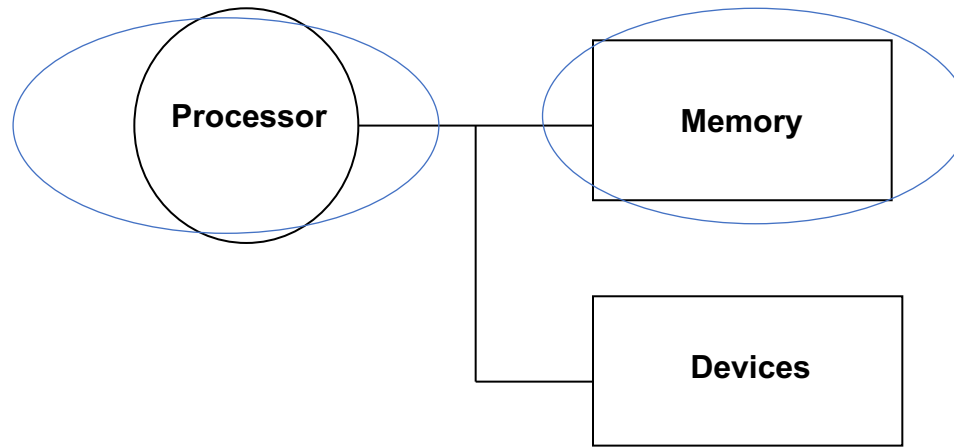
How to Design an Instruction Set?



Start thinking like a compiler writer

➔ What instructions are needed for each HLL construct?

Arithmetic/Logical Expressions



Start thinking like a compiler writer

➔ What instructions are needed for each HLL construct?

$c = a + b$

becomes

add c, a, b

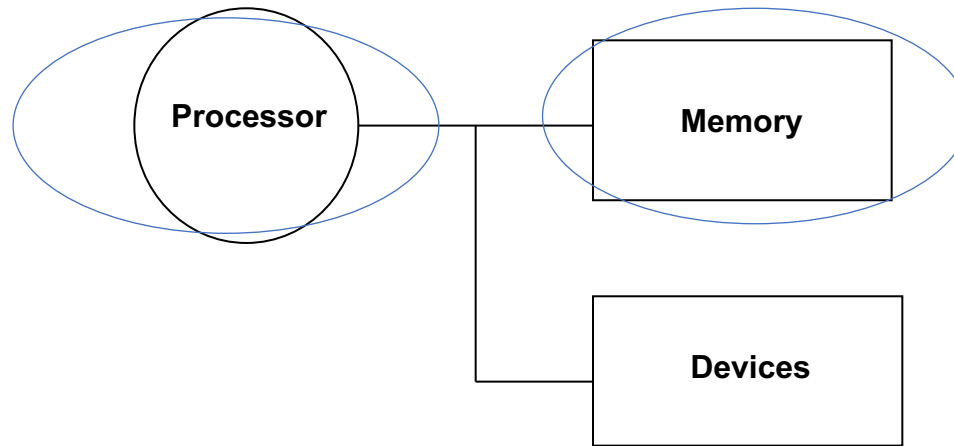
What do you call these?

➔ memory operands

Where are they?

➔ memory addressing mode

Arithmetic/Logical Expressions



Start thinking like a compiler writer

➔ What instructions are needed for each HLL construct?


$c = a + b$ ➔ add c, a, b

Keep adding to repertoire

$c = a - b$ ➔ sub c, a, b

$c = !(a \& b)$ ➔ nand c, a, b

Arithmetic/Logical Expressions



Start thinking like a compiler writer

➔ What instructions are needed for each HLL construct?


$c = a + b$ ➔ add c, a, b

Is there a downside
to operands in
memory?

How can we address
that?

A trip to memory is
EXPENSIVE!

Operands?




$c = a + b \rightarrow \text{add } c, a, b$

How about Load/Store instructions?

ld r_1, a

st c, r_2

Load/Store Instructions



ld r₁, a


st c, r₂

➔ We've got operands in registers

➔ Register addressing mode!

So how do we compile $c = a + b$ now?


Register Operands



Old way:
add c, a, b

New way:
ld r₁, a
ld r₂, b
add r₃, r₁, r₂
st c, r₃

Compiling with Register Operands



Old way:

add c, a, b } 1 instruction

New way:

ld r₁, a
ld r₂, b
add r₃, r₁, r₂
st c, r₃ } 4 instructions

This looks dumb!

Not really.
Why?


We can re-use the
values in registers!

Keep Frequently Used Tools Nearby!



...or, the principle of **locality**

Reusing Values



```
c = a + b  
d = a * b + c  
if (c == d) {  
    ...  
}
```

With operands left in registers, we can save three memory accesses here!



Why would we consider loading values into registers before computing with them?

Doesn't that use more instructions?

- 9% A. Yes it does, but it doesn't matter how many instructions it takes.
- 36% B. No it doesn't. You counted wrong.
- 27% C. Yes it does, but it saves memory accesses because we can re-use the values.
- 27% D. Yes it does and it's a terrible design trade-off.


Structs in HLL

```
struct {  
    int b;  
    int c;  
} a;
```

Elements of a struct are contiguous in memory

How do we load b and c into registers?

Let's say &a is already in register R₁



Accessing Struct Members

```
struct {  
    int b;  
    int c;  
} a;
```


Let's say &a is already in register R_1

To load b:

$R_2 \leftarrow \text{memory}[R_1 + 0]$

To load c:

$R_3 \leftarrow \text{memory}[R_1 + 4]$



Base + Offset Address Mode

```
struct {  
    int b;  
    int c;  
} a;
```

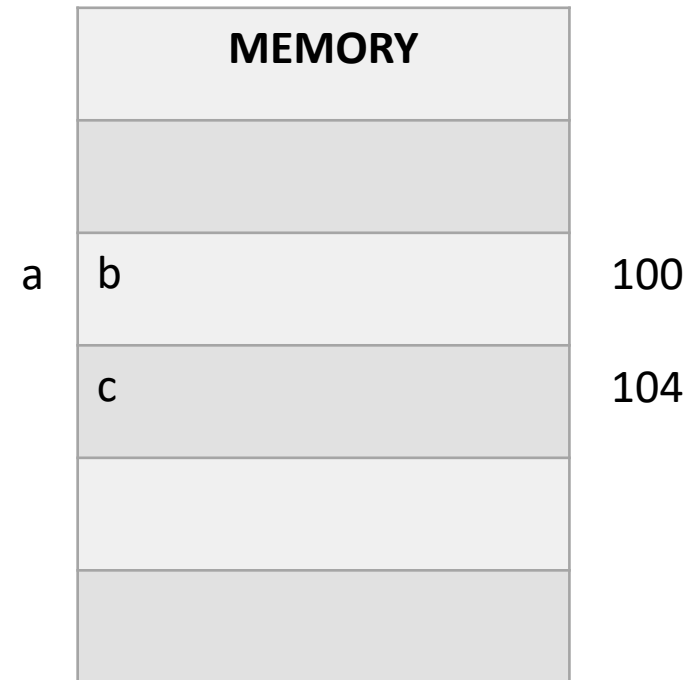
Let's say &a is already in register R_1

$\text{ld } R_i, \text{offset}(R_{\text{base}})$

To load b and c:

$\text{ld } R_2, 0(R_1)$

$\text{ld } R_3, 4(R_1)$



Operand Granularity

char → 8 bits → byte
short → 16 bits → half word
int → 32 bits* → word
long → 64 bits*

*depends on the word size of the architecture:
int=16, long=32 and others can happen

We need some instruction variants:


ldb, ldh, ldl, ...

similar for store instructions

Operand Alignment

```
struct {  
    char a;  
    char b[3];  
}
```


Dense Packing



- We're packing operands to save space.
- ISA may support ld/st with different precision!

A Different Struct


```
struct {  
    char a;  
    int b;  
}
```



If we have a 32-bit path to memory, how are we going to load b?

Why Alignment Rules Matter

```
struct {  
  char a;  
  int b;  
}
```



Now, how
to load b
is obvious

This is one of many space/time tradeoffs that ISA designers must address.

Accessing Array Operands

```
int a[100];
```

100		a[0]
104		a[1]
108		a[2]
112		a[3]
	.	
	.	
	.	
	.	
128		a[7]
132		a[8]
	.	
	.	
	.	
	.	

One approach, use base+offset

$$r_1 = 100$$

To load a[8]

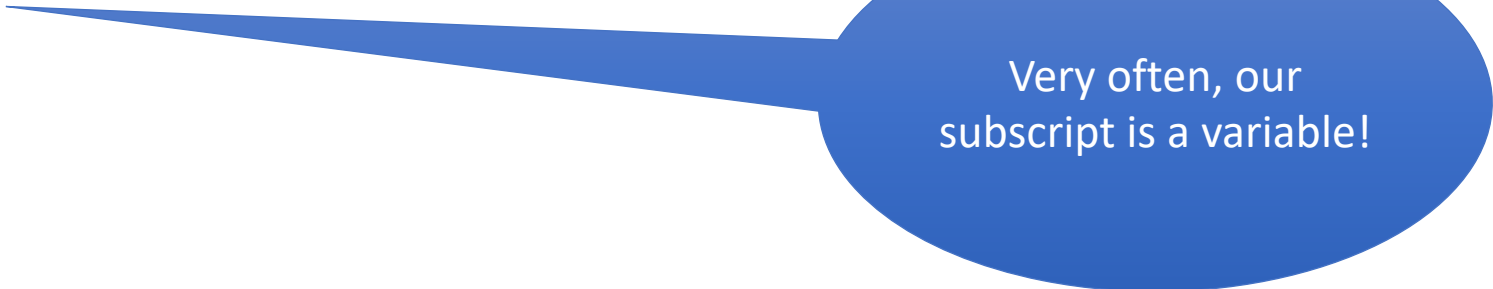
ld r₈, 32(r₁)

Is this the best we can do?

Typical Array Use

- How do we typically use arrays?

```
loop  
  c[i] = a[i]  
  ...  
  i = i + 1  
end loop
```



Very often, our subscript is a variable!

- Looks like it's time for a new addressing mode
- Let's implement base+index addressing

Accessing Array Operands

```
int a[100];
```

100		a[0]
104		a[1]
108		a[2]
112		a[3]
	.	
	.	
	.	
	.	
128		a[7]
132		a[8]
	.	
	.	
	.	
	.	

With base+index mode

$$r_1 = 100$$

$$r_2 = i * 4$$

Why?

To load a[8]

ld r₈, r₂(r₁)

Using base + index


Endianness

Say we have a 32-bit register, R1, that contains the value 1 in two's-complement. We store that register into an (aligned) memory location

```
ST R1, Mem[100]
```

The value is stored in addresses 100-103 (of course).

Which byte contains the 1? (The other 3 will be zero, right?)




Endianness

Tip: Endianness often shows up in byte-addressable memories because we can access individual bytes out of longer data types

Little Endian → addresses the LSB of the word

`int b = 0x11223344;`



Big Endian → addresses the MSB of the word

So What's the Difference

The difference only shows up when taking a “word” apart (in this case, loading 8 bits from a 32-bit integer)

ld**b** r₁, Mem[104]

Big Endian → loads **0x11** into r₁

Little Endian → loads **0x44** into r₁

+3		+2		+1	+0	
					a	100
b... 11	b... 22	b... 33	b... 44			104

So What's the Difference

- The difference shows up when taking a word apart
- Let's store 0x11223344 (a 32-bit integer) at location 104

	Little Endian Result	Big Endian Result
ldb r ₁ , Mem[104]	0x44	0x11
ldh r ₁ , Mem[104]	0x3344	0x1122
ldw r ₁ , Mem[104]	0x11223344	0x11223344

	107	106	105	104
	b...	b...	b...	b...
Big Endian	44	33	22	11
Little Endian	11	22	33	44

So, about the LC-3

- Was it Big Endian or Little Endian?
- Think carefully...
- You can't tell!
- There are no instructions that manipulate more or fewer than 16 bits
- So there isn't any way to show this implementation detail!
- Was this accidental or on purpose?

Recap

Software	Hardware
Expressions & assignments	ALU instructions
Variable reuse	register addressing mode ld/st instructions
Data abstraction <ul style="list-style-type: none">• struct• array	base + offset addr mode base + index addr mode
Granularity of operands	ldb/ldh/ldw instructions addressability (byte, word)
Packing operands	Memory alignment (space/time tradeoff)
Endianness 0x11223344	Little (first byte is 0x44) / Big (first byte is 0x11)

Review Question 1

An instruction set...

- 36% A. Serves as a level of abstraction between software and hardware.
- 18% B. Provides the details of the machine implementation.
- 32% C. Deals with the datapath and control of the processor.
- 14% D. None of the above.

Review Question 2

Addressing mode...

- 25% A. Refers to the kinds of opcodes supported in an architecture.
- 25% B. Refers to the way the operands are specified in an instruction.
- 20% C. Refers to the granularity of the memory element that can be addressed in an instruction.
- 10% D. Is a critical tool for USPS operations.
- 20% E. None of the above.

Review Question 3

Endianness of an architecture...

- 29% A. Is a key determinant of processor performance.
- 24% B. Is a key determinant of how the compiler lays out data structures in memory.
- 18% C. Matters if one declares a datatype of a particular granularity and accesses it at a different granularity.
- 6% D. Is the official name of the party held after completing this course.
- 24% E. None of the above.

What do we need for...


- Conditional statements
- Switch statements
- Loops
- Procedure calls
- Other considerations for ISA

Compiling Conditional Statements

- In what order are program statements normally executed?
- How do we know what instruction to execute next?
- How can we handle this high-level language construct:

```
if (x == y) z = 7;
```


What Do We Need to Do?



- Evaluate predicate
- Break the sequential flow of instructions
- Rejoin control path

Implementing a Conditional

- Evaluate predicate
 - ALU Op
- Break sequential flow
 - Need to know where we are
 - ➔ PC
 - Need a new instruction
 - ➔ BEQ $r1, r2, \text{offset}$
 - ➔ if $r1 == r2$ then $PC = PC + \text{offset}$
else do nothing
 - ➔ PC relative addressing mode!
- Rejoin control flow
 - ➔ need an unconditional jump



An Example

- C

```
if (a == b)
    c = d + e;
else
    c = f + g;
```

Assuming

r1 = a

r2 = b

r3 = c

r4 = d

r5 = e

r6 = f

r7 = g

- Assembly

```
        beq  r1, r2, then
        add  r3, r6, r7
        beq  r1, r1, skip*
then add  r3, r4, r5
skip ...
```

* Effectively an unconditional branch

Outcome of Conditional Statements

- Introduction of PC
- One new instruction
 BEQ r_1, r_2, offset
- One new addressing mode: PC-relative
- (optional) an Unconditional Jump
 J r_n ; $PC \leftarrow r_n$
- Do we really need an unconditional jump??

Compiling Switch Statements

```
if (n==0)
    x=a;
else if (n==1)
    x=b;
else if (n==2)
    x=c;
else
    x=d;
```

Do these produce
essentially equivalent
assembly code?

They can, but they
don't have to!

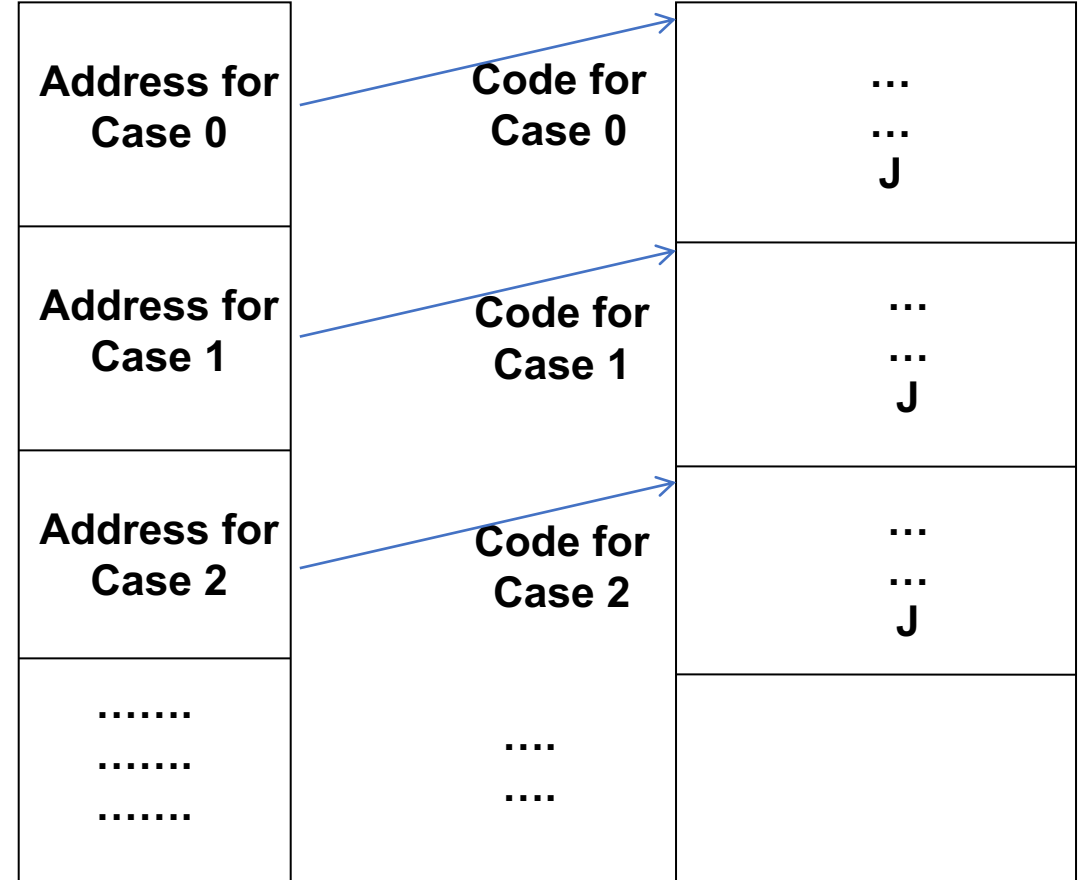
```
switch (n) {
    case 0:
        x=a;
        break;
    case 1:
        x=b;
        break;
    case 2:
        x=c;
        break;
    default:
        x=d;
}
```

Switch Can Use a Jump Table

- Think of a C array of pointers to the individual cases
- To do this we need an indirect addressing mode

J @r_i

➔ PC ← Mem[r_i]



Jump table

Loops

- Do we need anything new in the ISA?
- Not really.

Compiling Loops

- C

```
while(j != 0)
{
    /* loop body */
    t = t + a[j--];
}
```

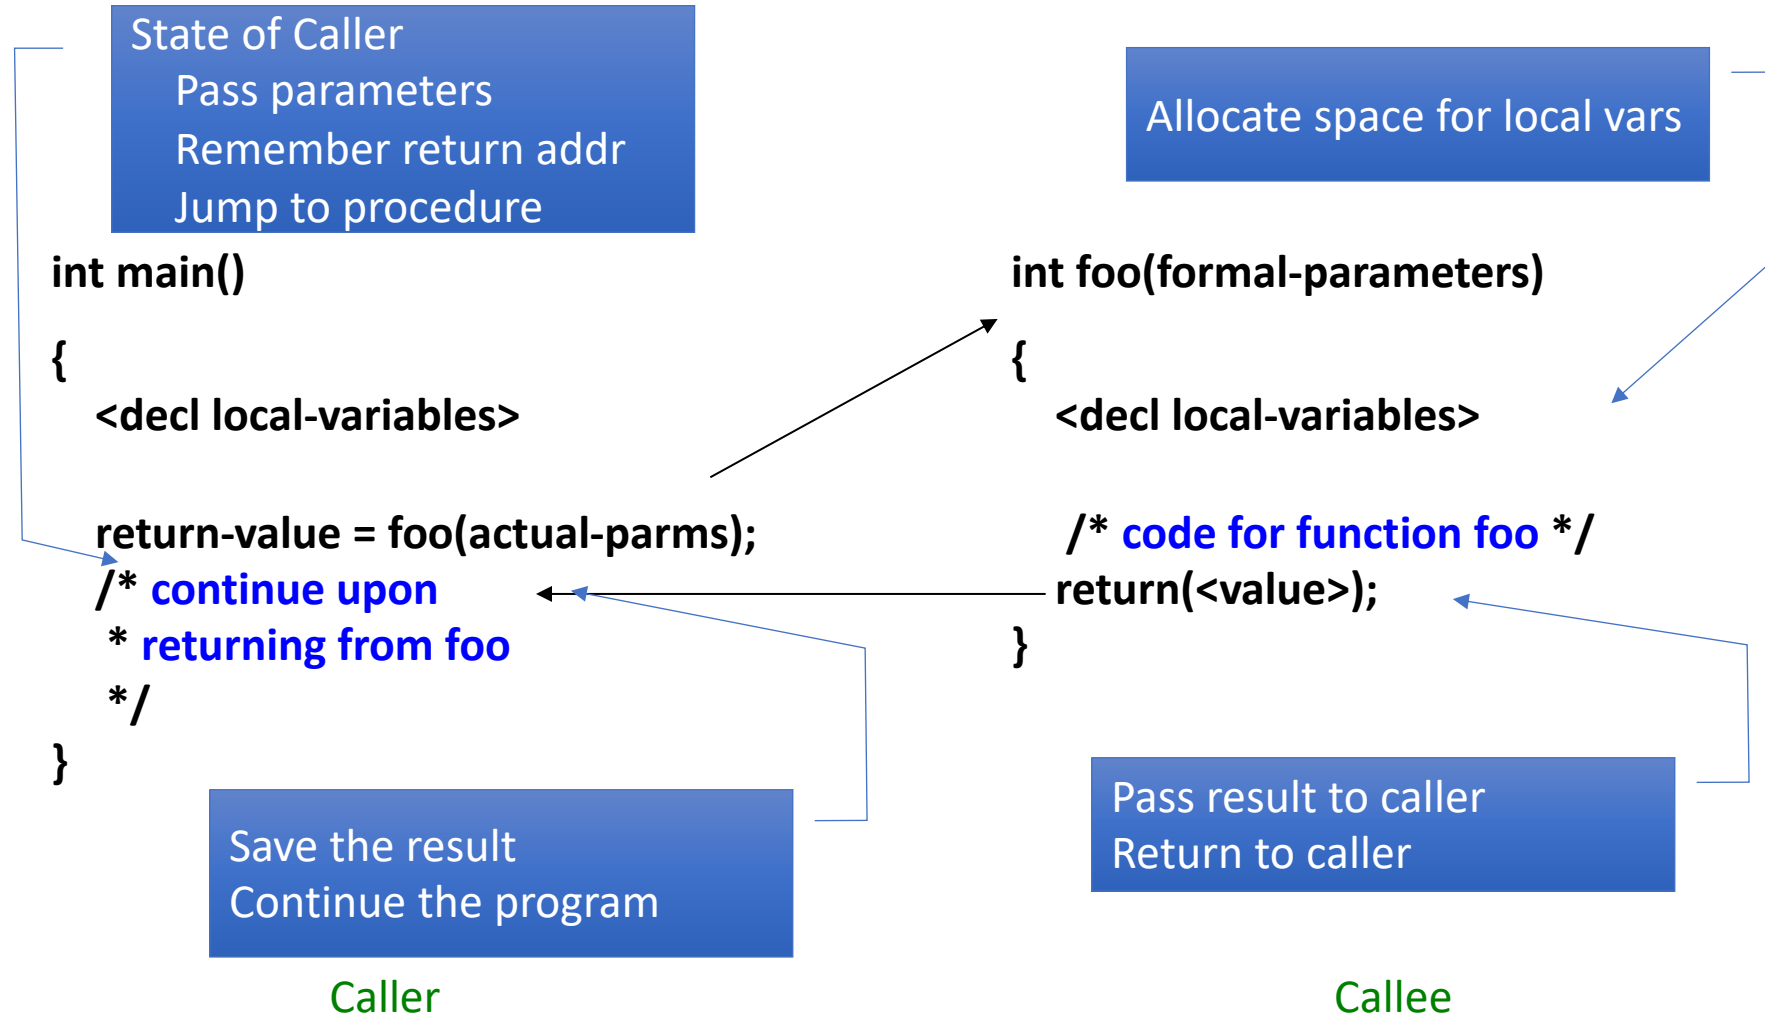
- Assembly

```
loop    beq r1,r0,done
        ; loop body
        ...
        beq r0, r0, loop
done    ...
```

Summary

Software	Hardware
Expressions & assignments	ALU instructions, LD/ST instructions
Data abstraction <ul style="list-style-type: none">• struct• array	register addr mode base + offset addr mode base + index addr mode
Conditionals & Loops	PC-relative addr mode branch/jump instruction (register or PC-relative) Indirect addr mode (optional)


How Do We Compile Function Calls?



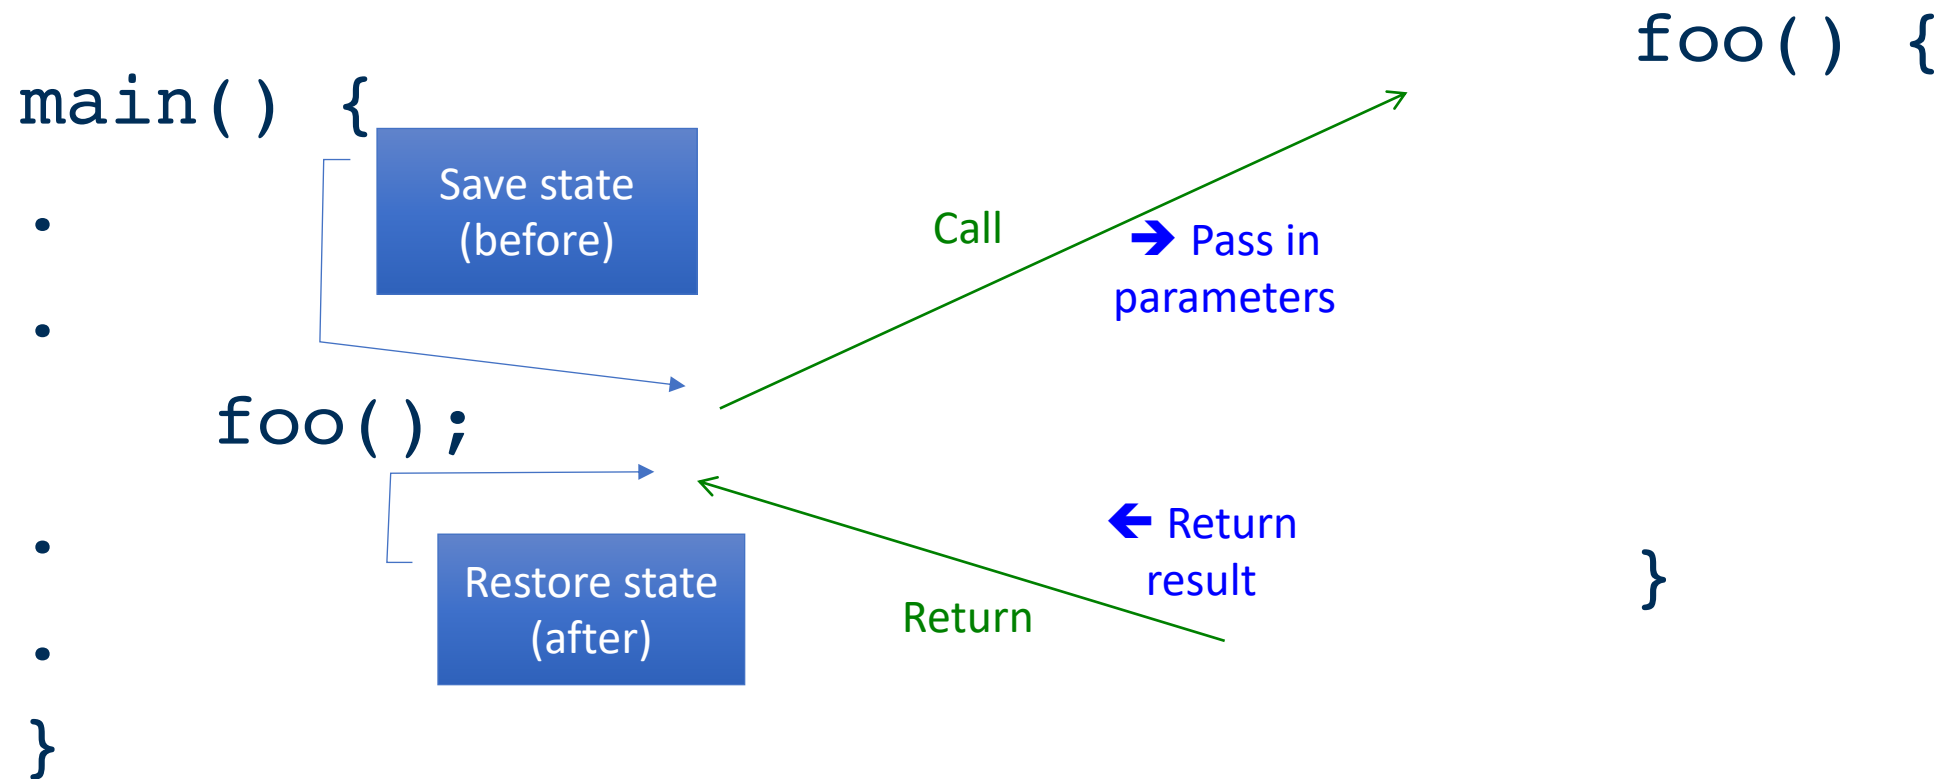
Remembering the Return Address

- Have we needed to do this before?
- Add a Jump & Link instruction
 - JALR $r_{\text{target}}, r_{\text{link}}$; $r_{\text{link}} \leq \text{PC}, \text{PC} \leq r_{\text{target}}$
- Recall
 - J r_{target} ; $\text{PC} \leq r_{\text{target}}$
- Do we need this instruction anymore?

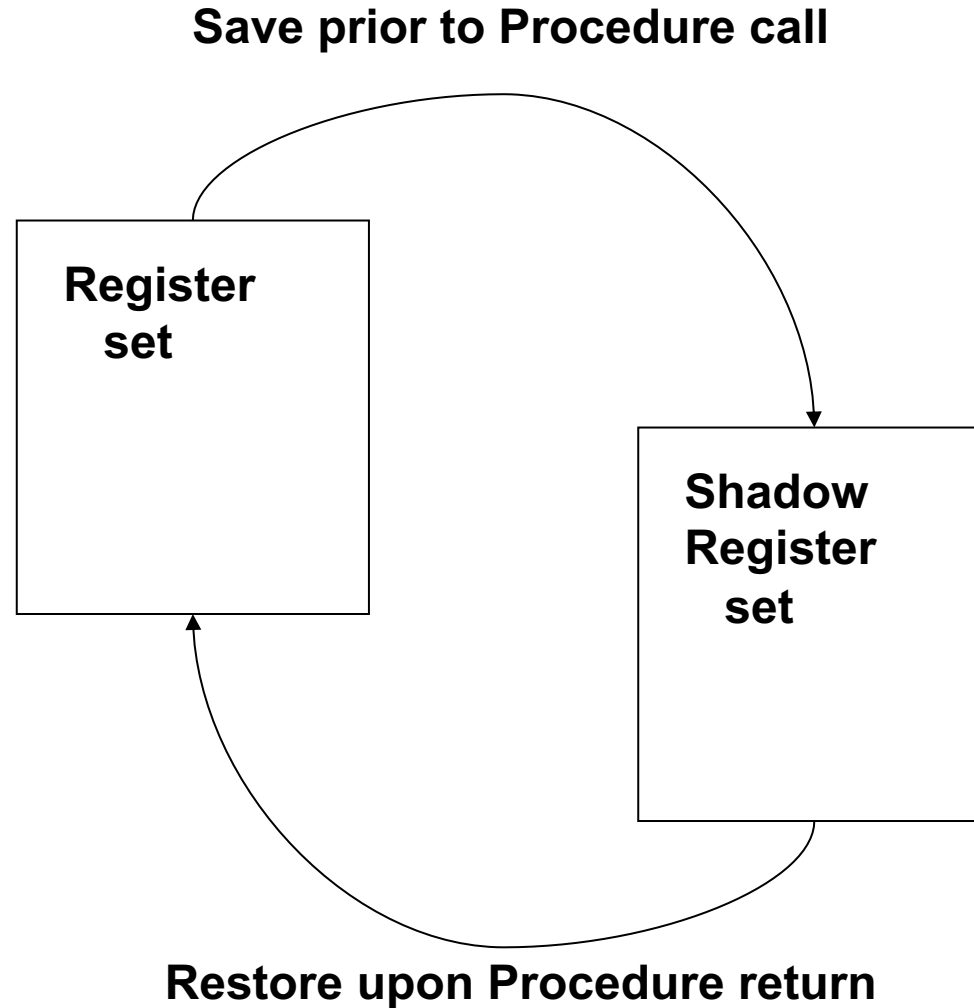
Control Flow



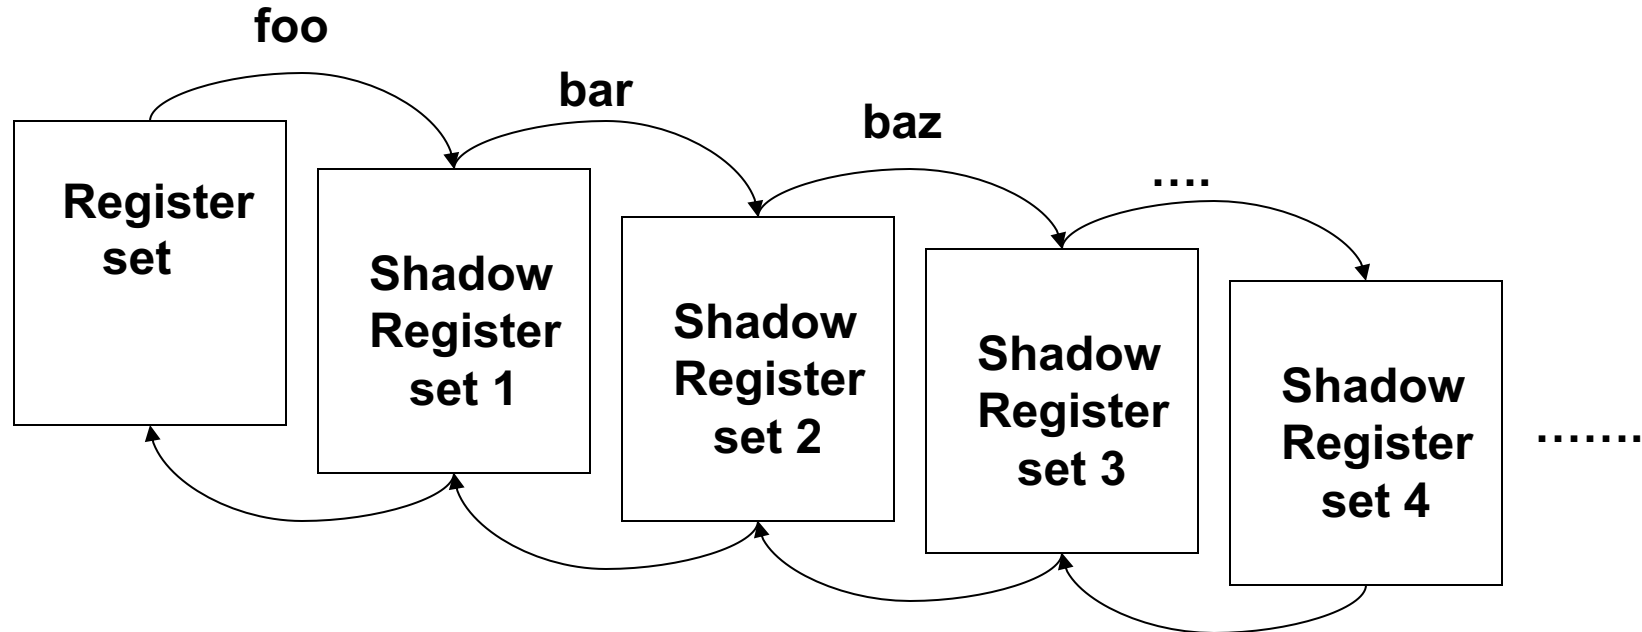
Control Flow



Another Way to Save State



Shadow Register Sets



- `foo()` calls `bar()` who calls `baz()`, etc.
- The Big Deal: No memory accesses!
(but we need lots of extra registers)
- Another form of this is called **register renaming**

Saving State

- If we don't have shadow registers, where are we going to save all that state?
- A stack

Where are we going to put the stack?

- In memory
- But in small cases, could we hold the state in a few extra registers?
(another space/time tradeoff)

Use a Stack to Communicate

```
main() {
```

```
•
```

```
•
```

```
    foo();
```

```
•
```

```
•
```

```
}
```

```
foo() {
```

```
- Save/restore state
```

```
- Pass parameters
```

```
- Return results
```

```
}
```

```
- What else is needed in ISA?
```

```
- Nothing new..
```

Saving Registers During a Procedure

- We can have the **caller** save all the registers
-or-
We can have the **callee** save all the registers
- What's wrong with those choices?
 - Not everything needs to be saved every time...

Saving Registers During a Procedure

- If we split the assignment of the registers, then most of the time, the **caller** and **callee** can each save fewer registers based on what they actually need to use
- In the LC-2200 case, we'll functionally divide the working register set
 - **s0-s2** registers which the **callee** must preserve if it wants to use them
 - **t0-t2** registers which the **caller** must preserve if it wants their values to persist over a function call
- This division of responsibility saves memory accesses.

Stack as Communication Area

- Saving/restoring state over a procedure call

Who does it?

➔ Split between Caller and Callee

Stack as Communication Area

- Returning results

Do we really need to put them on the stack?

➔ Use registers
(We'll call this register **v0**)

Stack as Communication Area

- Parameter Passing

Do we really need to put them on the stack?

➔ Use registers
(We'll call these registers **a0-a2**)

Stack as Communication Area

- Will we need the stack at all for parameters and results?
- What if we run out of registers?
- We use the stack if we run out
- Here we're trading time for complexity

Moral of the Story

- Use the stack sparingly
 - LD/ST instructions are expensive (i.e. memory access is slow)
- Software calling convention
 - Used by the compiler to keep track of the use of the stack and registers
 - Better have one!

Software Convention for LC-2200

Use:
Program
Data

- Registers s0-s2 are the caller's saved registers
- Registers t0-t2 are the temporary registers
- Registers a0-a2 are the parameter passing registers
- Register v0 is used for return value

Use:
Bookkeeping

- Register ra is used for return address (r_{link})
- Register at is used for target address (r_{target})
- Register sp is used as a stack pointer

Review Question 1

Saving and restoring of registers on a procedure call...

- 20% A. Is always done by the caller.
- 20% B. Is always done by the callee.
- 20% C. Is never done since hardware implicitly takes care of it.
- 20% D. Is done on a need basis partly by the caller and partly by the callee.
- 20% E. What is a caller/callee?

Review Question 2

On the LC-2200, how are actual parameters passed to a function?

- 20% A. On the stack.
- 20% B. On the heap.
- 20% C. Up to 3 in registers, the rest on the stack.
- 20% D. Up to 6 in registers, the rest on the stack.
- 20% E. None of the above.

Review Question 3

We store some values in registers during a procedure call...

- 25% A. Because we like to mix things up – variety is good!
- 25% B. Because it reduces memory references.
- 25% C. It makes the stack shorter so it reduces the danger of overflow.
- 25% D. It makes for prettier code.