Georgia Tech Computer Science

CS2200
Systems and Networks

Lecture 7: Interrupts

Alexandros (Alex) Daglis
School of Computer Science
Georgia Institute of Technology
adaglis@gatech.edu

Lecture slides adapted from Bill Leahy, Charles Lively of Georgia Tech

Today's agenda

- Finish up control path
  - Recap of basic ROM use
  - Couple more examples of instruction execution steps
  - Micro-sequencer

- Interrupts, Traps, Exceptions
  - Chapter 4

Control unit's ROM contents

	Drive Signals					Load Signals						Write Signals						
Current State	PC	ALU	Reg	MEM	OFF	PC	A	B	MAR	IR	Z	MEM	REG	func	Reg Sel	M	T	Next State
0000 0 00000	1						1		1									00001
0000 0 00001				1						1								00010
0000 0 00010		1				1								11		1		10000
…																		
0101 0 10000			1				1								00			10001
0101 0 10001			1					1							01			10010
0101 0 10010		1									1				10		1	10011
0101 0 10011	1						1		1									00001
…																		
0101 1 10011	1						1										1	10100
0101 1 10100					1			1									1	10101
0101 1 10101		1				1								00				00000

EXECUTE state: ADD instruction

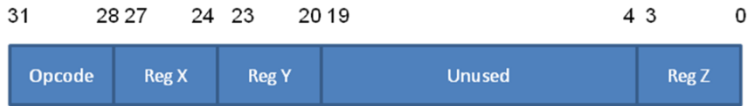

- $R_X \leftarrow R_Y + R_Z$

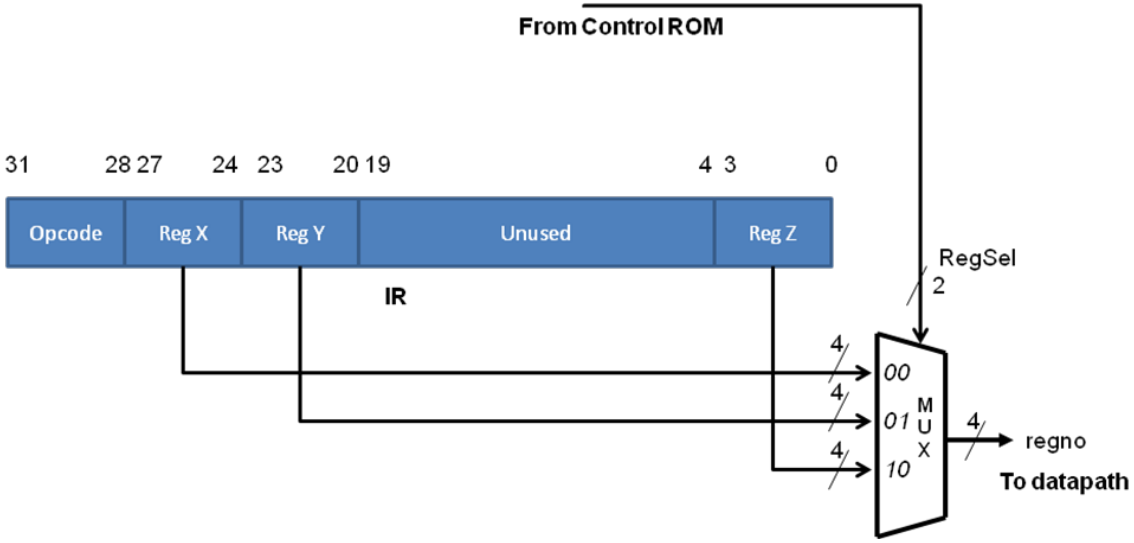

EXECUTE state: ADD instruction

add1

Ry → A

Control signals needed:

RegSel = 01

DrREG

LdA

add2

Rz → B

Control signals needed:

RegSel = 10

DrREG

LdB

add3

A+B → Rx

Control signals needed:

func = 00

DrALU

RegSel = 00

WrREG

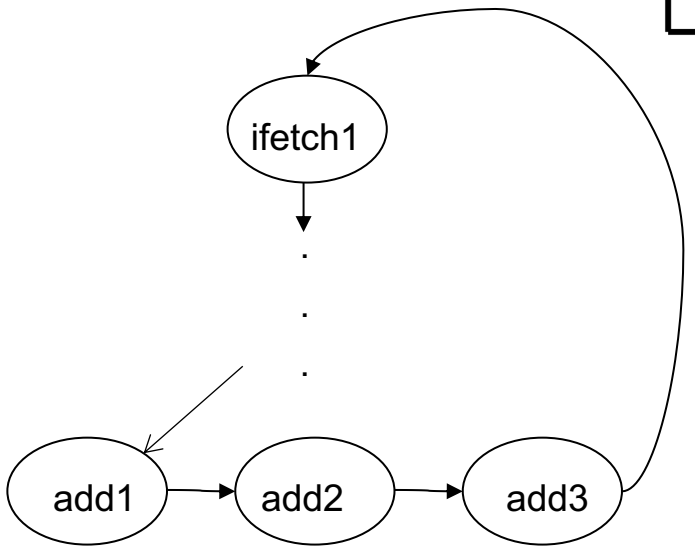

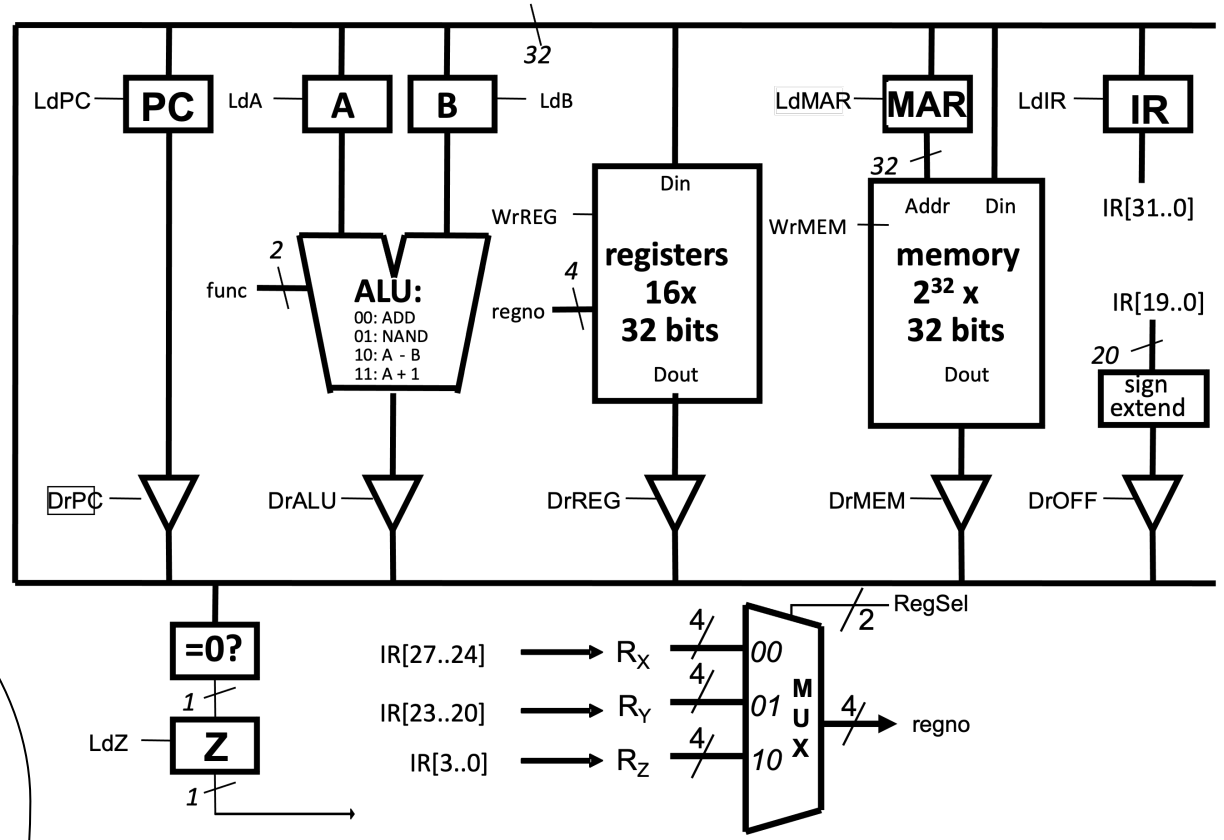

What must be changed in ADD to implement NAND?

EXECUTE state: JALR instruction

JALR instruction does the following:

$R_Y \leftarrow PC + 1$

$PC \leftarrow R_X$

jalr1

PC → Ry

Control signals needed:

DrPC

RegSel = 01

WrREG

jalr2

Rx → PC

Control signals needed:

RegSel = 00

DrREG

LdPC

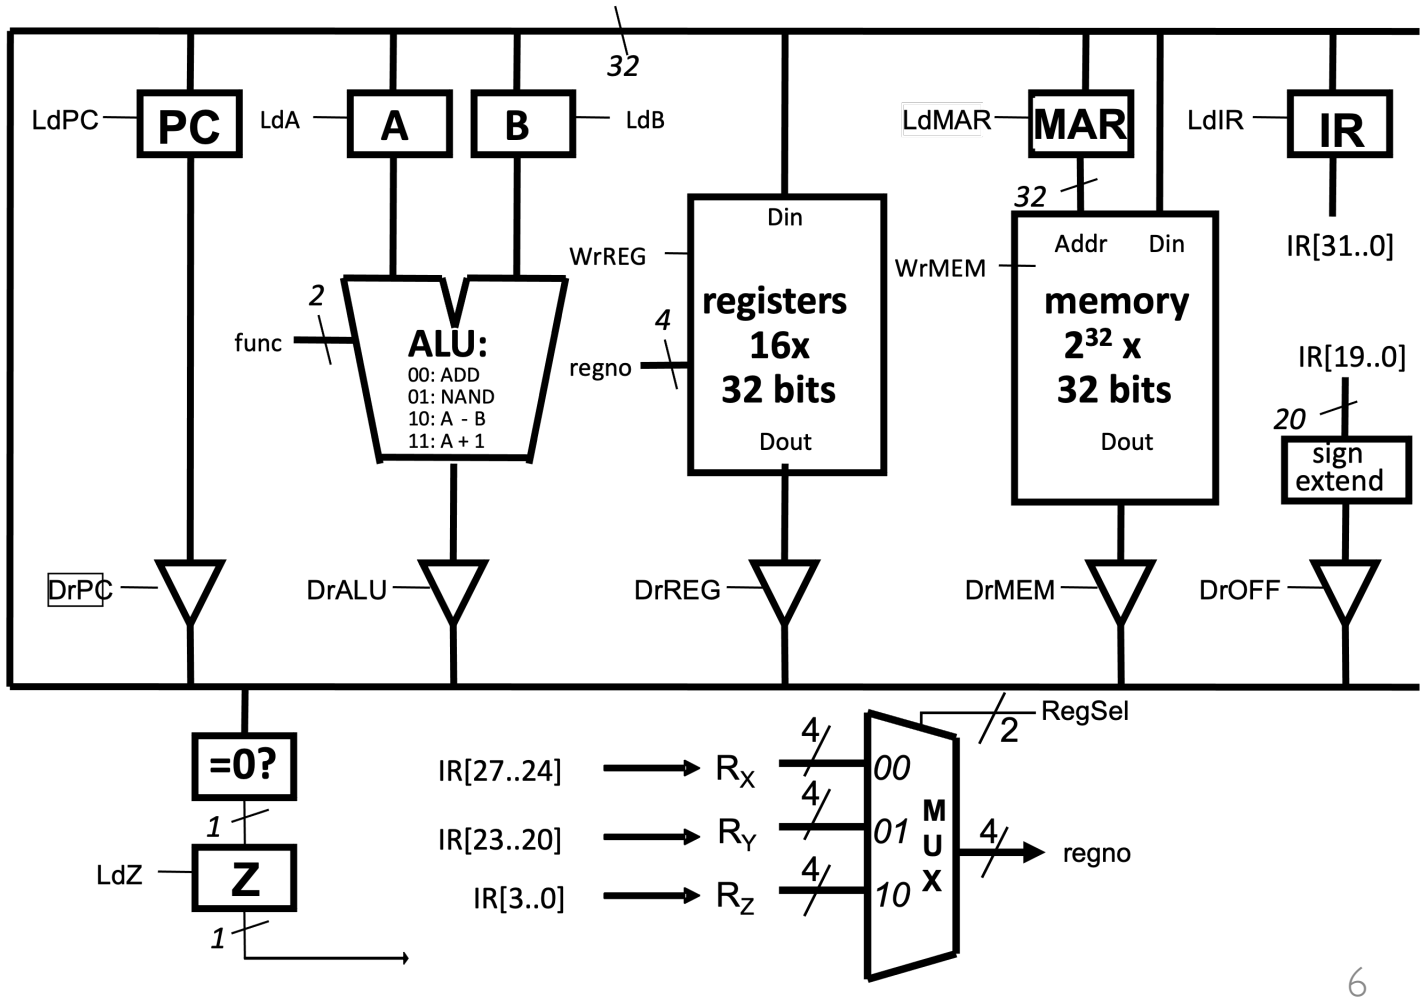

Session ID:
dag2200

Question

When all of the control signals are zero in the LC-2200 datapath, what value is being presented by the ALU to DrALU?

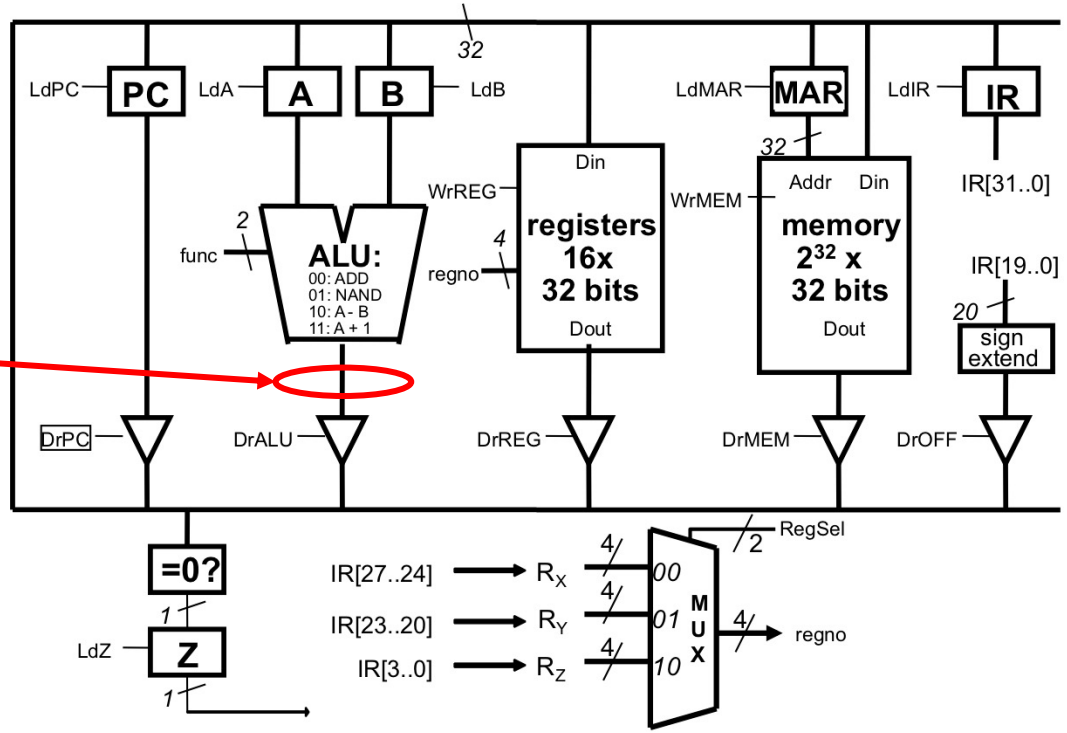

25% A. A + B

38% B. The value of one of the registers

25% C. Zero

13% D. 42

Alternative Style of Control Unit Design

A number of different approaches may be used to implement the Control Unit

Microprogrammed Control

- As presented our design works
- Problem: Too slow
  - Solution: Pre-fetch the next microinstruction
- Problem: Too much memory required
  - Solution: OR the opcode with the next state value 10000 instead of pre-pending it
  - Solution: Use more than one ROM and more sophisticated Decode/BEQ logic
    - One set of ROMs for which state comes next
    - One for what the control outputs should be in the state

3-ROM Microsequencer

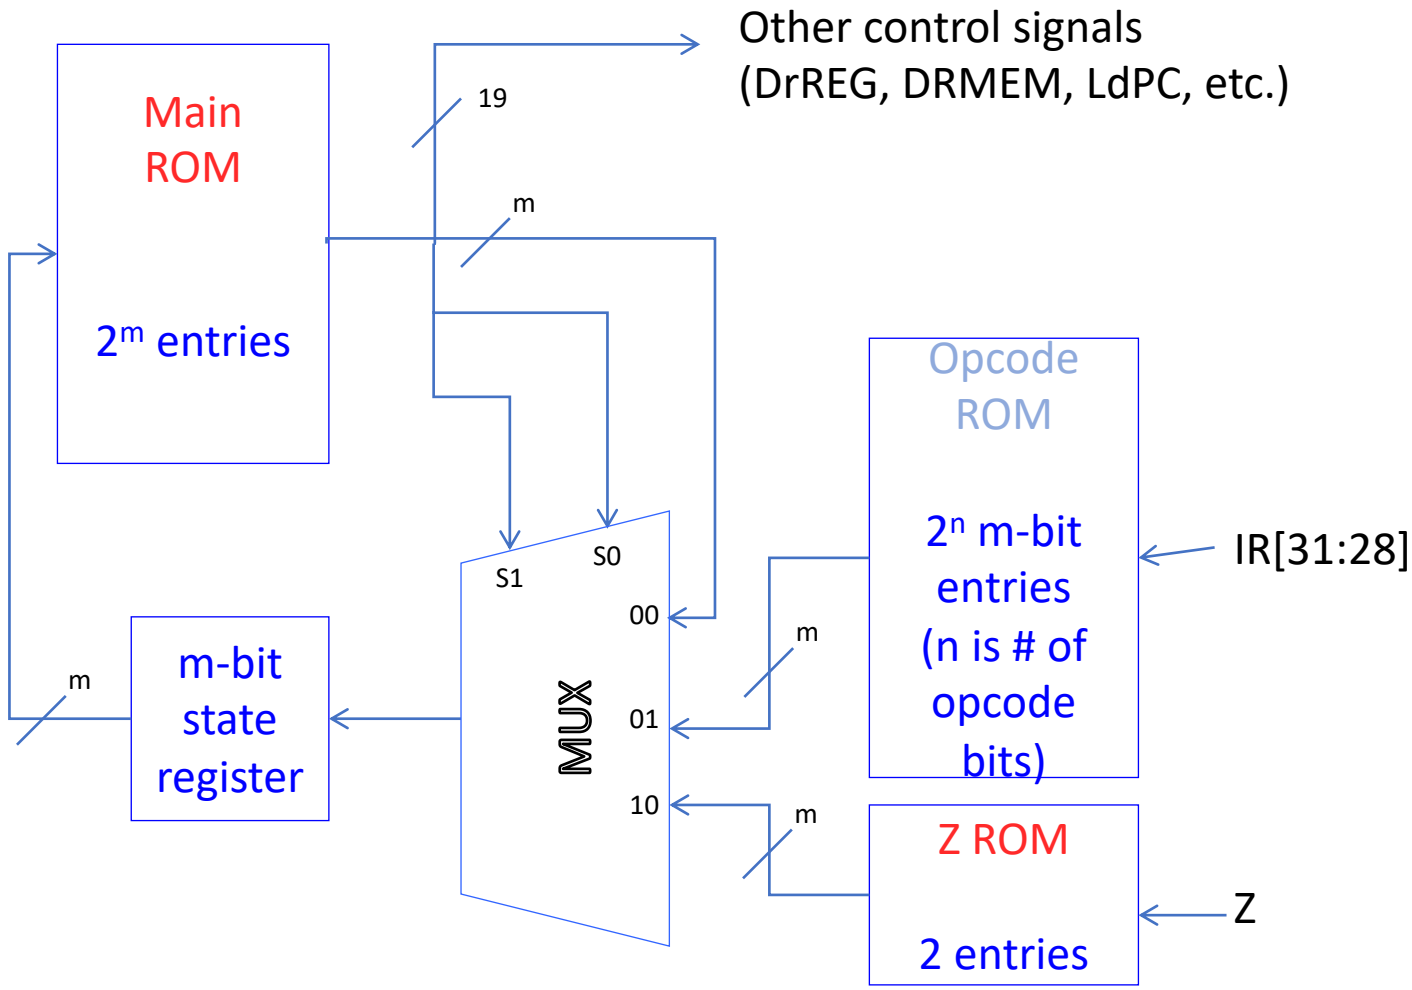

Space/Time Tradeoff

- Flat ROM
  - More space (since we increased the ROM by a factor of 32 for the occasional address modifiers, but have extra ROM space)
  - Faster since only one ROM access in each microinstruction

- Micro sequencer (3-ROM control unit)
  - Less space (main ROM much smaller than Flat ROM)
  - Slower since additional ROM access in every clock cycle

Hardwired Control

- State machine can be represented as sequential logic truth table
- Thus can be implemented using normal combinational logic or FPGA

- Can produce boolean function for each control signal
  - E.g., DrPC = ifetch1 + jalr1 + beq4 + …

Control Regime	Pros	Cons	Comment	When to use	Examples
Micro-programmed	Simplicity, maintainability, flexibility Rapid prototyping	Potential for space and time inefficiency	Space inefficiency may be mitigated with vertical microcode (micro-sequencer) Time inefficiency may be mitigated with prefetching	For complex instructions, and for quick non-pipelined prototyping of architectures	PDP 11 series, IBM 360 and 370 series, Motorola 68000, complex instructions in x86 architecture
Hardwired	Amenable for pipelined implementation Potential for higher performance	Potentially harder to change the design Longer design time	Maintainability can be increased with the use of structured hardware such as PLAs and FPGAs	For High performance pipelined implementation of architectures	Most modern processors including Intel Pentium series, IBM PowerPC, MIPS

Interrupts, Traps and Exceptions

- Interrupts, traps and exceptions are discontinuities in program flow

- Students asking a teacher questions in a classroom is a good analogy to the handling of discontinuities in program flow

Discontinuities in program execution

We must first understand

- **Synchronous** events: Occur at well defined points aligned with activity of the system
  - Making a phone call
  - Opening a file

- **Asynchronous** events: Occur unexpectedly with respect to ongoing activity of the system
  - Receiving a phone call
  - A user presses a key on a keyboard

Discontinuities in program execution

Definitions

- **Interrupts**: Asynchronous events usually produced by I/O devices which must be handled by the processor by interrupting execution of the currently running process
- **Traps**: Synchronous events produced by special instructions typically used to allow secure entry into operating system code
- **Exceptions**: Synchronous events usually associated with software requesting something the hardware can't perform i.e. illegal addressing, illegal op code, etc.

Discontinuities in program execution

Type	Sync/Async	Source	Intentional?	Examples
Exception	Sync	Internal	No	Overflow, Divide by zero, Illegal memory address
Trap	Sync	Internal	Yes and No	System call, Page fault, Emulated instructions
Interrupt	Async	External	Yes	I/O device completion

An example of a trap is...

A. Keyboard input
B. System call to open a file
C. Divide by zero
D. Arithmetic overflow
E. High speed I/O
F. ... the question itself is a trap!

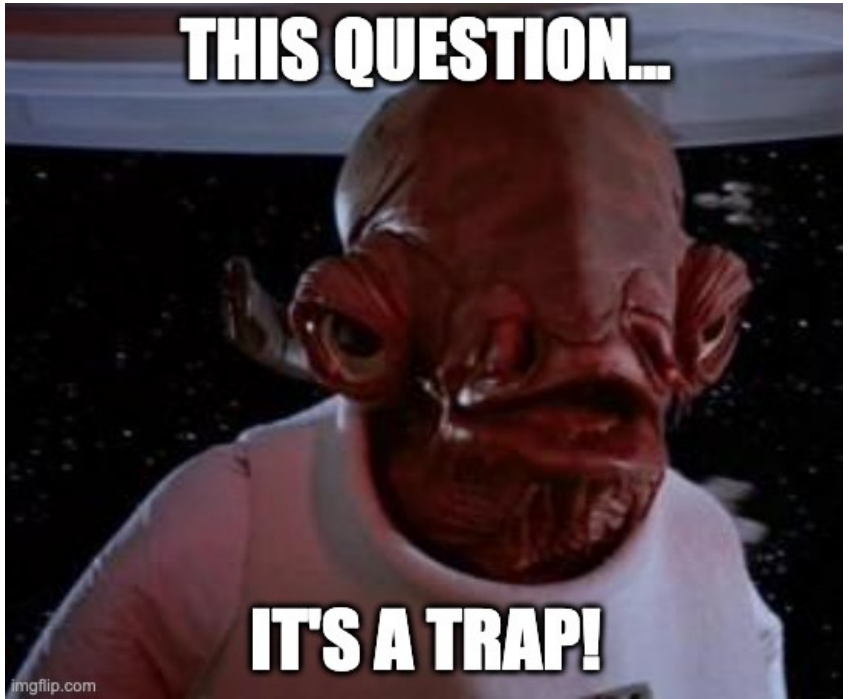

Execution path

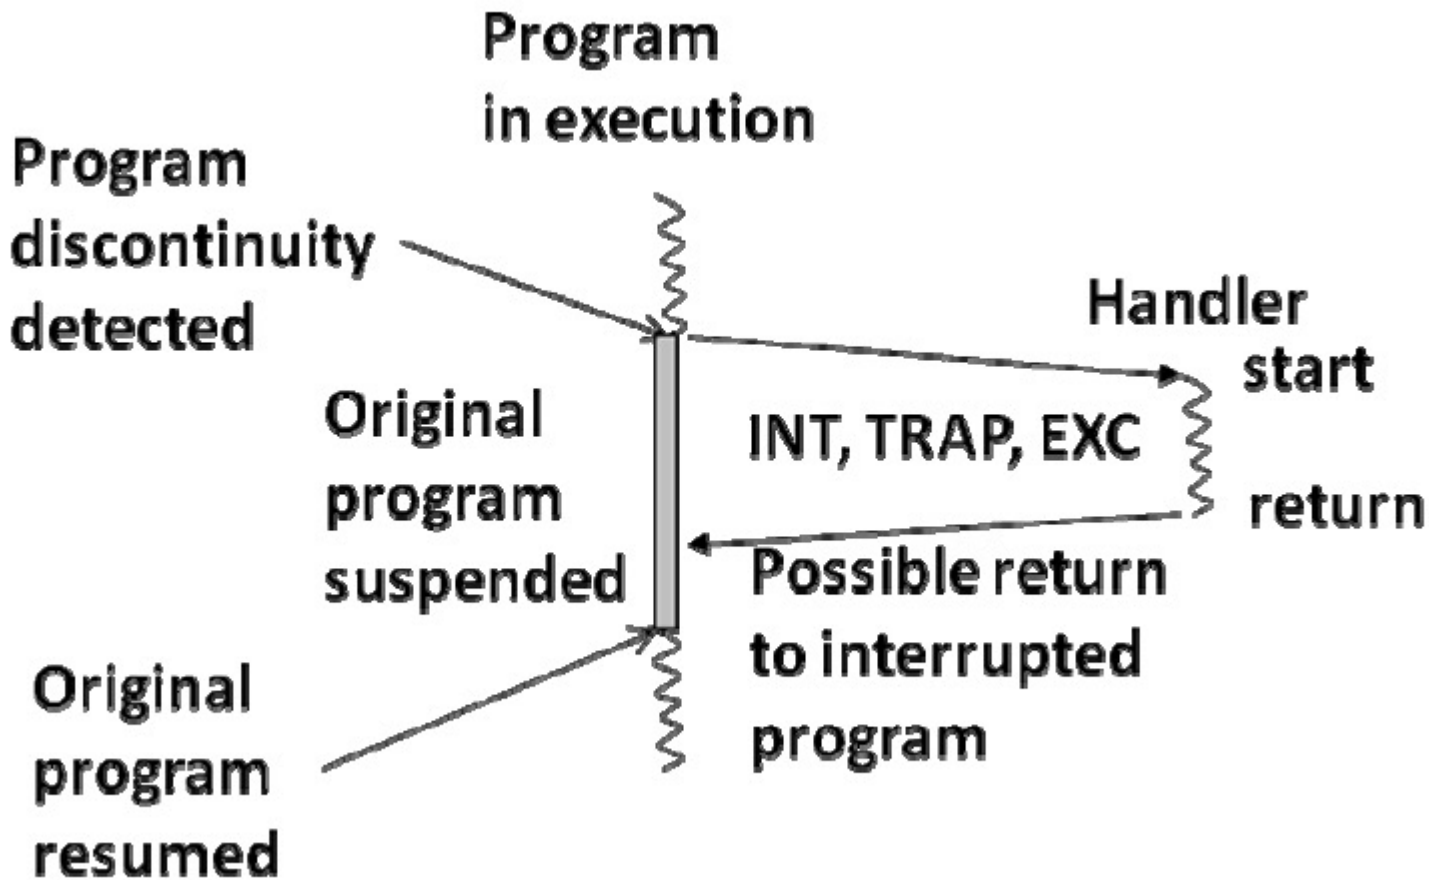

New internal processor register

Will contain a unique number stashed by the hardware to indicate the type of discontinuity

Interrupt Vector Table

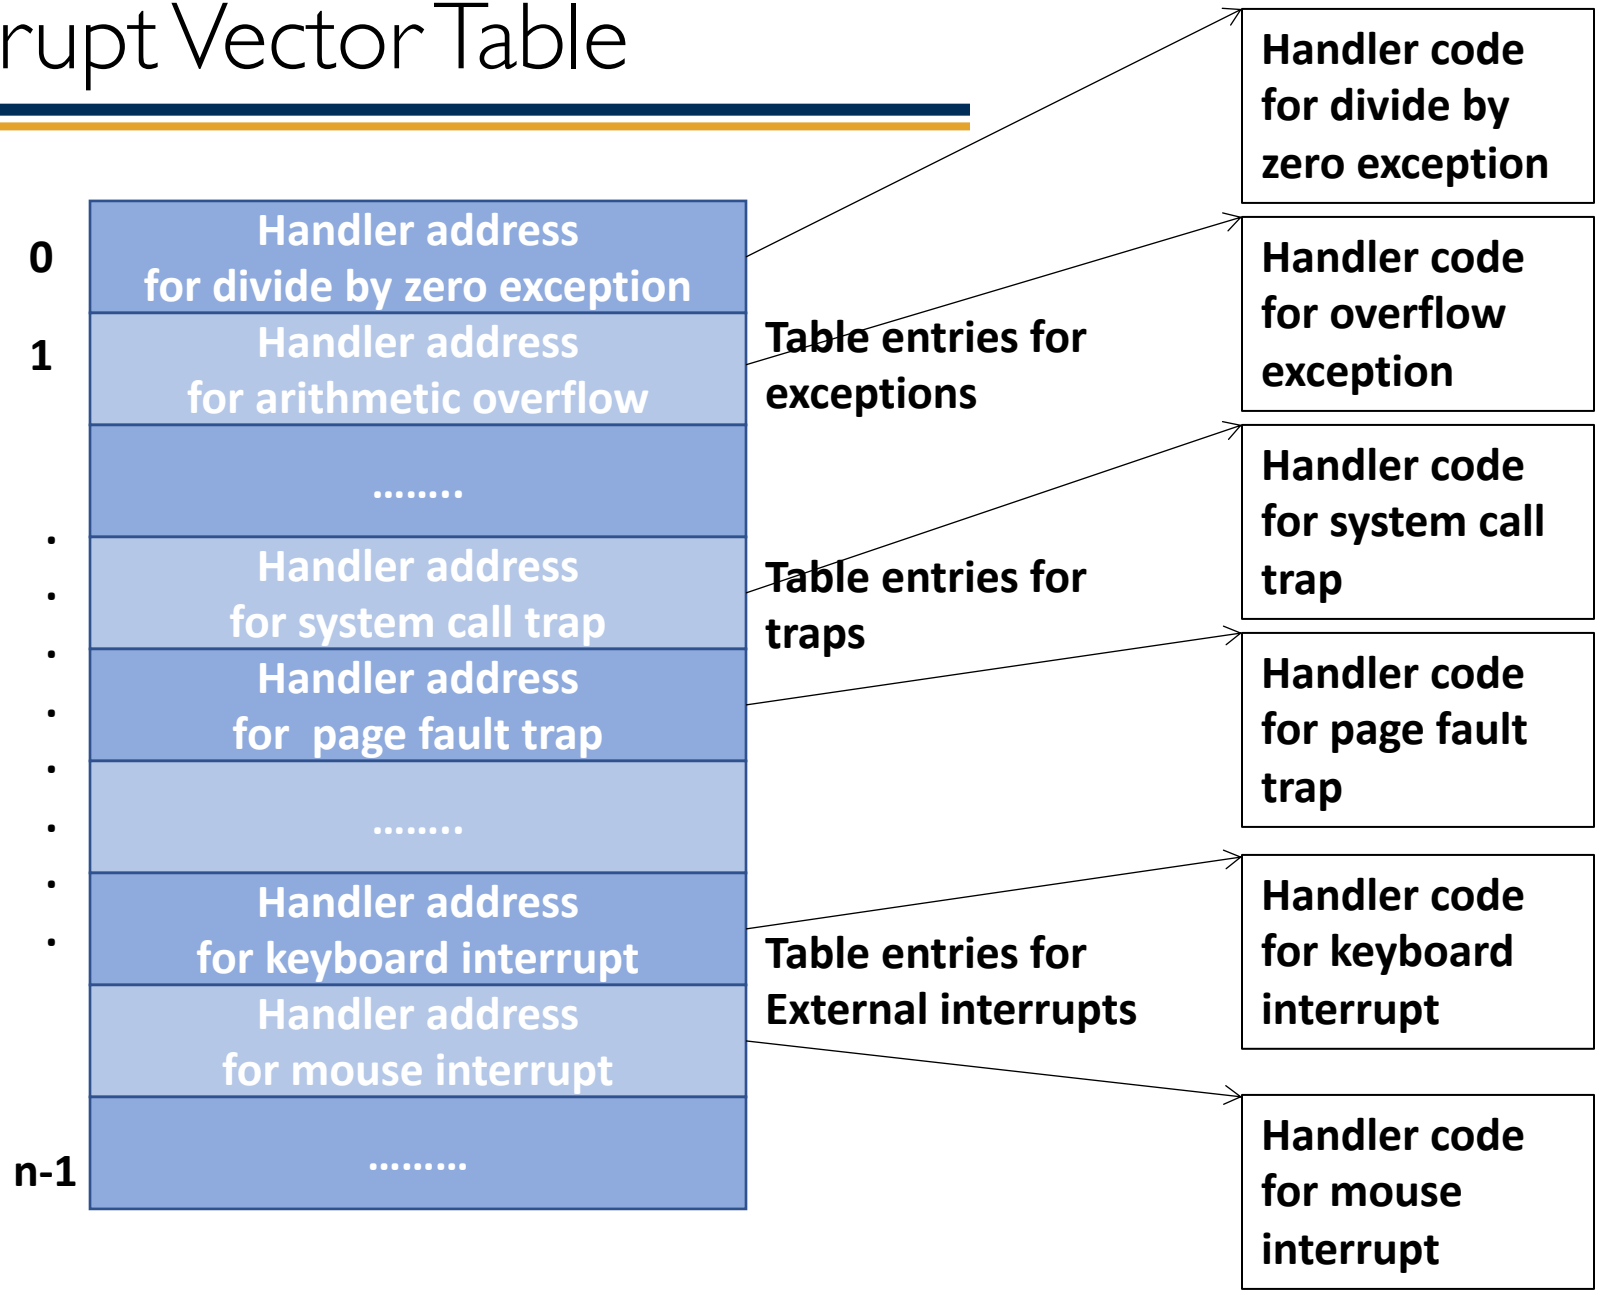

Dealing with program discontinuities

- Can happen anywhere even in the middle of an instruction execution.
- Unplanned for and forced by the hardware. Hardware has to save the program counter since we are jumping to the handler.
- Address of the handler is unknown. Therefore, hardware must manufacture an address.
- Since hardware saved the PC, handler has to discover where to return upon completion.

Architectural enhancements to handle program discontinuities

- When should the processor handle an interrupt?
- How does the processor know there is an interrupt?
- How do we save the return address?
- How do we manufacture the handler address?
- How do we handle multiple cascaded interrupts?
- How do we return from the interrupt?

Modifications to FSM

Where should we take an interrupt?

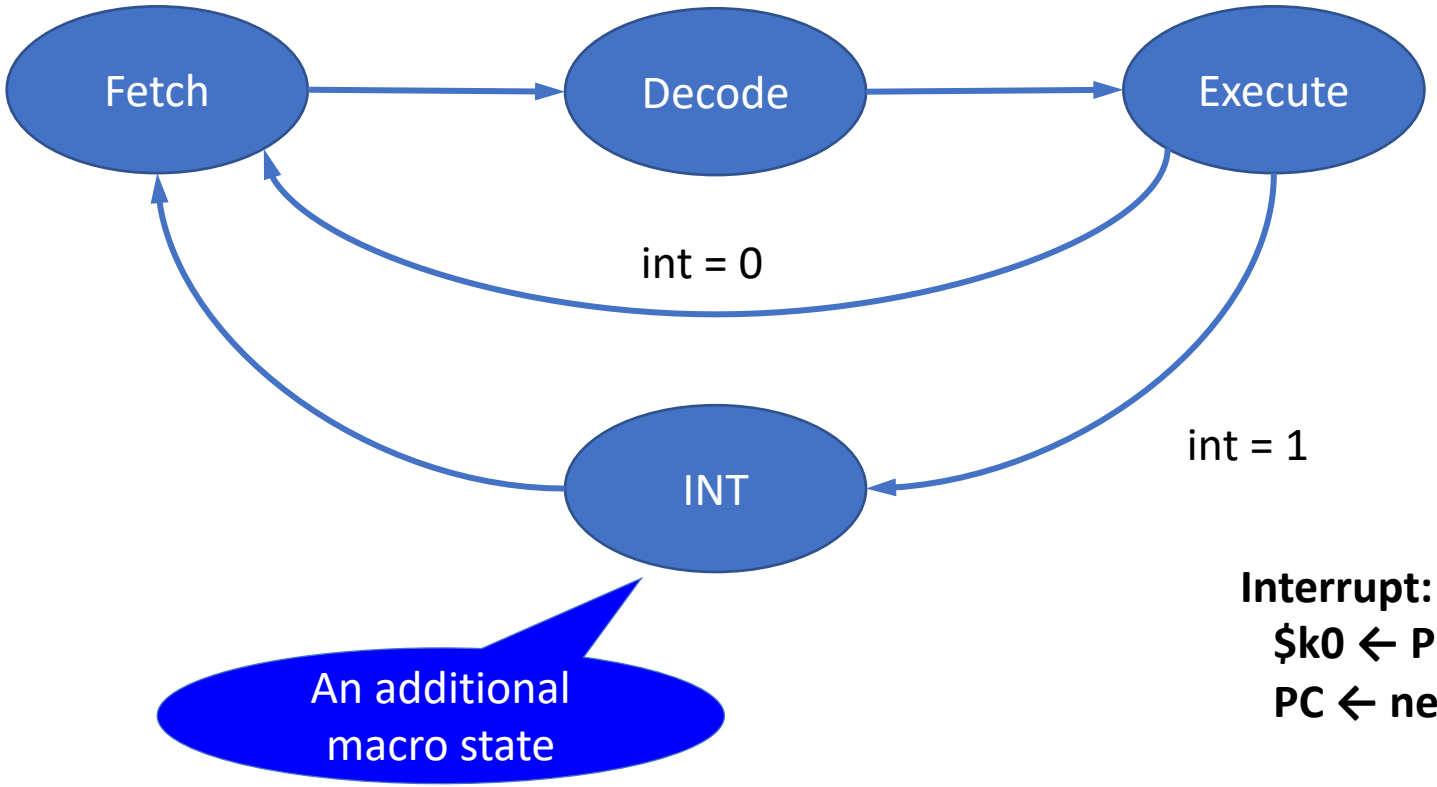

Interrupt:
$k0 ← PC
PC ← new PC

What needs to happen in software?

```
Handler:
  save processor registers;
  execute device code;
  restore processor registers;
  return to original program;
```

That's great, but…

- There are a couple of rubs.
- What happens when an interrupt handler takes an interrupt?

Handling cascaded interrupts

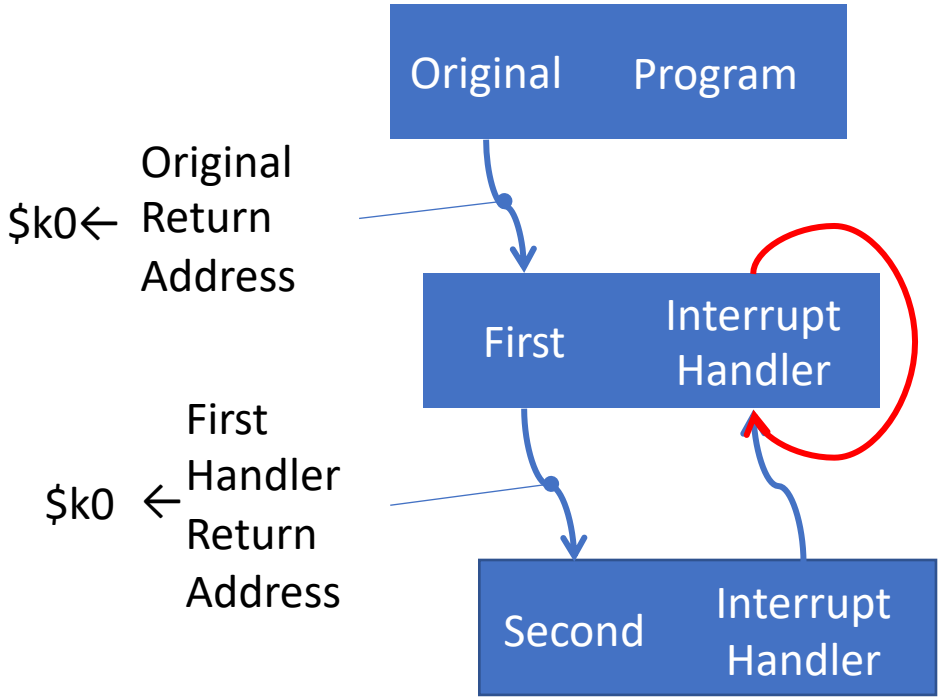

What needs to happen …

```
Handler:
    save processor registers
       (including $k0);
    execute device code;
    restore processor registers
       (including $k0);
    return to original program;
```

That's great, but…

- There are a couple of rubs.
- What happens when an interrupt handler takes an interrupt?
- OK. That's better. Save/restore $k0 in the handler.
- But one more little thing…
- What happens if the second interrupt hits before we save $k0?

What needs to happen …

```
Handler:
    save processor registers
       (including $k0);
    execute device code;
    restore processor registers
       (including $k0);
    return to original program;
```

What if an interrupt happens here?

No problem. We'll save $k0 first

How many instructions does it take to push a register on the stack?

What needs to happen …

```
Handler:
    save processor registers
       (including $k0);
    execute device code;
    restore processor registers
       (including $k0);
    return to original program;
```

Store $k0 first.

It takes
- Decrement $sp
- Store $k0,0($sp)

What if the interrupt happens after the Decrement!?!?

What are we lacking?

- We don't have a way to prevent an interrupt from happening between certain instructions
- In other words, we need for groups of machine instructions to behave **atomically** – i.e. as if they all were executed as a single instruction
- How could we do that?
- We could turn off interrupts between instructions?

The plan

- Create a new processor register, IE, that is 1 when interrupts are enabled
- For an interrupt to be recognized, i.e. for the microcode to advance to the INT macro state, an interrupt must be asserted AND IE must be 1
- In the INT macro state, turn off IE before fetching the first instruction in the handler
- We need two more instructions: EI and DI to respectively set IE to 1 and 0.
- Use EI after pushing $k0 on the stack

Handling cascaded interrupts

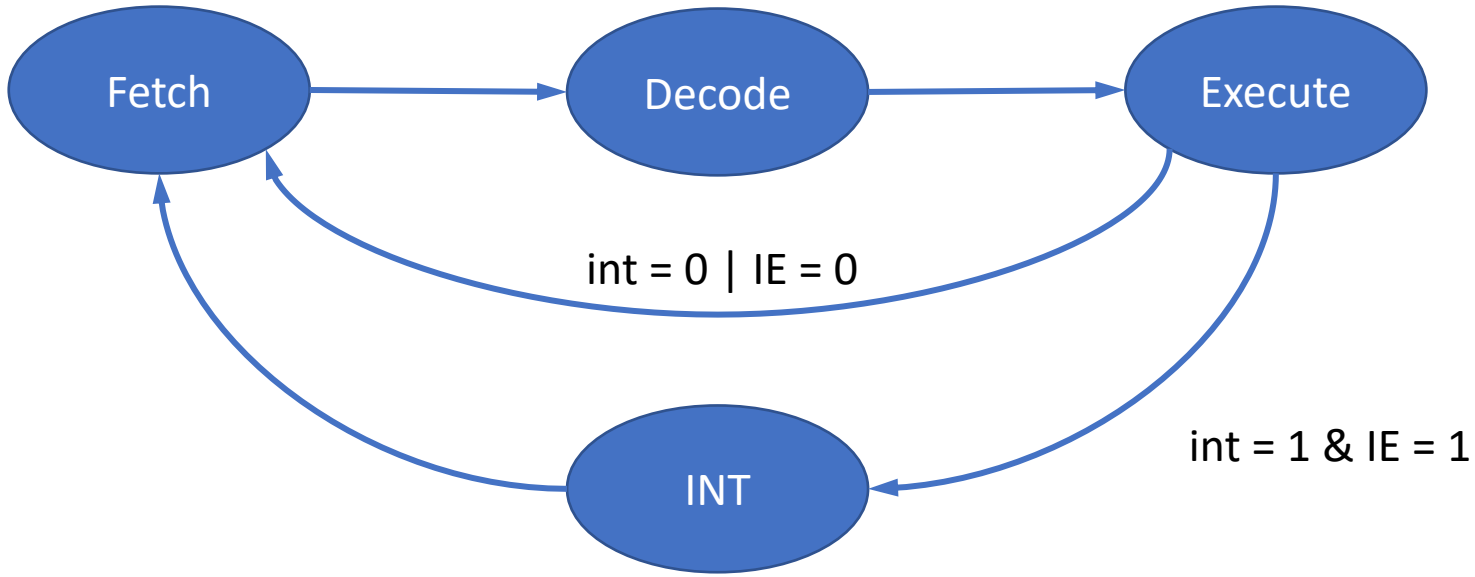

Add 2 new instructions
Enable Ints
Disable Ints

Interrupt:
disable Ints
$k0 ← PC
PC ← new PC

Yay! This will work perfectly!

```
Handler:
    save $k0;
    enable interrupts
    save processor registers;
    execute device code;
    restore processor registers;
    disable interrupts;
    restore $k0;
    enable interrupts;
    return to original program;
```

Or does it?

What if an interrupt occurs here?

Returning from the handler

- Returning involves jumping to the address in $k0 which can be accomplished with

```
jalr $k0, $zero
```

- But as we have just seen an interrupt at precisely the wrong moment would destroy $k0 and cause a failure
- What do we need?
- All this needs to be atomic, too!

```
restore $k0;
enable interrupts
return to original program;
```

Returning from the handler

- So we need another new instruction, RETI
- It atomically enables interrupts and sets the PC to return from the handler

- RETI:
  PC ← $k0
  EI ← 1

Handling cascaded interrupts

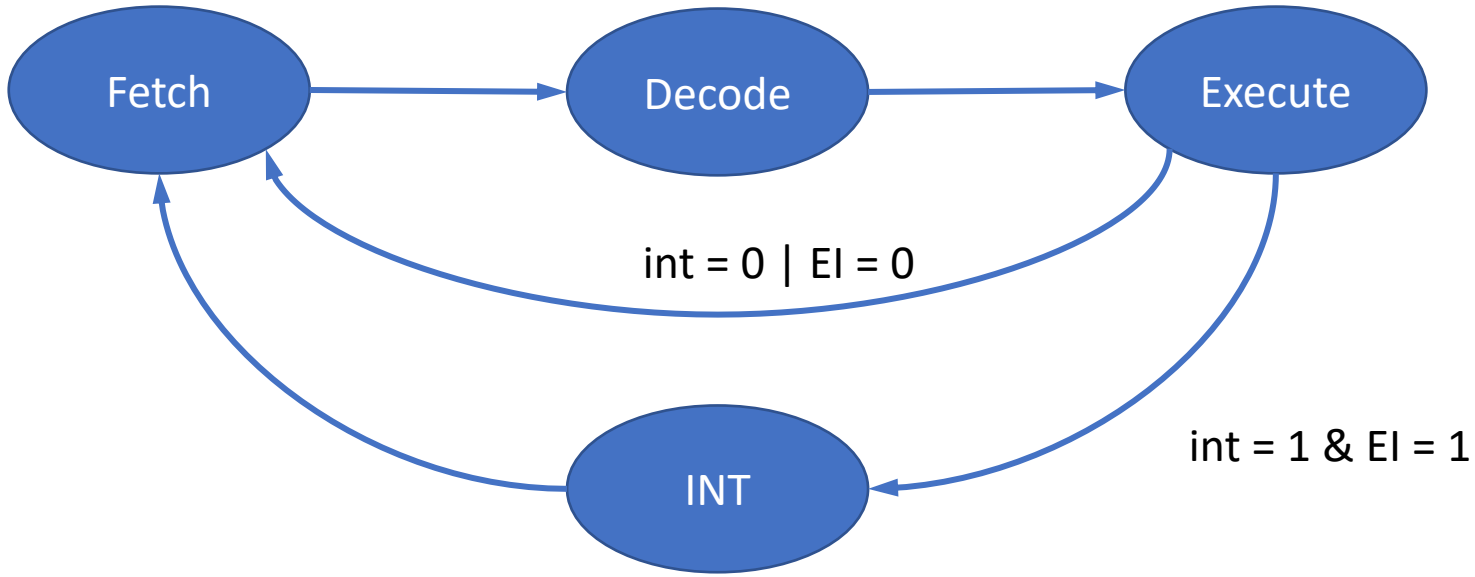

Add 3 new instructions
Enable Ints
Disable Ints
Return from interrupt

Interrupt:
Disable Ints
$k0 ← PC
PC ← new PC

Summary of architectural enhancements to LC-2200 to handle interrupts (so far)

- Three new instructions to LC-2200:
  - Enable interrupts
  - Disable interrupts
  - Return from interrupt

- Upon an interrupt, store the current PC implicitly into a special register $k0, disable interrupts, and set the PC to the address of the handler

- Upon returning from an interrupt (RETI), store $k0 into the PC and enable interrupts.

Session ID: dag2200

What instructions do we use to save $k0 on the system stack?

A.
```
addi    $sp,$sp,-1
sw      $k0,0($sp)
```

B.
```
sw      $k0,0($sp)
addi    $sp,$sp,-1
```

C.
```
sw      $k0,-1($sp)
addi    $sp,$sp,-1
```

D.
```
jalr    save,$ra
```

E.
```
nand    $r0,$r0,$zero
addi    $r0,$r0,1
add     $sp,$sp,$ro
sw      $k0,0($sp)
```

Why does the interrupt handler…

… disable and reenable interrupts during its execution?

A. To prevent overheating the interrupt data path
B. The device I/O operations won't function properly if interrupts are not disabled in the handler
C. It prevents the saving and restoring of all the processor registers from being interrupted.
D. It allows the interrupt handler to be interrupted
E. None of the above

Hardware details for handling external interrupts

- What we have presented thus far is what is required for interrupts, traps and exceptions
- What do we need specifically for external interrupts?

Wiring for external interrupts

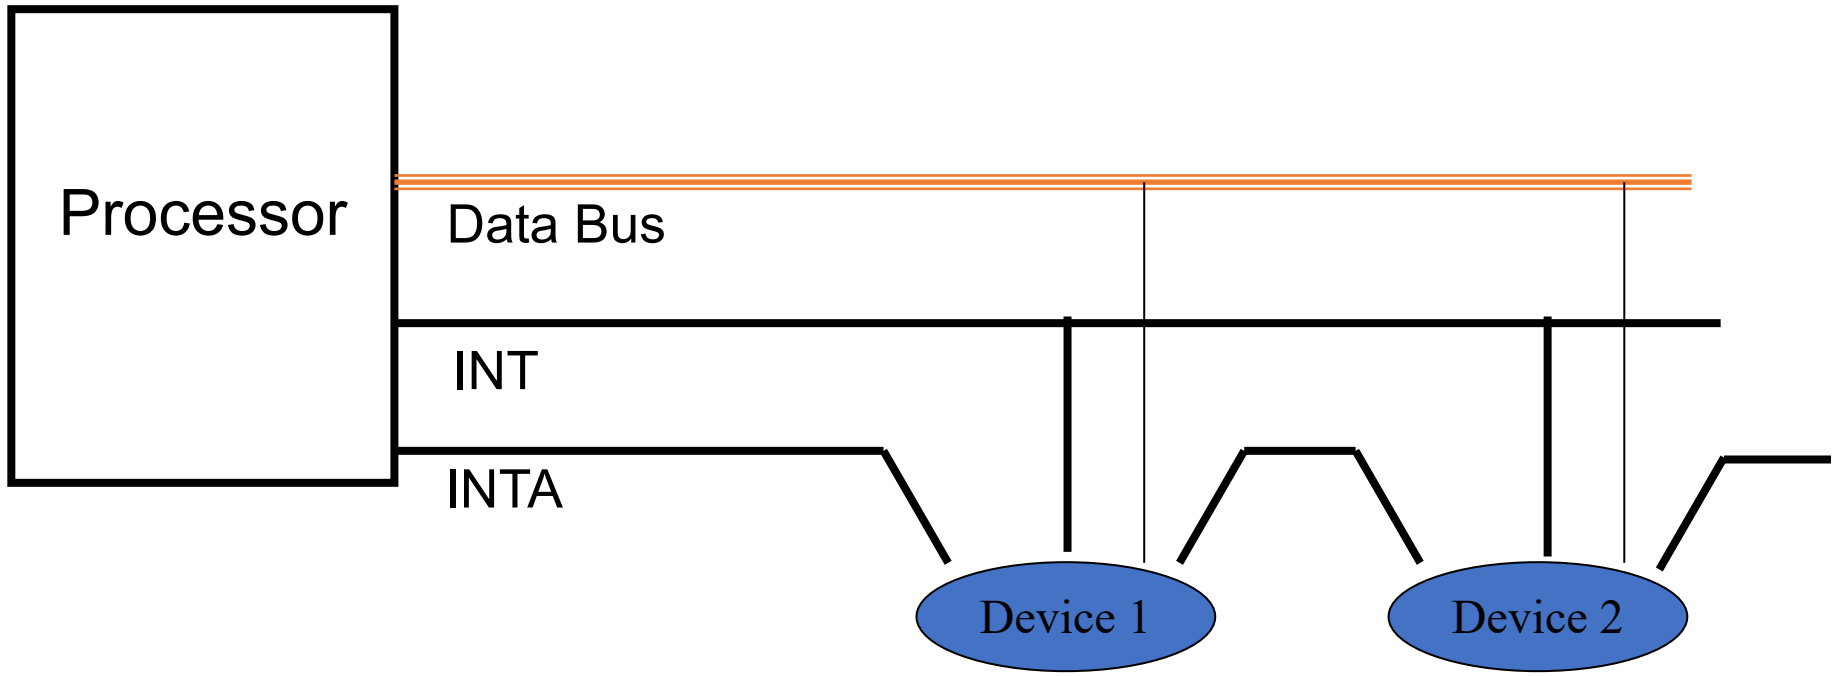

What happens at an interrupt?

- Device asserts the INT bus (it's wired so multiple devices can do this simultaneously)
- At the completion of the current instruction, CPU sees INT signal and microcode cycles into the INT macro state
- Microcode raises the INTA signal line
- Devices pass-through the INTA signal if they are not interrupting; otherwise the first interrupting device asserts its ID on the data bus
- Microcode reads the data bus and uses the ID as an index to determine which entry in the IVT to use to set the PC

Multiple interrupt priority levels

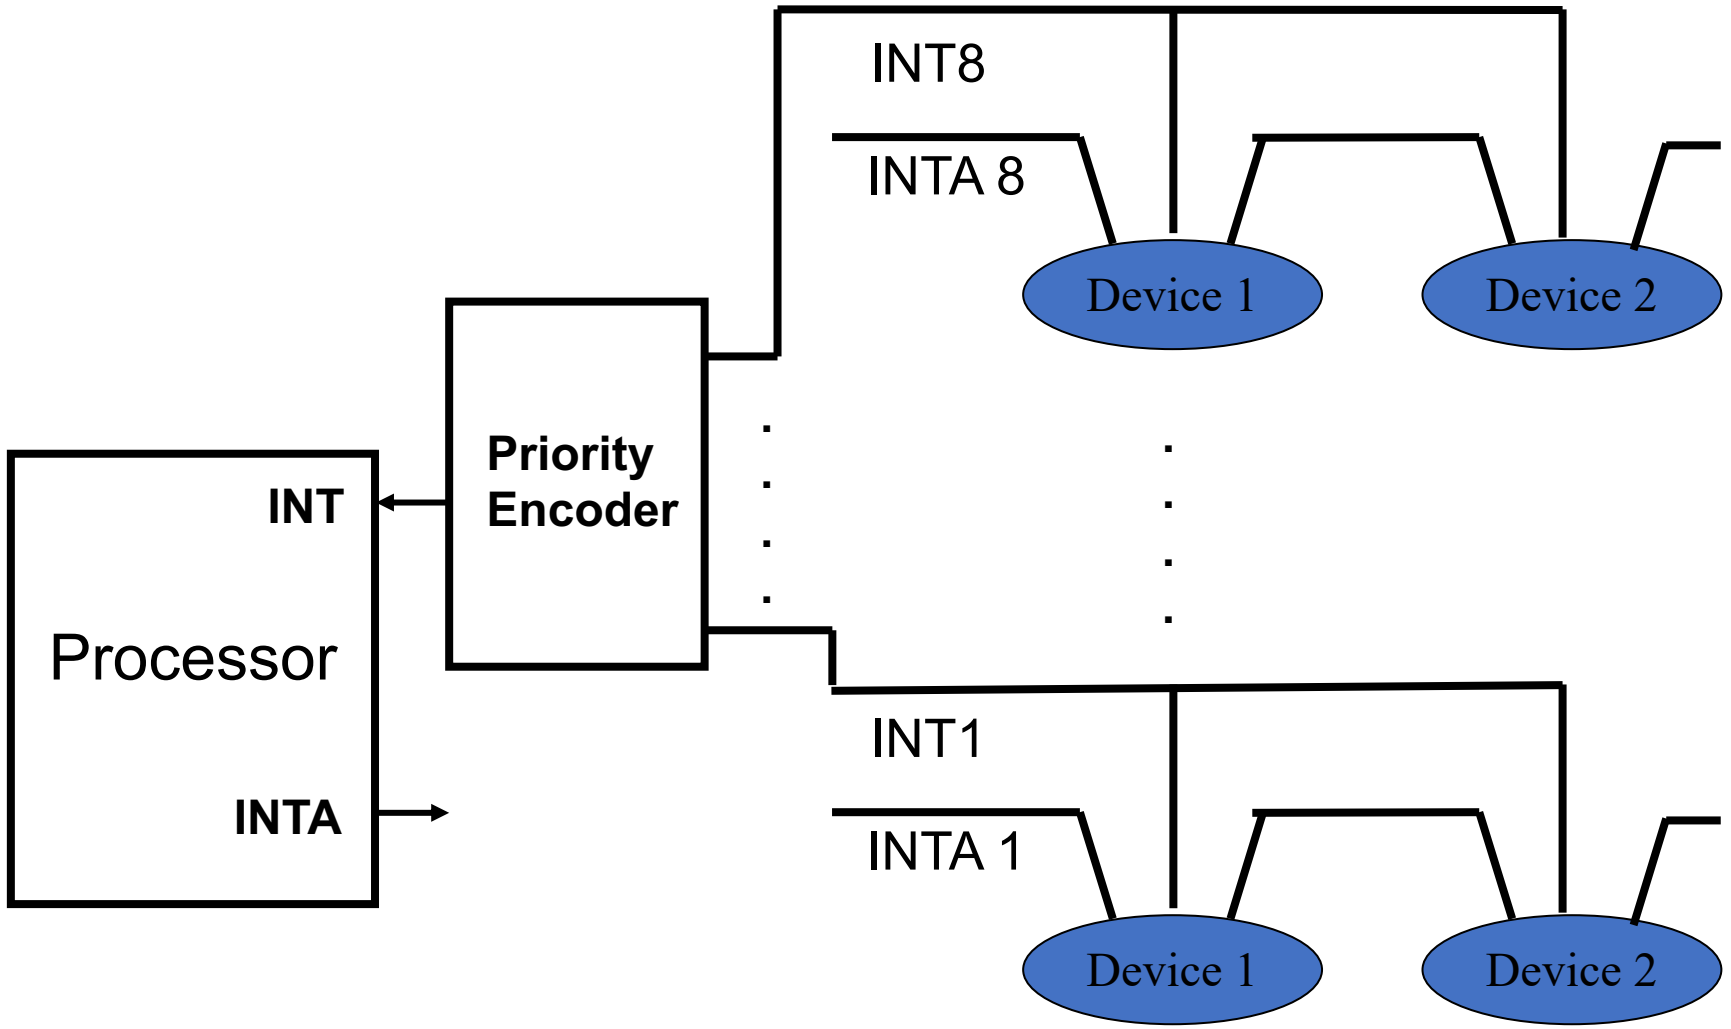

Priority encoder

- A priority encoder takes 2^n inputs and produces a 1-bit INT output and an n-bit ID output.
- If any of the input lines is high, the PE asserts the INT output
- The PE asserts the encoded value of the first high input line onto the ID output
  - E.g. if input 5 and 7 are high on a 3-bit PE, then it asserts INT and ID=101
  - If only input 7 is high, then it asserts INT and ID=111

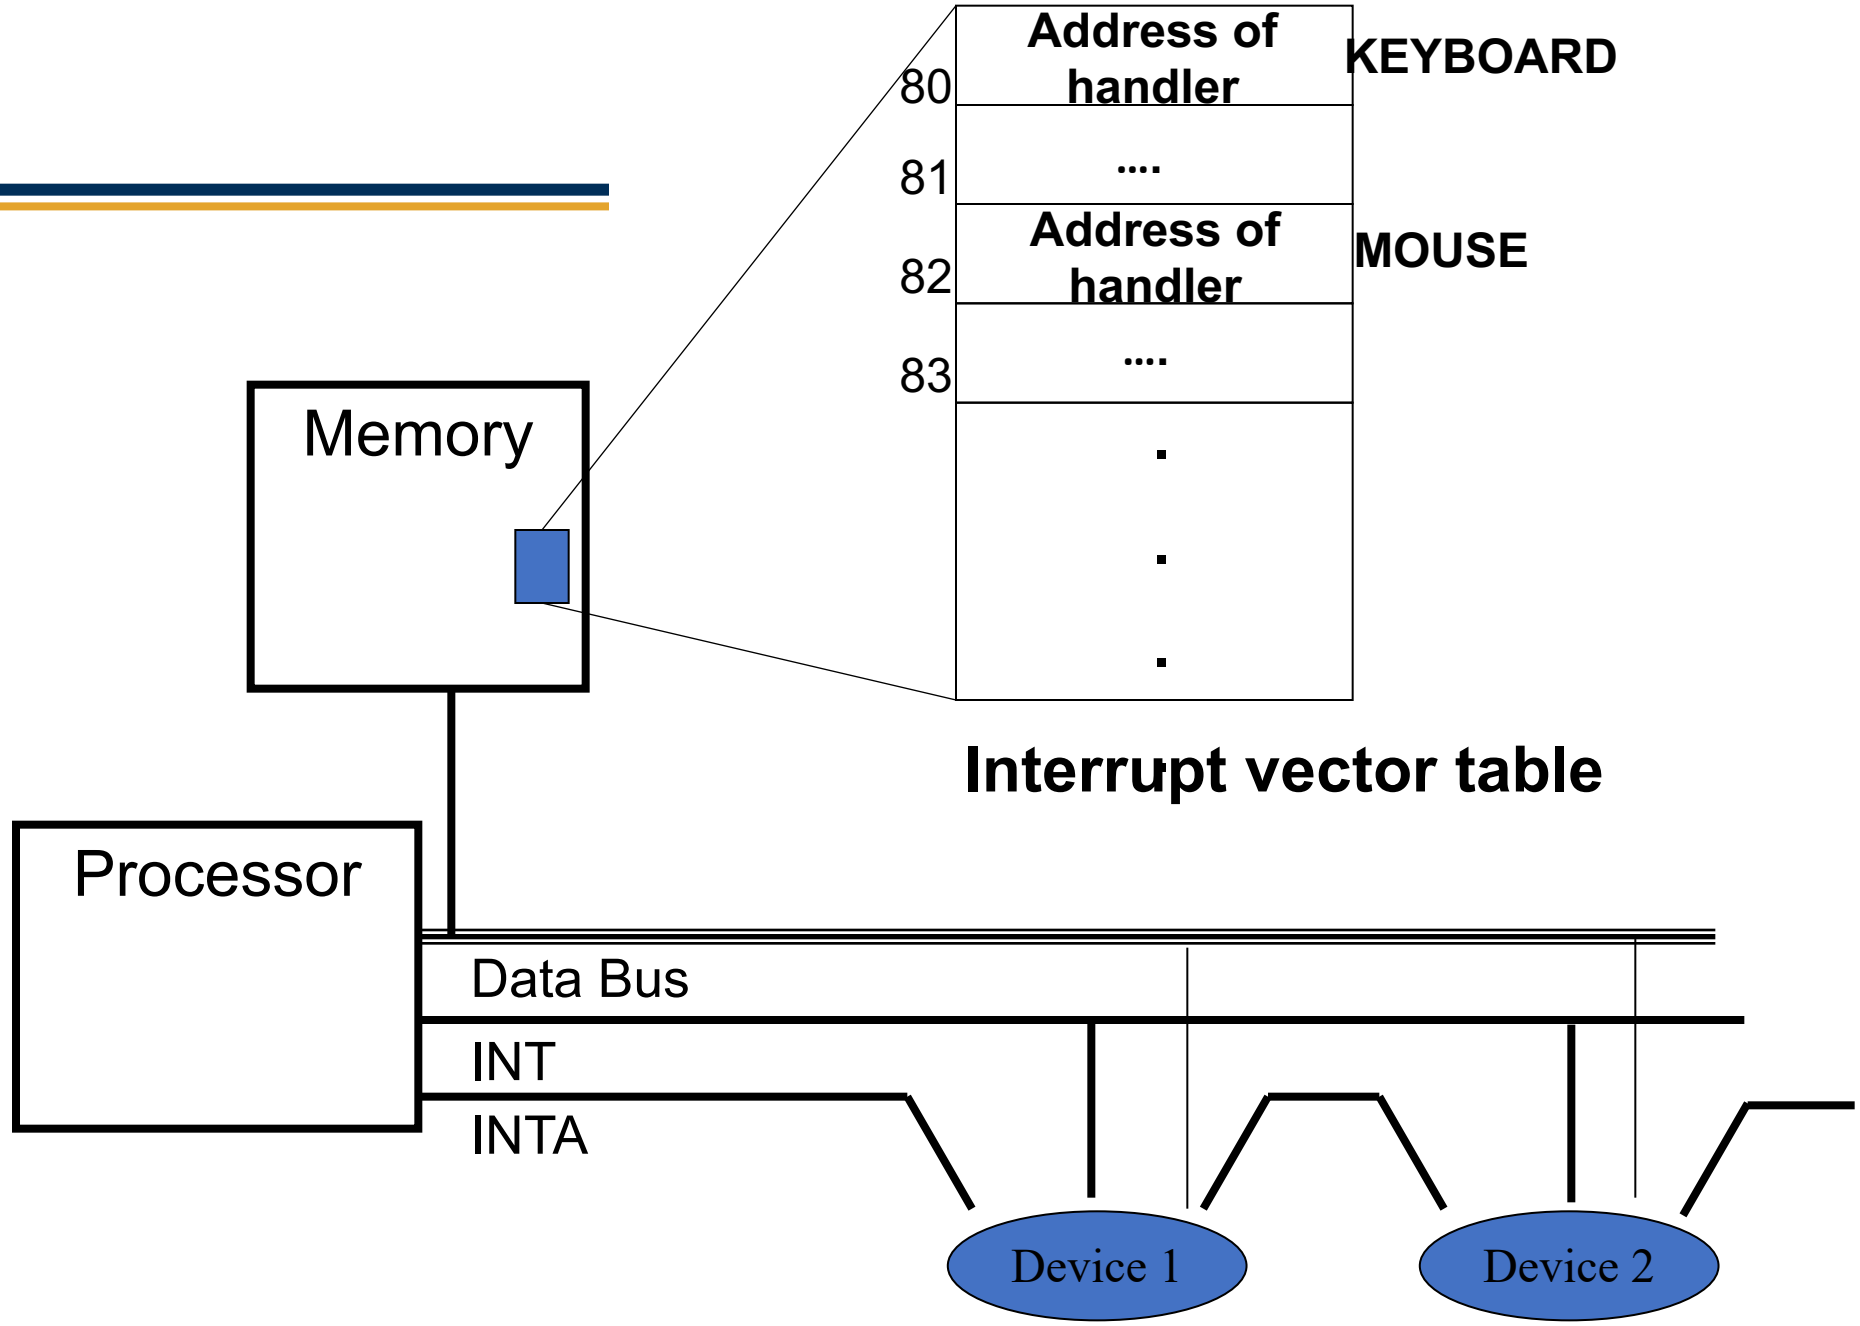
Address of handler
KEYBOARD
80
81
....
Address of handler
MOUSE
82
83
....
Memory
Interrupt vector table
Processor
Data Bus
INT
INTA
Device 1
Device 2

Where to save/restore CPU registers in the interrupt handler

- The user stack?
- Bad idea. The user doesn't even have to set $sp if he doesn't feel like it. Bad practice, but real possibility.
- Where, then?
- How about we let the OS have a system stack that we know is handled properly?

Stack for saving/restoring

- Hardware has no guarantee for stack behavior by user program (register/conventions)
- Equip processor with 2 saved stack pointers (User/System)
- On interrupt, save user stack pointer from $sp and restore the system stack pointer to $sp
- We'll need two more registers, USP and SSP

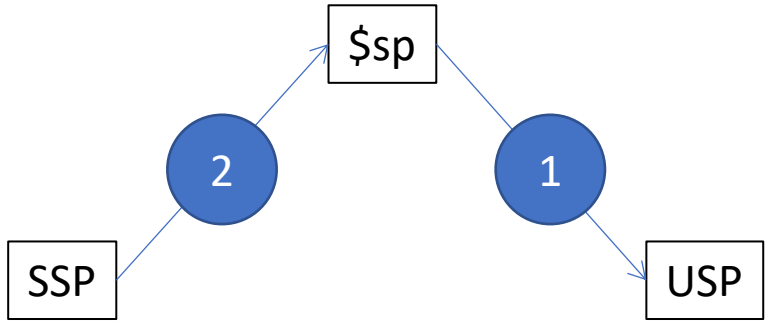

Stack for saving/restoring

- Use system stack for saving all necessary information
- Upon completion of interrupt restore registers, etc.
- The restore user stack pointer by reversing earlier swap
- Keep a user/kernel mode flag to record whether we're using the user or kernel stack

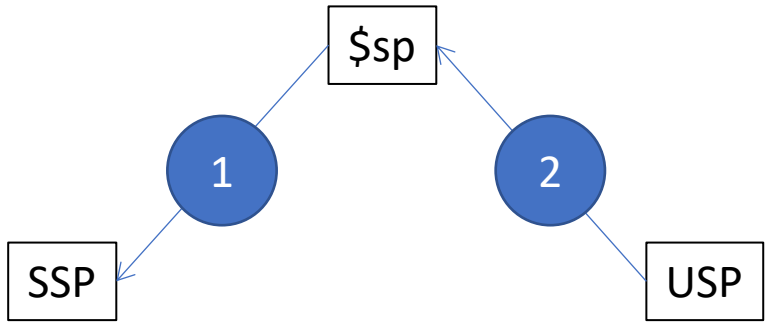

Stacks and modes during interrupts

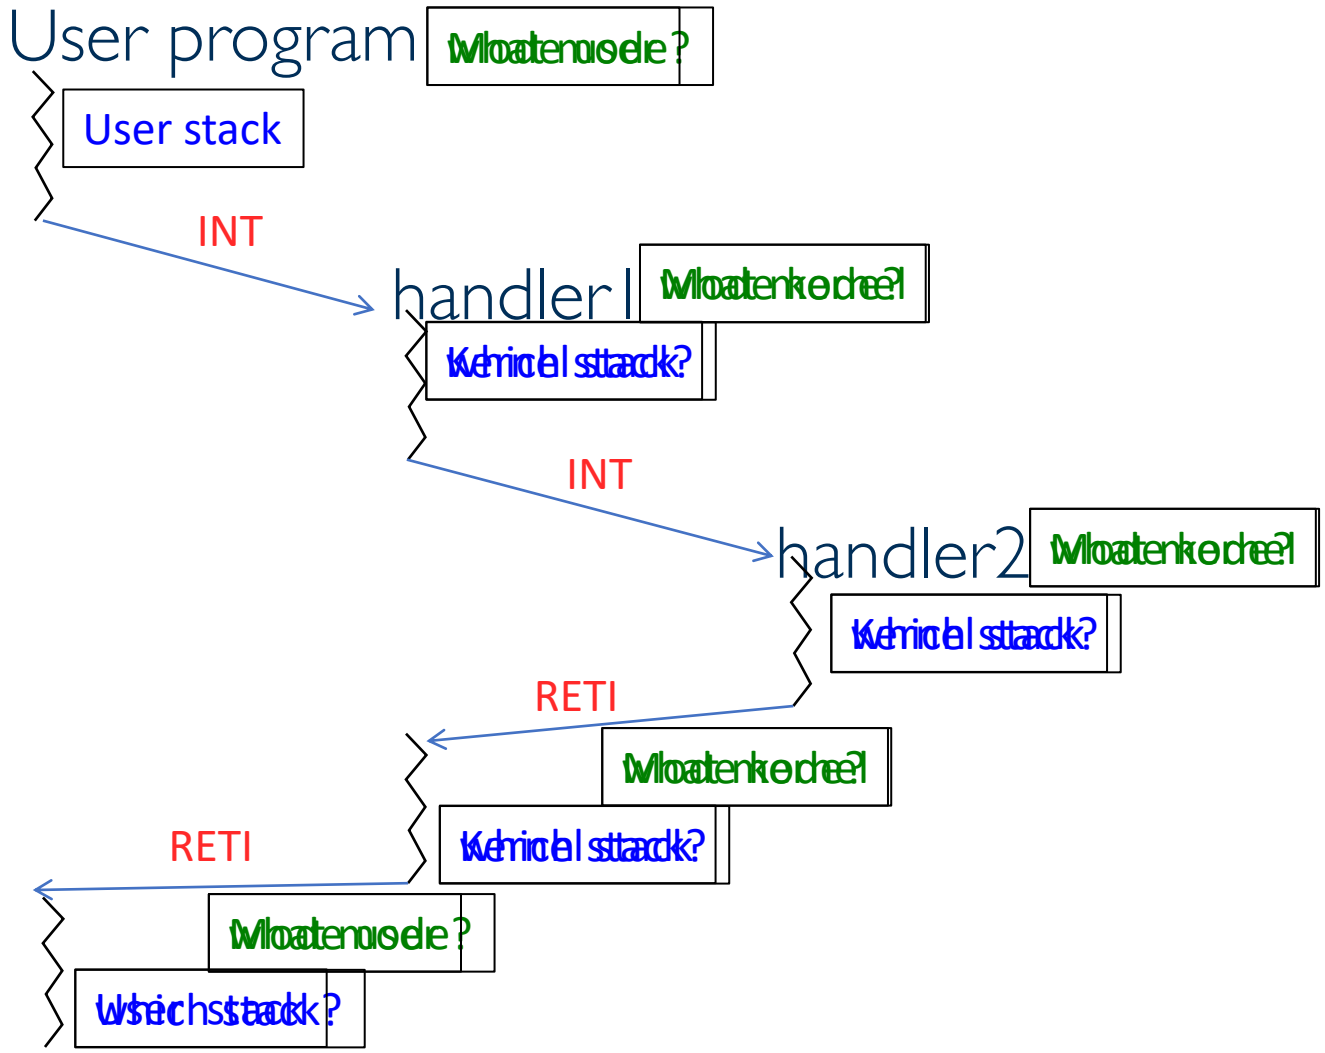

Summary of interrupt actions

INT macro state:

$k0 ← PC
Assert INTA to acknowledge interrupt
Receive IV (interrupt vector) from the device on the data bus
PC ← Mem[IV]
if user mode,
    USP ← $sp; $sp ← SSP
Push mode on stack
mode ← kernel
Disable interrupts

RETI instruction:

PC ← $k0
Pop mode from system stack
if user mode,
    SSP ← $sp; $sp ← USP
Enable interrupts

A working interrupt handler

```
Handler:
  // handler starts with interrupts disabled
  push $k0 onto system stack;
  enable interrupts;
  save processor registers to system stack;
  execute device code;
  restore processor registers from system stack;
  disable interrupts;
  pop $k0 from system stack;
  // handler ends with interrupts disabled
  return to original program using RETI;
```

Architecture enhancements to LC-2200 for interrupts

1. An interrupt vector table (IVT), to be initialized by the operating system with handler addresses.
2. An exception/trap register (ETR) that contains the vector for internally generated exceptions and traps.
3. A Hardware mechanism for receiving the vector for an externally generated interrupt.
4. User/kernel mode and associated mode bit in the processor.
5. User/system stack corresponding to the mode bit.
6. A hardware mechanism for storing the current PC implicitly into a special register $k0, upon an interrupt, and for retrieving the handler address from the IVT using the vector (either internally generated or received from the external device).
7. Three new instructions to LC-2200:
   Enable interrupts
   Disable interrupts
   Return from interrupt

Putting it all together

Executing instruction at 19999. The PC has already been incremented. Device signals interrupt in middle of instruction. $sp points to user stack

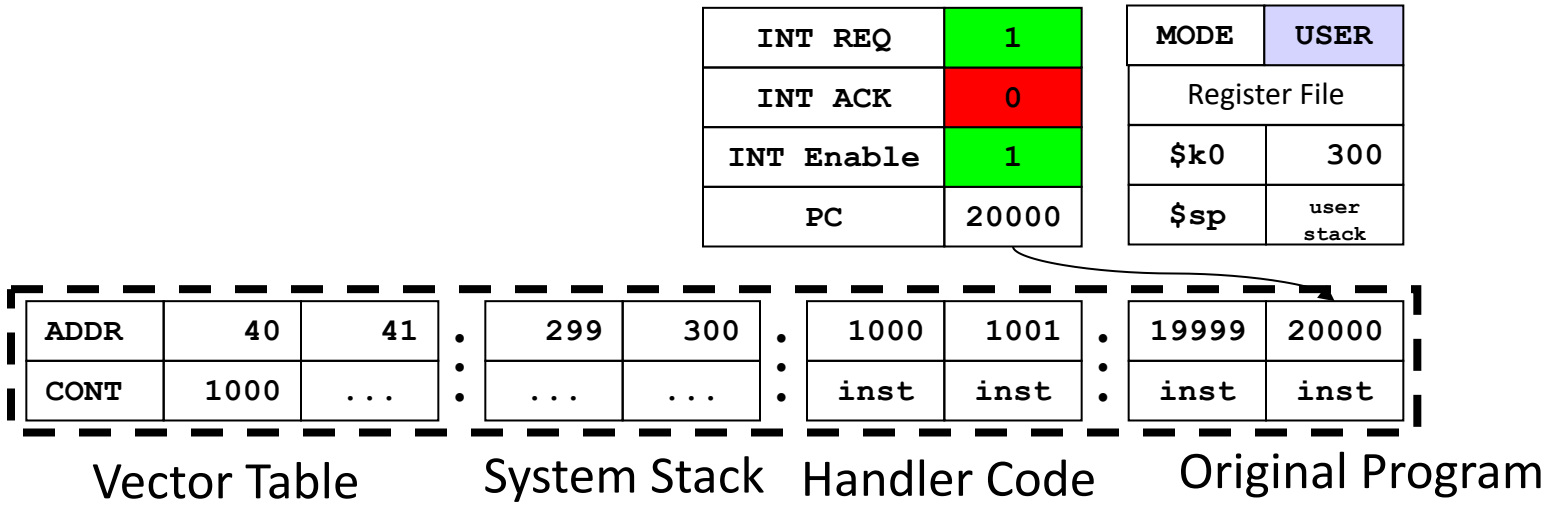

Putting it all together

Interrupt has been noticed.
$k0 gets PC.
Interrupts are disabled.
Interrupt is acknowledged.
Device puts vector on bus.

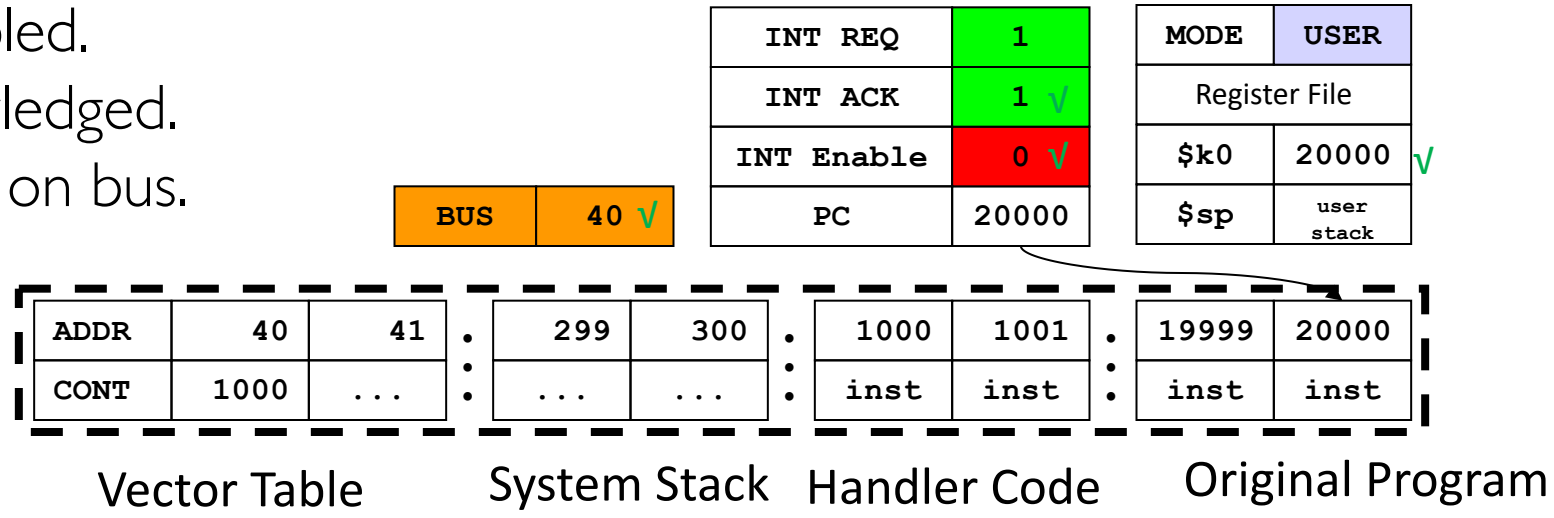

Putting it all together

Handler address is put into PC
$sp now points to system stack;
Current mode is saved in system stack;
New mode is set to kernel;
Interrupt code at 1000 will handle the interrupt.

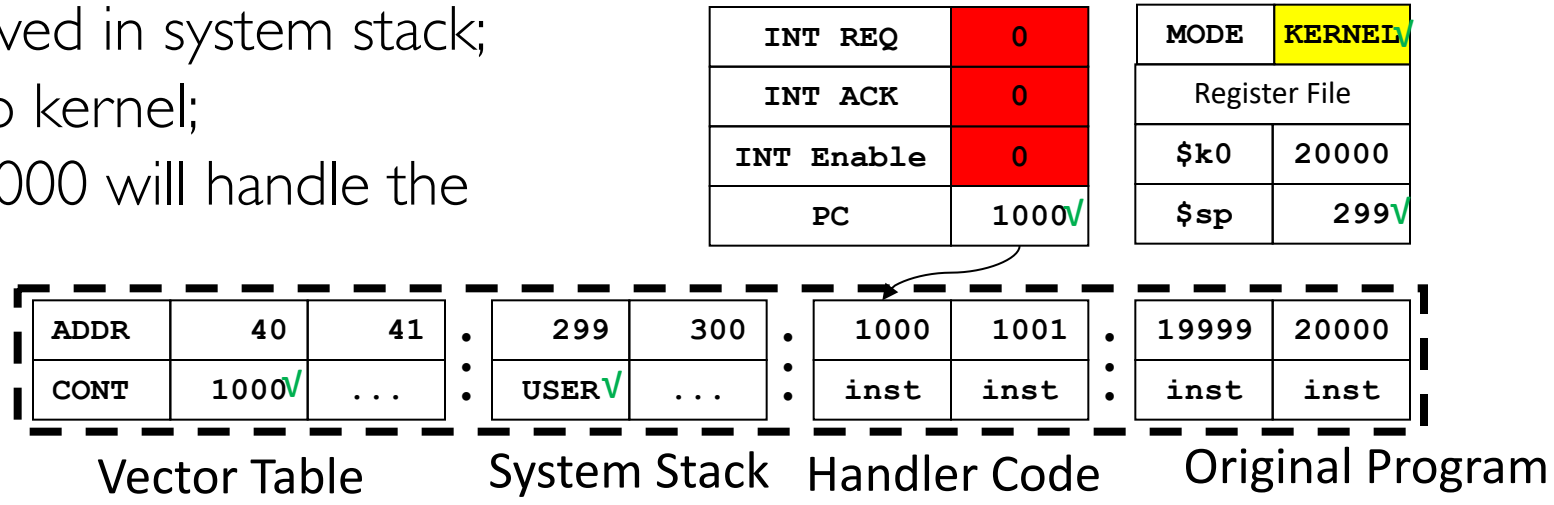

Putting it all together

Handler completes.
RETI instruction restores mode from system stack; since returning to user program in this example, sets Mode to User;
$sp now points to user stack;
copies $k0 into PC;
re-enables interrupts

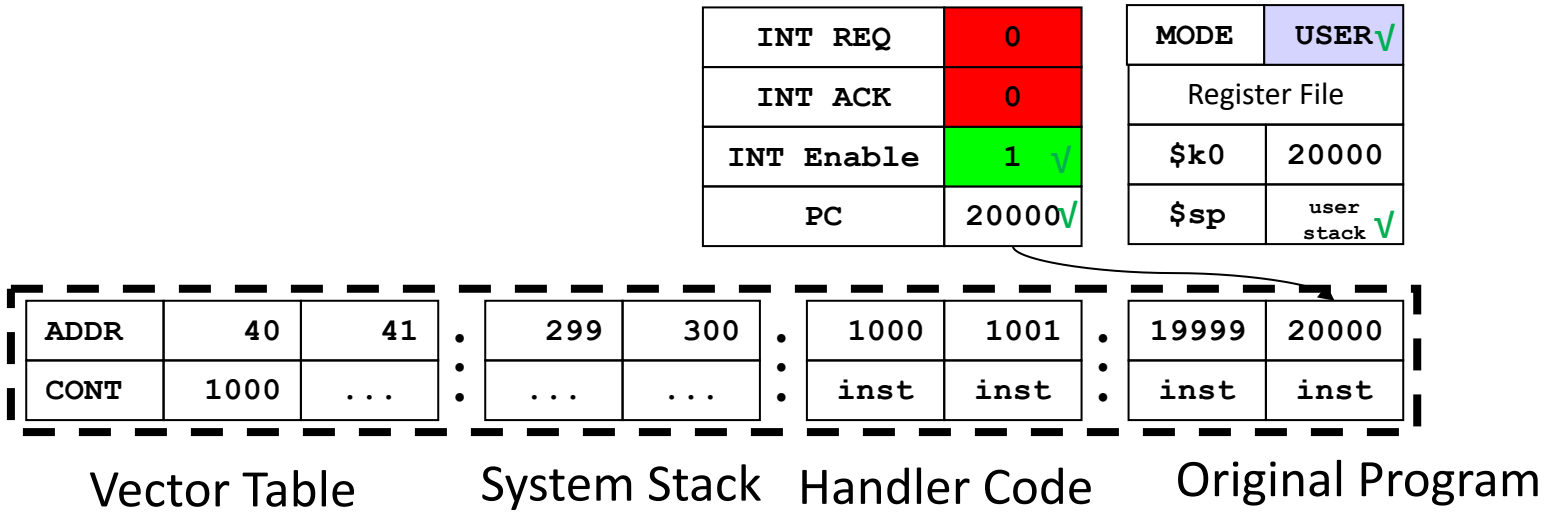

Summary

- Interrupts help a processor communicate with the outside world.
- An interrupt is a specific instance of program discontinuity.
- Processor/Bus enhancements included
  - Three new instructions
  - User stack and system stack pointers
  - Mode bit
  - INT macro state
  - Control lines called INT and INTA

Summary

- Software mechanism needed to handle interrupts; traps and exceptions are similar.
- Discussed how to write a generic interrupt handler that can handle nested interrupts.
- Intentionally simplified. Interrupt mechanisms in modern processors are considerably more complex. For example, modern processors categorize interrupts into two groups: *maskable* and *non-maskable*.
  - maskable: Interrupts that can be temporarily turned off
  - Non-maskable: Interrupts that cannot be turned off

Summary

- Presented mode as a characterization of the internal state of a processor. Intentionally simplistic view.
- Processor state may have a number of other attributes available as discrete bits of information (similar to the mode bit).
  - Modern processors aggregate all of these bits into one register called *processor status word (PSW)*.
  - Upon an interrupt and its return, the hardware implicitly pushes and pops, respectively, both the PC and the PSW on the system stack.
- The interested reader is referred to more advanced textbooks on computer architecture for details on how the interrupt architecture is implemented in modern processors.

Summary

- Presented simple treatment of the interrupt handler code to understand what needs to be done in the processor architecture to deal with interrupts. The handler would typically do a lot more than save processor registers.

- LC-2200 designates a register $k0 for saving PC in the INT macro state. In modern processors, there is no need for this since the hardware automatically saves the PC on the system stack.