

CS2200
Systems and Networks
Spring 2022

Lecture 9: Performance part deux

Alexandros (Alex) Daglis
School of Computer Science
Georgia Institute of Technology
adaglis@gatech.edu

Announcements

- Sign up for Project I demo (deadline Saturday)
- Any necessary rescheduling must be requested **24h in advance**
- Missing a demo without valid excuse incurs penalty (see Syllabus)

Please respect the TAs' time and schedule!

- Don't miss your demo slot
- Don't be late for your demo slot (arrive 5min early)

The “Iron Law” of Processor Performance

$$\text{Processor Performance} = \frac{\text{Time}}{\text{Program}}$$

$$= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

ISA & Compiler

Microarchitecture

Circuit

Improvement in execution speed

It takes me 8 minutes to walk to class from my office. I can run twice as fast as I walk. If I walk half the distance and run the remaining half, how much time will I take to reach class from my office?

$$8 = d/v$$

$$n = \frac{1}{2} d/v + \frac{1}{2} d/2v$$

$$8v = d$$

$$n = 4v/v + 4v/2v$$

$$n = 4 + 2$$

$$n = 6$$

Amdahl's Law

Amdahl's Law

$$\text{Time}_{\text{after}} = \frac{\text{Time}_{\text{affected}}}{\text{Amount of Improvement}} + \text{Time}_{\text{unaffected}}$$

My office walk: $6 = 4 / 2 + 4$

Improving an instruction

A processor spends 20% of its time on ADD instructions. An engineer proposes to improve the speed of the ADD instruction by 4 times. What is the speedup achieved by this modification?

The improvement only applies for the ADD instruction, so 80% of the execution time is unaffected by the improvement.

Original normalized execution time = 1

$$\begin{aligned}\text{New execution time} &= (\text{time spent in ADD}/4) + \text{remaining time} \\ &= 0.2 / 4 + .8 \\ &= 0.85\end{aligned}$$

$$\begin{aligned}\text{Speedup} &= \text{execution time before} / \text{execution time after} \\ &= 1 / 0.85 = 1.18 = 18\%\end{aligned}$$



Improving an instruction

An engineer is asked to improve the processor's overall performance by 2 times by optimizing the ADD instruction. The processor spends 20% of its time on ADD instructions. How much faster must the ADD instruction become?

- A. 2x
- B. 10x
- C. 100x
- D. That's impossible!!

Execution time?

CPI of Instruction Classes	Code 1	Code 2
R-type = 2	3	10
I-type = 10	3	1
J-type = 3	5	2
S-type = 4	2	3

13 inst
6+30
+15+8 =
59 cycles

16 inst
20+10+6
+12 =
48 cycles

- A. Code 1 is better than Code 2 since it has fewer instructions
- B. Code 2 is better than Code 1 since it has fewer instructions
- C. Code 1 is better than Code 2 since it takes fewer total clock cycles to execute
- D. Code 2 is better than Code 1 since it takes fewer total clock cycles to execute

Architecture change?

- We have a computer with three types of instructions that have the following CPI:

Type	CPI
A	2
B	5
C	1

- An architect determines that she can reduce the CPI for B to 3 but will need to slow the clock speed of the processor. What is the maximum permissible slowing of the clock that will make this change worthwhile?
- Assume that all the workloads for this processor use 30% of A, 10% of B, and 60% of C types of instructions

How do we answer that?

Execution time of the old machine:

$$ET_o = N * (F_{Ao} * CPI_{Ao} + F_B * CPI_{Bo} + F_C * CPI_{Co}) * C_o$$

(where F_x and CPI_x are the dynamic frequencies and CPIs of each type of instruction, respectively)

$$ET_o = N * (0.3 * 2 + 0.1 * 5 + 0.6 * 1) * C_o = N * 1.7 * C_o$$

Execution time for the new machine:

$$ET_n = N * (0.3 * 2 + 0.1 * 3 + 0.6 * 1) * C_n = N * 1.5 * C_n$$

For the design to be viable,

$$ET_n < ET_o$$

$$N * 1.5 * C_n < N * 1.7 * C_o$$

$$C_n < 1.7/1.5 * C_o$$

$$C_n < 1.13 * C_o$$

Maximum permissible increase in clock cycle time = 13%

Type	CPI
A	2
B	5
C	1

30% of A,
10% of B,
60% of C

Combining two instructions?

Instruction	CPI
Add	2
Shift	3
Others	2
Add/Shift	4

If the sequence ADD followed by SHIFT appears in 20% of the dynamic frequency of a program, what is the speedup of the program with all {ADD, SHIFT} replaced by the combined instruction?

[HINT: For every 10 instructions in the original program, 2 instructions are the ADD/SHIFT combo. Thus the number of instructions in the new program shrinks to 90% of the original program.]

A solution

$$\begin{aligned} ET_o &= N * (F_{ADD} * 2 + F_{SHIFT} * 3 + F_{others} * 2) * C \\ &= N * (0.1 * 2 + 0.1 * 3 + 0.8 * 2) * C \\ &= 2.1 * N * C \end{aligned}$$

With the combo instruction replacing {ADD SHIFT}, the number of instructions in the new program shrinks to 0.9N in the new program. The frequency of the combo instruction is 1/9 and the other instructions are 8/9.

$$\begin{aligned} ET_n &= (0.9 * N) * (F_{COMBO} * 4 + F_{others} * 2) * C \\ &= (0.9 * N) * (1/9 * 4 + 8/9 * 2) * C \\ &= 2 * N * C \end{aligned}$$

$$\begin{aligned} \text{Speedup} &= \text{old execution time} / \text{new execution time} \\ &= 2.1NC / 2NC \\ &= 1.05 \end{aligned}$$

Instruction	CPI
Add	2
Shift	3
Others	2
Add/Shift	4

Benchmarks

- **Benchmarks** are a set of programs that are **representative** of the workload for a processor.
- The key difficulty is to be sure that the benchmark program selected **really** are representative of the prospective workload.
- Standard benchmark suites (SPEC, LINPACK, Whetstone, Dhrystone, and many more) facilitate “apples-to-apples” comparisons by summarizing performance *across a set of programs*
 - E.g., SPEC uses perl, gcc, AI apps, compression, imaging apps, modeling apps, etc. to represent a common workload
- A radical new design is hard to benchmark because there may not yet be a compiler or much code.

Using a benchmark

Some caveats:

- Testing a single system component (e.g., the processor) only gives a limited view: e.g., memory organization and memory—processor—bus bandwidths are also key
- Some processors do well on certain benchmark programs and other do well on other programs
- A composite index can be useful when we want to compare two processors without knowing the exact kind of workload they are going to run, but composite functions come with caveats, so we must be very cautious.
 - More on this soon

Reasons to be skeptical of a benchmark

- The vendor gave you the benchmark results (i.e., conflict of interest)
- The vendor wrote the benchmark suite
- The benchmark suite doesn't seem to have any elements that represent your workload (e.g., you run web server farms and the benchmark represents only computationally intensive scientific calculations)
- The equipment being benchmarked is different from the equipment you want to evaluate (maybe a little different, maybe a lot different)
- The benchmark uses a different compiler suite than you plan to use
- The cost of mistakenly choosing the wrong equipment is very high



#Dieselgate

Comparing Multiple Programs

	Computer A	Computer B	Computer C
Program 1 (secs)	1	10	20
Program 2 (secs)	1000	100	20
Program 3 (secs)	1001	110	40

A is 10 times faster than B for program 1
B is 10 times faster than A for program 2
A is 20 times faster than C for program 1
C is 50 times faster than A for program 2
B is 2 times faster than C for program 1
C is 5 times faster than B for program 2

Each statement above is correct...

...but I just want to know which machine is the best?

Need a composite metric



Let's Try a Simpler Example

Two machines timed on two benchmarks

	<u>Machine A</u>	<u>Machine B</u>
Program 1	2 seconds	4 seconds
Program 2	12 seconds	8 seconds

How much faster is Machine A than Machine B?

Attempt 1: ratio of runtimes, normalized to Machine A runtimes

program1: $4/2$

program2 : $8/12$

- Machine A ran 2 times faster on program 1, $2/3$ times faster on program 2
- On average, Machine A is $(2 + 2/3) / 2 = 4/3$ times faster than Machine B

It turns this “averaging” stuff can fool us; watch...

Example (cont'd)

Two machines timed on two benchmarks

	<u>Machine A</u>	<u>Machine B</u>
Program 1	2 seconds	4 seconds
Program 2	12 seconds	8 seconds

How much faster is Machine A than B?

Attempt 2: ratio of runtimes, normalized to Machine B runtimes

program 1: $2/4$ program 2 : $12/8$

- Machine A ran program 1 in $1/2$ the time and program 2 in $3/2$ the time
- On average, $(1/2 + 3/2) / 2 = 1$
- Put another way, Machine A is 1.0 times faster than Machine B

Example (cont'd)

Two machines timed on two benchmarks

	Machine A	Machine B
Program 1	2 seconds	4 seconds
Program 2	12 seconds	8 seconds

How much faster is Machine A than B?

Attempt 3: ratio of aggregated runtimes, norm. to A

- Machine A took 14 seconds for both programs
- Machine B took 12 seconds for both programs
- Therefore, Machine A takes $14/12$ of the time of Machine B
- Put another way, Machine A is $6/7$ faster than Machine B

Which is Right?

Question:

- How can we get three different answers?

Answer:

- Because, while they are all reasonable calculations...

...each answers a different question

Need to be more precise in understanding and posing these performance & metric questions

Arithmetic and Harmonic Mean

Average of the execution time that tracks total execution time is the arithmetic mean

$$\frac{1}{n} \sum_{i=1}^n \textit{Time}_i$$

This is the definition for “average” you are most familiar with

If performance is expressed as a rate, then the average that tracks total execution time is the harmonic mean

$$\frac{n}{\sum_{i=1}^n \frac{1}{\textit{Rate}_i}}$$

This is a different definition for “average” you are probably less familiar with

Geometric Mean

- Used for relative rate (i.e., ratio) or normalized performance

$$\textit{Relative_Rate} = \frac{\textit{Rate}}{\textit{Rate}_{ref}} = \frac{\textit{Time}_{ref}}{\textit{Time}}$$

- Geometric mean

$$\sqrt[n]{\prod_{i=1}^n \textit{Relative_Rate}_i} = \frac{\sqrt[n]{\prod_{i=1}^n \textit{Rate}_i}}{\textit{Rate}_{ref}}$$

Why does the choice of the mean matter?

Benchmark	Ops (millions)	Computer 1	Computer 2	Speedup (C2 vs C1)
<i>Absolute performance (Time)</i>				
Program 1	100	1	20	
Program 2	100	1000	20	
Total time		1001	40	25
Avg (arith mean)		500	20	25

Why does the choice of the mean matter?

Benchmark	Ops (millions)	Computer 1	Computer 2	Speedup (C2 vs C1)
<i>Absolute performance (Time)</i>				
Program 1	100	1	20	
Program 2	100	1000	20	
Total time		1001	40	25
Avg (arith mean)		500	20	25
<i>Performance in MFLOPS (Rate)</i>				
Program 1		100	5	
Program 2		0.1	5	
Arith. mean		50.1	5	0.1
Geom. mean		3.2	5	1.6
Harm. mean		0.2	5	25



Quiz

A car drives 30 mph for first 10 miles, 90 mph for next 10 miles

What is the car's average speed?

$$\text{Average speed} = (30 + 90)/2 = 60 \text{ mph}$$

A. True

B. False

Quiz

- E.g., 30 mph for first 10 miles, 90 mph for next 10 miles

- What is the average speed?

Average speed = $(30 + 90)/2 = 60$ mph (Wrong!)

- Correct answer:

Average speed = total distance / total time
 $= 20 / (10/30 + 10/90)$
 $= 45$ mph

For rates use Harmonic Mean!

Problems with Arithmetic Mean

- Applications do not have the same probability of being run
- Longer programs weigh more heavily in the average

For example, two machines timed on two benchmarks

	<u>Machine A</u>	<u>Machine B</u>
Program 1	2 seconds (20%)	4 seconds (20%)
<u>Program 2</u>	<u>12 seconds (80%)</u>	<u>8 seconds (80%)</u>

- If we do arithmetic mean, Program 2 “counts more” than Program 1
An X% improvement in Program 2 changes the average more than an X% improvement in Program 1
- But perhaps Program 2 is 4 times more likely to run than Program 1

Weighted Execution Time

Often, one runs some programs more often than others. Therefore, we should *weight* the more frequently used programs' execution time

$$\sum_{i=1}^n \text{Weight}_i \times \text{Time}_i$$

Weighted Harmonic Mean

$$\frac{1}{\sum_{i=1}^n \frac{\text{Weight}_i}{\text{Rate}_i}}$$

Using a Weighted Sum (or weighted average)

	<u>Machine A</u>	<u>Machine B</u>
Program 1	2 seconds (20%)	4 seconds (20%)
Program 2	12 seconds (80%)	8 seconds (80%)
Total	10 seconds	7.2 seconds

Allows us to determine relative performance $10/7.2 = 1.39$

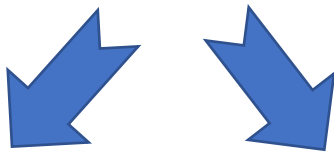
→ Machine B is 1.39 times faster than Machine A

(instead of $14/12 = 1.16x$ faster without weighting)

What if we only know normalized runtimes?

Normalize runtime of each program to a reference

	<u>Machine A (ref)</u>	<u>Machine B</u>
Program 1	2 seconds	4 seconds
Program 2	12 seconds	8 seconds



	<u>Machine A (norm to B)</u>	<u>Machine B (norm to A)</u>
Program 1	0.5	2.0
Program 2	1.5	0.666
Average?	1.0	1.333

- When we normalize A to B and average, it looks like A & B are the same.
- But when we normalize B to A and average, it looks like A is better!

Using Geometric Mean

	Machine A (norm to B)	Machine B (norm to A)
Program 1	0.5	2.0
Program 2	1.5	0.666
Geometric Mean	0.866	1.154

Note that $1.154 = 1/0.866$

Drawbacks:

- Does not reflect actual runtime because it normalizes
- Each application now counts equally

When is geomean useful?

Geometric mean of ratios is not proportional to total time

Use to compare machines when

- Relative performance on each program is known
- Relative runtime/weights of different programs is not known
- E.g., to aggregate speedups on set of programs

Rule of thumb: Use AM for times, HM for rates, GM for ratios

Summary of metrics

Name	Notation	Units	Comment
Memory footprint	-	Bytes	Total space occupied by the program in memory
Execution time	$(\sum \text{CPI}_j) * \text{clock cycle time, where } 1 \leq j \leq n$	Seconds	Running time of the program that executes n instructions
Arithmetic mean	$(E_1 + E_2 + \dots + E_p)/p$	Seconds	Average of execution times of constituent p benchmark programs
Weighted Arithmetic mean	$(f_1 * E_1 + f_2 * E_2 + \dots + f_p * E_p)$	Seconds	Weighted average of execution times of constituent p benchmark programs
Geometric mean	$p^{\text{th}} \text{ root } (E_1 * E_2 * \dots * E_p)$	Seconds	p^{th} root of the product of execution times of p programs that constitute the benchmark
Static instruction frequency		%	Occurrence of instruction i in compiled code
Dynamic instruction frequency		%	Occurrence of instruction i in executed code
Speedup (M_A over M_B)	E_B/E_A	Number	Speedup of Machine A over B
Improvement in Exec time	$(E_{\text{old}} - E_{\text{new}})/E_{\text{old}}$	Number	New Vs. old
Amdahl's law	$\text{Time}_{\text{after}} = \text{Time}_{\text{unaffected}} + \text{Time}_{\text{affected}}/x$	Seconds	x is amount of improvement
Iron Law	$\text{Time} = \text{Insns/program} * \text{cycles/insn} * \text{time/cyc}$	Seconds	Three terms to improve performance