

CS2200

Systems and Networks


Spring 2022

Lecture 10: Pipelining

Alexandros (Alex) Daglis
School of Computer Science
Georgia Institute of Technology
adaglis@gatech.edu

Going faster...

Latency vs Throughput



Bill's Sandwich Shop



station I
(place order)

station II
(select bread)

station III
(cheese)

station IV
(meat)

station V
(veggies)



"THE LC-2200"
VOTED BEST SANDWICH
IN MIDTOWN

Bill's Sandwich Shop



station I
(place order)

station II
(select bread)

station III
(cheese)

station IV
(meat)

station V
(veggies)



Bill's Sandwich Shop



station I
(place order)

station II
(select bread)

station III
(cheese)

station IV
(meat)

station V
(veggies)



Bill's Sandwich Shop



station I
(place order)

station II
(select bread)

station III
(cheese)

station IV
(meat)

station V
(veggies)



Bill's Sandwich Shop



station I
(place order)

station II
(select bread)

station III
(cheese)

station IV
(meat)

station V
(veggies)

A yellow square containing a brown sandwich shape. Inside the sandwich shape, the text "THE LC-2200" VOTED SLOWEST SANDWICH EVER" is written in bold, brown, sans-serif capital letters.

**"THE LC-2200"
VOTED SLOWEST
SANDWICH EVER**



Net result: One complete sandwich every five cycles.

Bill's Mega-Sandwich Shop



station I
(place order)

station II
(select bread)

station III
(cheese)

station IV
(meat)

station V
(veggies)

"What will You Have?"




Manager


Net result: One complete sandwich every cycle! (Once you fill the pipeline.)
A 5x speedup!




Interpreting an instruction




Interpreting an instruction




Interpreting an instruction




Interpreting an instruction



Interpreting an instruction




Interpreting an instruction





One at a time...

instructions




Three at a time?

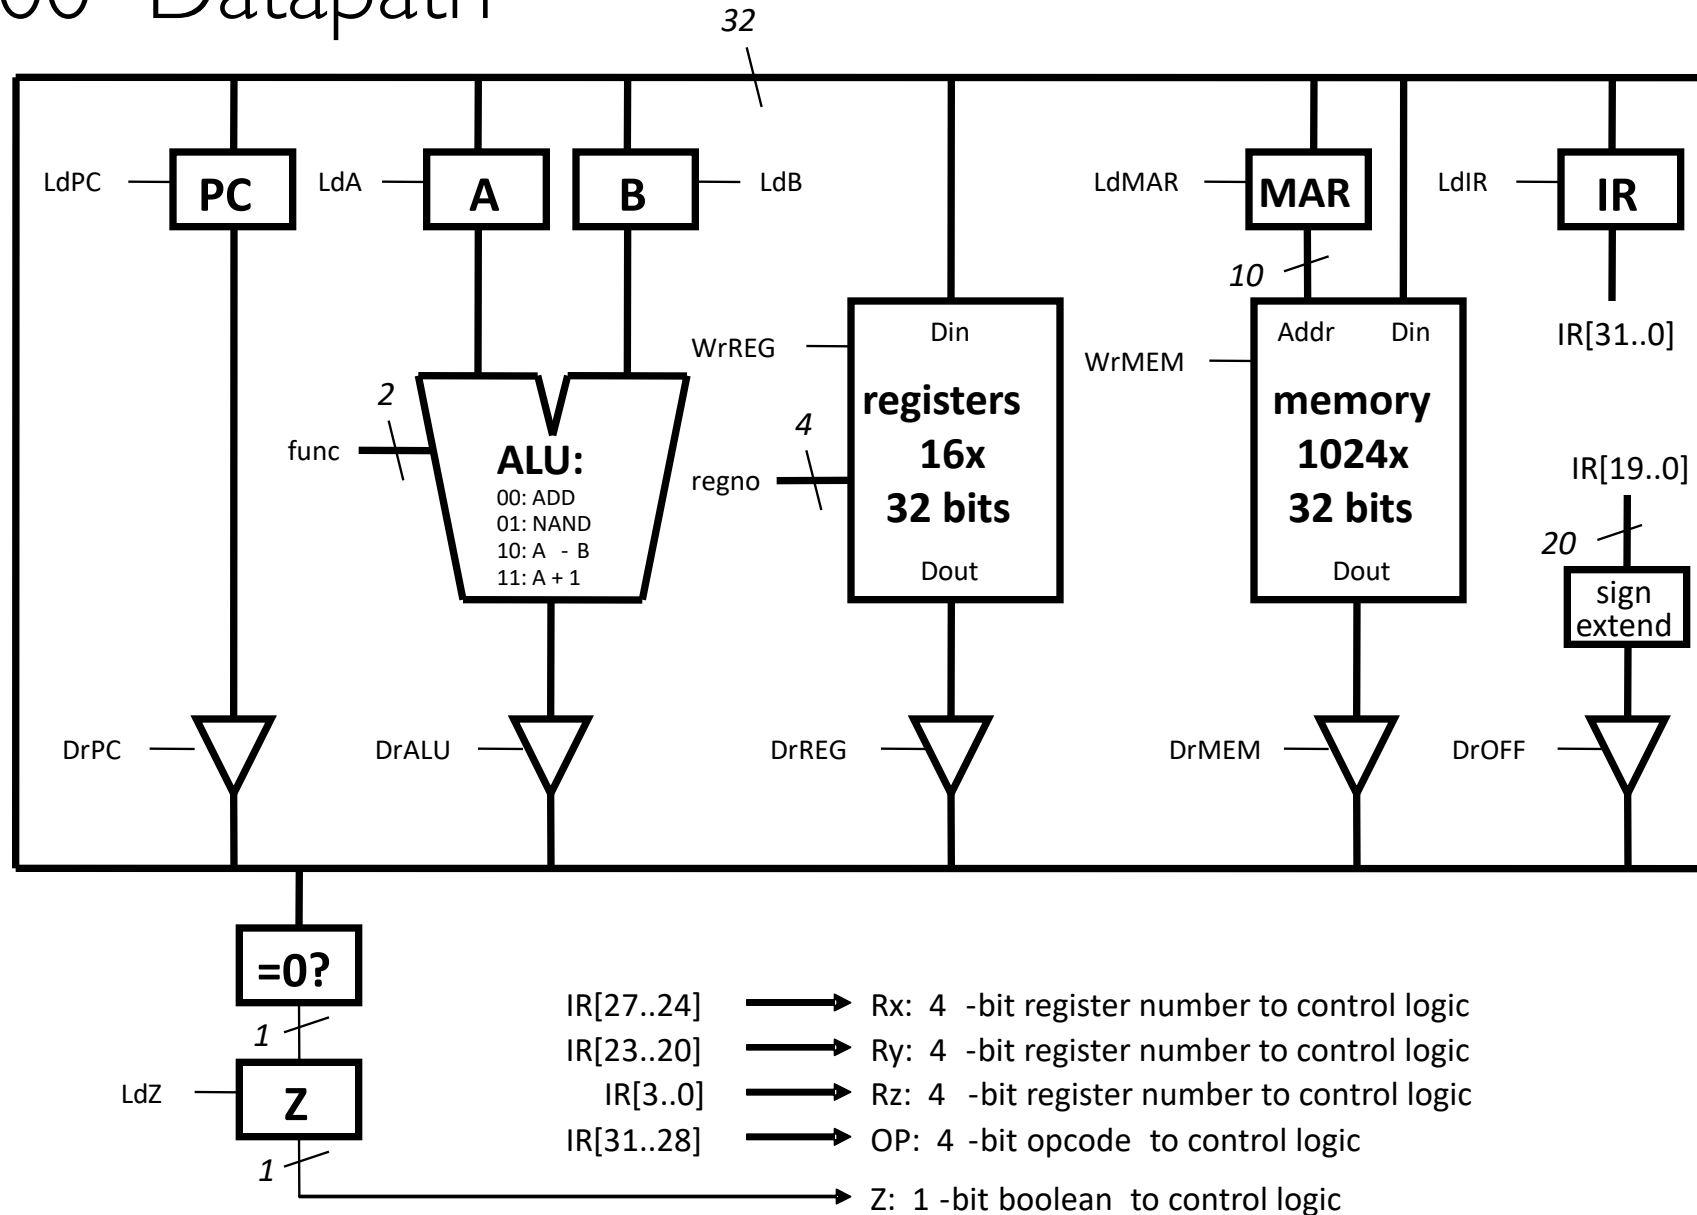




What do we need in the datapath?



LC-2200 Datapath



What units are used?

Macro State

FETCH

DECODE

EXECUTE (ADD)

EXECUTE (LW)

Units in Use

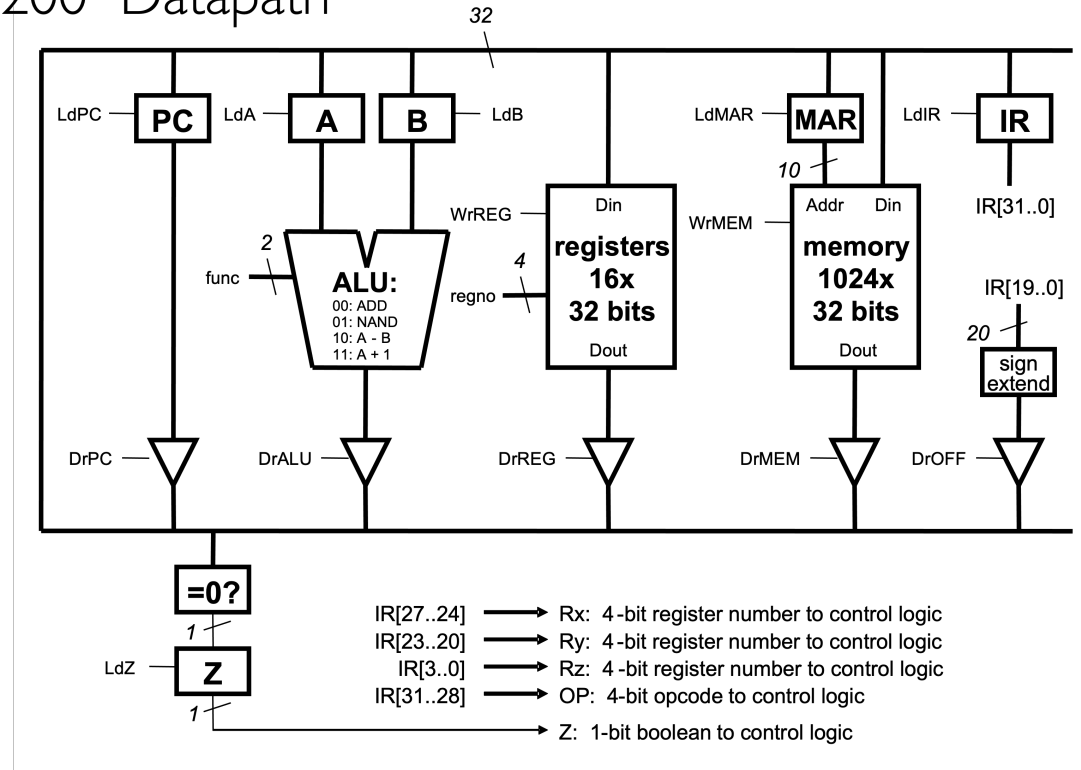
IR, ALU, PC, MEM

IR

IR, ALU, Reg-file

IR, ALU, Reg-file, MEM, Sign extender

LC-2200 Datapath



Imitation: The Sincerest Form of Flattery



station I
(place order)
New (5th order)



station II
(select bread)
4th order



station III
(cheese)
3rd order



station IV
(meat)
2nd order




station V
(veggies)
1st order

Points to note:

- Work done by each station is roughly equal
- Every sandwich goes through every station, regardless of what you want or don't want
- Order form and partially assembled sandwich is passed from one station to the next
- Each station does part of the work for assembling a sandwich

How do we mimic sandwich assembly?



- Unequal division of labor
- Violates first principle in sandwich assembly line

Bill's Cost-Cutting Measure



station I
(place order)
New (5th order)



station II
(select bread)
4th order



station III
(cheese)
3rd order



station IV
(meat)
2nd order



station V
(veggies)
1st order



Bill's Cost-Cutting Measure



Station 1
(place order)
3rd order



station II
(select bread)
2nd order



station III
(cheese, meat, veggies)
1st order



Disaster!


- Increased waiting times
- One sandwich every 3 cycles, not one every cycle!
- Two workers who only work about a third of the time



Bill originally used 5 stages rather than 3 in his sandwich pipeline because...

- A. He copied the instruction pipeline
- B. He could only find 5 workers
- C. He is solving the unemployment problem
- D. He wants to make sure that the amount of work per stage is roughly the same
- E. He is a nice guy
- F. You are making us hungry...

How do we mimic sandwich assembly?



- Get to the basics!
- What needs to happen for **every** instruction?

Each step is
roughly the
same amount
of work

- 1 ■ Fetch into IR and increment PC
- 2 ■ Decode instruction and read register contents
- 3
 - Perform arithmetic/logic (maybe)
 - Perform address computation (maybe)
- 4 ■ Fetch/store memory operand (maybe)
- 5 ■ Write to register (maybe)

How can we know what registers to read?

- What does the “bread guy” care about?
- How does the “bread guy” know **what** bread?
- It's the order form!



- What's the equivalent of the order form in the instruction pipeline?
- IR

LC-2200 Instruction set

- R-type instructions (add, nand):

bits 31-28: opcode;

bits 23-20: reg Y;

bits 3-0: reg Z

bits 27-24: reg X;

bits 19-4: unused (should be all 0s);

- I-type instructions (addi, lw, sw, beq):

bits 31-28: opcode;

bits 23-20: reg Y;

bits 27-24: reg X

bits 19-0: Imm. Offset

- J-type instructions (jalr):

bits 31-28: opcode;

bits 23-20: reg Y;

bits 27-24: reg X


bits 19-0: unused

- O-type instructions (halt):

bits 31-28: opcode;

bits 27-0: unused


Every instruction goes through



- Some stages don't do anything for some instructions
- Each stage works on a different instruction


IF	- fetch instruction into IR and increment PC
ID/RR	- decode and read register contents
EX	- perform arithmetic/logic (maybe) - perform address computation (maybe)
MEM	- fetch/store memory operand (maybe)
WB	- write to register (maybe)

What moves between stages?




- Think of our sandwich assembly line
- Each station passes two things to the next
 - The order form
 - A partially assembled sandwich
- Analogies with our data path
 - Is the IR equivalent to the order form?
 - What is a “partially assembled data sandwich”?

Electrical isolation between stages




- Registers act as a **wall** to isolate the actions of each combinational logic cloud!
- On each clock tick, the registers send their output to drive the next stage of combinational logic

What's in these pipeline boxes?




- How do we separate our pipeline boxes?
- The boxes are really just combinational logic clouds
- So we already know how to separate them...

This is where buffers come from...



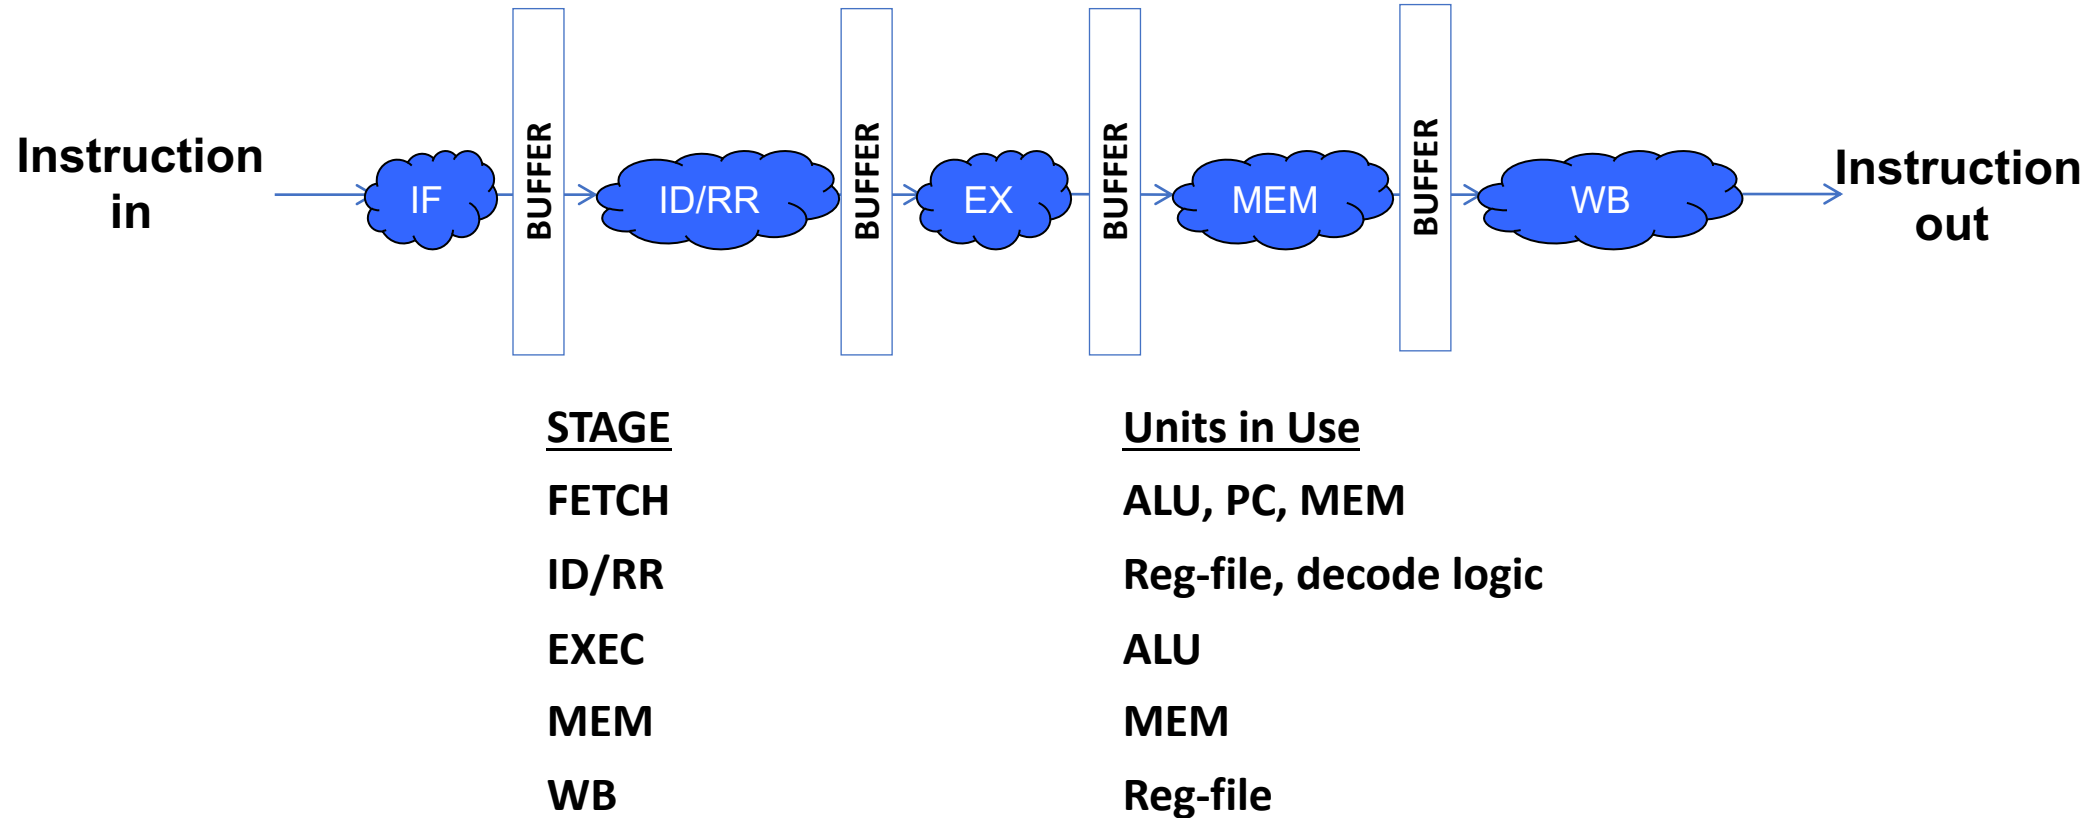
- There is combinational logic in each cloud for performing the actions of that stage
- The buffers are merely clocked registers for passing the partially assembled results

Again: What does each stage do?




- IF - fetch instruction into IR and increment PC
- ID/RR - read register contents
- EX
 - perform arithmetic/logic (maybe)
 - perform address computation (maybe)
- MEM - fetch/store memory operand (maybe)
- WB - write to register (maybe)

Again: What does each stage need?



- Conclusion: we need a little datapath in each stage

A shiny new data path!



STAGE

FETCH

ID/RR

EXEC

MEM

WB

Units in Use

ALU, PC, MEM


Reg-file, decode logic

ALU

MEM

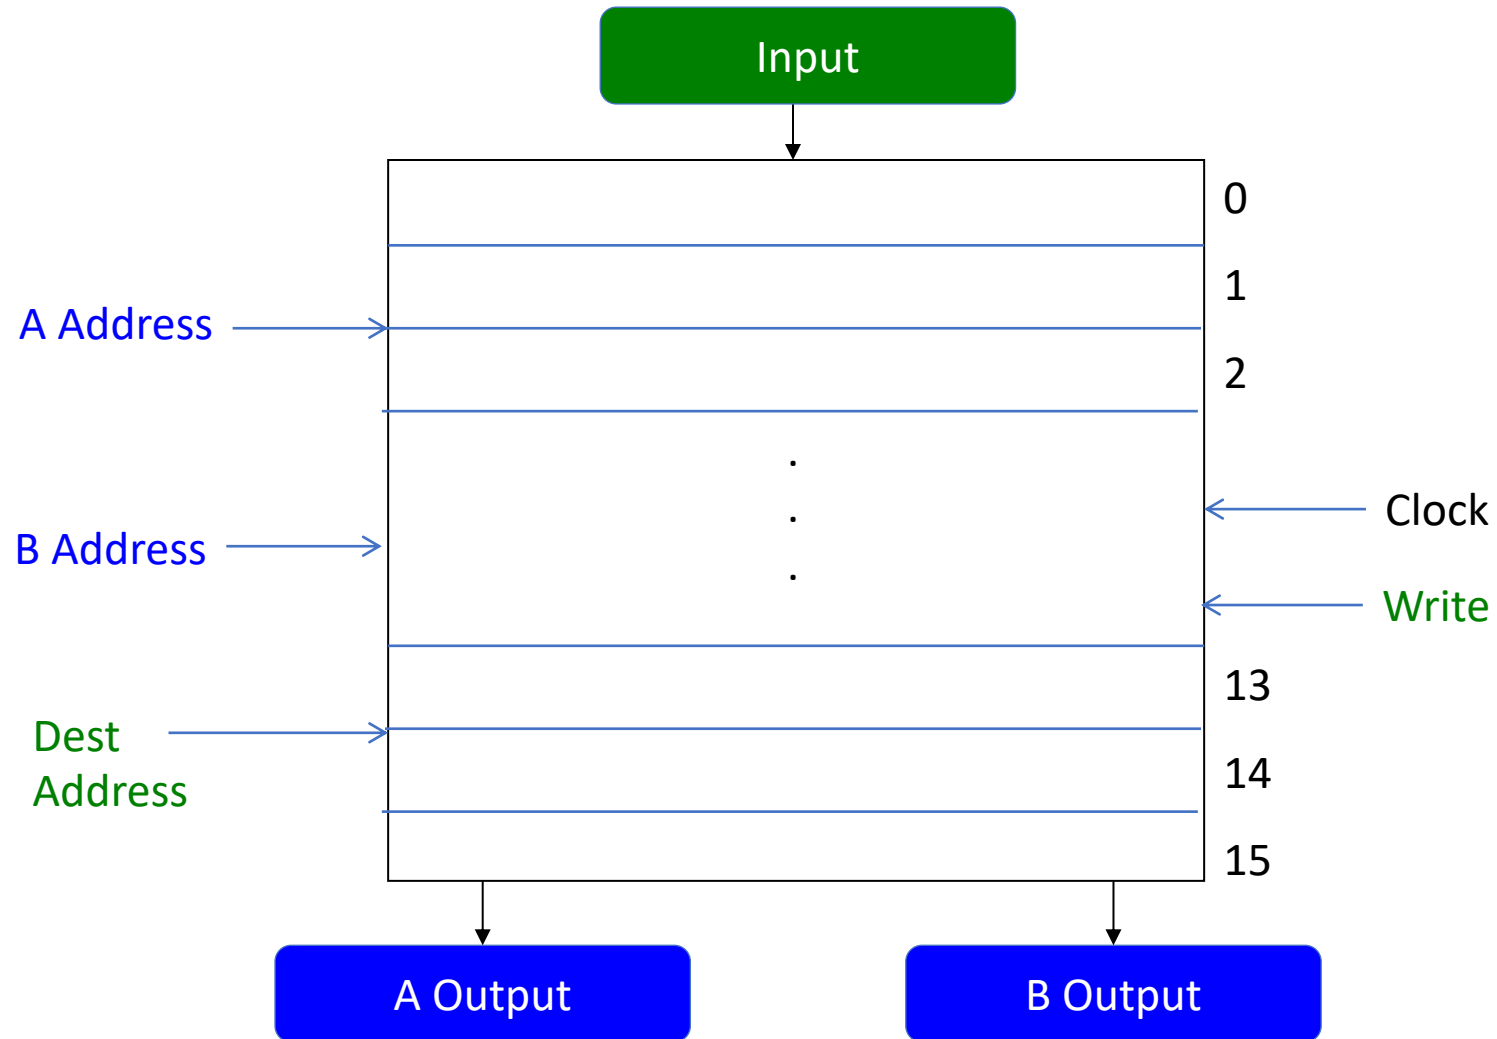
Reg-file

Two register files? Not.




- We use a register file in two stages

Do you remember this device?

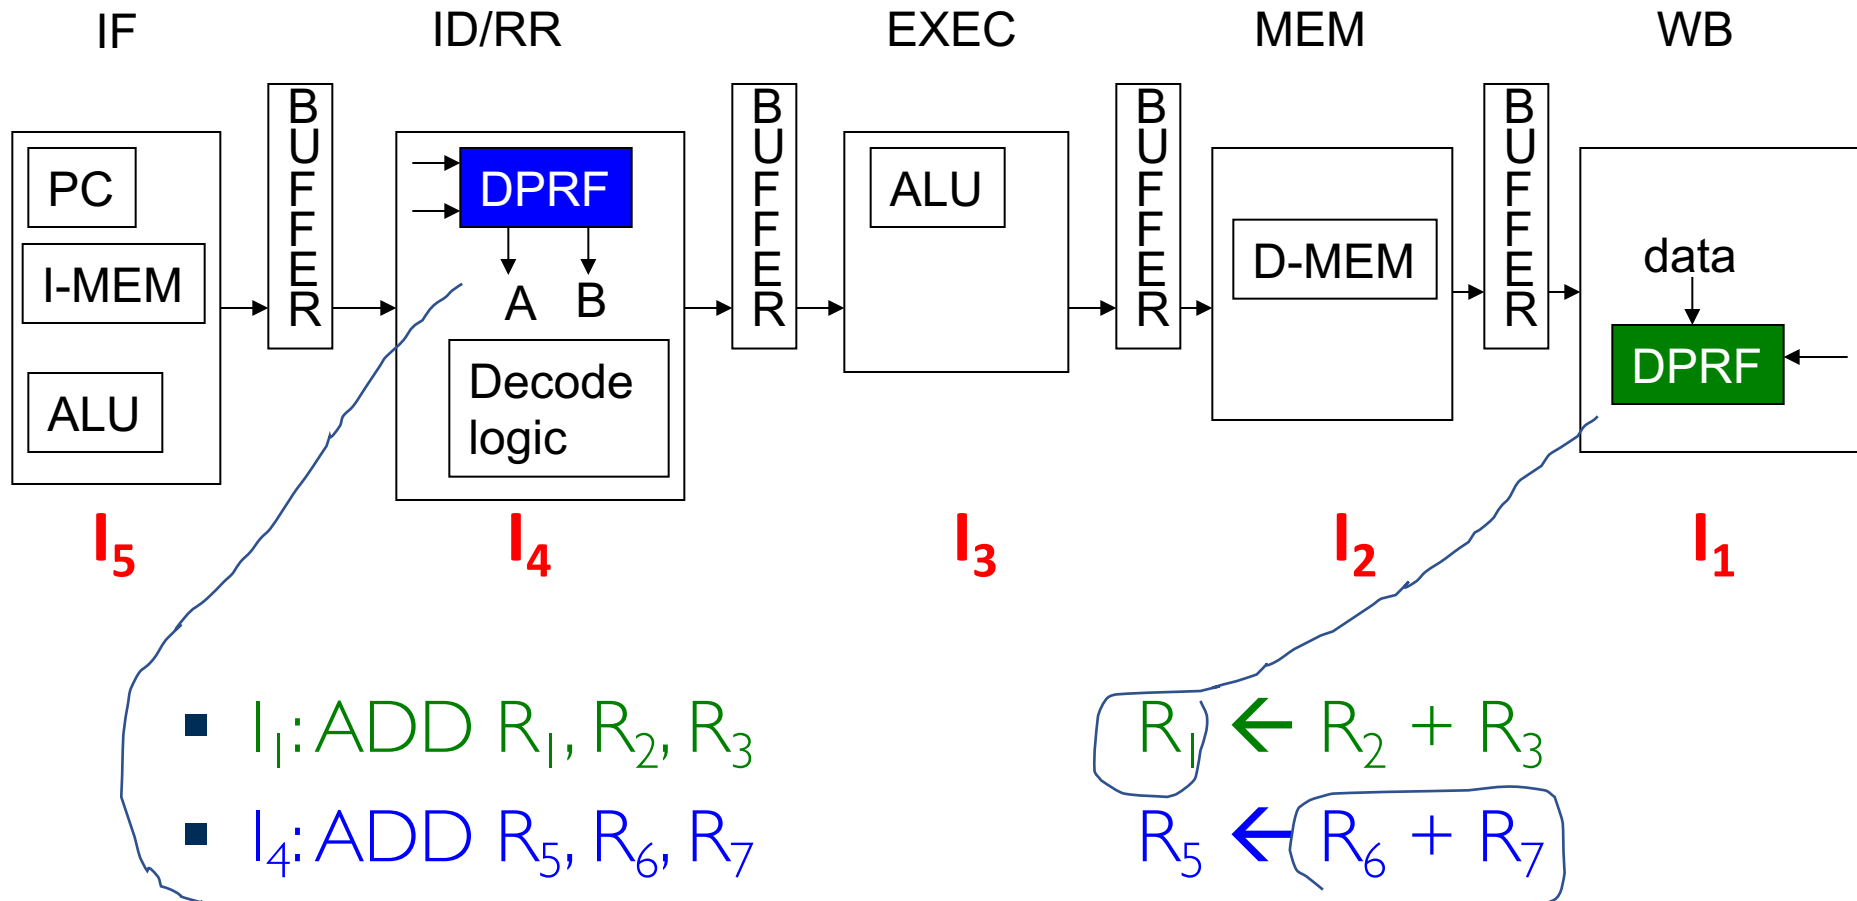


Two register files? Not.

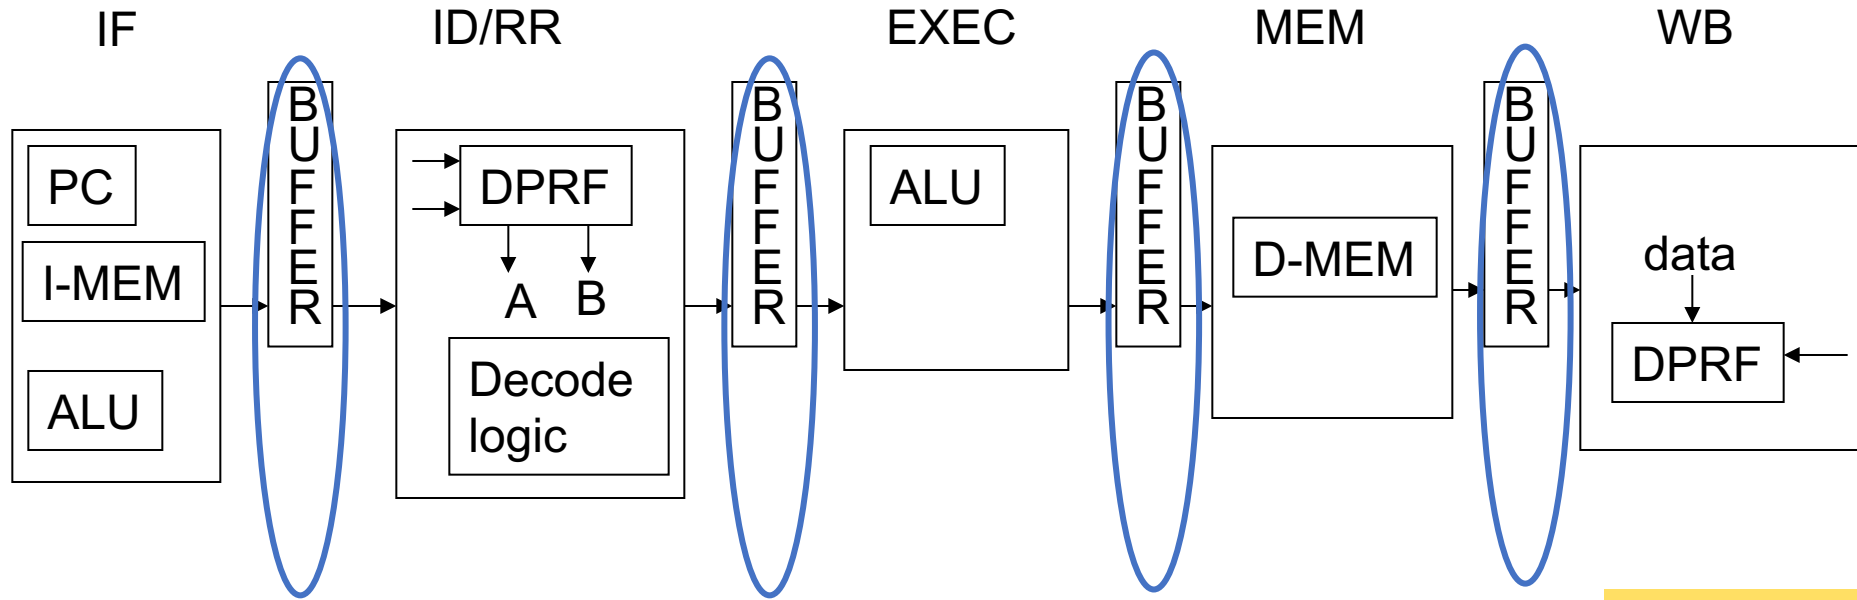


- We use a register file in two stages, but conveniently
 - ID/RR uses read ports (A addr, B addr B, A output, B output)
 - WB uses write port (Dest addr, Input)
- So we can happily share the same register file on each clock cycle

Example of sharing the register file



But what about those buffers?



- What goes in those buffers? Same? Different?
- The different stages seem to need different information
- What was that sandwich analogy?






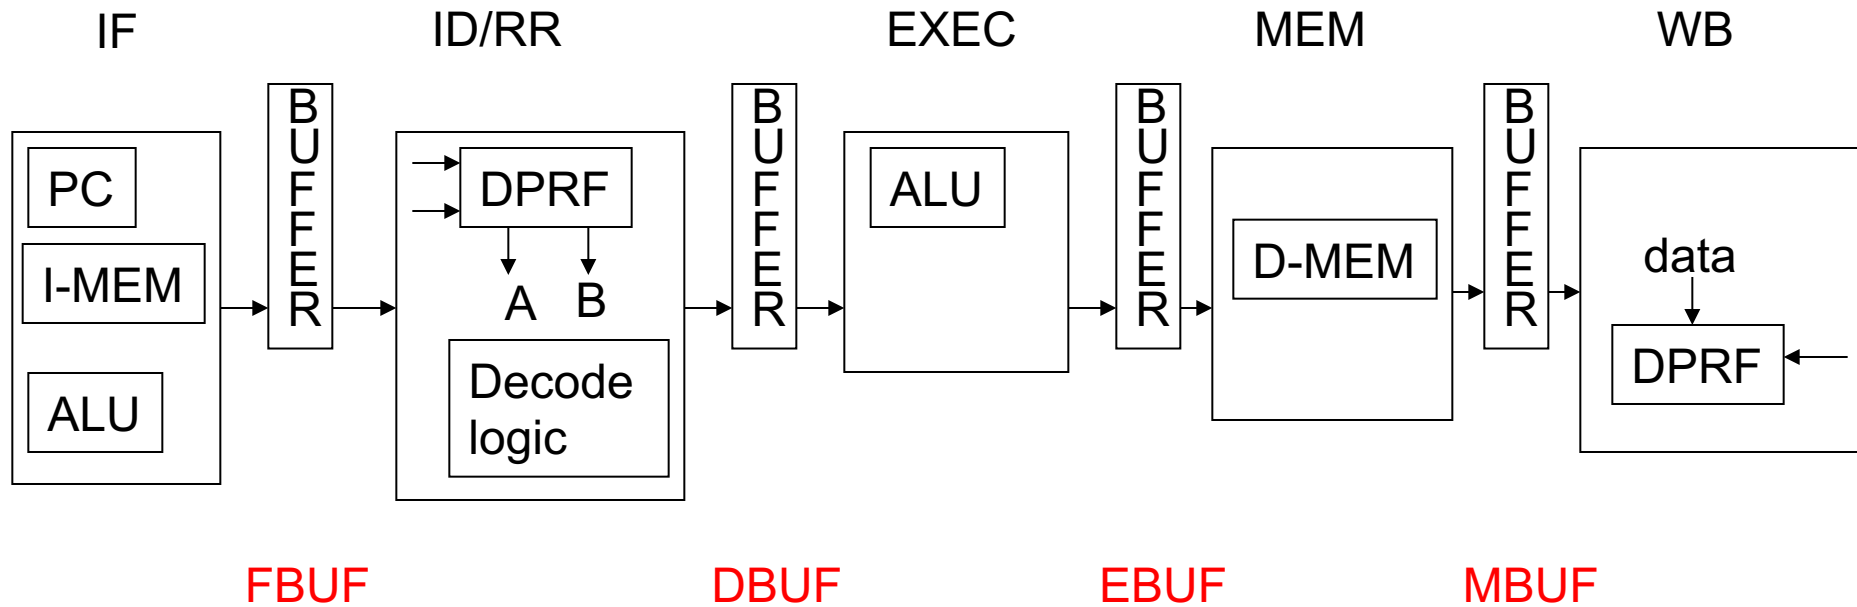
Comparing the instruction pipeline...

...to Bill's sandwich pipeline, the buffers between stages serve the same function as

- A. A partially assembled sandwich
- B. Order form passed between stations
- C. A combination of the above two
- D. No comparison
- E. Buffer? You mean buffet?
- F. You are making us hungry... again!

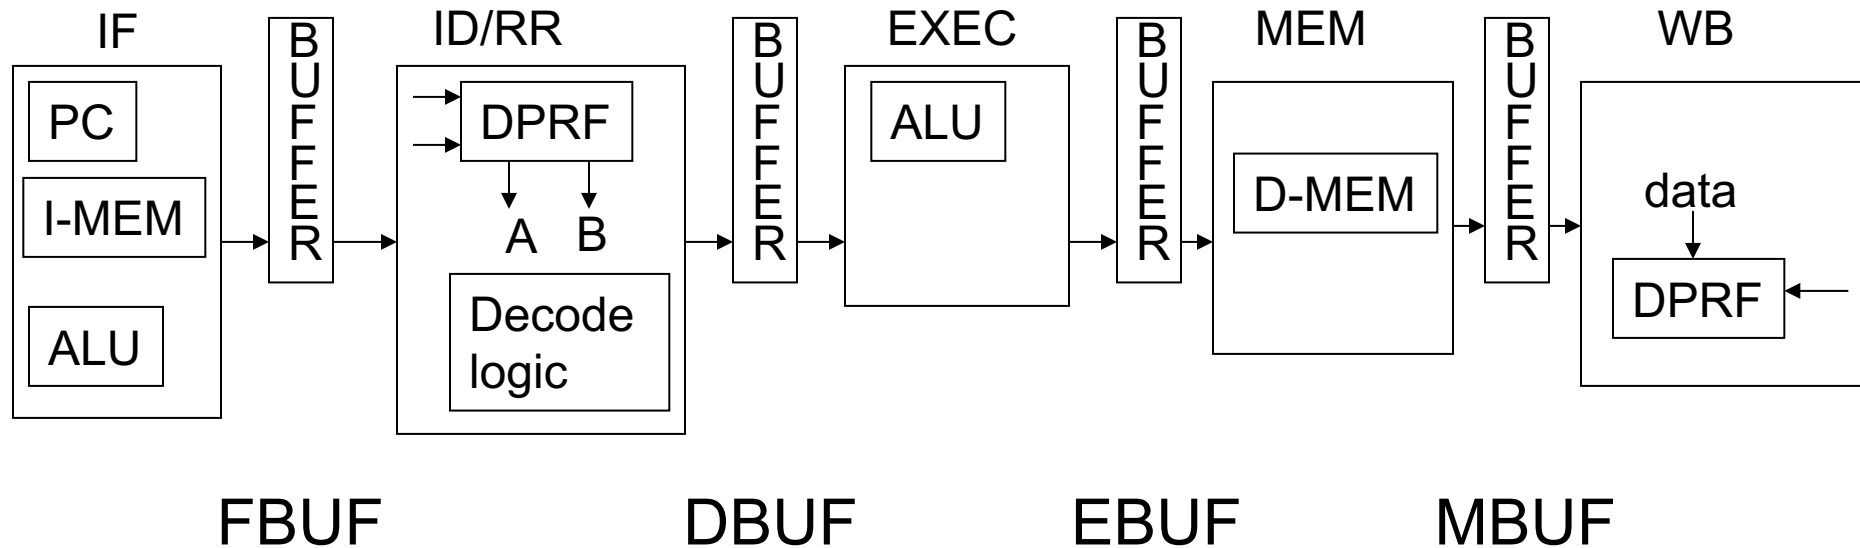


So what makes up our partially assembled data sandwich?




- The IR (or parts of it) is the order form
- The data sandwich is the partial execution results
- Do the buffers have to all carry the same information? Or can they be different?

What does each stage need in the buffers?



Name	Output of Stage	Contents
FBUFF	IF	Primarily contains instruction read from memory
DBUFF	ID/RR	Decoded IR and values read from register file
EBUFF	EX	Primarily contains result of ALU operation plus other parts of the instruction depending on the instruction specifics
MBUFF	MEM	Same as EBUFF if instruction is not LW or SW; If instruction is LW, then buffer contains the contents of memory location read

Our non-pipelined LC-2200



At each point, the processor is in exactly one macro state



What “macro state” is a pipelined processor in at any point in time?

- A. Fetch
- B. Decode
- C. Execute
- D. All of the above
- E. We need a new “sandwich” state

Anatomy of an ADD instruction

- IF stage (cycle 1)
- ID/RR stage (cycle 2)
- EX stage (cycle 3)
- MEM stage (cycle 4)
- WB stage (cycle 5)

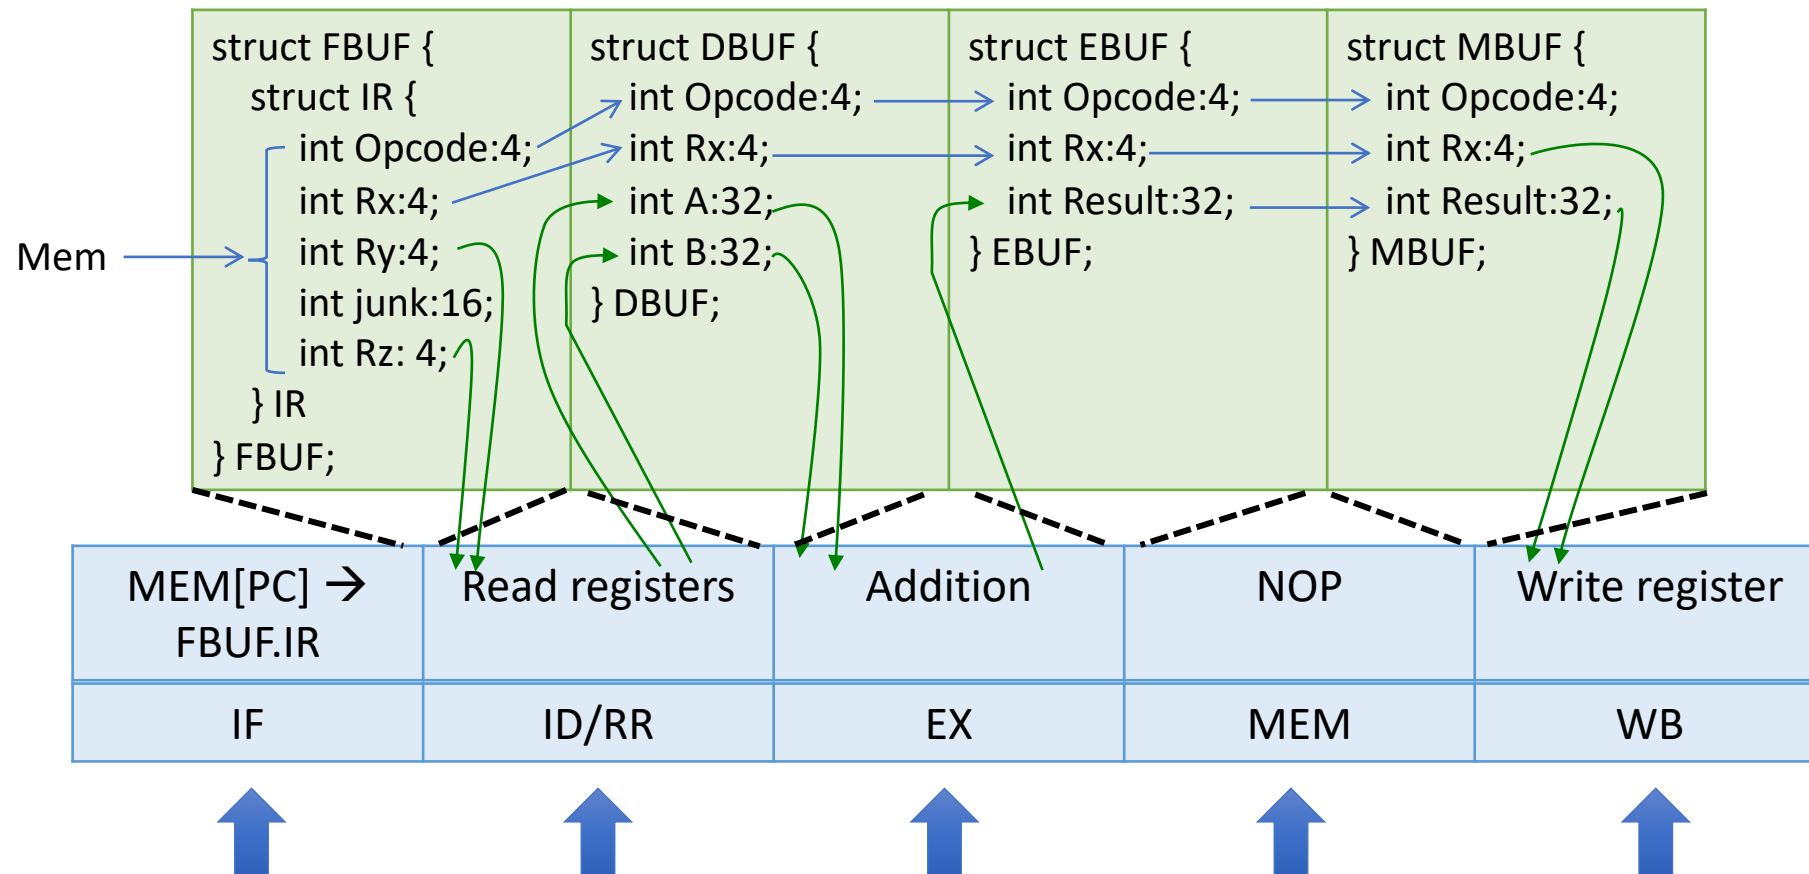
Actions for an ADD instruction

Stage	Actions	Comments
IF	$\text{MEM}[\text{PC}] \rightarrow \text{FBUF.IR}$ $\text{PC} + 1 \rightarrow \text{PC}$	Instruction is fetched from memory and placed in FBUF (which is essentially the IR)
ID/RR	$\text{DPRF}[\text{FBUF.IR.Ry}] \rightarrow \text{DBUF.A}$ $\text{DPRF}[\text{FBUF.IR.Rz}] \rightarrow \text{DBUF.B}$ $\text{FBUF.IR.Opcode} \rightarrow \text{DBUF.Opcode}$ $\text{FBUF.IR.Rx} \rightarrow \text{DBUF.Rx}$	Read register Ry into DBUF.A Read register Rz into DBUF.B Copy opcode from FBUF to DBUF Copy Rx from FBUF to DBUF
EX	$\text{DBUF.A} + \text{DBUF.B} \rightarrow \text{EBUF.Result}$ $\text{DBUF.OPCODE} \rightarrow \text{EBUF.Opcode}$ $\text{DBUF.Rx} \rightarrow \text{EBUF.Rx}$	Perform addition Copy opcode from DBUF to EBUF Copy Rx from DBUF to EBUF
MEM	$\text{EBUF} \rightarrow \text{MBUF}$	Copy EBUF to MBUF
WB	$\text{MBUF.Result} \rightarrow \text{DPRF}[\text{MBUF.Rx}]$	Write the result of addition into register Rx

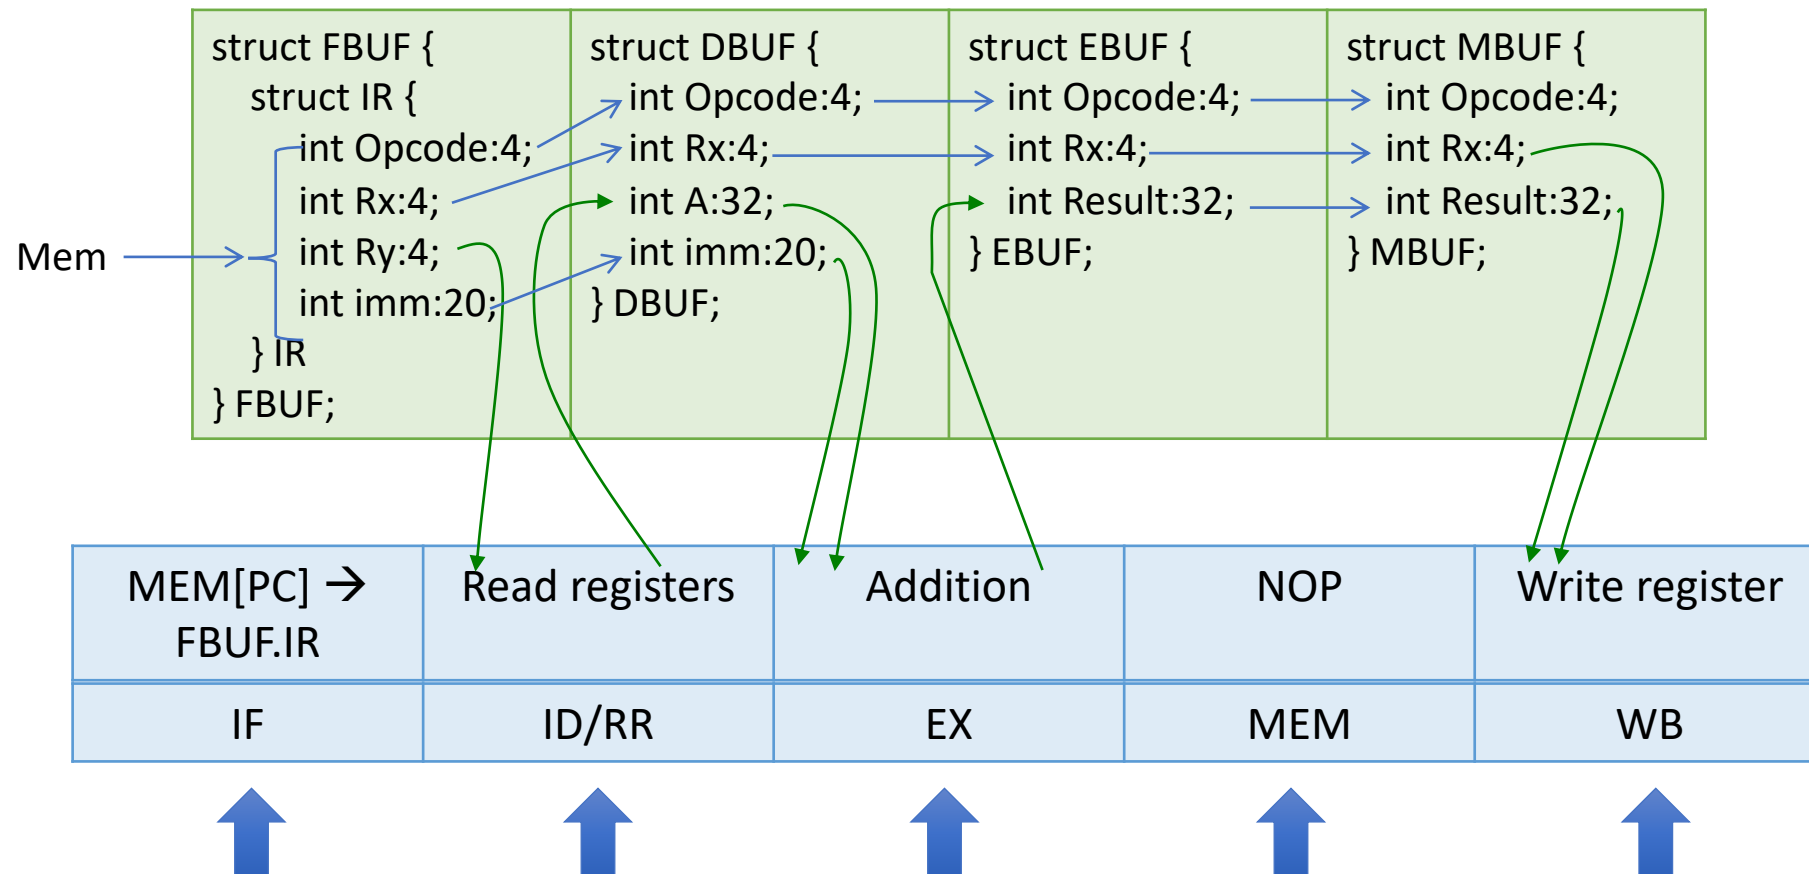
Actions for an ADD instruction

Stage	Actions	Comments
IF	MEM[PC] \rightarrow FBUF.IR PC + 1 \rightarrow PC	Instruction is fetched from memory and placed
ID/RR	DPRF[FBUF.IR.Ry] \rightarrow DBUF.A DPRF[FBUF.IR.Rz] \rightarrow DBUF.B FBUF.IR.Opcode \rightarrow DBUF.Opcode FBUF.IR.Rx \rightarrow DBUF.Rx	Considering only the Add instruction, quantify the sizes of the various buffers between the stages of the pipeline.
EX	DBUF.A + DBUF.B \rightarrow EBUF.Result DBUF.OPCODE \rightarrow EBUF.Opcode DBUF.Rx \rightarrow EBUF.Rx	<div> FBUF? Addition DBUF? Data from DBUF to EBUF EBUF? Data from DBUF to EBUF MBUF? </div>
MEM	EBUF \rightarrow MBUF	Copy EBUF to MBUF
WB	MBUF.Result \rightarrow DPRF[MBUF.Rx]	Write the result of addition into register Rx

How does this pipeline work for ADD?




How does this pipeline work for ADDI?



How to generalize design of pipeline register?

- Look at each instruction in the ISA
- Decide what needs to be carried over from one stage to the next



Design the DBUF register for LC-2200

- Don't optimize by overloading the different fields of the register
- Decide what fields need to be in DBUF for each instruction in the ISA

ADD/NAND	ADDI	BEQ	And so on...
<pre>struct DBUF { int Opcode:4; int Rx:4; int A:32; int B:32; } DBUF;</pre>	<pre>struct DBUF { int Opcode:4; int Rx:4; int A:32; int imm:20; } DBUF;</pre>	<pre>struct DBUF { int Opcode:4; int Rx:4; int imm:20; int PC:32; } DBUF;</pre>	

Now take the UNION of the requirements

ADD/NAND	ADDI	BEQ	And so on...
<pre>struct DBUF { int Opcode:4; int Rx:4; int A:32; int B:32; } DBUF;</pre>	<pre>struct DBUF { int Opcode:4; int Rx:4; int A:32; int imm:20; } DBUF;</pre>	<pre>struct DBUF { int Opcode:4; int Rx:4; int imm:20; int PC:32; } DBUF;</pre>	



```
struct DBUF {  
    int Opcode:4;  
    int Rx:4;  
    int A:32;  
    int B:32;  
    int imm:20;  
    int PC:32;  
} DBUF;
```

Now take the UNION of the requirements

```
struct DBUF {  
    int Opcode:4;  
    int Rx:4;  
    int A:32;  
    int B:32;  
    int imm:20;  
    int PC:32;  
} DBUF;
```

- In the ID/RR stage, **fill** in fields depending on the instruction in FBUF.IR
- Leave unneeded fields **unfilled**
- What fields would we fill in for
JALR \$at, \$ra # \$ra \leftarrow PC; PC \leftarrow \$at



For the BEQ instruction, the fields in DBUF *not filled in...*

BEQ Rx, Ry, offset ; If (Rx == Ry) $PC \leftarrow PC + \text{offset}$

- A. A (contents of Ry)
- B. B (contents of Rx)
- C. PC
- D. Immediate value
- E. Rx specifier

```
struct DBUF {
    int Opcode:4;
    int Rx:4;
    int A:32;
    int B:32;
    int imm:20;
    int PC:32;
} DBUF;
```


IF

ID/RR

EX

MEM

WB

