

CS2200  
Systems and Networks  
Spring 2022

# Lecture 13: Processor Scheduling

Alexandros (Alex) Daglis  
School of Computer Science  
Georgia Institute of Technology  
[adaglis@gatech.edu](mailto:adaglis@gatech.edu)

# Announcements

---

- Starting Chapter 6
- Assignment deadlines extended – check Canvas announcements

# Topics

---

- Process as an abstraction
- Scheduler
- Process, job, thread, task
- States of process
- Process control block (PCB)
- Types of scheduler
- Algorithms
- Metrics & Evaluation

# What is an operating system?

---

- An operating system is a special program that manages access of user programs to hardware resources
- The OS provides a number of abstractions, e.g.
  - Multiple processes (how can one physical processor run multiple programs at the same time?)
  - Memory permissions to protect other processes and the OS
  - Shared access to I/O devices, e.g. networks, file systems
  - Resource sharing (processors, memory, I/O, ...)
  - Additional instructions (read, write, exit, get more memory, change permissions, ...) implemented through traps

# Levels of Abstraction

---

Application (Algorithms expressed in High Level Language)
System Software (Compiler, OS, etc.)
Computer Architecture (ISA)
Machine Organization (Datapath and Control)
Sequential and Combinational Logic Elements
Logic Gates
Transistors
Solid-State Physics (Electrons and Holes)

# Levels of Abstraction

---

Application (Algorithms expressed in High Level Language)
System software (Compiler, OS, etc.)
Computer Architecture (ISA)
Machine Organization (Datapath and Control)
Sequential and Combinational Logic Elements
Logic Gates
Transistors
Solid-State Physics (Electrons and Holes)

# Levels of Abstraction

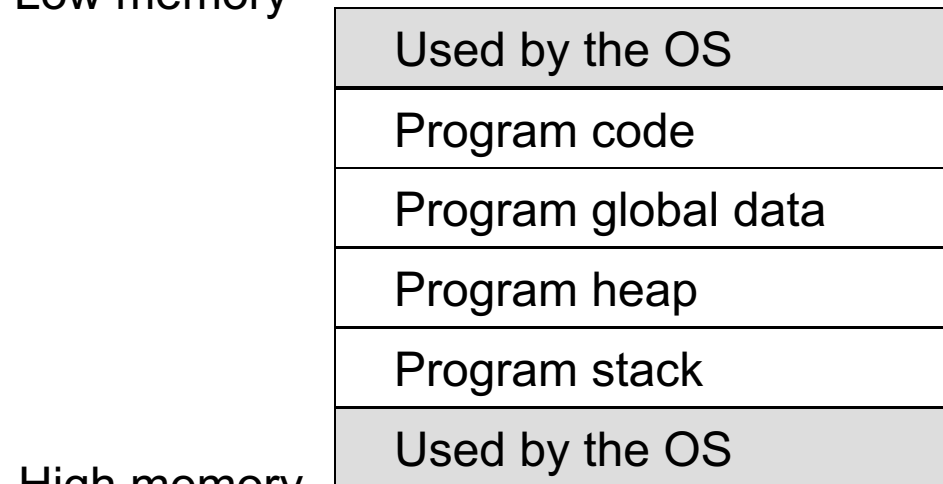
---

Application (Algorithms expressed in High Level Language)	
Compiler	
➡ Operating System	
Computer Architecture (ISA)	
Machine Organization (Datapath and Control)	
Sequential and Combinational Logic Elements	
Logic Gates	
Transistors	
Solid-State Physics (Electrons and Holes)	


# A program in memory

---

Low memory



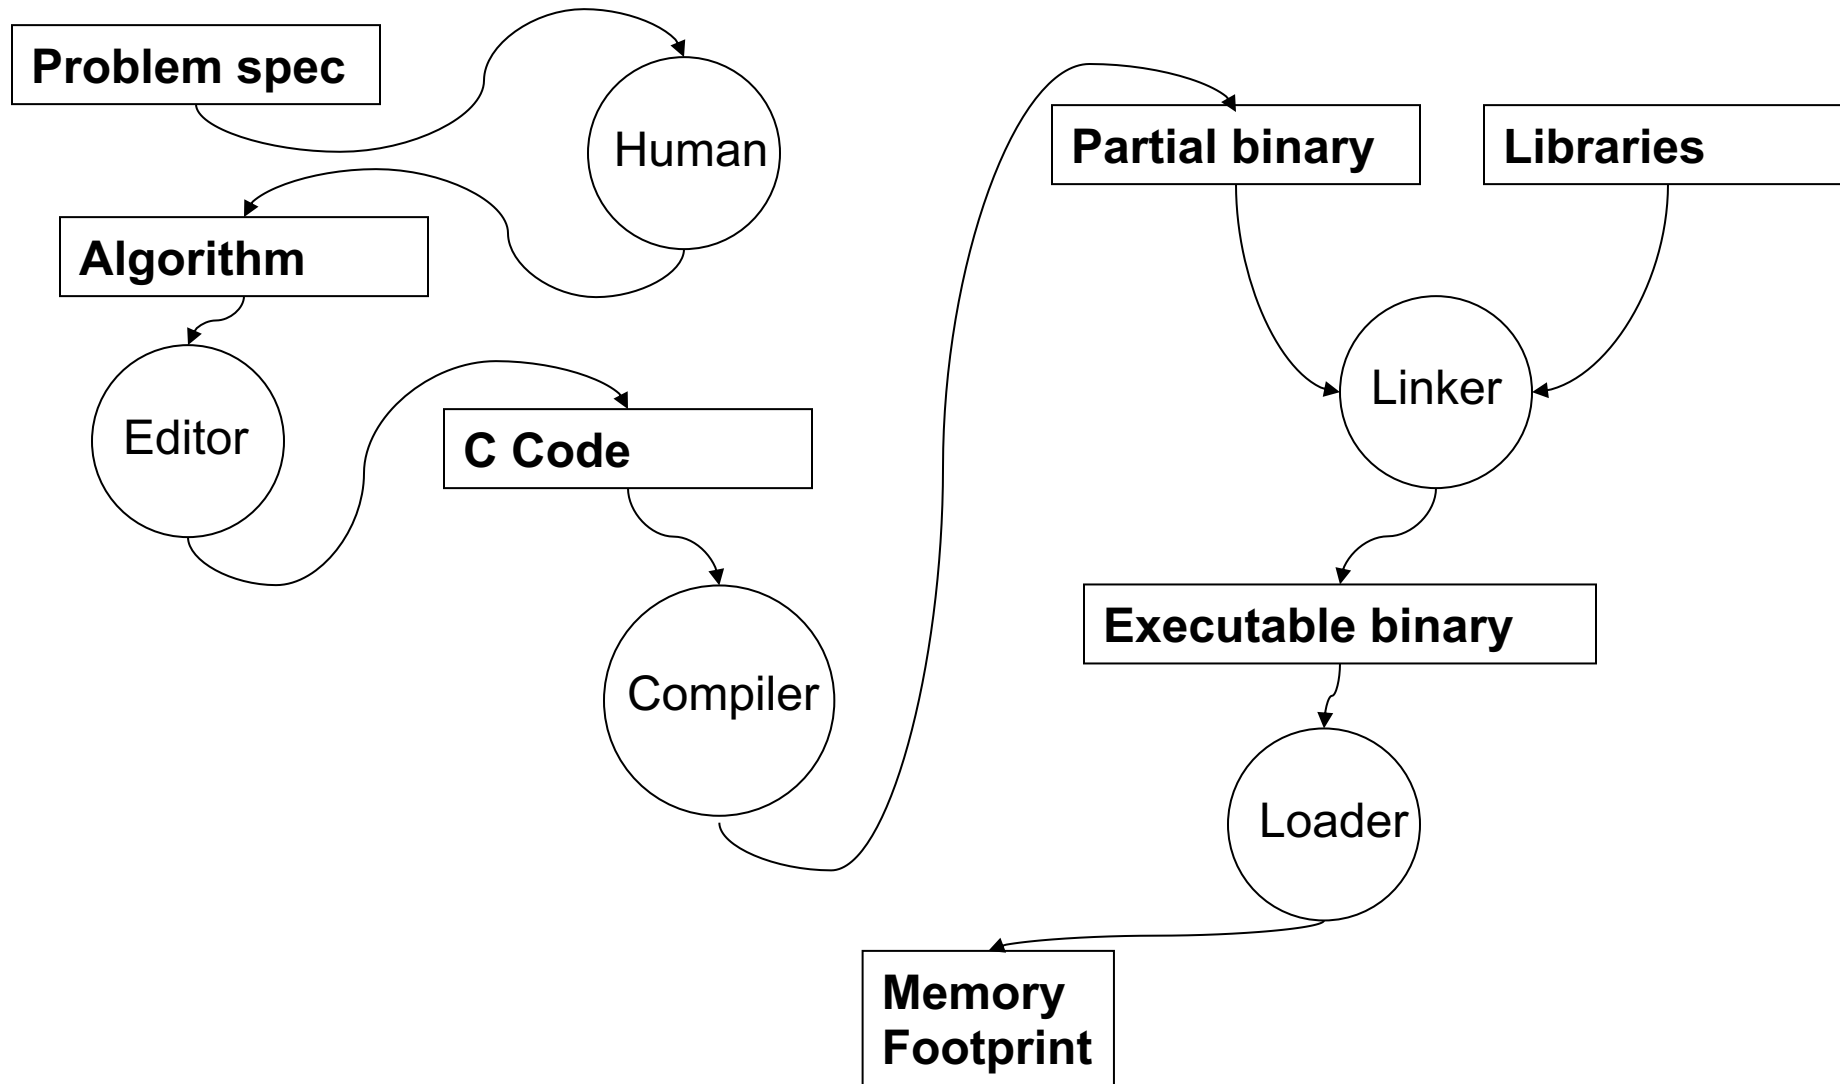
High memory



Memory footprint  
of  
User program



# How do we create a program?



# More Than 1 Program in Memory?

---

<b>Program 1</b>
<b>Program 2</b>
.
.
.
<b>Program n</b>
<b>OS Data Structures</b>
<b>OS routines</b>

- What's the difference between a process and a program?
- A process is a program AND all of the state that represents its execution, e.g.
  - Registers
  - Memory
  - PC
  - Stack
- So far we've been talking about computers that have **one** single process with the state maintained in hardware
- Interrupts were our first clue that it doesn't have to be this way



# Process differs from program in that

- 0% A. There is no difference
- 11% B. A process is a machine language representation of a program
- 68% C. A process is a program in execution
- 0% D. A process is a program that executes correctly
- 21% E. A process is a special kind of program that is part of the operating system
- 0% F. No clue

# Concepts

---

- Multitasking
  - Shortest Job First
  - Priority
  - Round Robin
  - Preemption
  - First Come First Served

# Concepts

- Multitasking
- Shortest Job First
- Priority
- Preemption
- Round Robin
- First Come First Served

In the context of your life

Your To Do List:

- Laundry
- Prepare food and eat
- Call Mom
- Prepare for tests

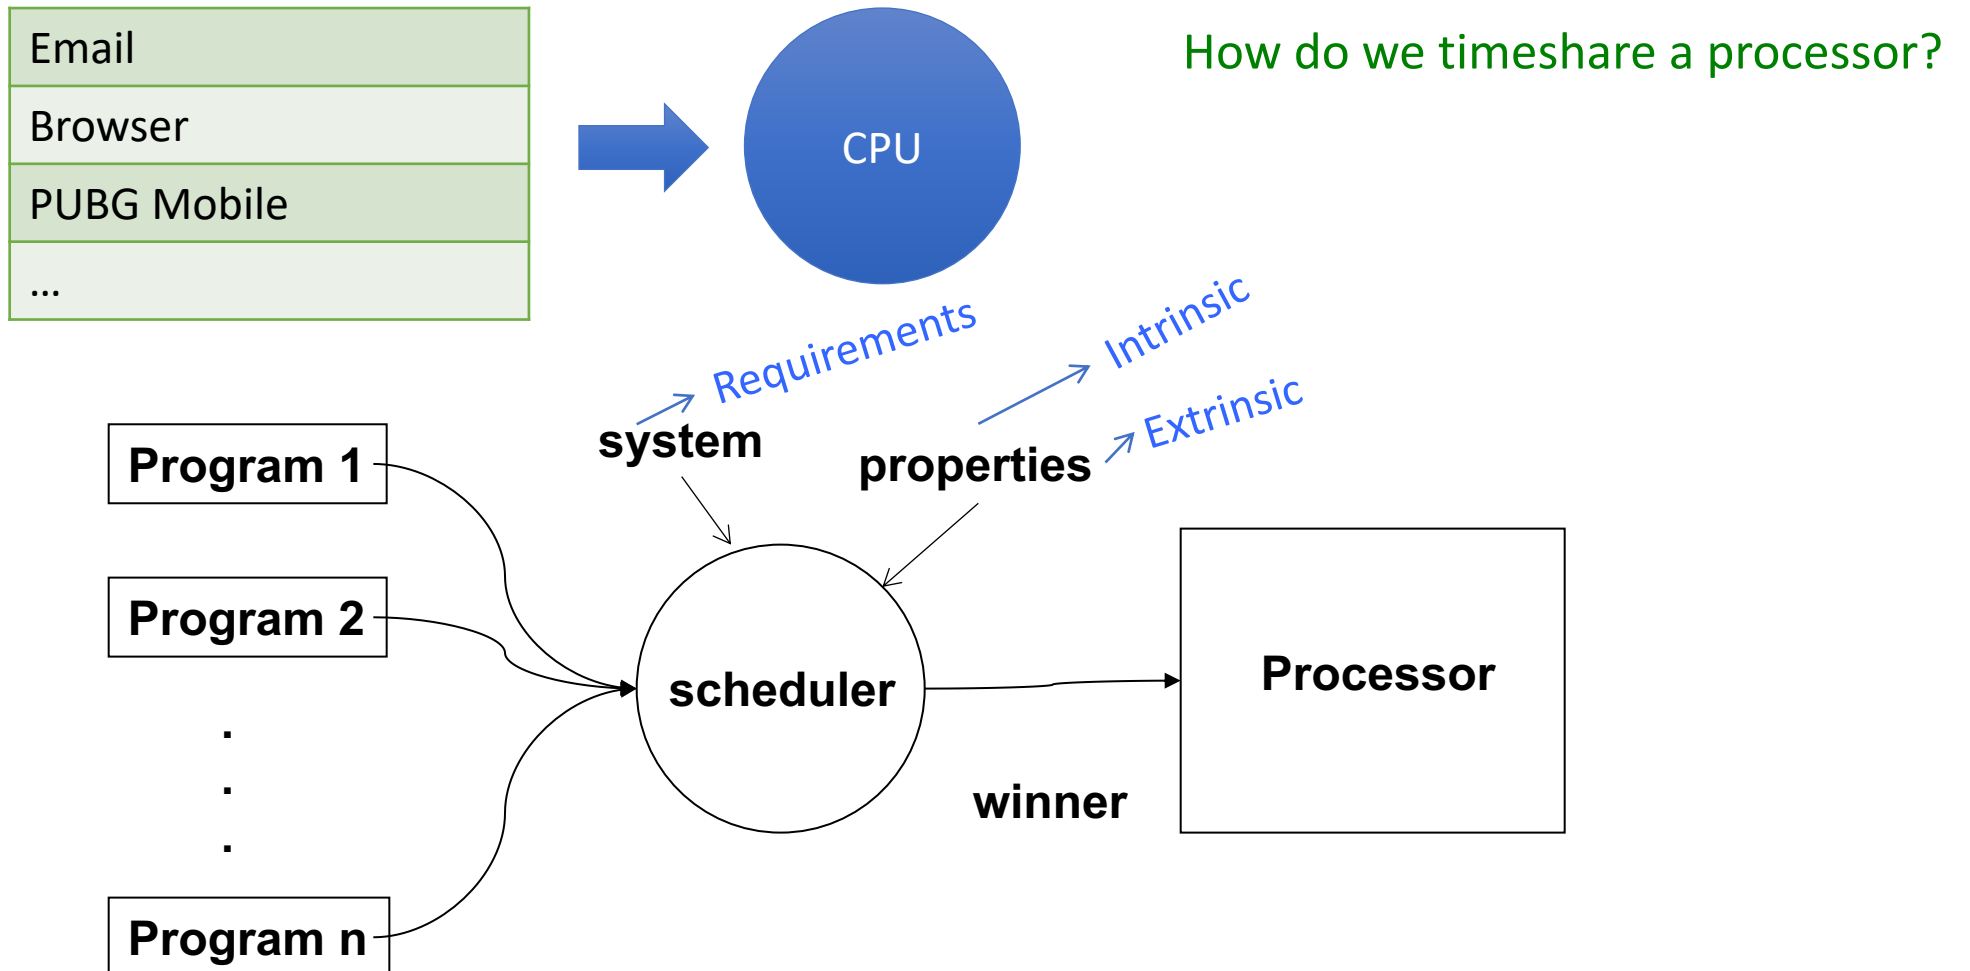
You:



How do you timeshare yourself?!

Laundry	-----
Prep & Eat	-----
Call Mom	-----
Test Prep	-----


# Multitasking in the computer world



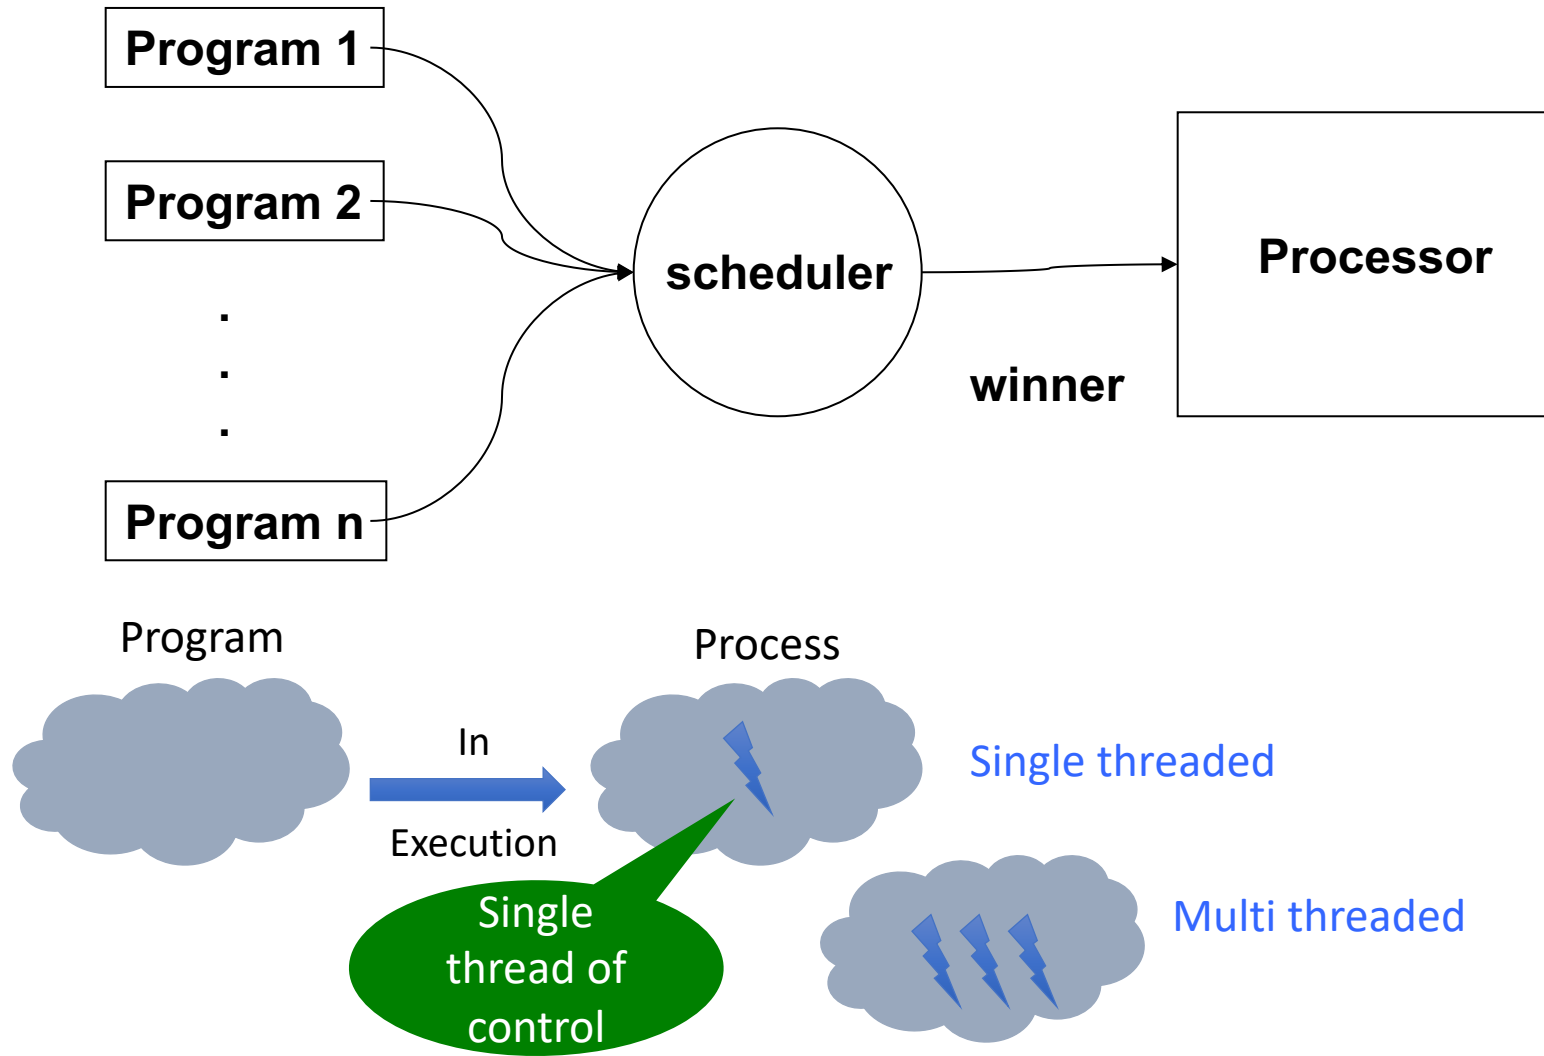


# A scheduler

- A. Is part of the operating system that makes decisions to allocate the processor resources to processes
- B. Is a piece of hardware
- C. Is an instruction in the ISA of the processor
- D. Tells me when to call Mom



# Quick aside: Threads



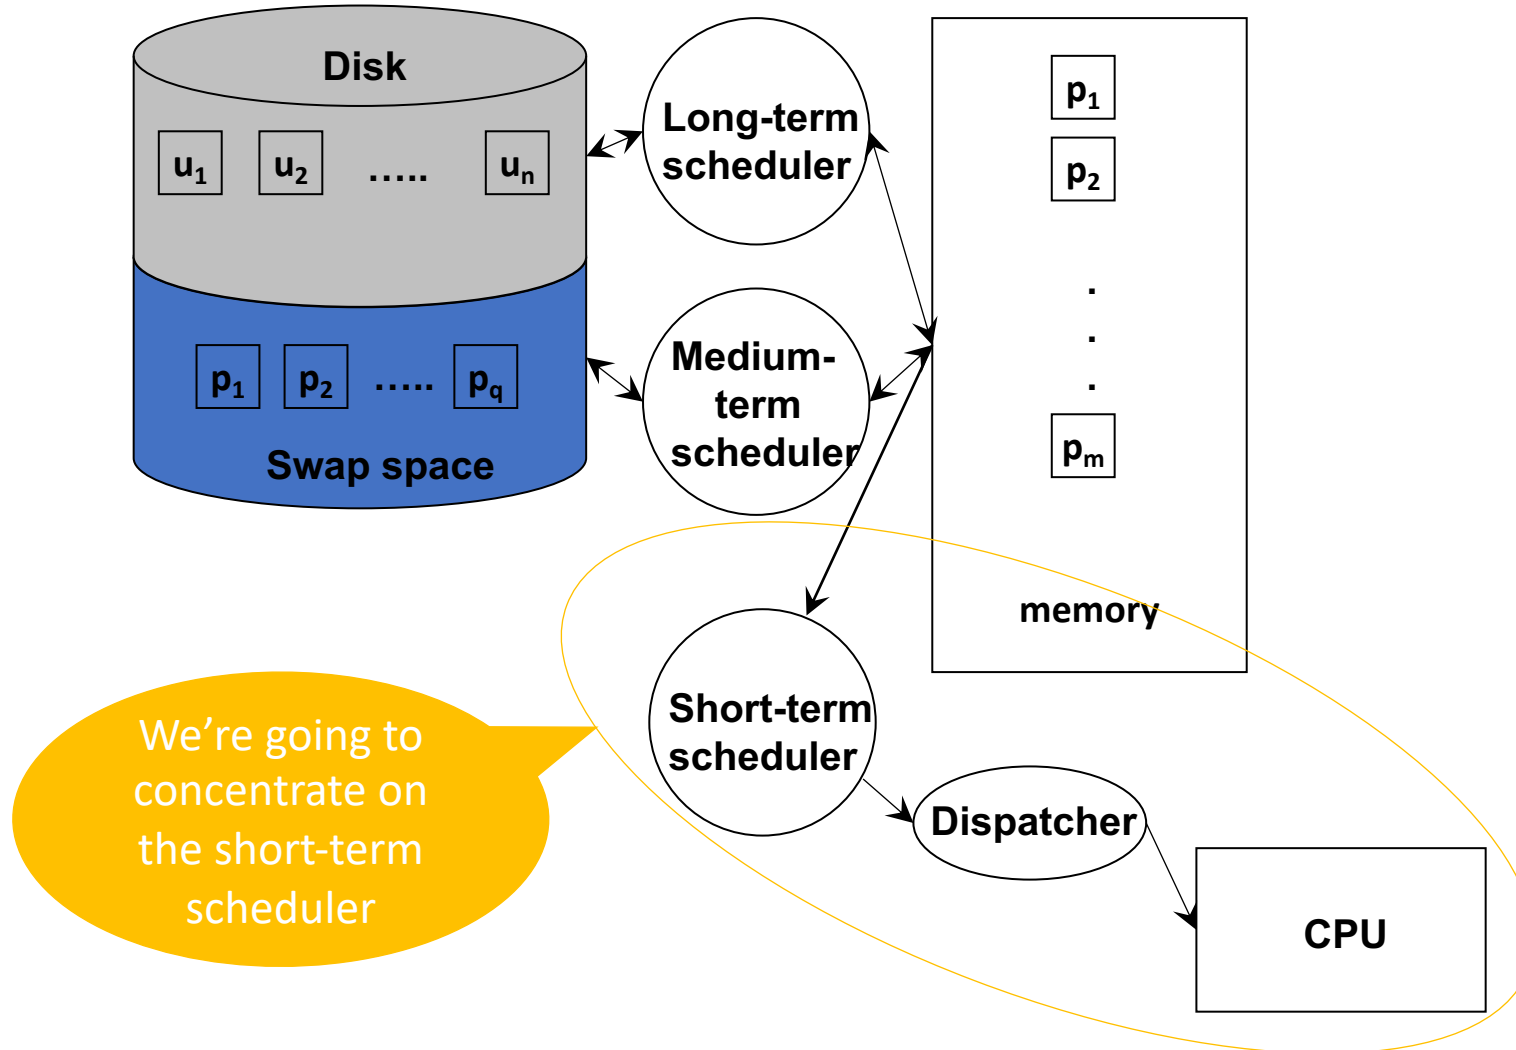


# Terminology

---

Name	Usual Connotation	Use in this chapter
Job	Unit of scheduling	Synonymous with process
Process	Program in execution; unit of scheduling	Synonymous with job
Thread	Unit of scheduling and/or execution; contained within a process	Not used in the scheduling algorithms described in this chapter
Task	Unit of work; unit of scheduling	Not used in the scheduling algorithms described in this chapter, except in describing the scheduling algorithm of Linux

# Schedulers




# Terminology

Name	Environment	Role
Loader	In every OS	Load user program from disk into memory
Long term scheduler	Batch oriented OS	Control the job mix in memory to balance use of system resources (CPU, memory, I/O)
Medium term scheduler	Every modern OS (time-shared, interactive)	Balance the mix of processes in memory to avoid thrashing
Short term scheduler	Every modern OS (time-shared, interactive)	Schedule the memory resident processes on the CPU
Dispatcher	In every OS	Populate the CPU registers with the state of the process selected for running by the short-term scheduler

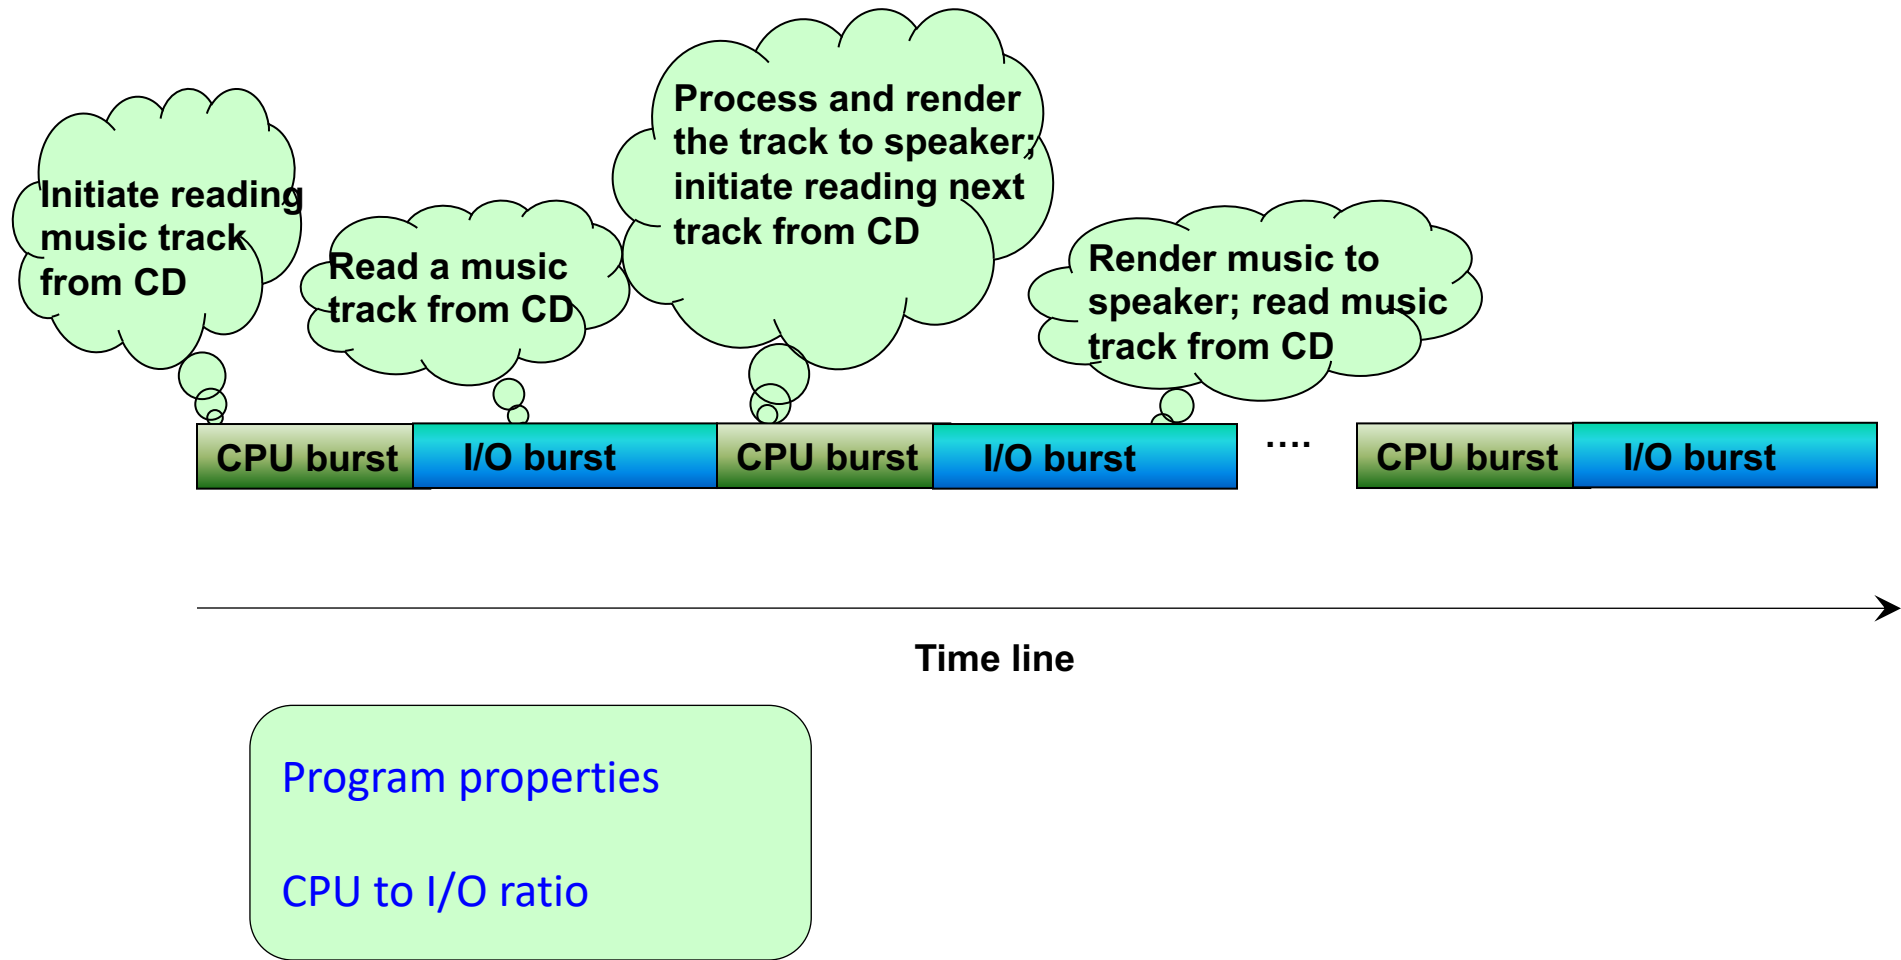


# In most modern-day operating systems

- A. There is exactly one program in residence in memory at any point in time
- B. There are multiple programs in residence in memory at any point in time
- C. There is exactly one program but multiple processes in residence in memory at any point in time
- D. I'll just choose this wrong answer

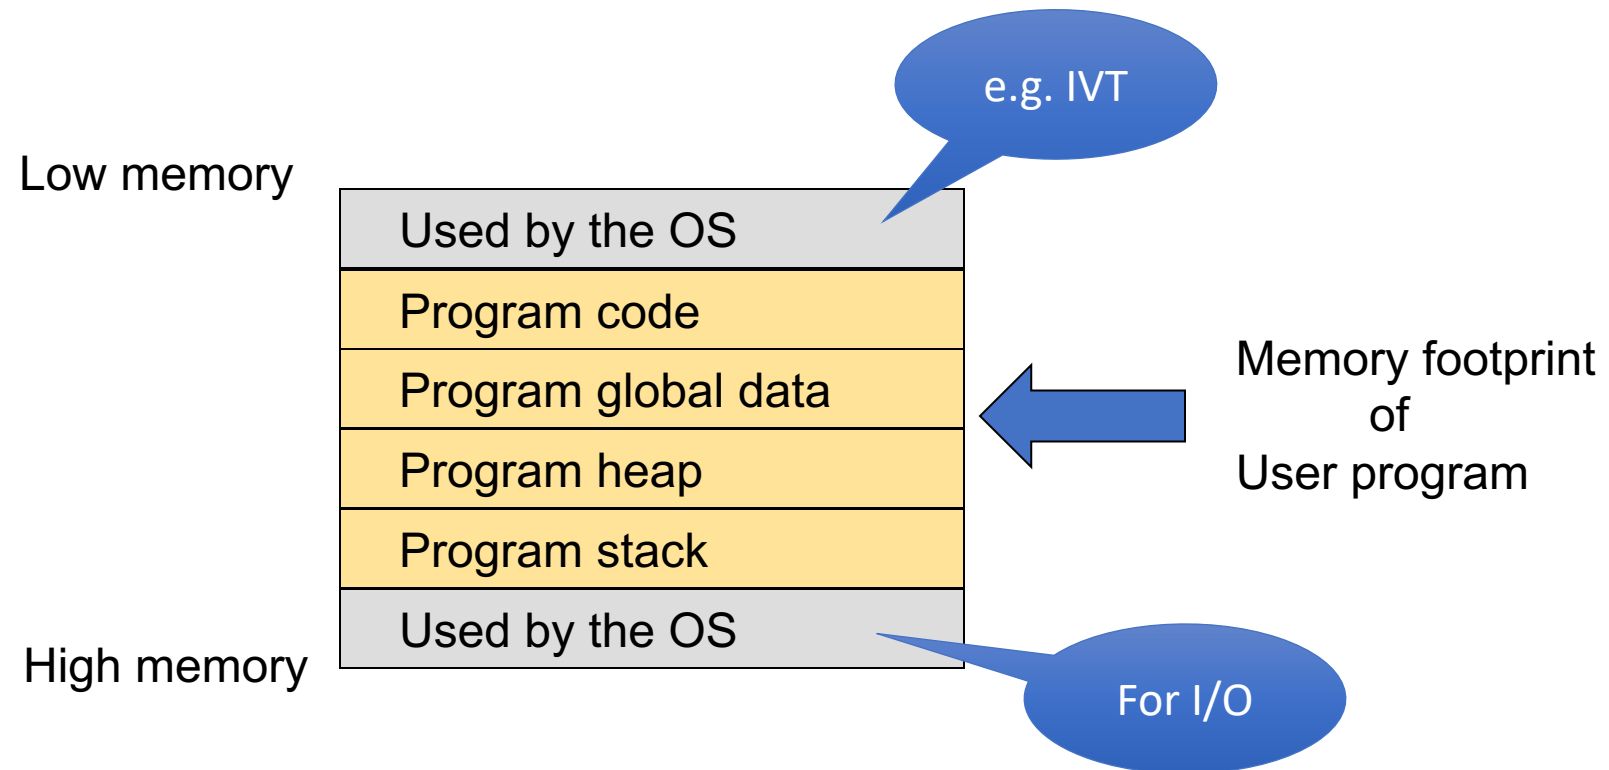


# What's the process doing?



# Recall the memory footprint of a program

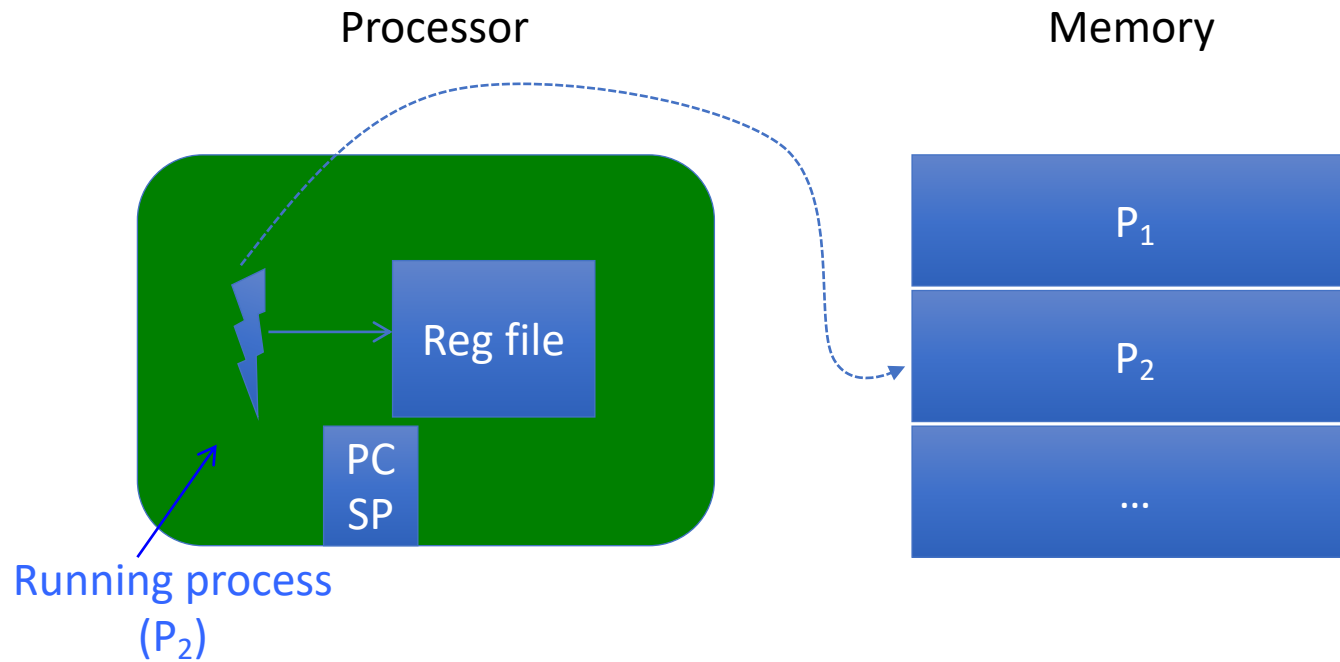
---



# Process state

Process<sub>2</sub> is running. The program is in memory. Where is its state?

Where will its state be when some other process is running?



# PCB – Process Control Block

---

We need a data structure to represent a process' state

```
enum  state_type {new, ready, running,
                  waiting, halted};
typedef struct control_block_type {
    enum state_type state;
    address PC;
    int reg_file[NUMREGS];
    struct control_block *next_pcb;
    int priority;
    address memory_footprint;
    ...
    ...
} control_block;
```





# One of these is **NOT** part of the PCB

- 20% A. General purpose registers that are visible to the instruction set
- 9% B. Program counter and the register that represents the stack pointer
- 31% C. Layout of the program in memory
- 34% D. Internal registers in the datapath of the processor
- 6% E. Priority information

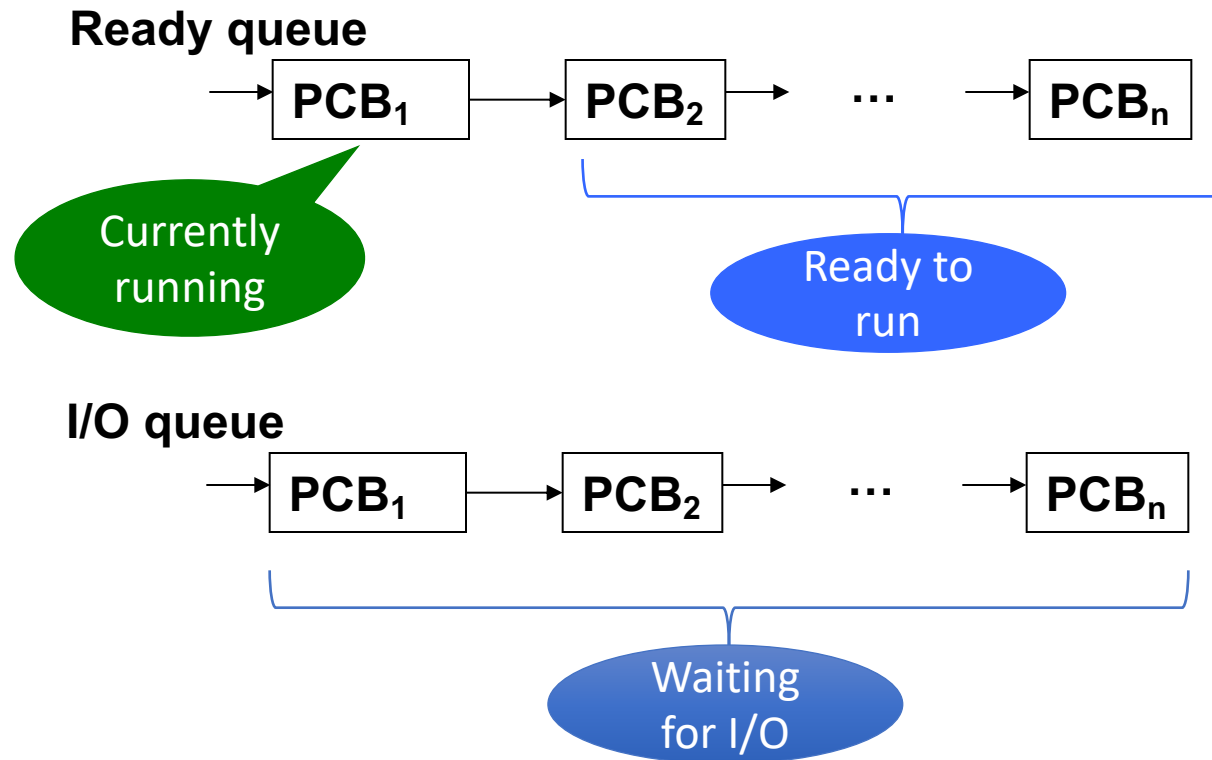
# PCB – Process Control Block

---

```
enum state_type {new, ready, running,
                 waiting, halted};
typedef struct control_block_type {
    enum state_type state;
    address PC; // where to resume
    int reg_file[NUMREGS]; // reg contents
    struct control_block *next_pcb;
    int priority; // extrinsic attribute
    address memory_footprint; // memory occupancy
    ...
    ...
} control_block;
```

} volatile state

# Data structure used by scheduler



# Steps in scheduling

---

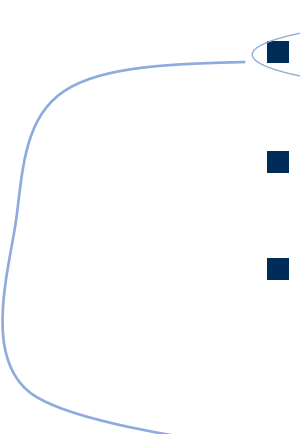
- Grab the attention of the processor
- Save the state of the current process
- Select a new process to run
- Dispatch the selected process
- Preemptive vs. non-preemptive

External interrupt  
(e.g. timer)

System call (trap), I/O  
request, process exit

# Steps in scheduling

---

- Grab the attention of the processor
  - Save the state of the current process
  - Select a new process to run
  - Dispatch the selected process
  - Dump the “state” (PC, registers) into PCB of currently running process
- 
- A blue curved arrow originates from the right side of the second list item, "Save the state of the current process", which is also enclosed in a light blue oval. The arrow curves downwards and to the right, ending with a small arrowhead pointing to the first list item of the fifth step, "Dump the 'state' (PC, registers) into PCB of currently running process".

# Steps in scheduling

---

- Grab the attention of the processor
- Save the state of the current process
- Select a new process to run
- Dispatch the selected process

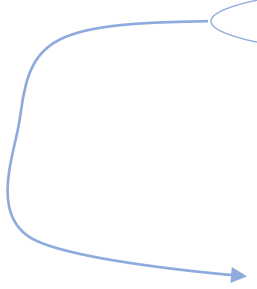


■ This is the short-term scheduling algorithm result → select a PCB to “dispatch”

# Steps in scheduling

---


- Grab the attention of the processor
- Save the state of the current process
- Select a new process to run
- Dispatch the selected process

- 
- What is “dispatch”?
    - ➔ load “state” of the selected PCB into processor registers (PC, reg file)

# Steps in scheduling

---

This whole  
process is  
called a  
“context  
switch”

- 
- Grab the attention of the processor
  - Save the state of the current process
  - Select a new process to run
  - Dispatch the selected process




Name	Description
CPU burst	Continuous CPU activity by a process before requiring an I/O operation
I/O burst	Activity initiated by the CPU on an I/O device
PCB	Process context block that holds the state of a process (i.e., program in execution)
Ready queue	Queue of PCBs that represent the set of memory resident processes that are ready to run on the CPU
I/O queue	Queue of PCBs that represent the set of memory resident processes that are waiting for some I/O operation either to be initiated or completed
Non-Preemptive algorithm	Algorithm that allows the currently scheduled process on the CPU to voluntarily relinquish the processor (either by terminating or making an I/O system call)
Preemptive algorithm	Algorithm that forcibly takes the processor away from the currently scheduled process in response to an external event (e.g. I/O completion interrupt, timer interrupt)
Thrashing	A phenomenon wherein the <b>dynamic memory usage</b> of the processes currently in the ready queue <b>exceed</b> the total <b>memory capacity</b> of the system



# A preemptive scheduling algorithm requires

- 40% A. A trap instruction
- 54% ☒ B. An external interrupt
- 3% C. The currently running process to terminate
- 3% D. The currently running process to make an I/O request

# Metrics



$w_i$ ,  $e_i$ , and  $t_i$ , are respectively the wait time, execution time, and the elapsed time (turnaround time) for a job  $j_i$


Throughput?

Avg. Turnaround Time?

Avg. Wait Time?

Response time?

# Metrics



$w_i$ ,  $e_i$ , and  $t_i$ , are respectively the wait time, execution time, and the elapsed time (turnaround time) for a job  $j_i$

System Centric	Throughput?	$3 / t_3$ jobs/sec
	Avg. Turnaround Time?	$(t_1 + t_2 + t_3) / 3$ sec
	Avg. Wait Time?	$(w_1 + w_2 + w_3) / 3$ sec
User Centric	Response time?	$R_{P1} = t_1, R_{P2} = t_2, R_{P3} = t_3$

Name	Notation	Units	Description
CPU Utilization	-	%	Percentage of time the CPU is busy
Throughput	$n/T$	Jobs/s	System-centric metric quantifying the number of jobs $n$ executed in time interval $T$
Avg. Turnaround time ( $t_{avg}$ )	$(t_1 + t_2 + \dots + t_n)/n$	Secs	System-centric metric quantifying the average time it takes for a job to complete
Avg. Waiting time ( $w_{avg}$ )	$(w_1 + w_2 + \dots + w_n)/n$	Secs	System-centric metric quantifying the average waiting time that a job experiences
Response time	$t_i$	Secs	User-centric metric quantifying the turnaround time for a specific job $I$
Variance in Response time	$E[(t_i - t_{avg})^2]$	Secs <sup>2</sup>	User-centric metric that quantifies the statistical variance of the actual response time ( $t_i$ ) experienced by a process ( $P_i$ ) from the expected value ( $t_{avg}$ )
Starvation	-	-	User-centric qualitative metric that <b>signifies denial of service</b> to a particular process or a set of processes due to some <b>intrinsic property</b> of the <b>scheduler</b>
Convoy effect	-	-	User-centric qualitative metric that results in a <b>detrimental effect</b> to some set of processes due to some <b>intrinsic property</b> of the <b>scheduler</b> [This often appears as a “convoy” of short jobs waiting for the completion of a long job; non-preemptive FCFS is the convoy effect’s native habitat.]



# The most user-centric metric of a scheduler is...

- 6% A. Throughput
- 28% B. Average waiting time
- 11% C. Average turnaround time
- 3% D. CPU utilization
- 50% E. Response time
- 3% F. None of the above