

CS2200  
Systems and Networks  
Spring 2022

Lecture 16:  
Virtual Memory pt 2

Alexandros (Alex) Daglis  
School of Computer Science  
Georgia Institute of Technology  
[adaglis@gatech.edu](mailto:adaglis@gatech.edu)

# Announcements

---

- Don't forget to sign up for Project 2 demo

# Recap

---

Hardware	Software
Fence register	User/kernel split
Bounds register	Static relocation
Base + limit register	Dynamic relocation

## Next

Allocation policies

Paging

# Memory allocation by OS

---


- Fixed size partition
- Variable size partition
- Both use the hardware base + limit registers

# Fixed size partitions

**OS memory manager allocation table**

Occupied bit	Partition Size	Process
0	5K	---
0	8K	---

memory




```
Struct AT_entry {  
    int occupied;  
    int size;  
    int pid;  
};
```

# P1 needs 6K of memory

**Allocation table**

Occupied bit	Partition Size	Process
0	5K	---
1	8K	P1

**Memory**




# P1 needs 6K of memory

**Allocation table**

**Table size is fixed = number of partitions**

Occupied bit	Partition Size	Process
0	5K	---
1	8K	P1

**Memory**



Needs only  
6K


# Internal fragmentation

**Internal fragmentation** = size of partition – actual used

**Allocation table**

Occupied bit	Partition Size	Process
0	5K	---
1	8K	P1

**Memory**



Needs only  
6K

Wasted  
(internal  
fragmentation)



# Another process needs 6K memory

**Do we have it?**


**Memory manager has only a 5K partition...**

**Not possible**

**Allocation table**

Occupied bit	Partition Size	Process
0	5K	---
1	8K	P1

**Memory**




# External fragmentation

**Consider P3 that needs 4K of memory**  
**Is it possible to allocate?**  
**Memory is available, but not contiguous**

**Allocation table**

Occupied bit	Partition Size	Process
1	1K	P1
0	2K	---
1	3K	P2
0	2K	---

**Memory**



External fragmentation

=  $\sum$  All non-contiguous free memory partitions

And we have 4K of non-continuous memory here

Which gives us 4K of **external fragmentation**

# Fragmentation

---

- Internal fragmentation
  - Size of partition – actual memory used
- External fragmentation
  - **$\Sigma$  All non-contiguous free partitions**

# Fixed size partition memory management

---

## Pros

- Simplicity

## Cons

- Fragmentation
  - Internal
  - External




# Total internal fragmentation is...

**Allocation table**

Occupied bit	Partition Size	Process
1	5K	P2 (need 2k)
1	8K	P1 (need 6K)

**Memory**



- A. 2K
- B. 3K
- C. 5K
- D. 8K




# Total external fragmentation is...

**Allocation table**

Occupied bit	Partition Size	Process
1	5K	P2 (need 2k)
1	8K	P1 (need 6K)

**Memory**



- A. 0K
- B. 2K
- C. 5K
- D. 8K

# Overcoming internal fragmentation

---

- Allocate exactly what is needed
- Variable size partitions

# Variable size partitions


---

**Memory manager allocation table**

<b>Start address</b>	<b>Size</b>	<b>Process</b>
<b>0</b>	<b>13K</b>	<b>FREE</b>

```
Struct AT_entry {  
    int start;  
    int size;  
    int pid;  
};
```

**memory**






# Partition table a little while later

Grows and shrinks as partitions get created and released  
→ No internal fragmentation!

**Allocation table**

Start address	Size	Process
0	2K	P1
2K	6K	P2
8K	3K	P3
11K	2K	FREE

**Memory**




# P1 exits

Allocation table

Start address	Size	Process
0	2K	P1 → FREE
2K	6K	P2
8K	3K	P3
11K	2K	FREE

Memory



External fragmentation




New request:  
P4 needs 4K  
Possible?

# P2 exits

Allocation table

Start address	Size	Process
0	2K	FREE
2K	6K	P2 → FREE
8K	3K	P3
11K	2K	FREE

Memory




# Coalescing two free areas

**Allocation table**

Start address	Size	Process
0	8K	FREE
8K	3K	P3
11K	2K	FREE

**Memory**




# Reducing external fragmentation

- Best fit algorithm
  - A little better memory utilization
- First fit algorithm
  - A little quicker but less space efficient in the average case

**Allocation table**

Start address	Size	Process
0	8K	FREE
8K	3K	P3
11K	2K	FREE

**Memory**




# Compaction

- Request for 9K

**Allocation table**

Start address	Size	Process
0	8K	FREE
8K	3K	P3
11K	2K	FREE

**Memory**




# Compaction

- Relocate P3
- Create contiguous space
- Note this is a rather expensive action

**Allocation table**

Start address	Size	Process
0	3K	P3
3K	10K	FREE

**Memory**





# With variable size partition memory management there is ...

---

- A. No external fragmentation
- B. No internal fragmentation
- C. No fragmentation
- D. Both internal and external fragmentation



# External fragmentation with variable size partitions


---

- Can limit full usage of memory resources
- Compaction is too expensive

# How might we solve the external fragmentation problem?

---

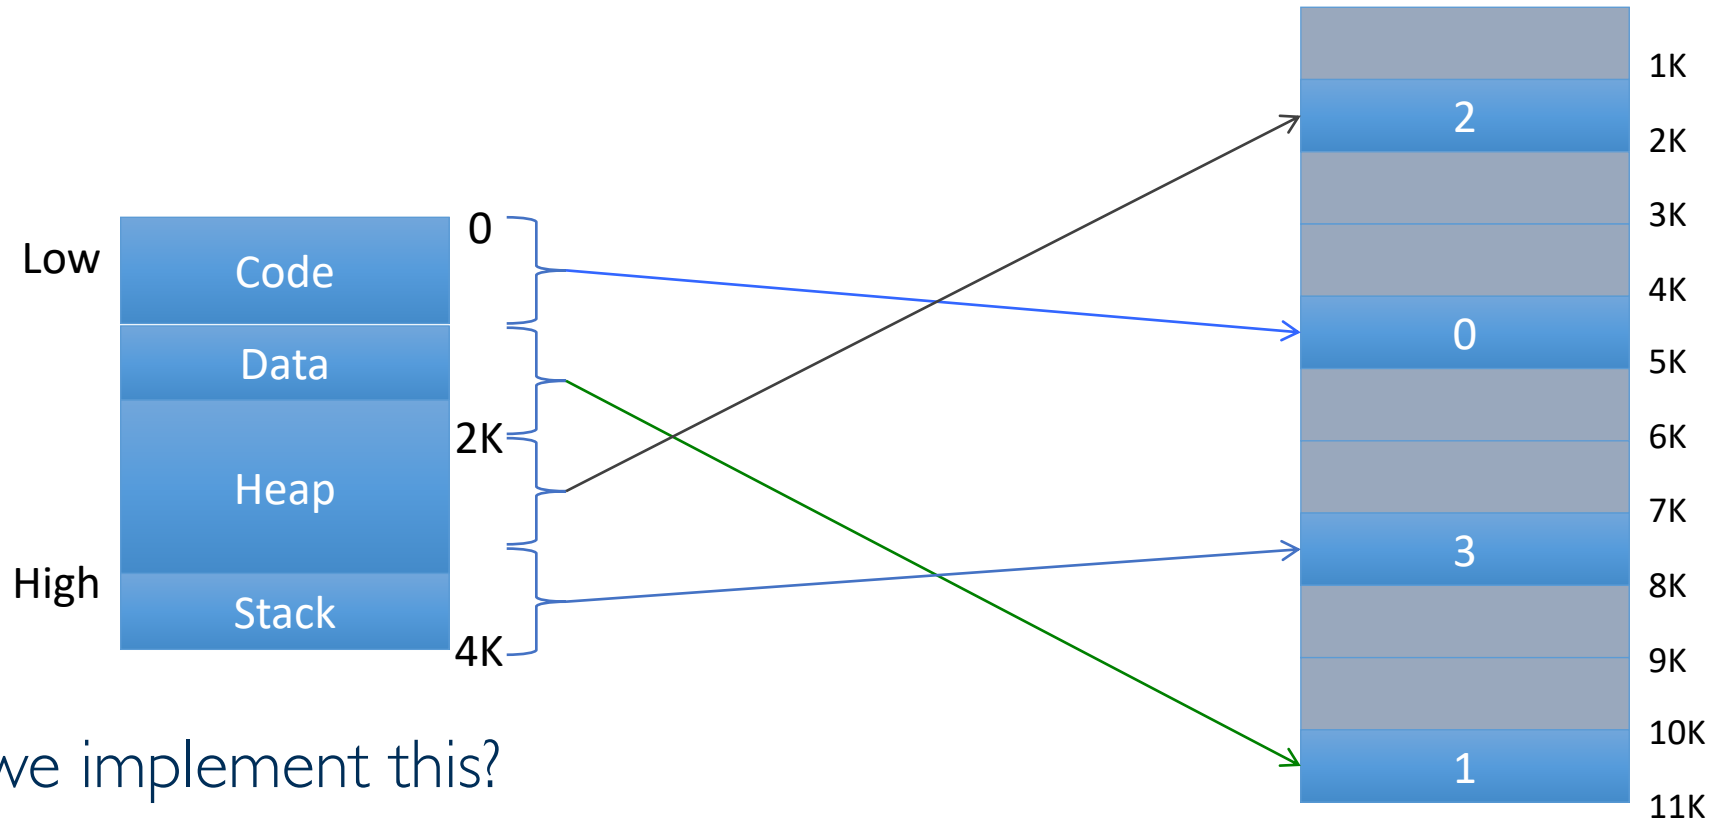
- Our process' memory footprint looks like this



- What's the main limiting assumption?
- That the process address space is contiguous!


# How might we solve the external fragmentation problem?

- What if we could store our memory footprint in non-contiguous memory locations?



- But how can we implement this?

# Use more sophisticated broker between CPU and memory




# Broker

---

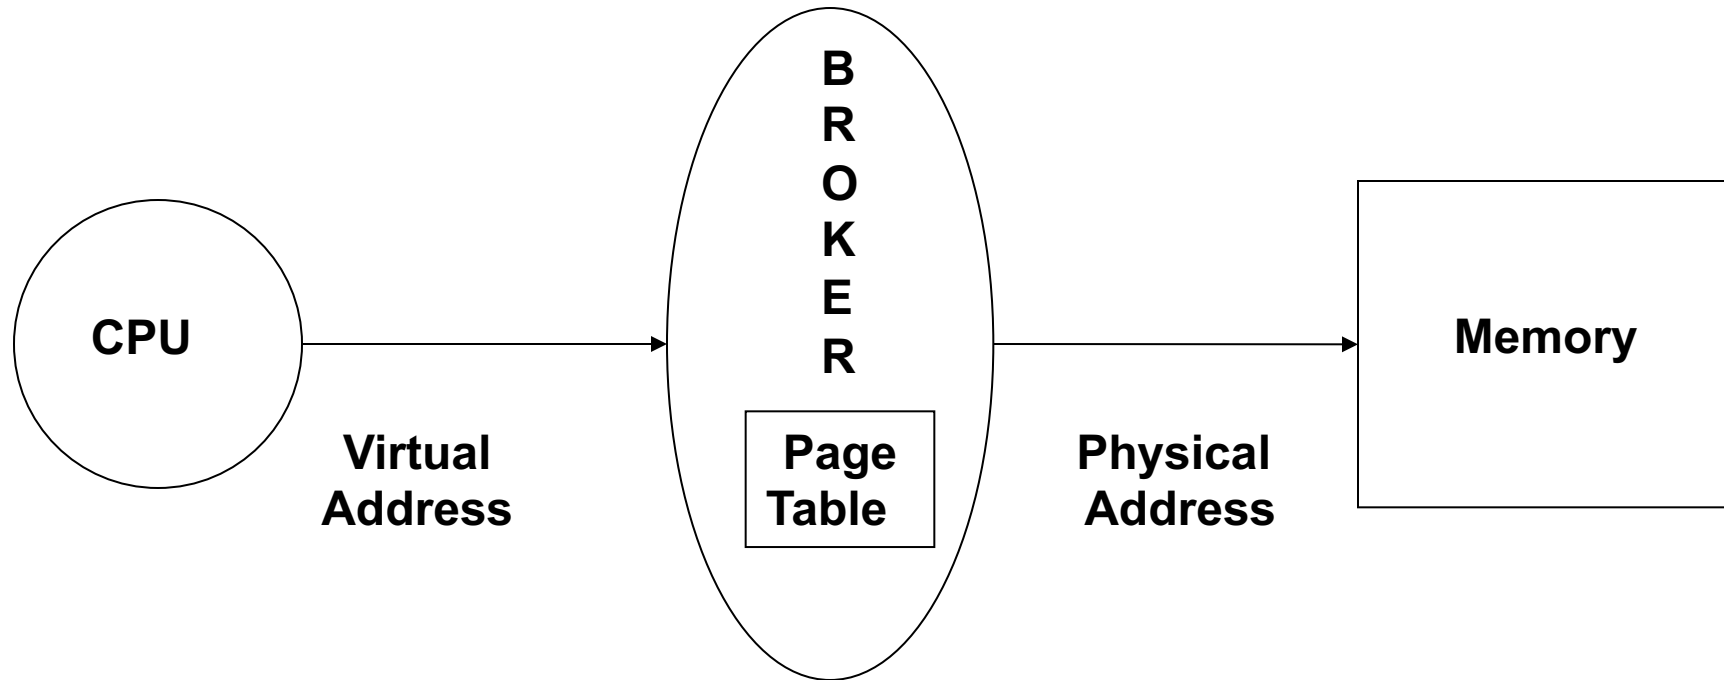
This broker maps

- Virtual address (VA) from the CPU  
to
- Physical address (PA) in memory



# Broker

- How does Broker map VA to PA?
- Perhaps like a phone directory?
  - Who sets it up? ➔ **The OS**
  - Who looks it up? ➔ **The hardware, on every access**



# How big is this table?

---

- At the lower limit, we could map the whole program
  - There would be only one entry in the page table
  - Isn't that the same as Base + Limit?
- At the upper limit, we could map every word
  - The table (# of entries) would be the size of the address space divided by the word size
  - Not practical at all
- So to strike a balance, we choose a **page size** to map
  - Bigger pages get us more **internal fragmentation** (average is  $\frac{1}{2}$  of the last page)
  - Smaller pages get us a bigger page table and take more CPU time to manage it

# Choosing a page size


---

- When memory was expensive (and small)
  - Page sizes were 512 to 2048 bytes
- These days
  - 4 KB up to 1 GB
  - Often it's configurable per process
- Page size is always a power of 2
  - The power of two allows us to split the VA into **virtual page number** (VPN) and **offset** within the page at a bit boundary
  - If the page size is  $2^n$ , the lower **n** bits are the **offset** within the page  
and bits **n** and up are the **virtual page number**




# Splitting up a virtual address

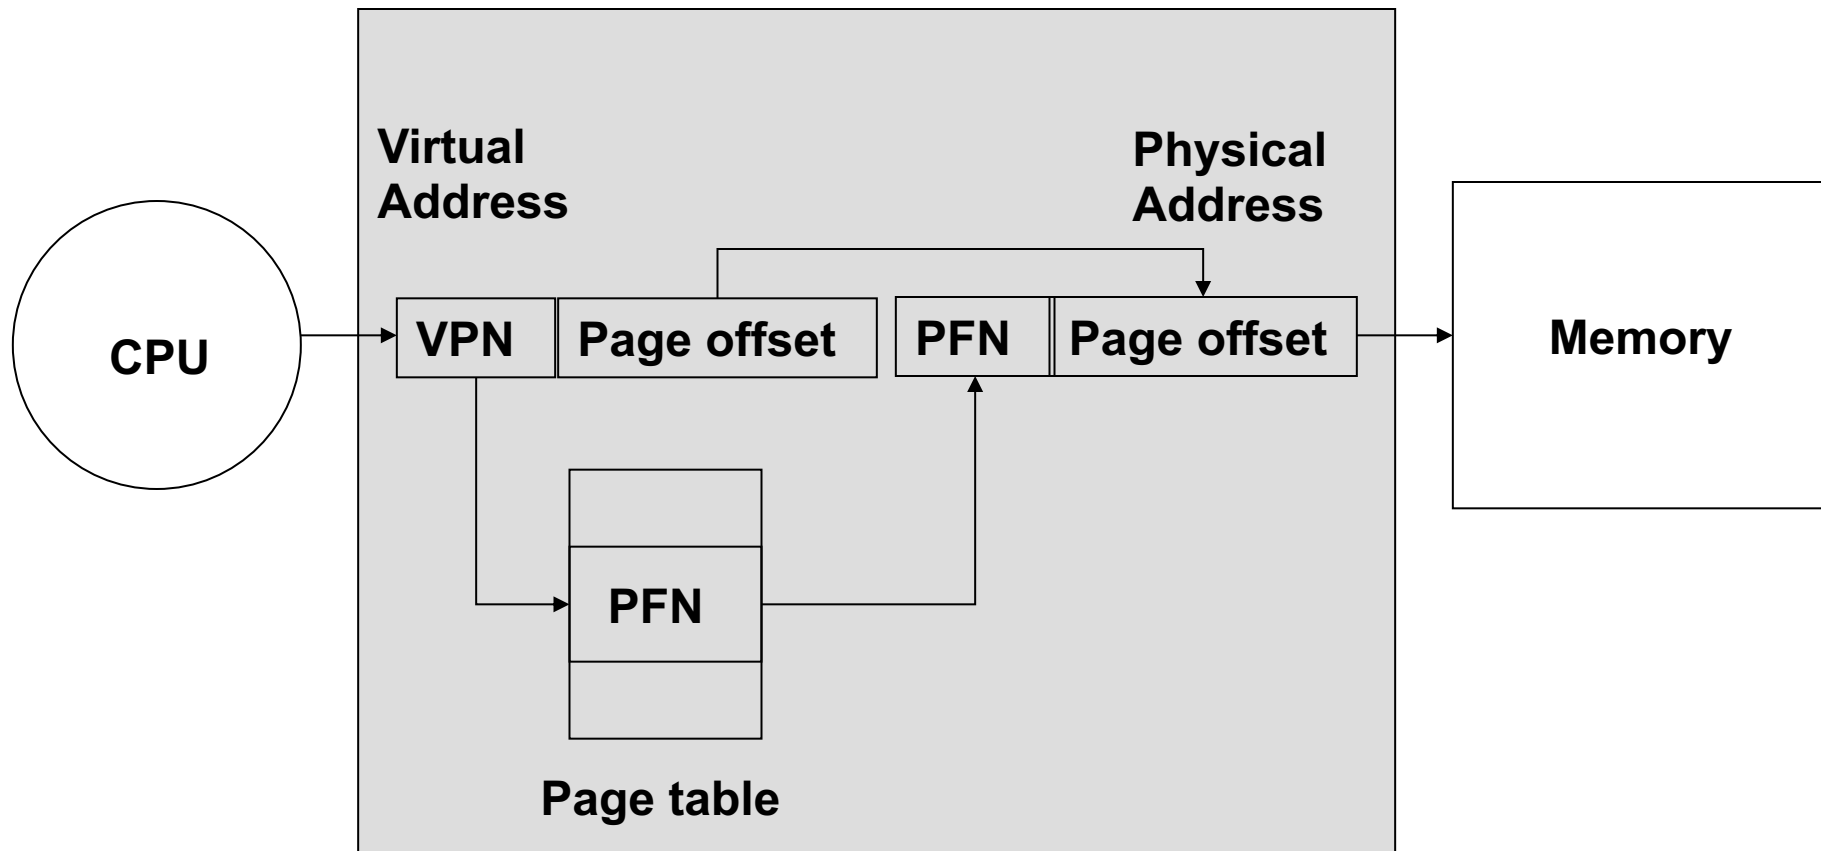
- Say we have a 4KB page size and a 32 bit virtual address space
  - 4KB is  $2^{12}$ , so the bottom 12 bits are offset and the top 20 bits are VPN
- For example, for virtual address 0x00004FFF:



# Page Table in Use




# Address translation



# Examples

- Consider a memory system with a 32-bit virtual address. Let us assume the pagesize is 8K Bytes.
- How big is the page table (# of entries)?
- $8k = 2^{13}$



- VPN is 19 bits, so page table is  $2^{19}$  or 524,288 entries

# Examples

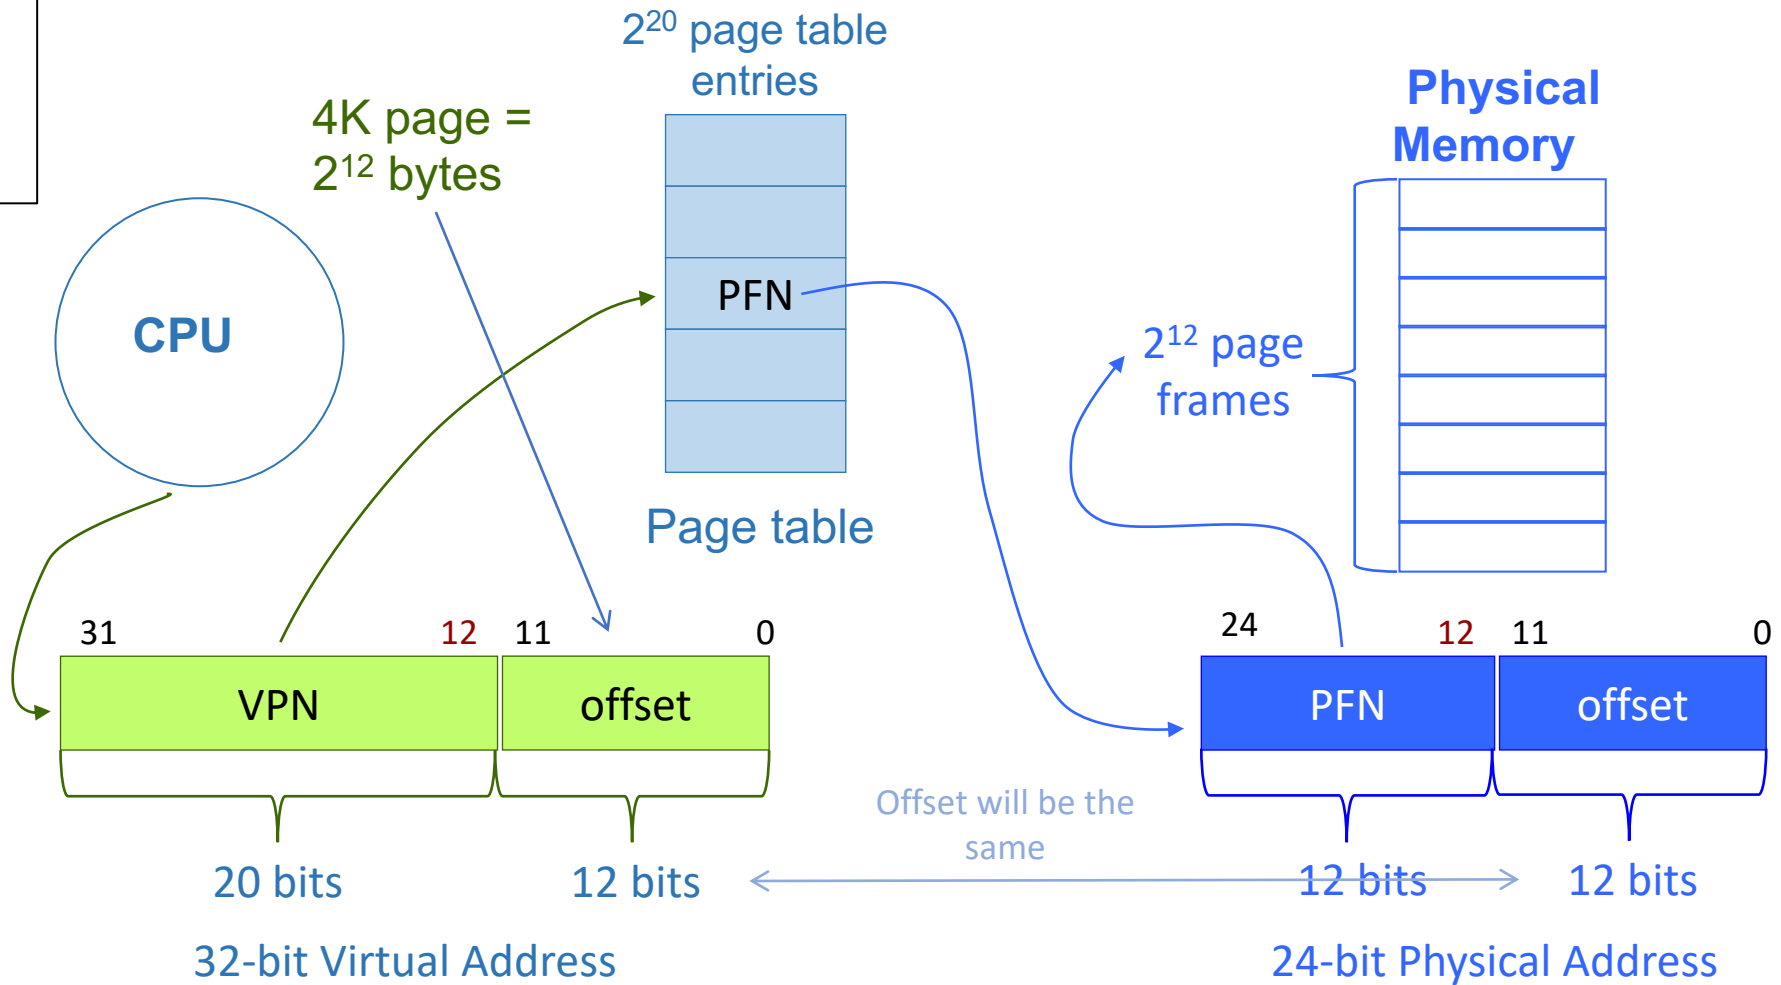
---

Consider a memory system with **32-bit virtual addresses** and **24-bit physical memory addresses**. Assume that the pagesize is 4K Bytes.

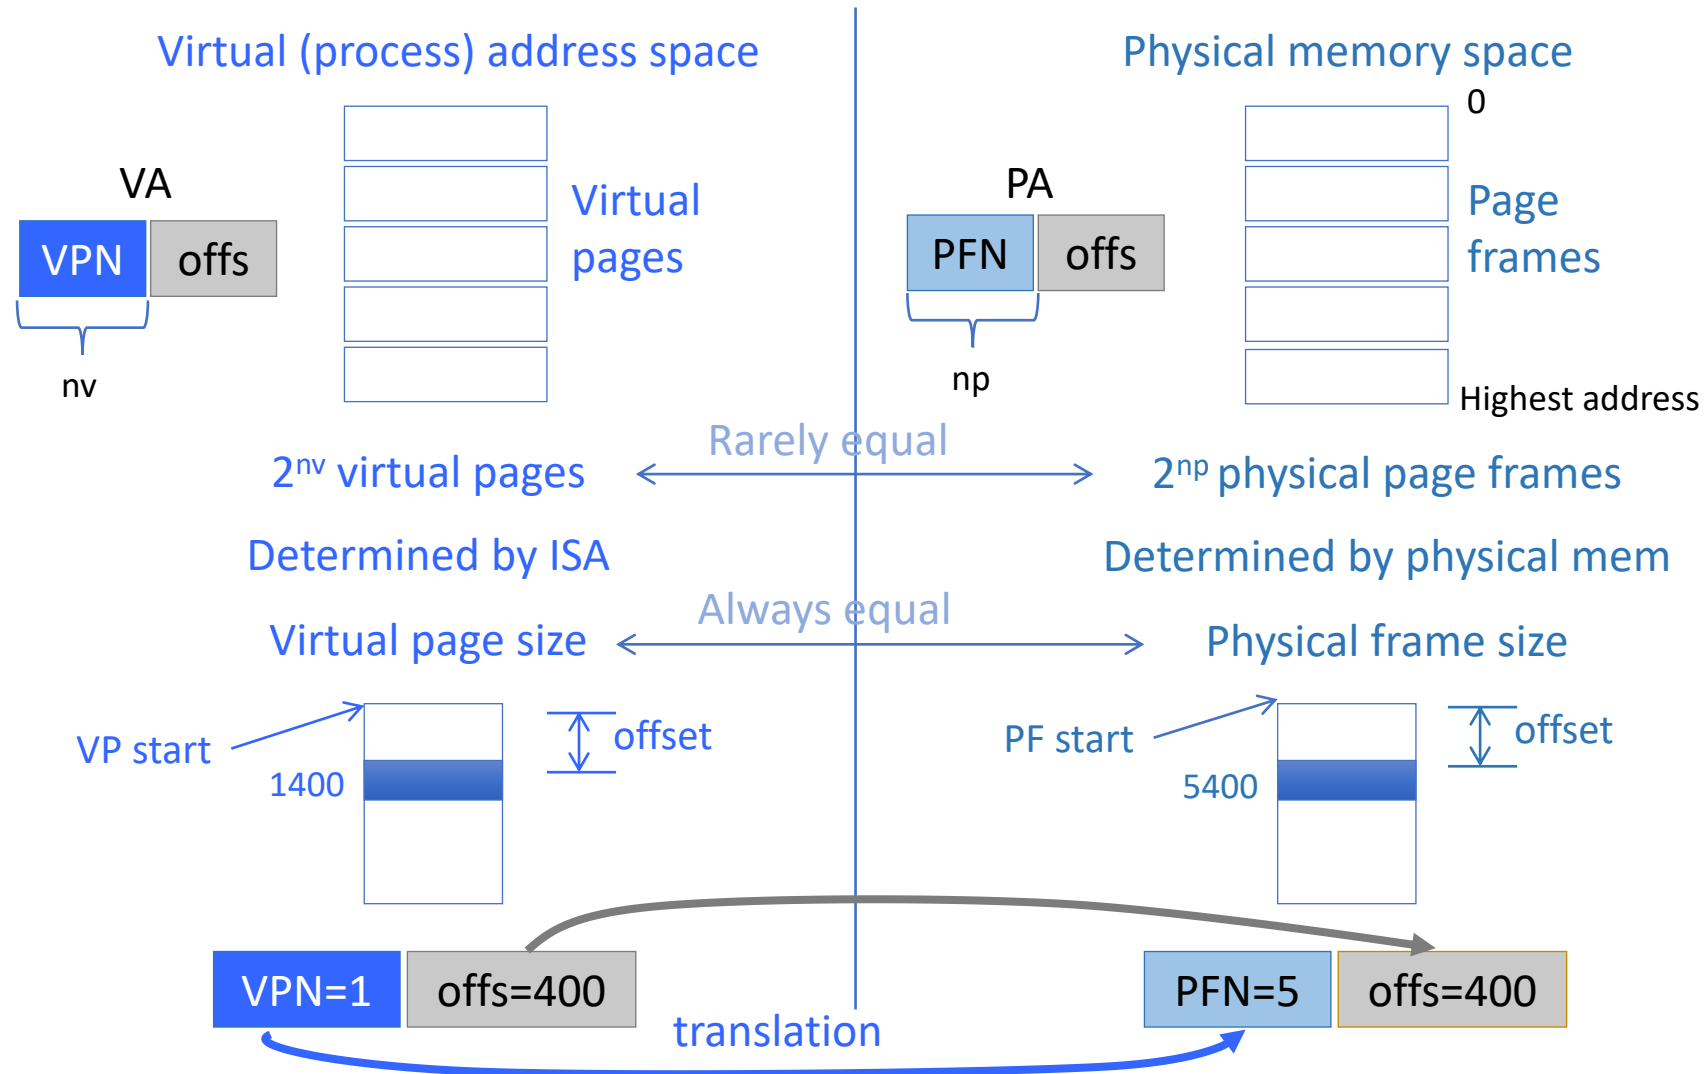
- How many page frames in memory?
- How big is the page table?
- How many page frames are there in this memory system?

# Example

32-bit VA  
24-bit PA  
4KB page size




# Important facts about paging



# So exactly where is the page table?

---



Hardware!


But it's not a special device

It's in physical memory!



# How many page tables are there?

---

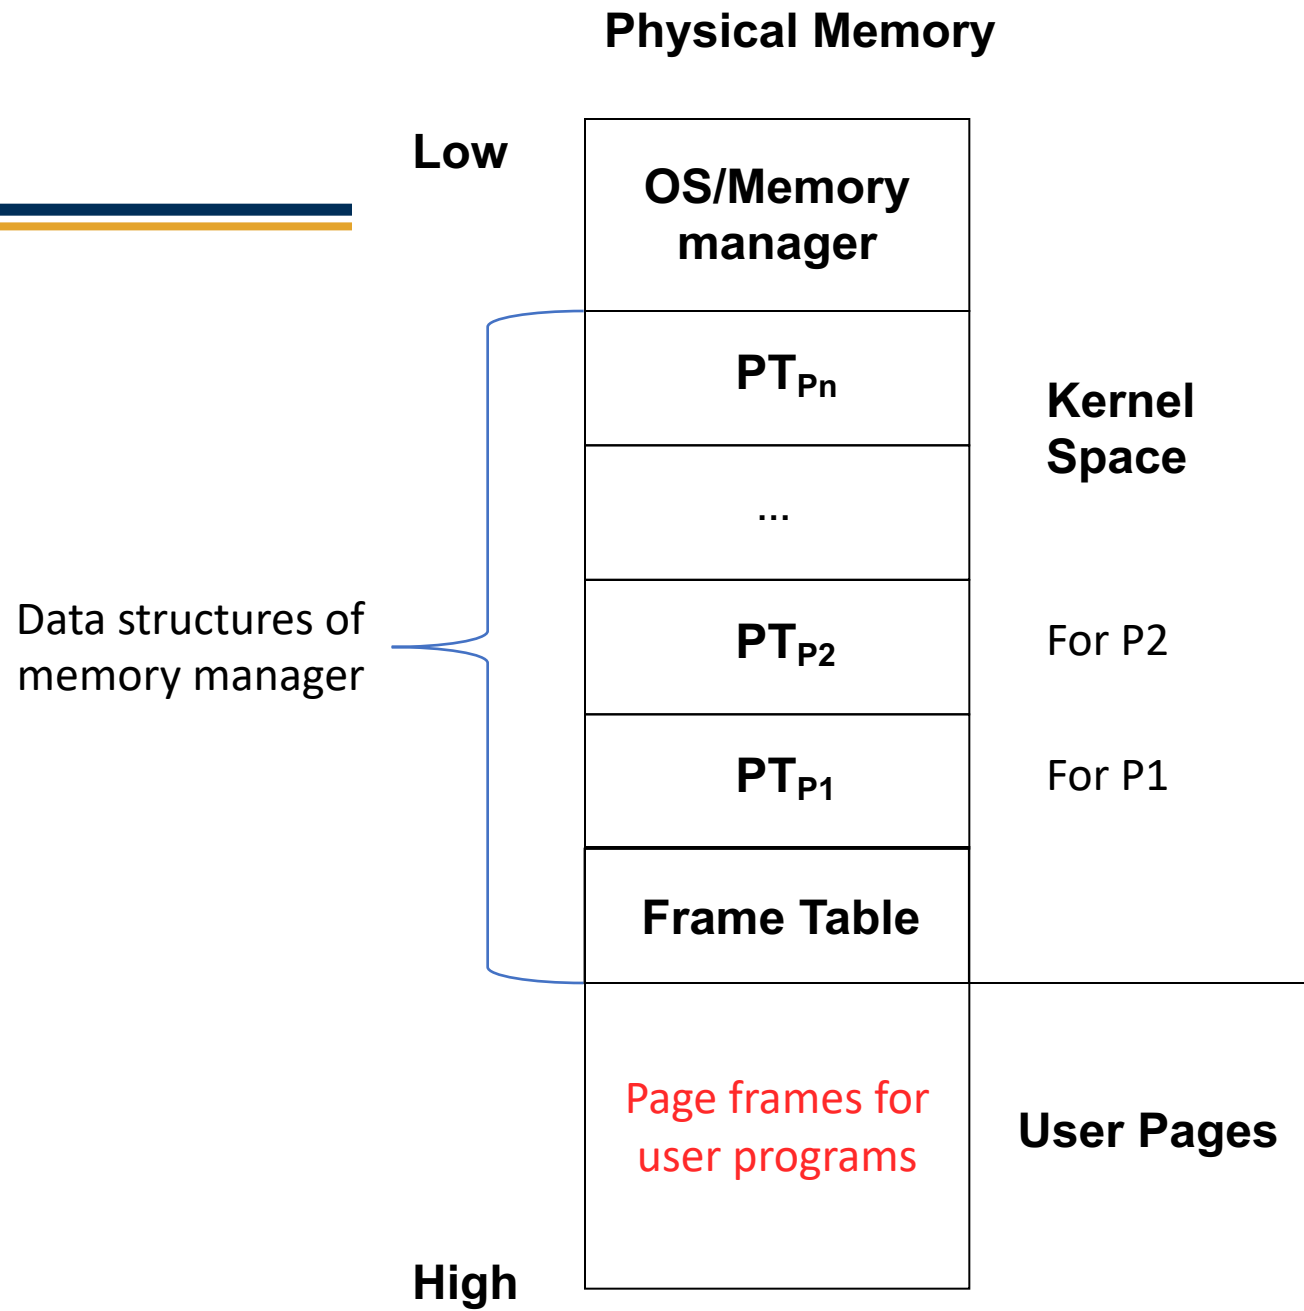


Process 1, 2, 3, ..., n in memory

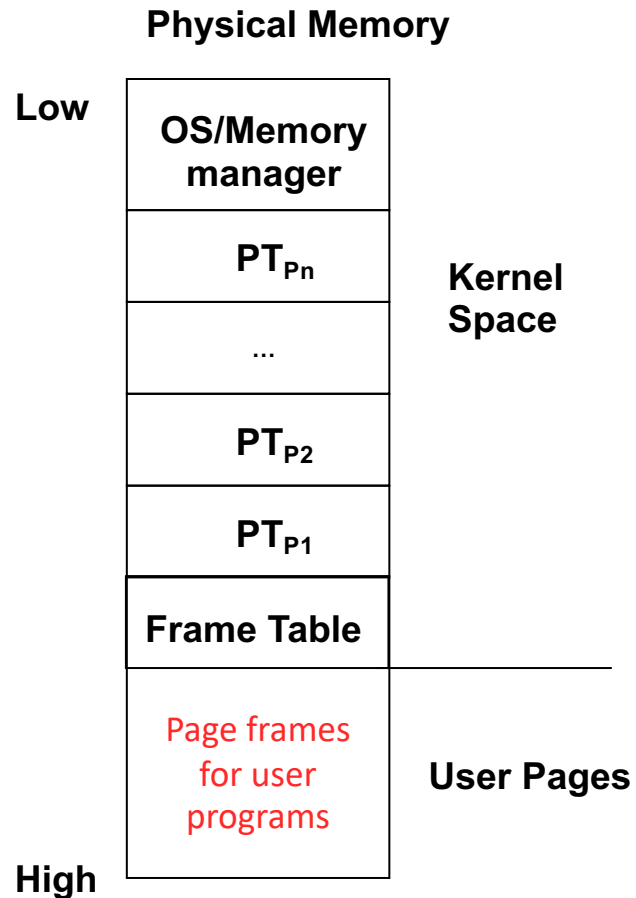
We'll need  $n$  page tables!

We need as many page tables as the number of processes!

# Physical memory layout



# What hardware assist does LC-2200 need?



- Just one new register: PTBR
- This holds the base physical address of the page table for the running process
- And of course hardware to look up the PFN from the page table for each memory reference

# PCB

---

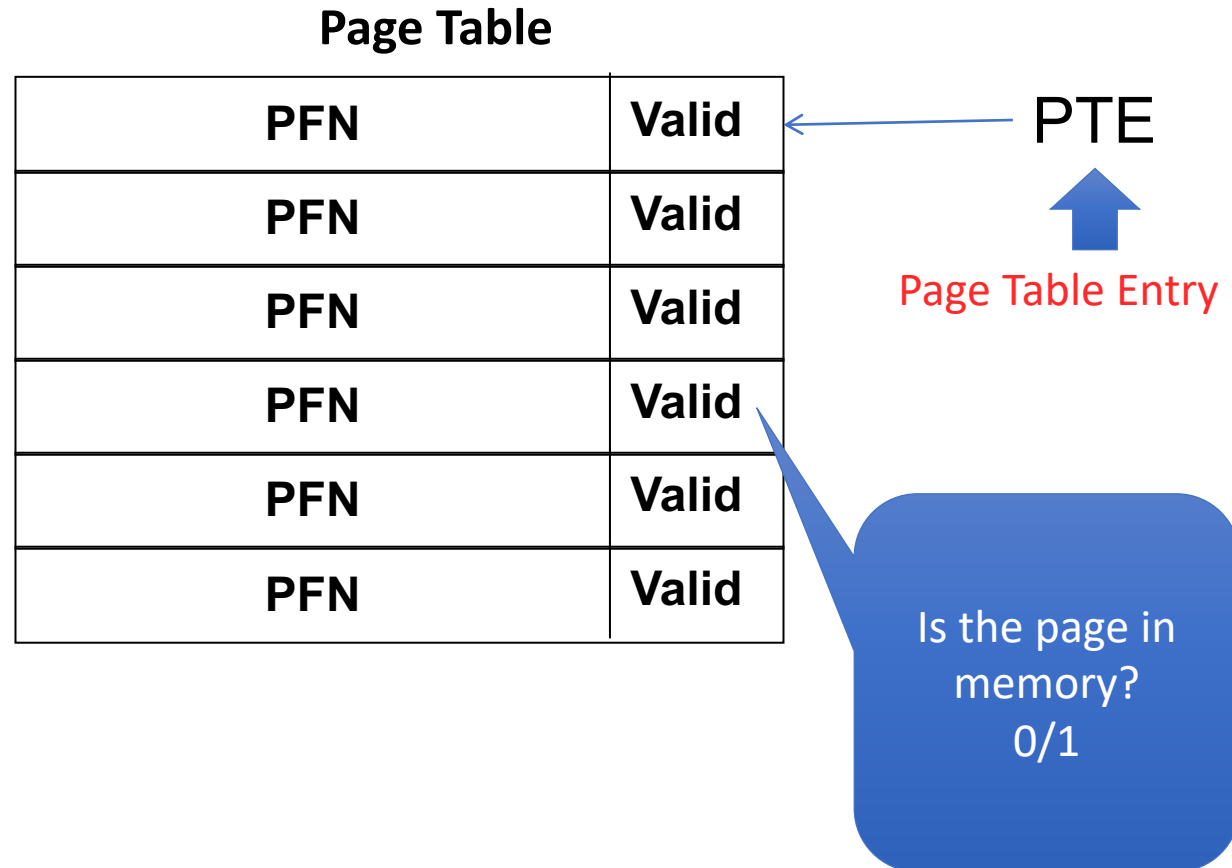
```
enum    state_type {new, ready, running,
                  waiting, halted};
typedef struct control_block_type {
    enum state_type state;
    address PC;
    int reg_file[NUMREGS];
    struct control_block *next_pcb;
    int priority;
    address PTBR;
    ...
    ...
} control_block;
```

# Paged memory allocation

---

- Allocate all at once at load time? → Not good: slow
- Allocate on demand? → Better utilization

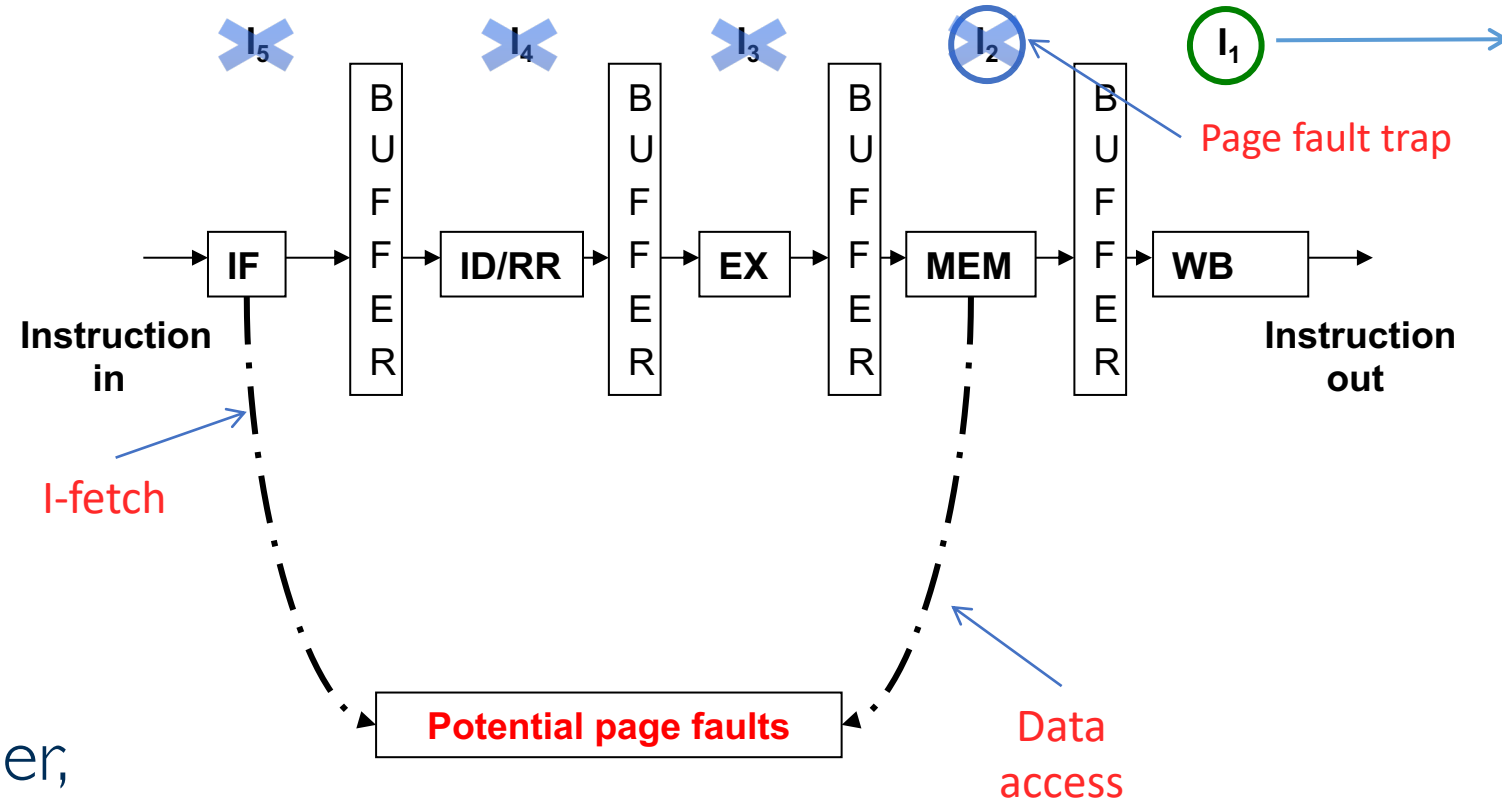
# Demand paging



When VPN's corresponding entry not valid: **Page Fault!**

# Ramification of demand paging

Where can a page fault hit us?



Let I<sub>1</sub> complete  
Squash I<sub>2</sub>-I<sub>5</sub>  
Trap to page fault handler,  
saving PC of I<sub>2</sub> for restart  
after page fault is serviced

As you may have guessed, we'll need to make the original PC value part of the pipeline buffers

# Page fault handler

---


1. Find a free page frame
2. Load the faulting virtual page from disk into the page frame
3. Give up the CPU while waiting for the paging I/O to complete
4. Update the page table entry for the faulting page
5. Place the PCB of the process back in the ready\_q of the scheduler
6. Call the scheduler



# Three Data structures for page fault handler

## (1/3): Freelist

```
struct pframe {  
    address PFN;  
    ...  
    pframe *next;  
};
```



# (2/3): Frame Table

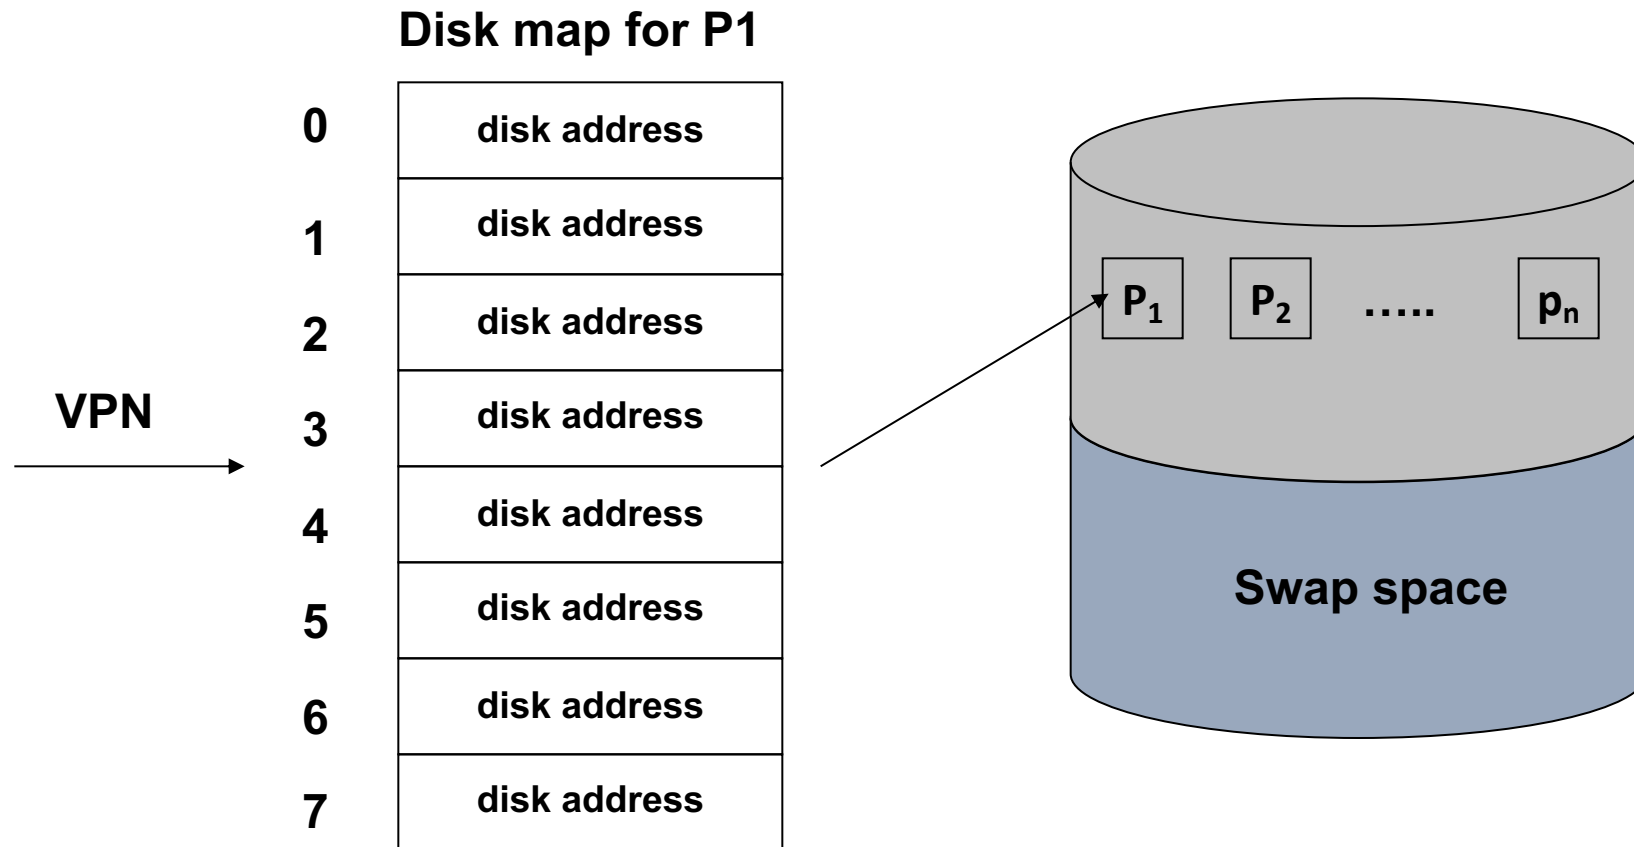
PFN →	0	<P2, 20>	<PID, VPN> →
	1	free	
	2	<P5, 15>	
	3	<P1, 32>	
	4	free	
	5	<P3, 0>	
	6	<P4, 0>	
	7	free	

Number of entries =  
number of physical  
memory page frames

Reverse mapping compared to PT

We need this for evicting pages

# (3/3): Disk Map



# Virtual memory manager data structures

---

Per process	PCB	Holds saved PTBR register
	Page table	VPN → PFN mapping Dual role: Memory manager uses it for setup Hardware uses on each memory access
	Disk map	VPN to disk block mapping needed for bringing missing pages from disk to physical memory
Per system	Free list	Free page frames in physical memory
	Frame table	PFN to <PID, VPN> mapping needed for evicting pages from physical memory

# PCB

---

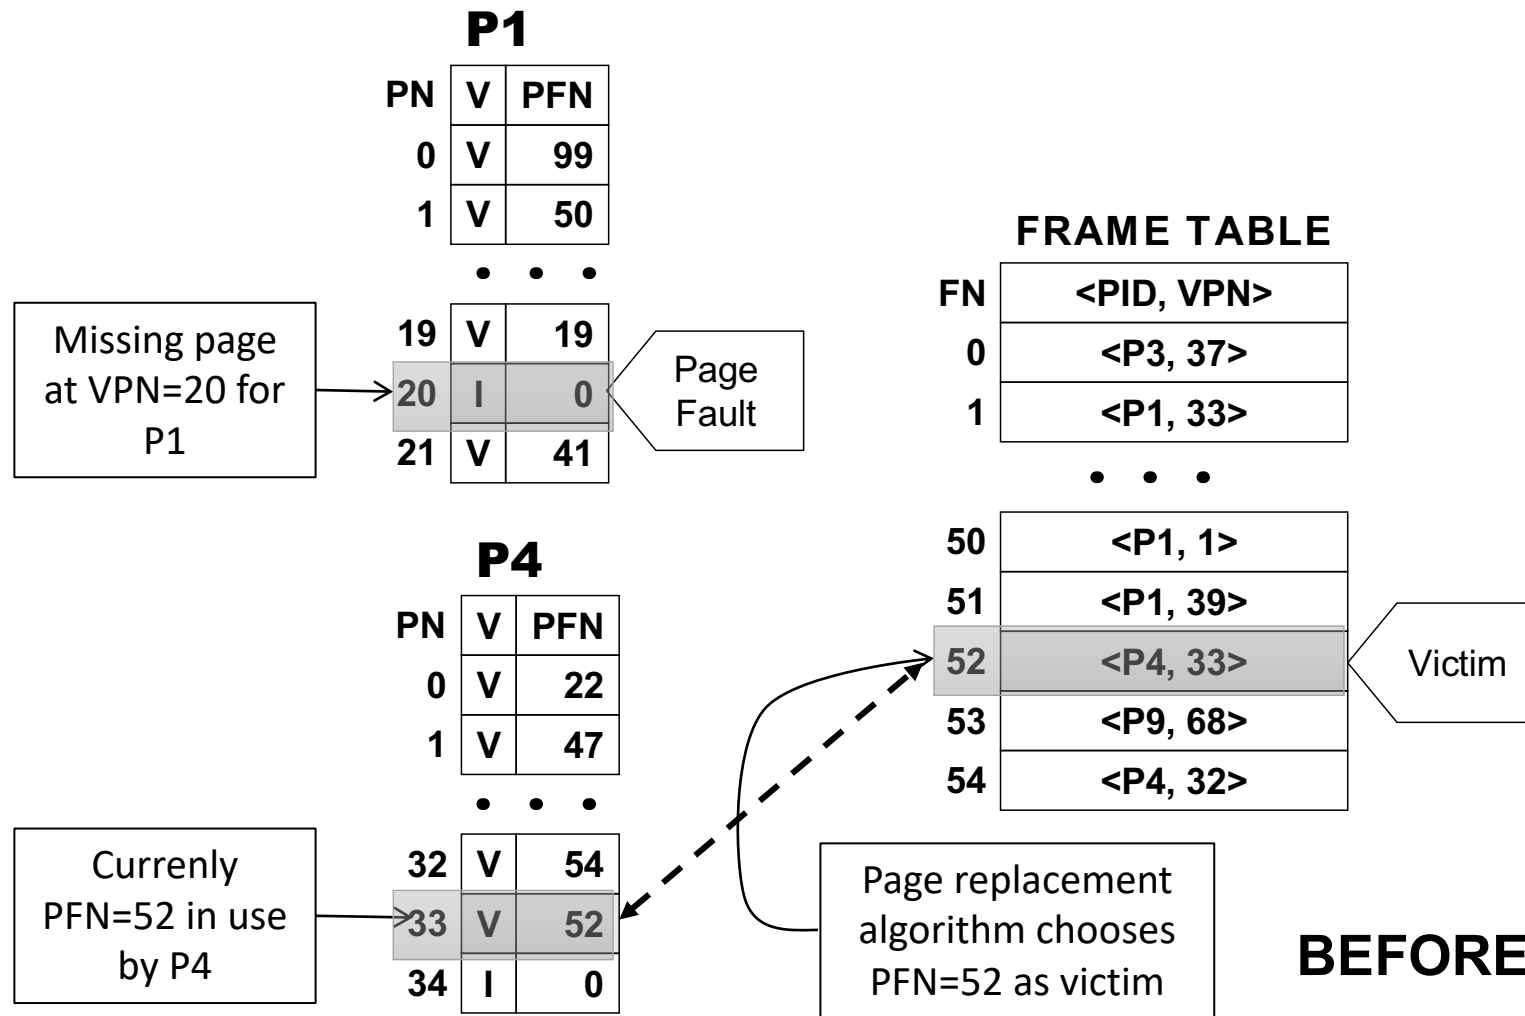
```
enum    state_type {new, ready, running,
                    waiting, halted};
typedef struct control_block_type {
    enum state_type state;
    address PC;
    int reg_file[NUMREGS];
    struct control_block *next_pcb;
    int priority;
    address PTBR;
    disk_address *disk_map;
    ...
    ...
} control_block;
```

# Example


---

- process P1 page fault at  $\text{VPN} = 20$ .
- free-list is empty.
- selects page frame  $\text{PFN} = 52$  as the victim.
- frame currently houses  $\text{VPN} = 33$  of process P4

# Before



# After








# With paged memory management there can be...

---

- A. External fragmentation
- B. Internal fragmentation
- C. No fragmentation
- D. Both internal and external fragmentation

# FIFO Page Replacement

- Manage the Frame Table like a queue




# FIFO example

Consider a sequence of page references by a process:

Reference number: 1 2 3 4 5 6 7 8 9 10 11 12 13

-----  
Virtual page number: 9 0 3 4 0 5 0 6 4 5 0 5 4

Assume there are 3 physical frames.




Ref: 1 2 3 4 5 6  
VPN: 9 0 3 4 0 5


---

---


**Reference #1 (PF)**




**Reference #2 (PF)**




**Reference #3 (PF)**




**Reference #4 (PF)**



**Reference #5 (HIT)**



**Reference #6 (PF)**





# Size of the FIFO queue?

---

- A. 42
- B. Number of physical page frames
- C. Number of virtual pages
- D. No clue

# Belady's Min

Consider a string of page references by a process:

Reference number: 1 2 3 4 5 6 7 8 9 10 11 12 13

Virtual page number: 9 0 3 4 0 5 0 6 4 5 0 5 4

Look into the future!

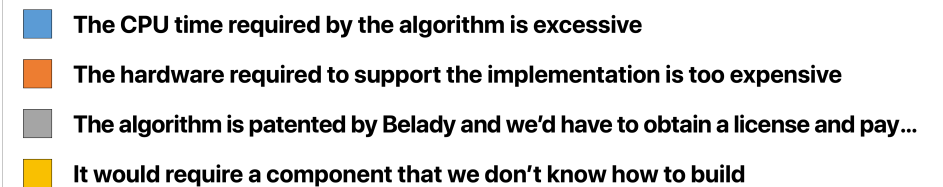



- Theoretically the best algorithm
- As the victim page, choose the page with the longest time to its next reference
- Merely requires us to predict the future
- If we had pages 0-9 in memory, which pages should we evict first?
- (1, 2, 7, 8) can go first. Then 6, 5, 4, 3, 0, 9



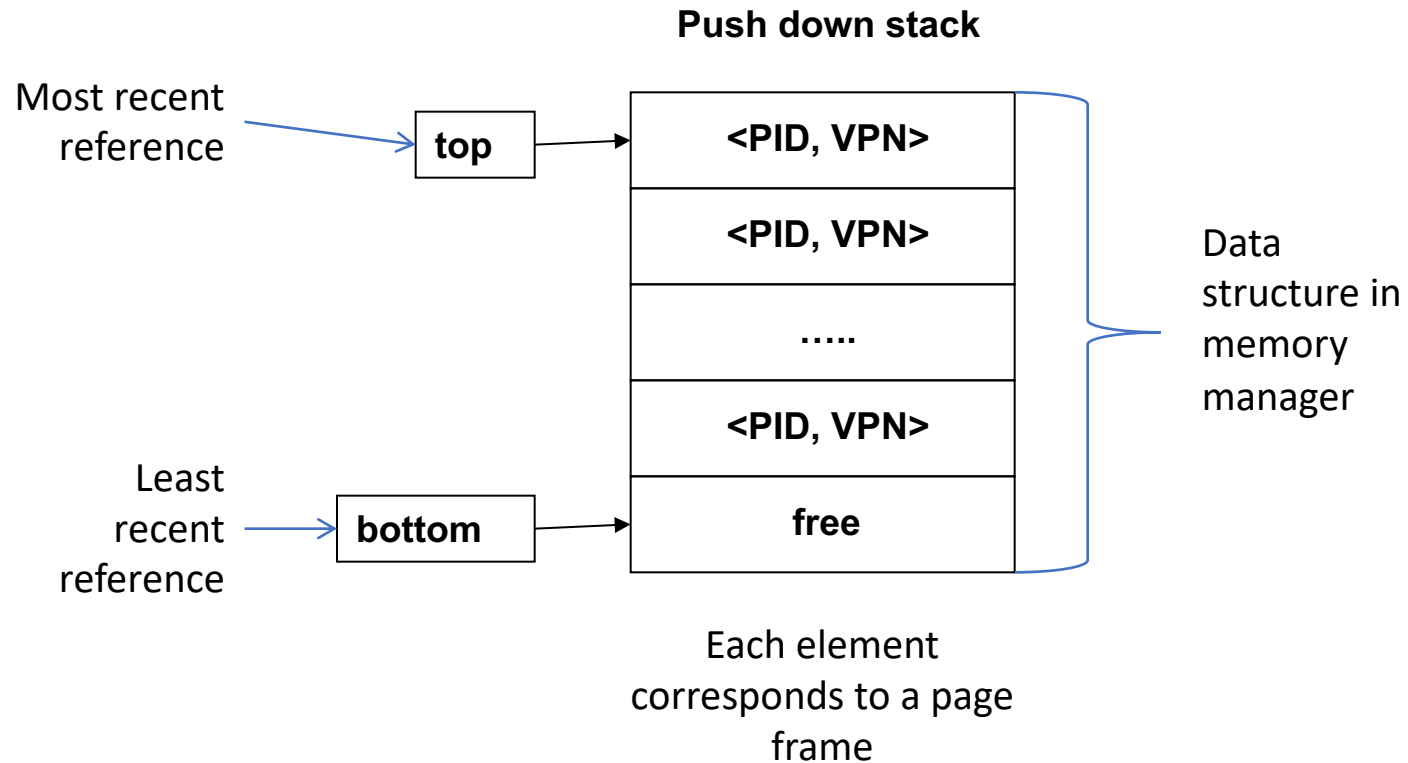
# Why not implement Belady's Min as a page replacement algorithm?

- A. The CPU time required by the algorithm is excessive
- B. The hardware required to support the implementation is too expensive
- C. The algorithm is patented by Belady and we'd have to obtain a license and pay royalties
- D. It would require a component that we don't know how to build



# LRU

- Use the past as a predictor of the future.






# LRU example

Consider a string of page references by a process:

Reference number: 1 2 3 4 5 6 7 8 9 10 11 12 13

-----  
Virtual page number: 9 0 3 4 0 5 0 6 4 5 0 5 4


Assume there are 3 physical frames.




Ref: 1 2 3 4 5 6  
VPN: 9 0 3 4 0 5

---


Reference #1 (PF)




Reference #2 (PF)




Reference #3 (PF)




Reference #4 (PF)



Reference #5 (HIT)



Reference #6 (PF)





# Size of LRU “stack”

---

- A. 42
- B. Number of virtual pages
- C. Number of physical frames
- D. No clue

# Problems with LRU

---

- Memory references are known to the hardware, but memory management (i.e. victim selection) is in software
- One possibility: make the stack shared by HW & SW
  - Implement stack in hardware
  - Let hardware update stack on each reference
  - Let software (OS) use this stack as a data structure
- Will it work?
- Still, no. The size of the stack is the number of page frames (i.e., quite large), so a memory write is required for each memory reference

# Ways to approximate LRU


---

- Use a small hardware register stack
  - Remember the last 16 or 32 references?
- Reference bit per physical frame
  - Add a “referenced” bit to each PTE
  - Set it each time the hardware uses the PTE to translate a memory reference
  - If it’s already set, don’t set it again.

**Page Table**

<b>PFN</b>	<b>Ref</b>	<b>Valid</b>
<b>PFN</b>	<b>Ref</b>	<b>Valid</b>
<b>PFN</b>	<b>Ref</b>	<b>Valid</b>

# Approximate LRU with ref bits



- Keep ref bits in PT
  - Set bit when page referenced → done by hardware
- Paging daemon → background OS process
  - Flush ref bits to software “counters” periodically  
 $\text{counter} = (\text{ref} \ll 31 \mid \text{counter} \gg 1)$
  - Clear ref bits
- Victim? → The page with the counter that has the lowest value

# “Second chance” page replacement using reference bits

---

1. Initially clear all the reference bits
2. As the process runs, set referenced bits on each page referenced
3. If a page has to be evicted, the memory manager selects a page in a FIFO manner
4. If the chosen victim's referenced bit is set, the manager clears the referenced bit and moves to the next page
5. The victim is the first page that doesn't have the referenced bit set

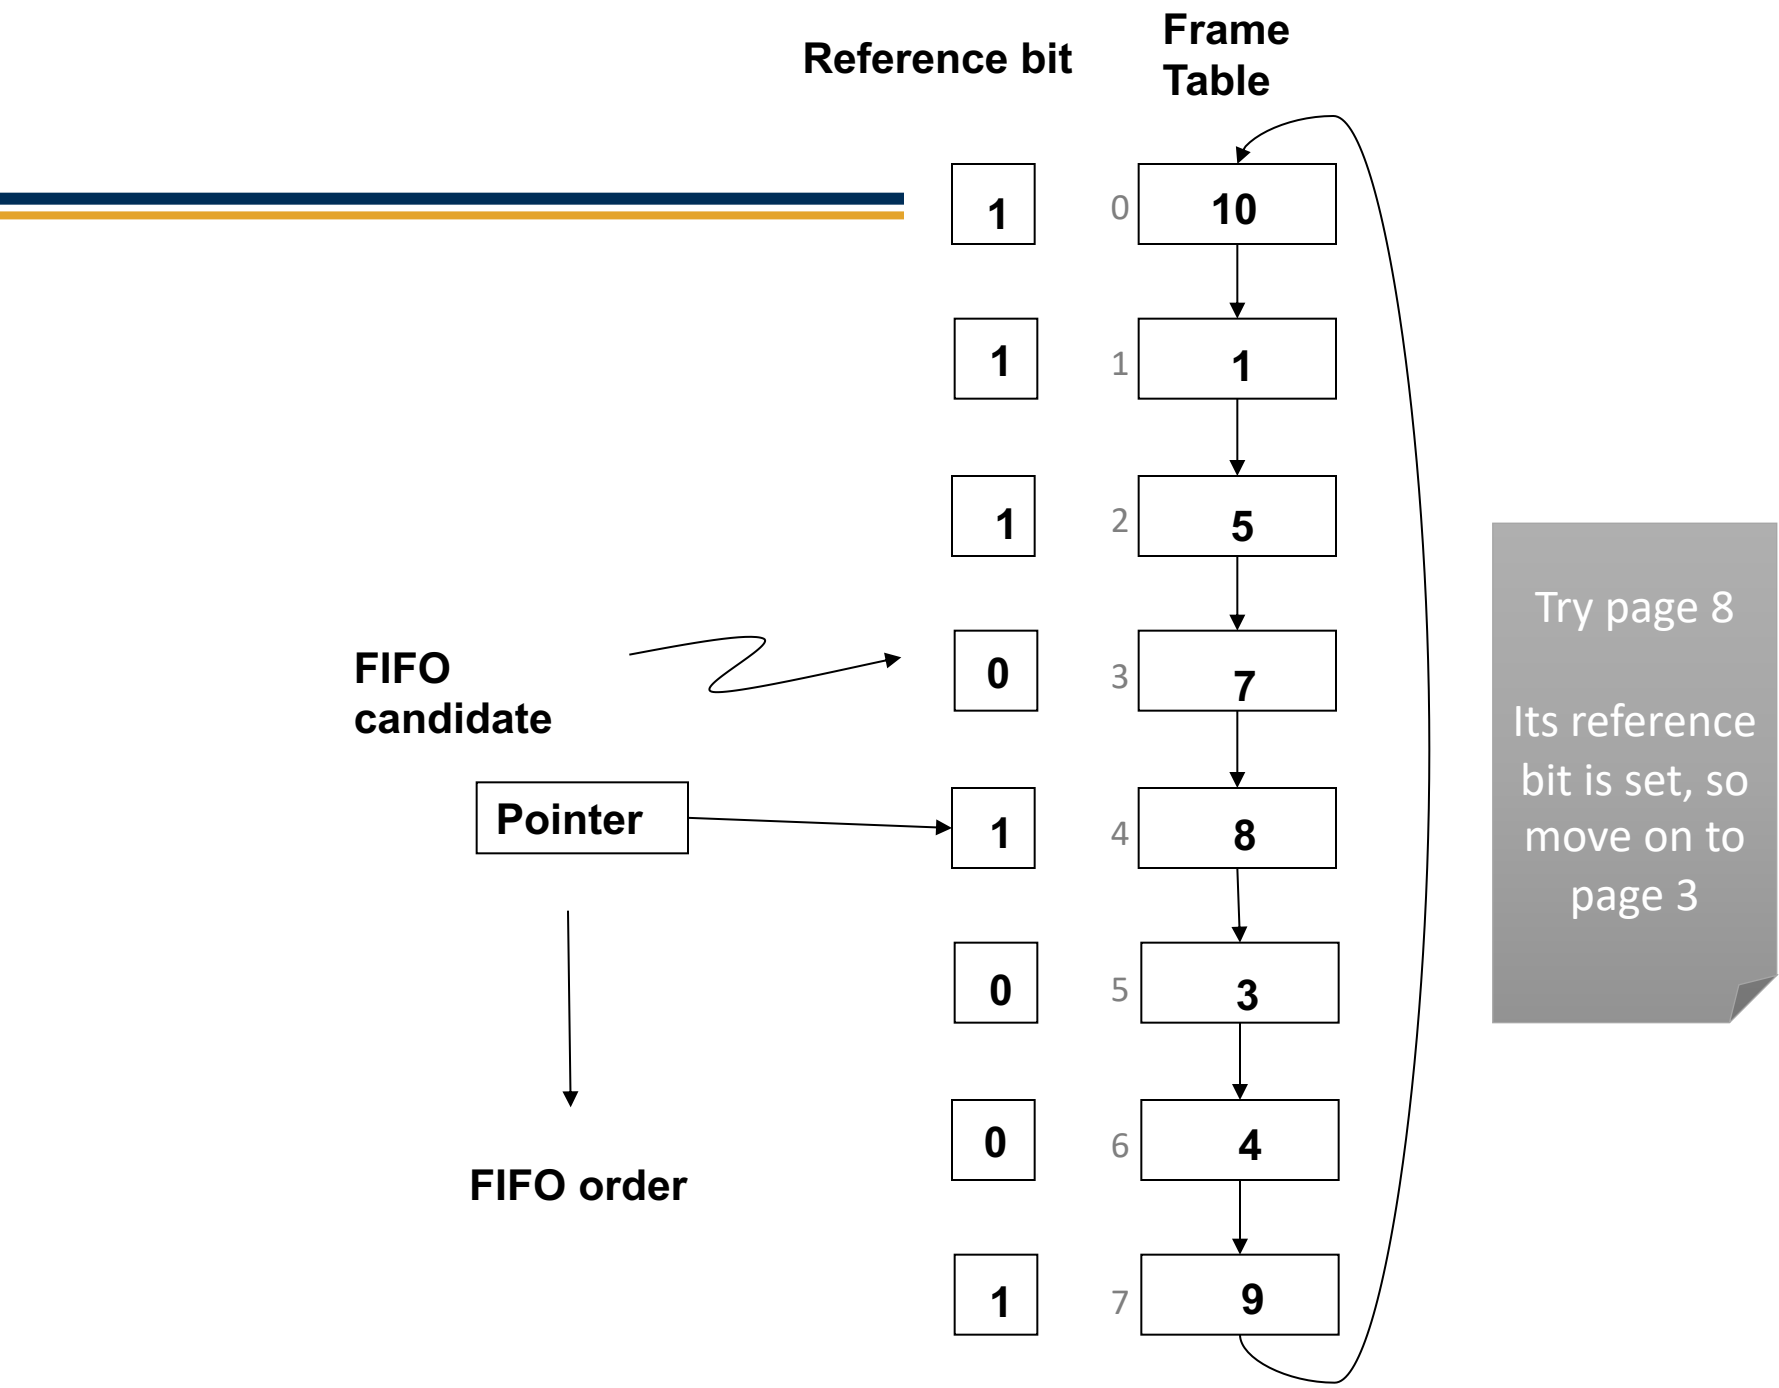
Try page 7

Its reference bit is set, so move on to page 8

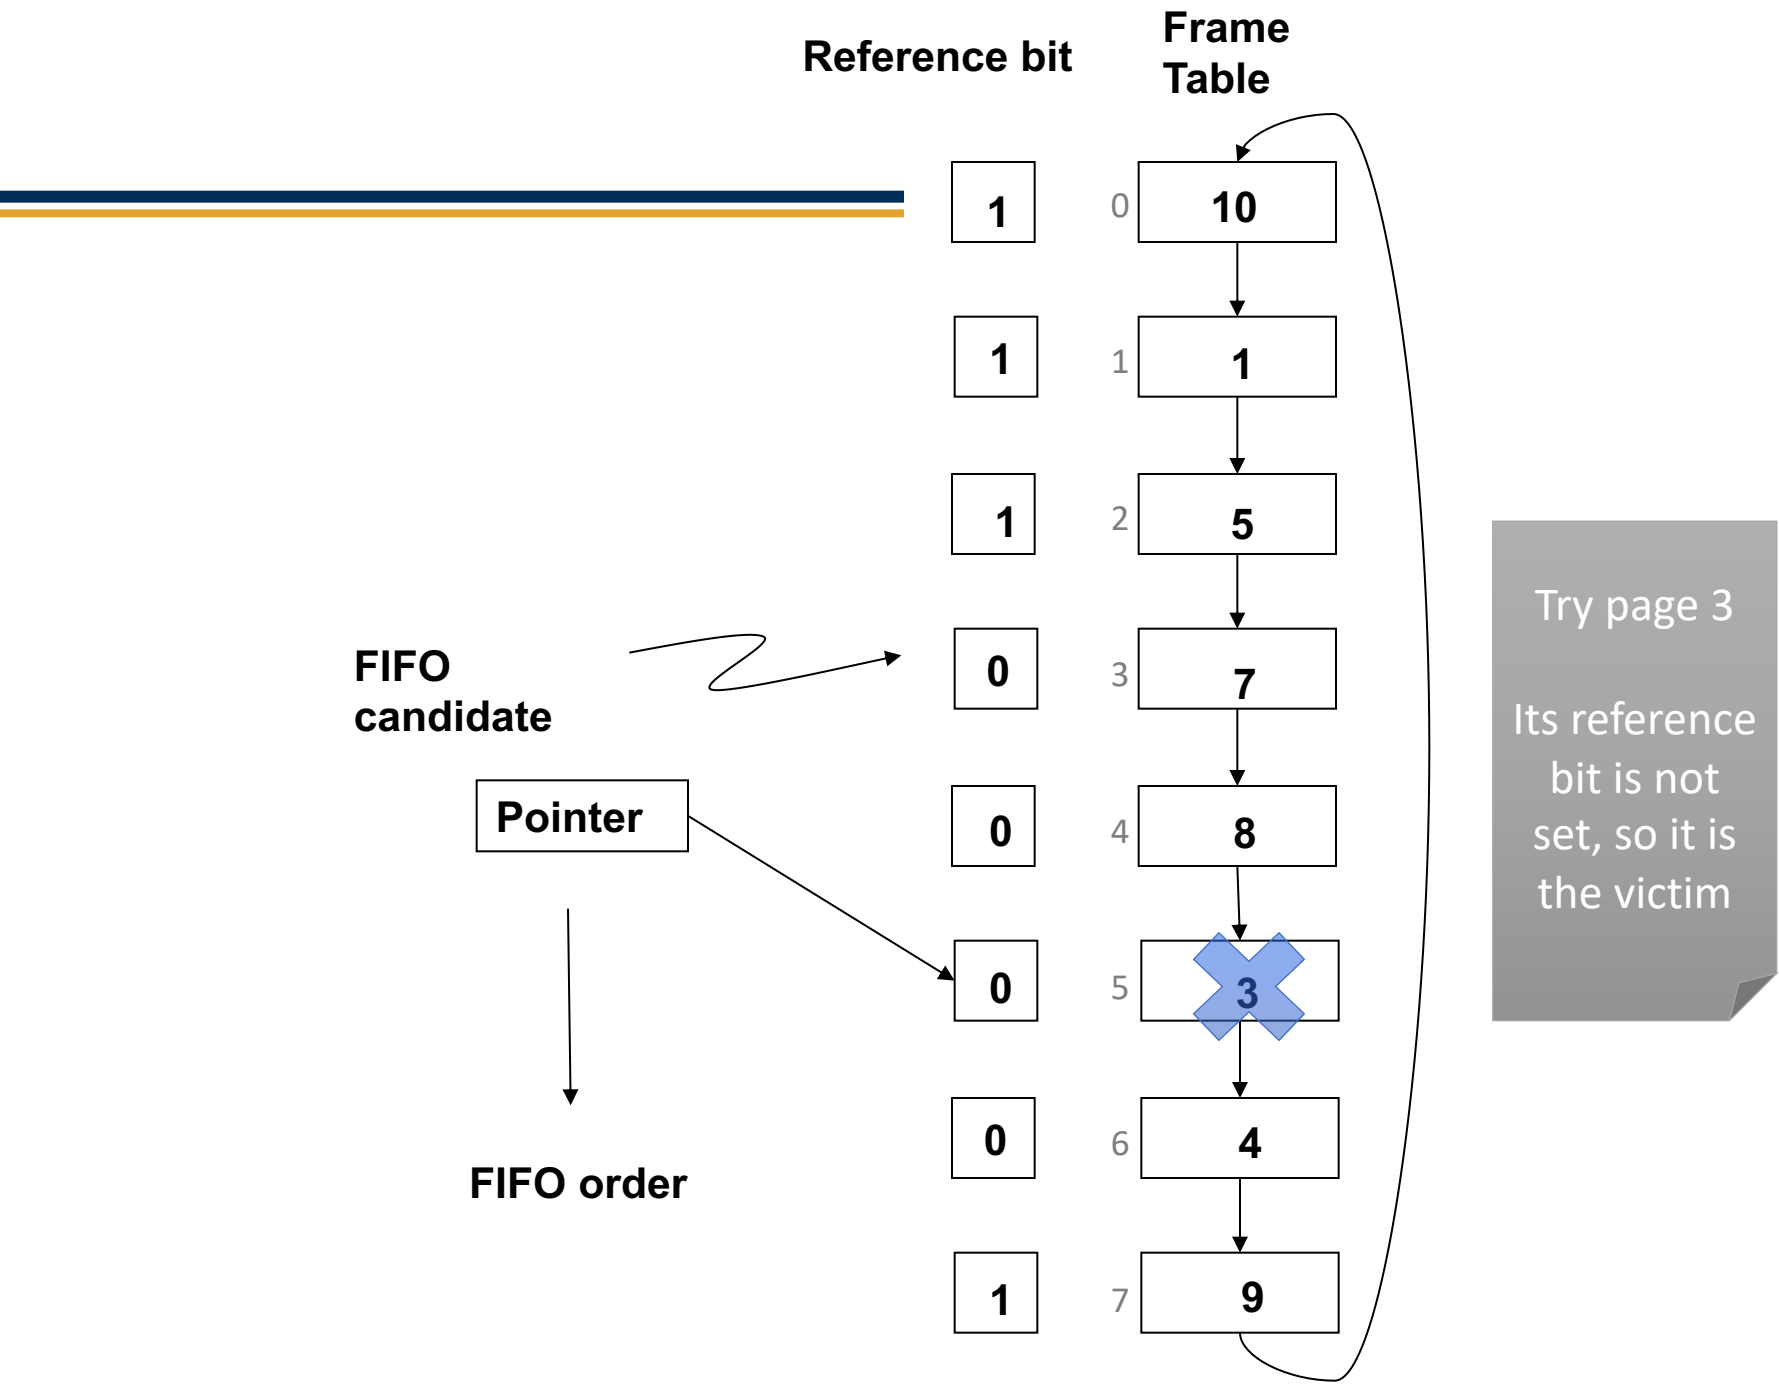




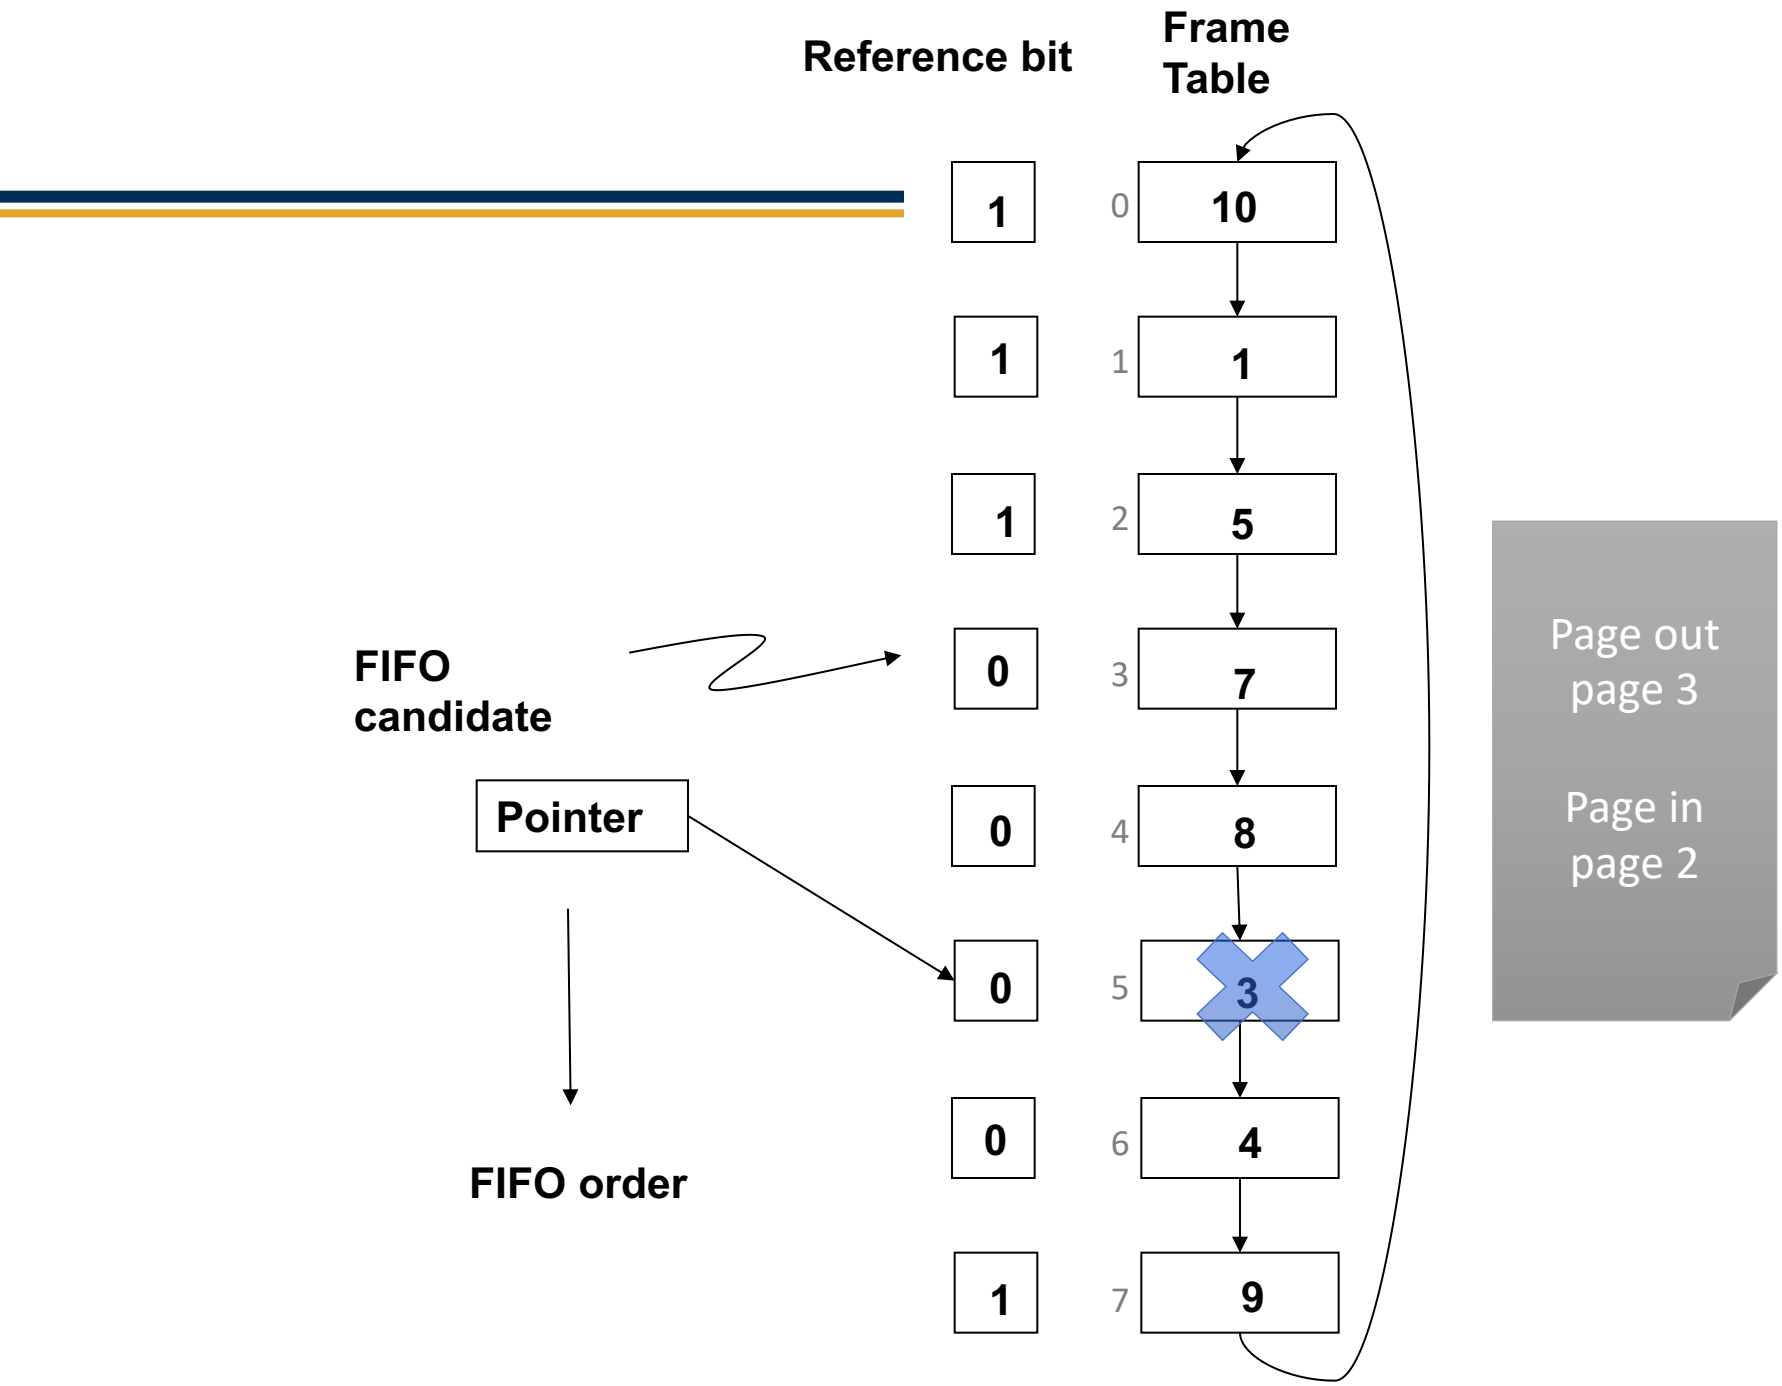
# Second chance replacement



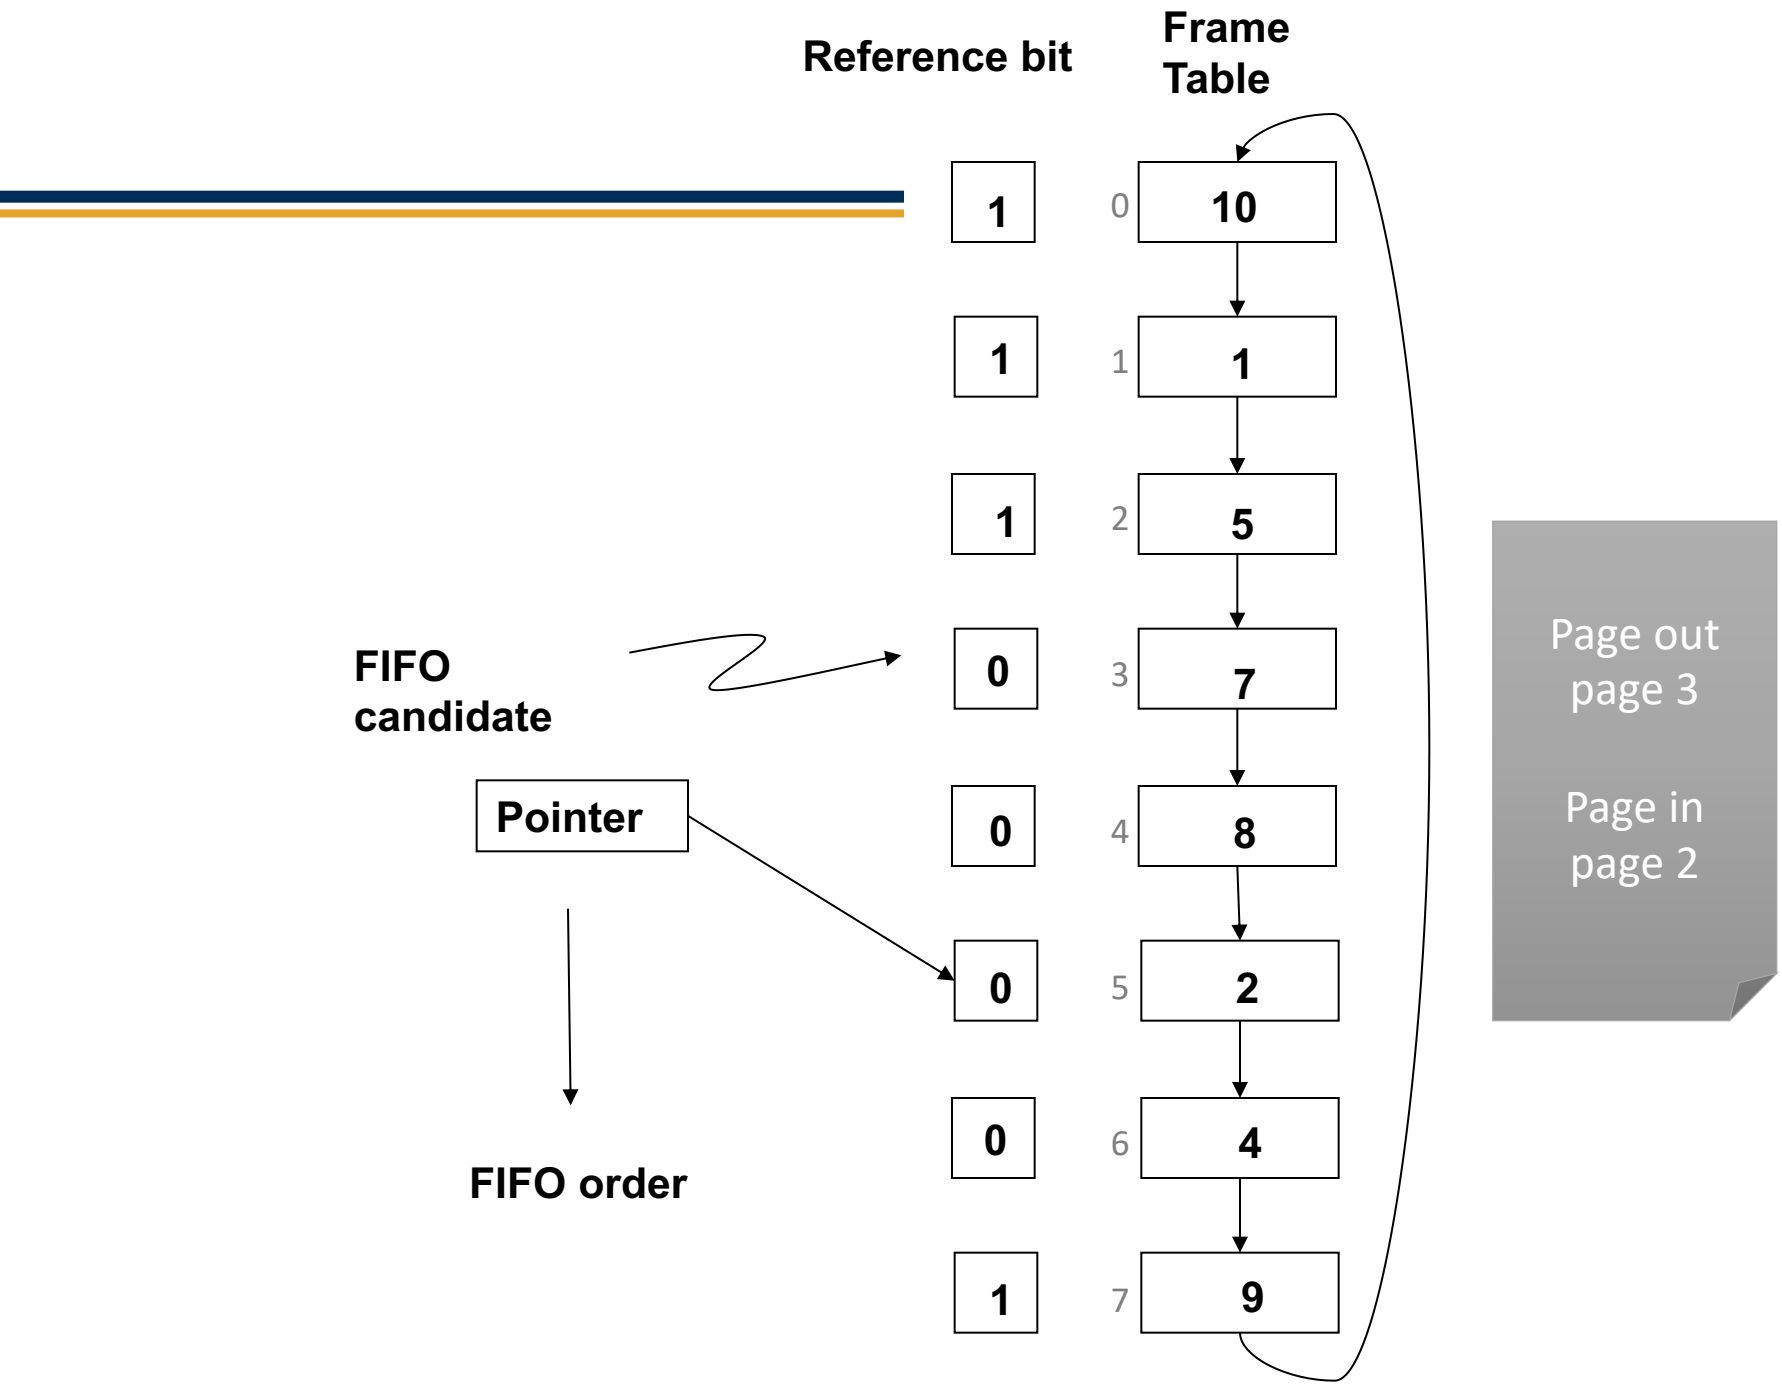
# Second chance replacement



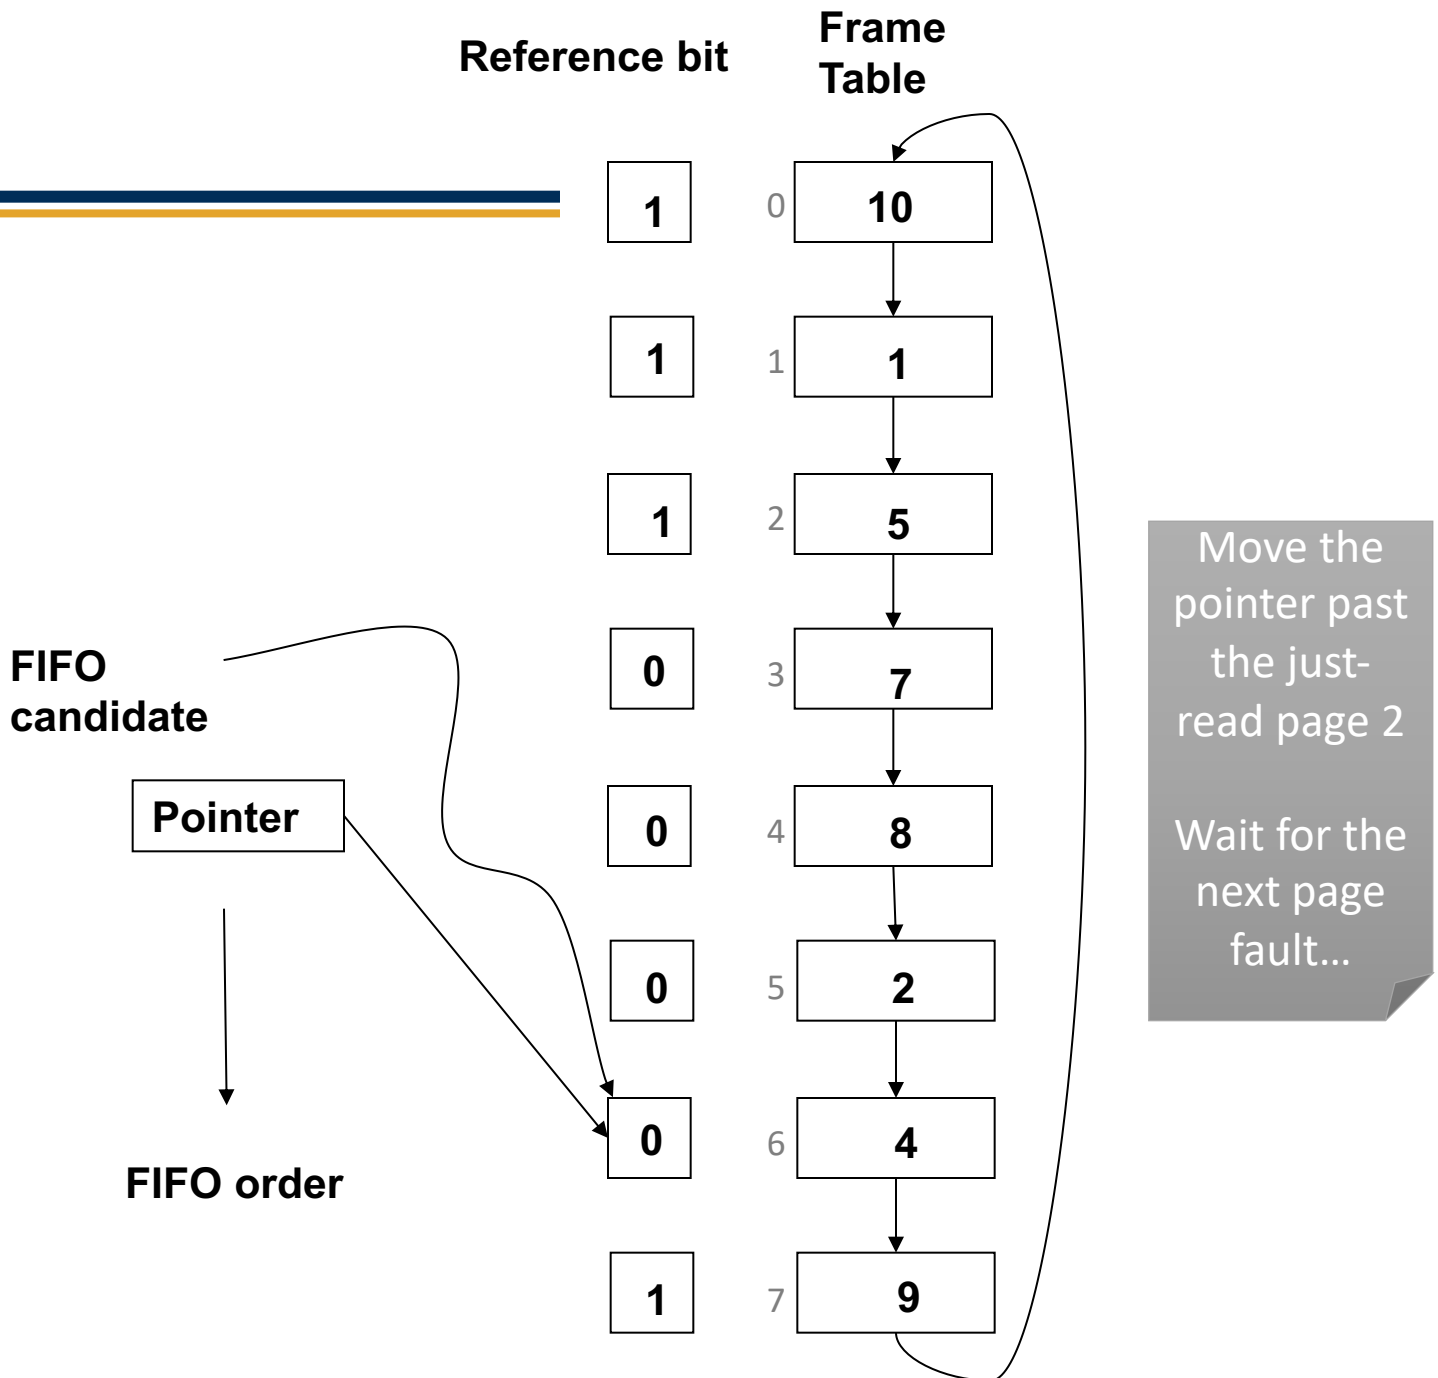
# Second chance replacement



# Second chance replacement



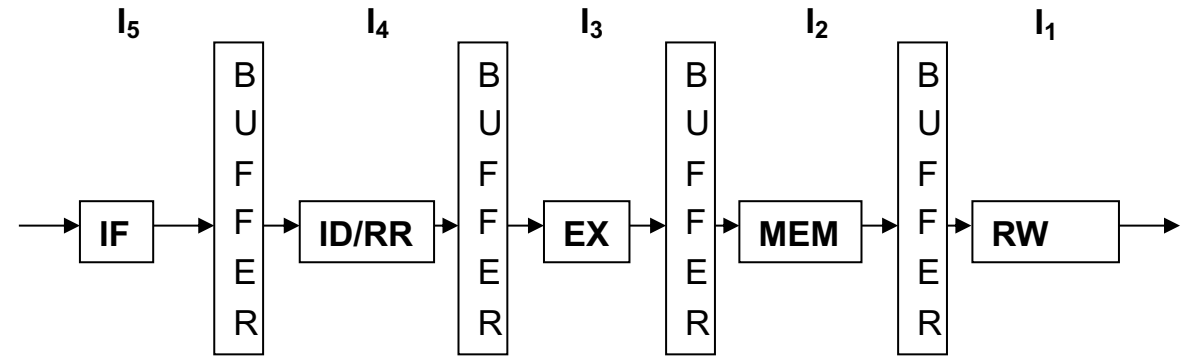
# Second chance replacement



# Page replacement algorithms

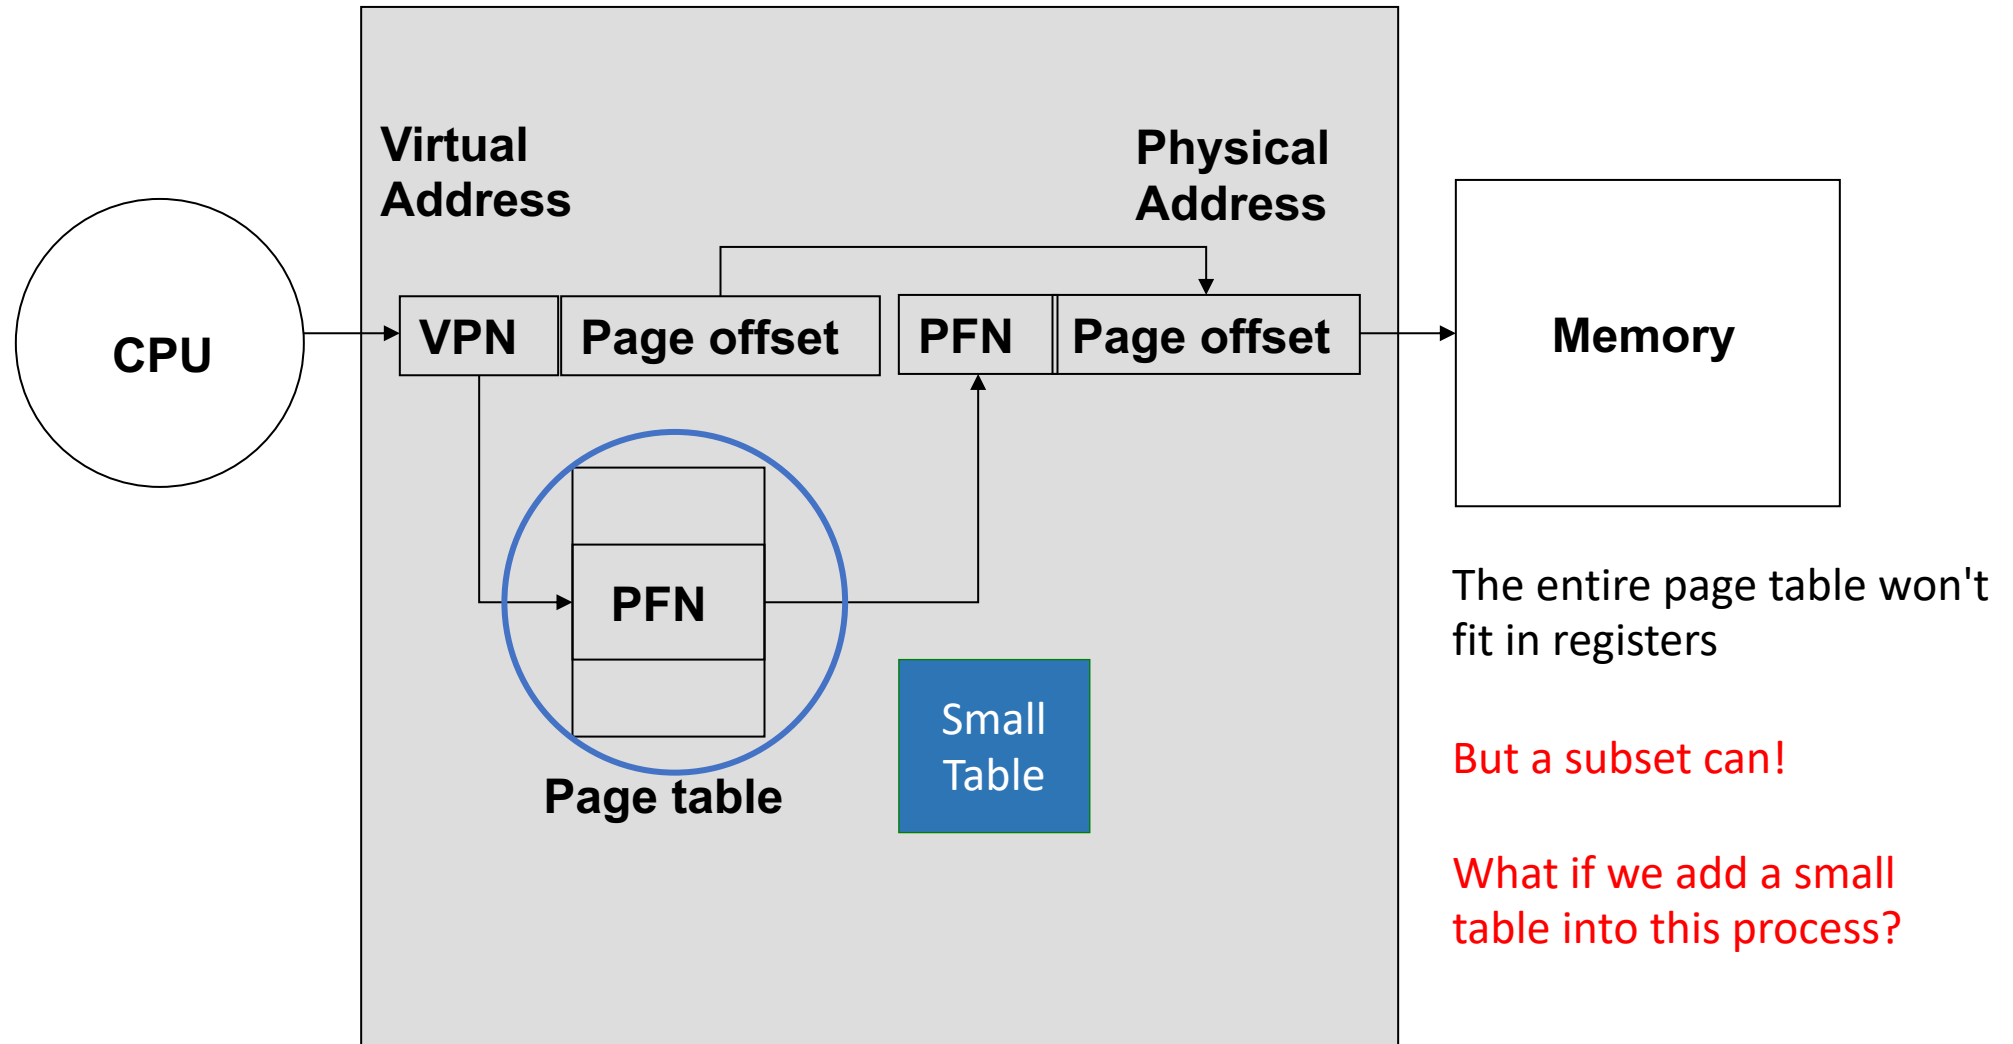
ALGORITHM	HARDWARE ASSIST	COMMENTS
FIFO	None	Could lead to performance anomalies
Belady's MIN	An oracle	Provably optimal; not realizable in hardware; useful as a standard
True LRU	Push down stack	Expected performance close to optimal; infeasible
Approximate LRU #1	A small hardware stack	Expected performance close to optimal; worst-case performance may be similar to FIFO
Approximate LRU #2	Reference bit per page	Expected performance close to optimal; moderate hardware complexity
Second chance replacement	Reference bit per page	Expected performance better than FIFO; memory manager implementation simplified compared to LRU schemes

# Back to our pipelined processor



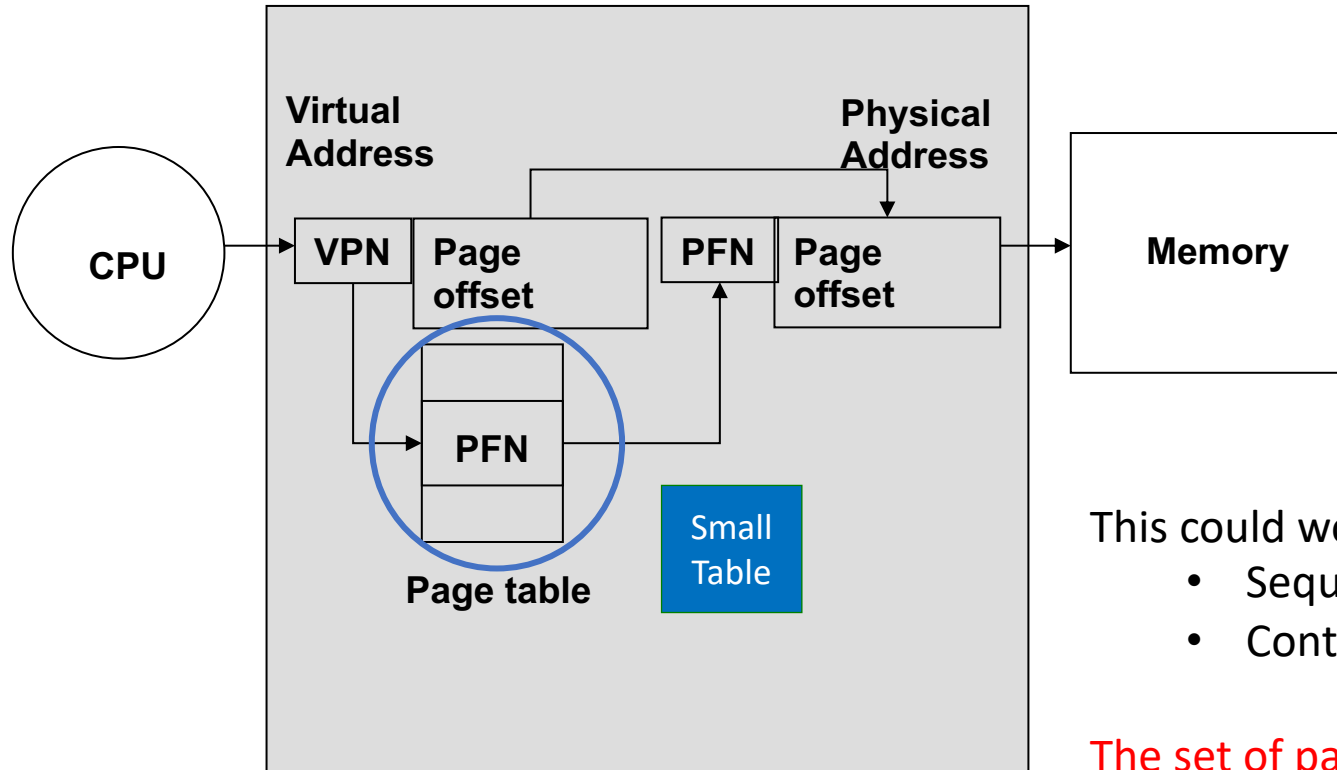
- With virtual memory...
- Every memory access requires two memory accesses!
  - PTE
  - Actual memory word
- This is bad news for the pipeline
- At least one bubble for every instruction

# Speeding up address translation





# Why will this work?



This could work because of the **locality** of a program

- Sequential instructions
- Contiguous data structures

The set of pages being referenced at a given time is called the **Working Set**

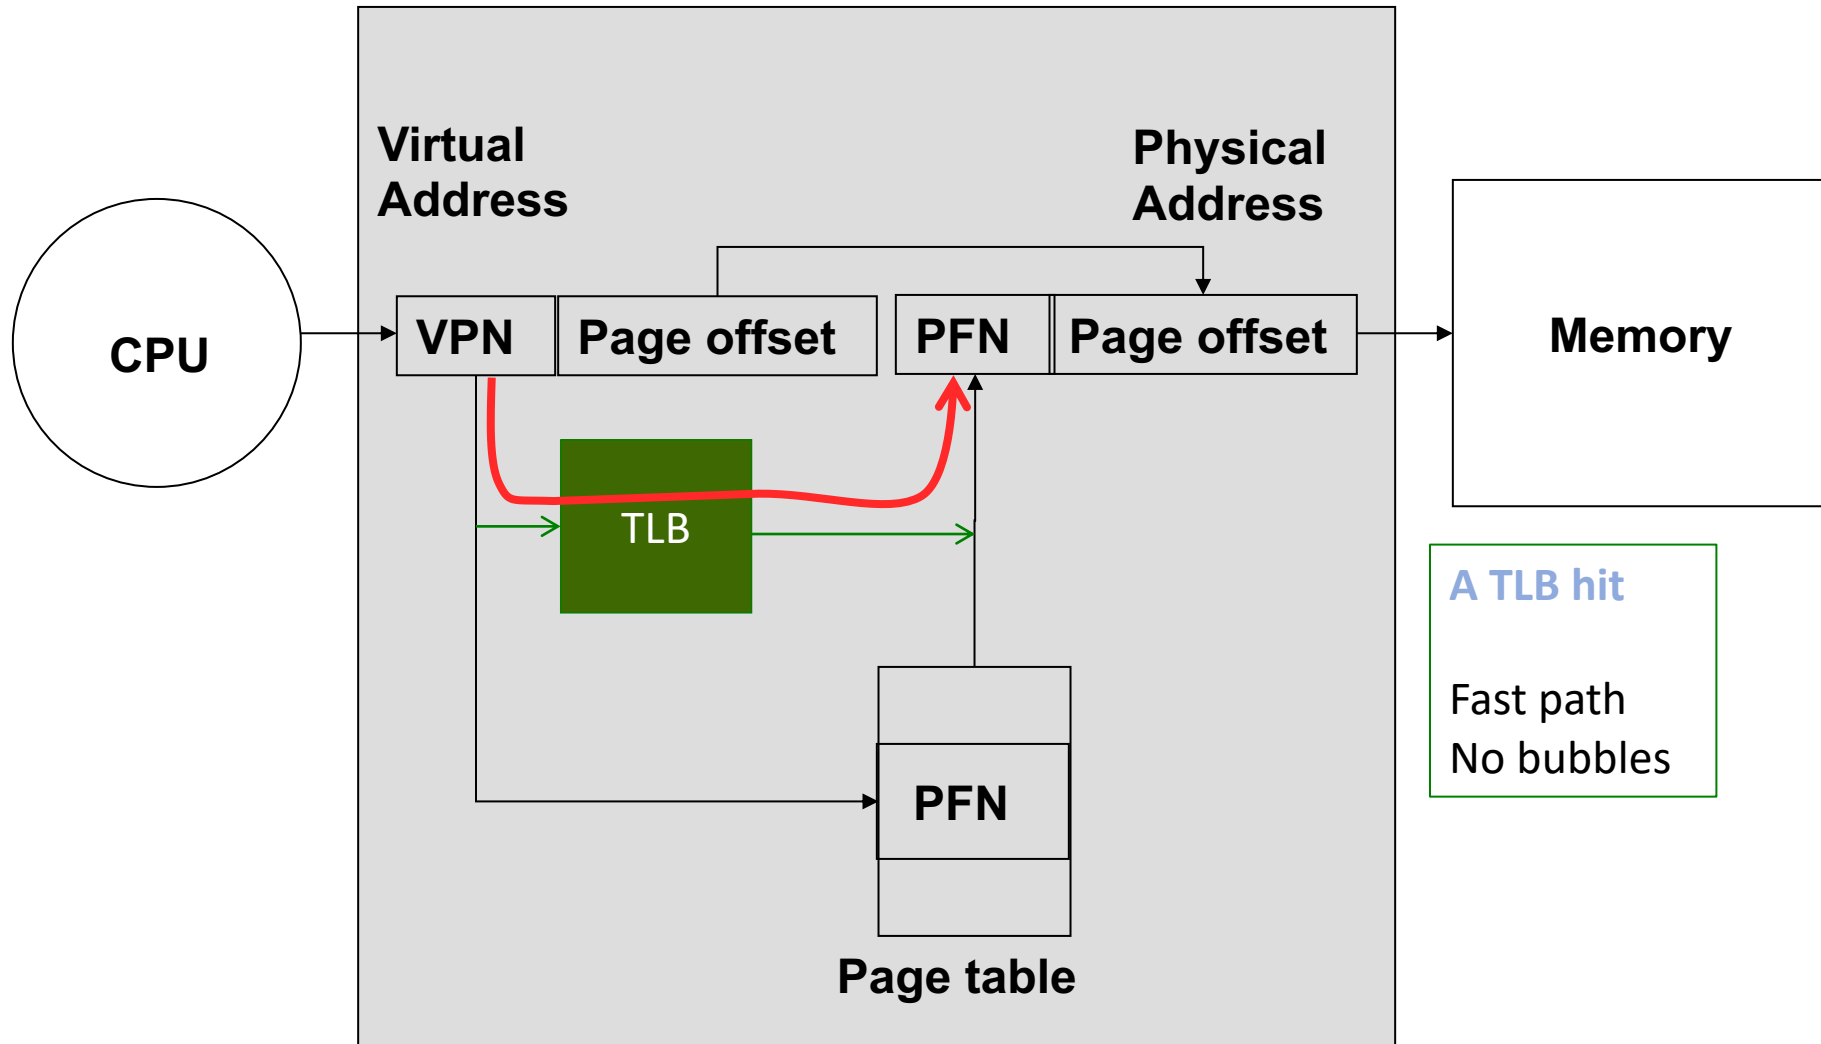
# TLB (translation lookaside buffer)

---

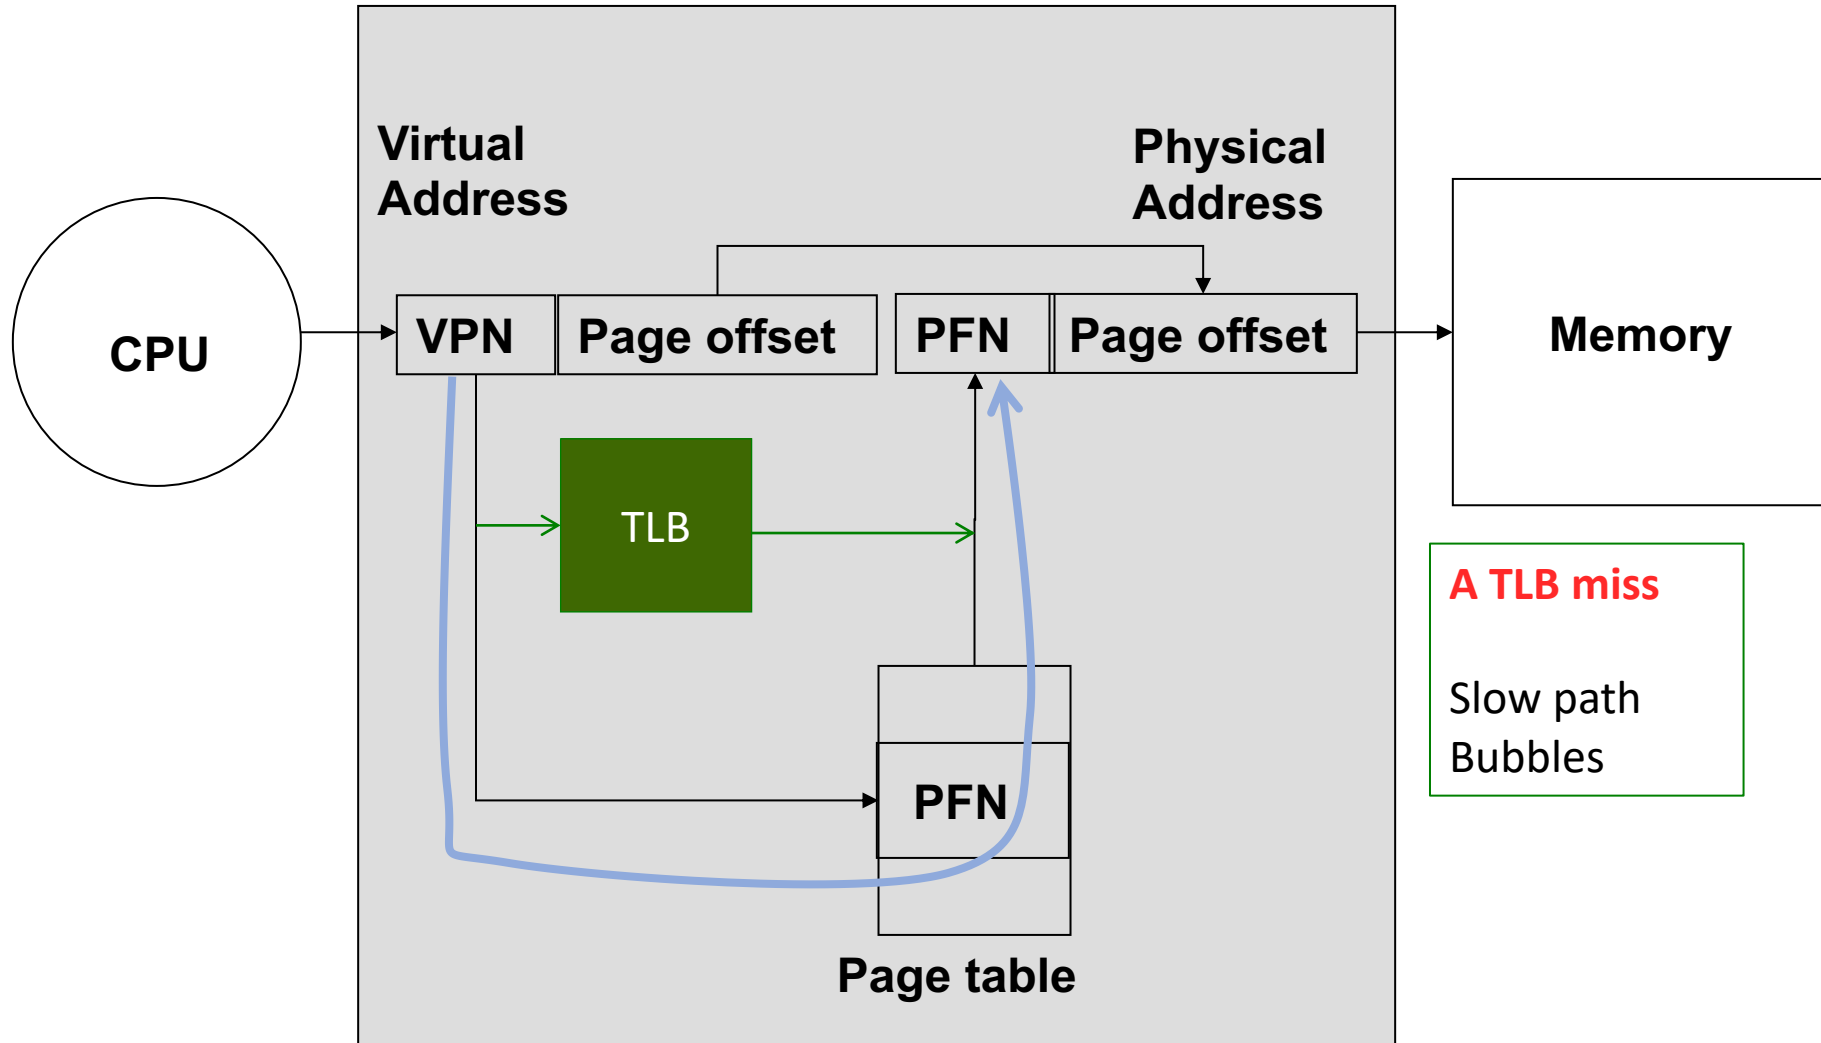
- It will look like the following table
- It is an **associative** memory: it can search on a match on the first two columns and output the corresponding last two columns in one cycle

USER/KERNEL	VPN	PFN	VALID/INVALID
U	0	122	V
U	XX	XX	I
U	10	152	V
U	11	170	V
K	0	10	V
K	1	11	V
K	3	15	V
K	XX	XX	I

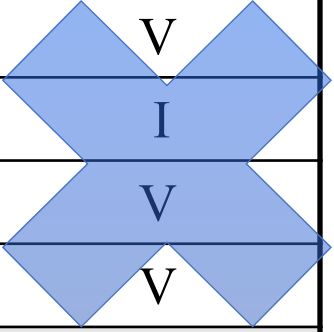
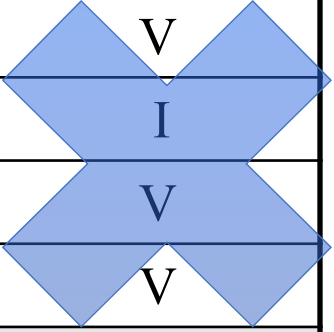
# Speeding up address translation



# Speeding up address translation



# TLB

		USER/KERNEL	VPN	PFN	VALID/INVALID
Specific to each process		U	0	122	I V
		U	XX	XX	I I
		U	10	152	I V
		U	11	170	I V
Same for all processes		K	0	10	V
		K	1	11	V
		K	3	15	V
		K	XX	XX	I

What's the implication of the U entries for a context switch?

→ We'll need to flush the U entries on context switch

# Another new instruction

---

- The LC-2200 is going to need
  - Purge TLB
  - or TLB flush
- Can only be executed by the kernel




# Upon a context switch...

---


- A. The entire TLB must be flushed
- B. Only the kernel portion of the TLB must be flushed
- C. Only the user portion of the TLB must be flushed
- D. Leave the poor TLB alone!

# Given the nature of a process

---




- We want to increase multiprogramming to keep the CPU busy doing useful work
- This is what we want to see:





# Extending the utilization curve

CPU Utilization



Thrashing

Nobody is getting any useful work done


- Too many processes asking for memory
- This is overcommitment of memory
- We have a bigger working set than physical memory can hold

Paging is implicit I/O on behalf of a process  
→ System became I/O bound  
→ Throughput drops precipitously

Degree of multiprogramming

# Working set of a program


**Working set** is the set of pages needed to keep a process from paging



**Working set size:** number of page frames needed to hold working set

# Working set of a program

The working set changes over time



$$WS_{t3-t4} = \{ p1, p2 \}$$

$$WS_{t7-t8} = \{ p4, p1 \}$$


# Memory pressure

---

$$\text{memory.pressure} = \sum_{i=1}^n WSS_i$$


- $P_1, P_2, P_3, \dots$  are processes in memory each with a working set  $WSS_i$
- The count of active processes signifies the degree of multiprogramming
- How do we control the degree of multiprogramming
  - **$\sum WSS > \text{total physical memory}$** 
    - swap out some processes
  - **$\sum WSS < \text{total physical memory}$** 
    - increase degree of multiprogramming

# Controlling thrashing



# Page faults are disruptive...

---

- ... from a process point of view
  - ➔ implicit I/O
- ... from a CPU-utilization perspective
  - ➔ overhead that doesn't contribute to work
- We need to limit impact of page faults to improve system performance



# We can tell a system is thrashing if

- A. It has too few page faults per second
- B. It has too many page faults per second
- C. The ratio of I/O operations to CPU operations is not optimal
- D. The combined working set of all processes is greater than the number of available page frames

If only it were as easy as B! Thrashing implies too many page faults, but too many page faults don't always imply thrashing! Applications can be changing the pages in use without changing their working set size, for instance.

In reality, to diagnose thrashing, you'd look for a high paging rate, low CPU utilization, and several processes waiting on paging I/O for several seconds. Those metrics together are a good clue.



# We can reduce thrashing by

---

- A. Using a medium-term scheduler that suspends processes until the condition improves
- B. Reducing the physical memory size
- C. Adding additional processes to increase the multiprogramming factor
- D. Reducing the page size


Of course this begs the question of how the medium-term scheduler is going to figure out that the system is thrashing...



# Paging Optimizations


Accelerate page-in by removing functionality from critical path

- Keep a (small) pool of free frames
  - Don't wait to start page replacement algorithm on a page fault
- Page replacement
  - Background activity of OS when CPU is not in use
  - If I/O is not busy, write out a “dirty” page which makes it “clean”



# Reverse mapping to page table

- Gives a “third” chance for reuse of a page before being kicked out
- P4 is running and page faults on VPN=33
- No need to go to disk!
- Just remap PFN=52 into PT for P4, VPN=33 and take it out of the freelist



# Linux VM and kswapd

```
$ free -h
```

	total	used	free	shared	buffers	cached
Mem:	15G	7.1G	8.5G	164K	703M	2.4G
...						
-/+ buffers/cache:		4.0G	11G			
Swap:	2.0G	26M	1.9G			


## ■ Kswapd

- Paging daemon
- Runs when “free” memory is low (about 2% of memory)
- Uses a modified version of second-chance replacement
- Links victim pages into the free list and sets their “invalid” PTE bits

## ■ Page fault handler:

- If the target page is still on the free list, it is reclaimed by removing it from the free list, marking its PTE bit “valid”, and writing it out if it’s dirty
- Otherwise, the first frame in the Free List is removed, the target page is read into it, and the target page’s PTE is modified to point to it and “valid” is set

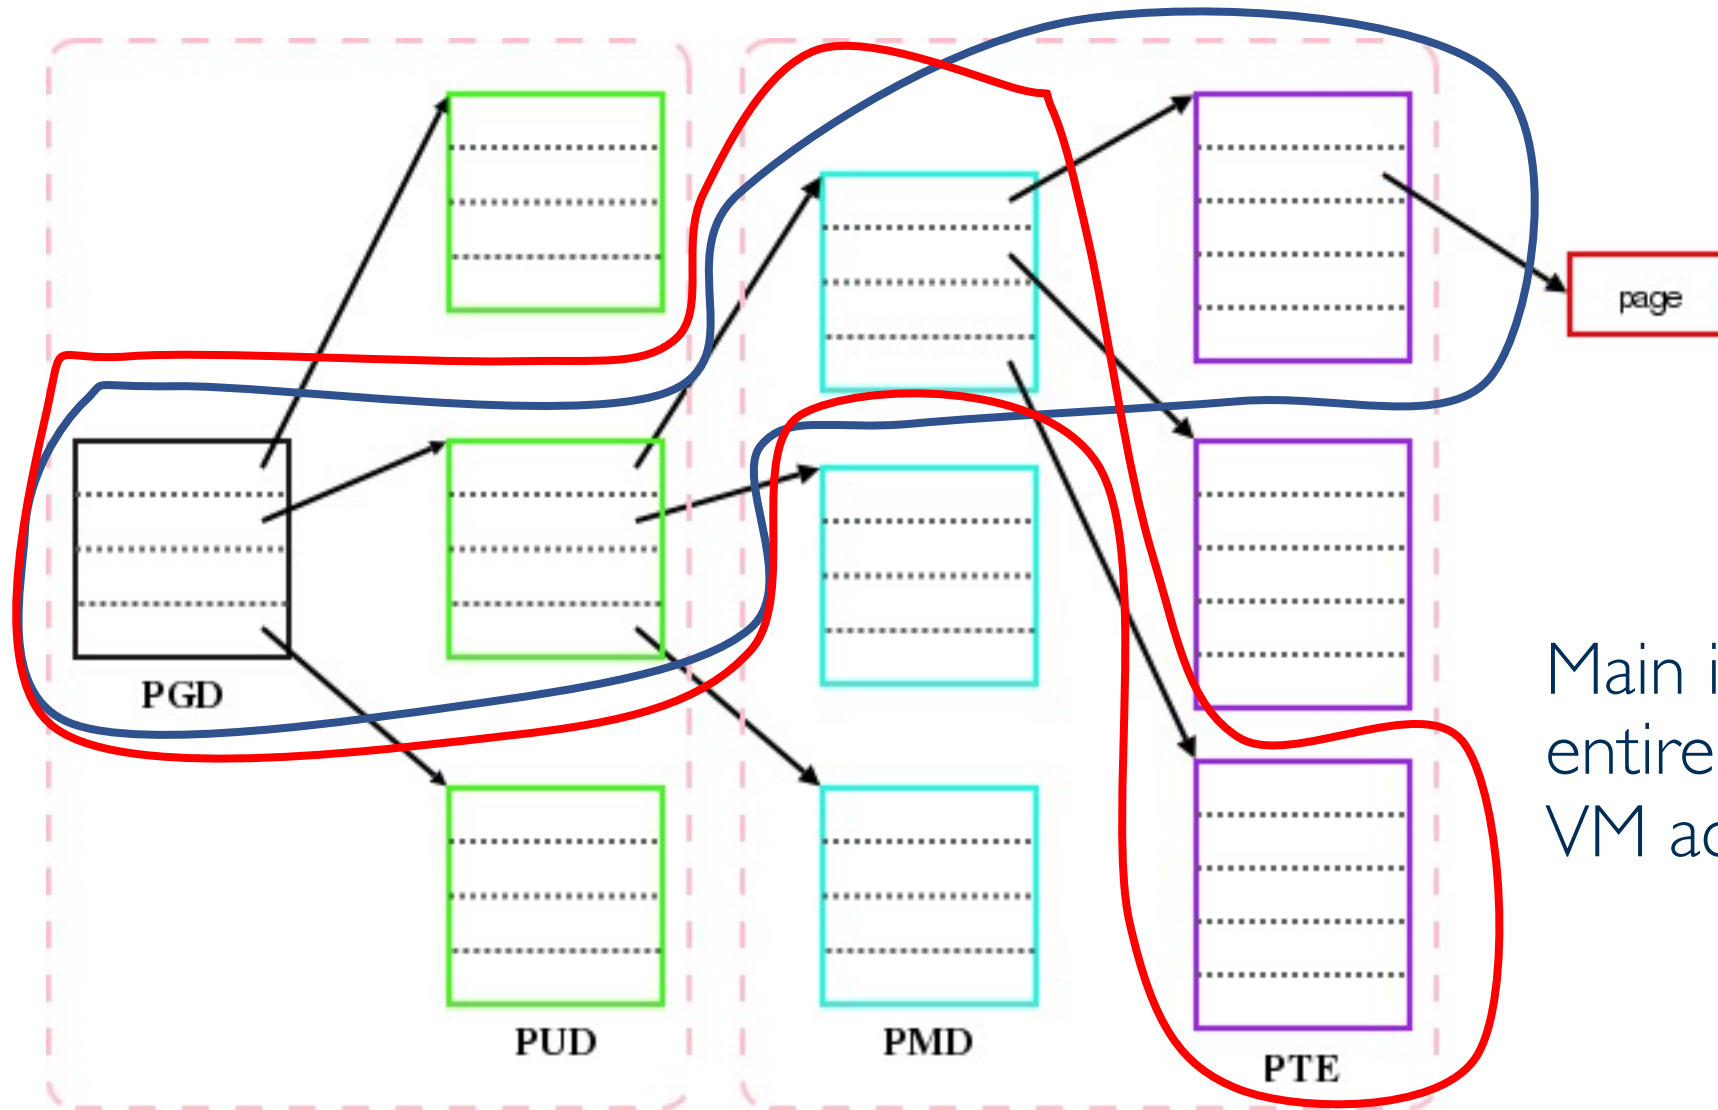
# How big are page tables?



With 48-bit virtual address space, 4KB pages...


- Need  $2^{36}$  PT entries (per process!)
- That's several GBs of memory just for the PT!

# Teaser: Hierarchical Page Table



Main idea: only load a fraction of entire page table, analogous to VM actually used by the process

# Intel's X86-64 4-level Page Tables



More about this in cs4290