

CS2200

Systems and Networks

Spring 2022

Lecture 21: Parallel Systems

Alexandros (Alex) Daglis
School of Computer Science
Georgia Institute of Technology
adaglis@gatech.edu


Abstractions

- Program
 - A static image loaded into memory
- Process
 - A program in execution
- In other words, $\text{process} = \text{program} + \text{state}$
 - State evolves as the program executes
- Thread
 - We split the idea of state up into data (memory) and threads of control (PC and registers)
 - One memory space, but one or more contexts of control

So then, what's a thread?

In general:


Program



n threads of
control ($n > 0$)



Program

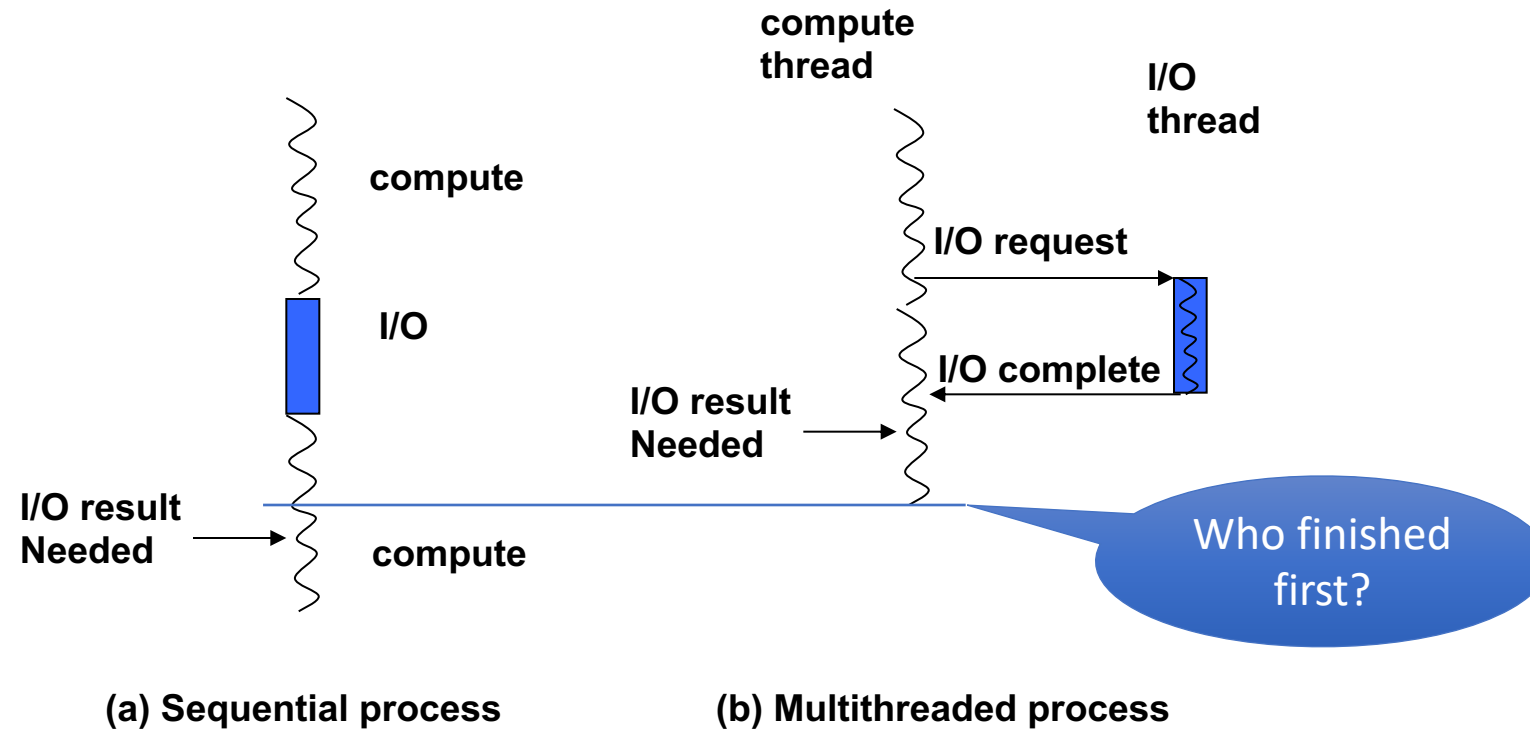


One thread
of control


Second
thread of
control

Process = program + state of all threads executing in the program

Example use of threads - I



Example use of threads - 2



Where are we headed?

- Programming support for threads
- Synchronization and communication between threads
- Architecture and OS support for threads

-
- Multithreaded program
 - Main and subroutines (procedures)
 - Thread starts executing in some top-level procedure upon “thread create”
 - Thread terminates
 - When main() terminates or
 - When the thread's top-level procedure terminates
 - We are going to need synchronization and communication among threads

Programming Support for Threads

- creation
 - `pthread_create`(top-level procedure, args)
- termination
 - `return` from top-level procedure
 - explicit `kill`
- rendezvous
 - creator can `wait` for children
 - `pthread_join`(child_tid)
- synchronization
 - `mutex`
 - `condition` variables

main thread



thread_create(foo, args)

(a) Before thread creation

main thread



thread_create(foo, args)

foo thread



(b) After thread creation

Sample program – thread create/join

```
int foo(int n)
{
    ....
    return 0;
}

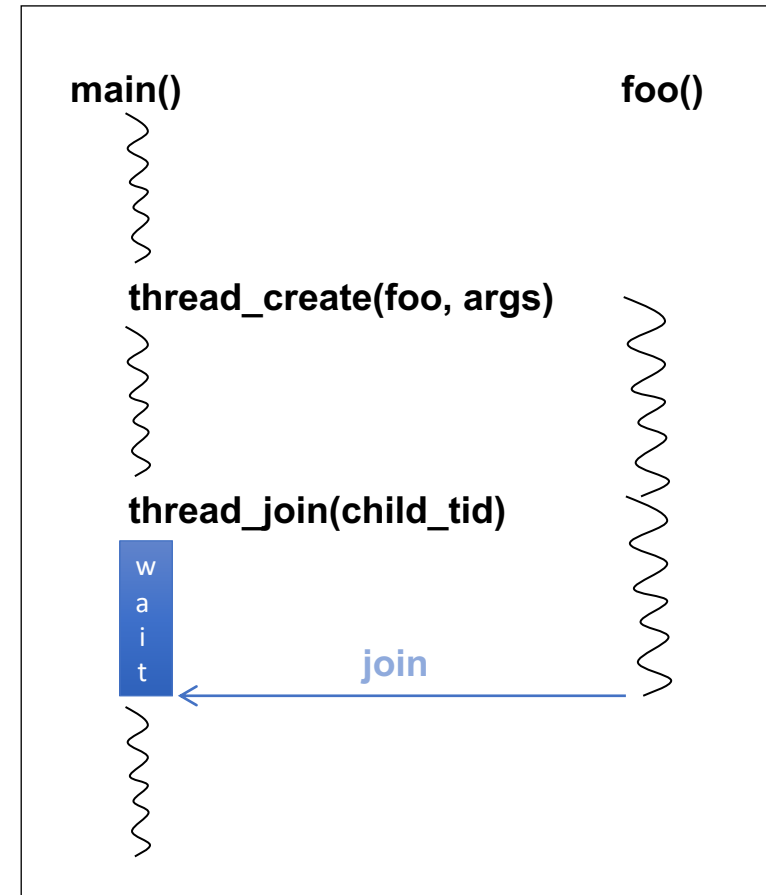
int main()
{
    int f;
    thread_type child_tid;

    ....

    child_tid = thread_create (foo, &f);

    ....

    thread_join(child_tid);
}
```



Threads within the same process



A thread...

7%

- A. I just want the participation credit
- B. ...is the same as a process
- C. ...is usually part of a process
- D. ...has nothing to do with a process
- E. ...usually refers to a set of processes
- F. ...is part of the memory hierarchy
- G. ...often involves a needle

A thread starts its execution ...

- 67% A. I just want the participation credit
- 0% B. In main()
- 0% C. At some top-level procedure that is part of the same program
- 33% D. At some top-level procedure that is part of a different program
- E. None of the above (B to D)



A thread

11% A. I just want the participation credit

22% B. ... lives forever

22% C. ... terminates **ONLY** when the top-level procedure where it started returns

11% D. ... terminates **ONLY** when main() terminates


11% E. ... terminates when **EITHER** the top-level procedure where it started or main() returns

22% F. ... terminates at the first context switch


Programming with Threads

- synchronization
 - for coordination of the threads
- communication
 - for inter-thread sharing of data
 - threads can be executing on different processors
 - how to achieve sharing?

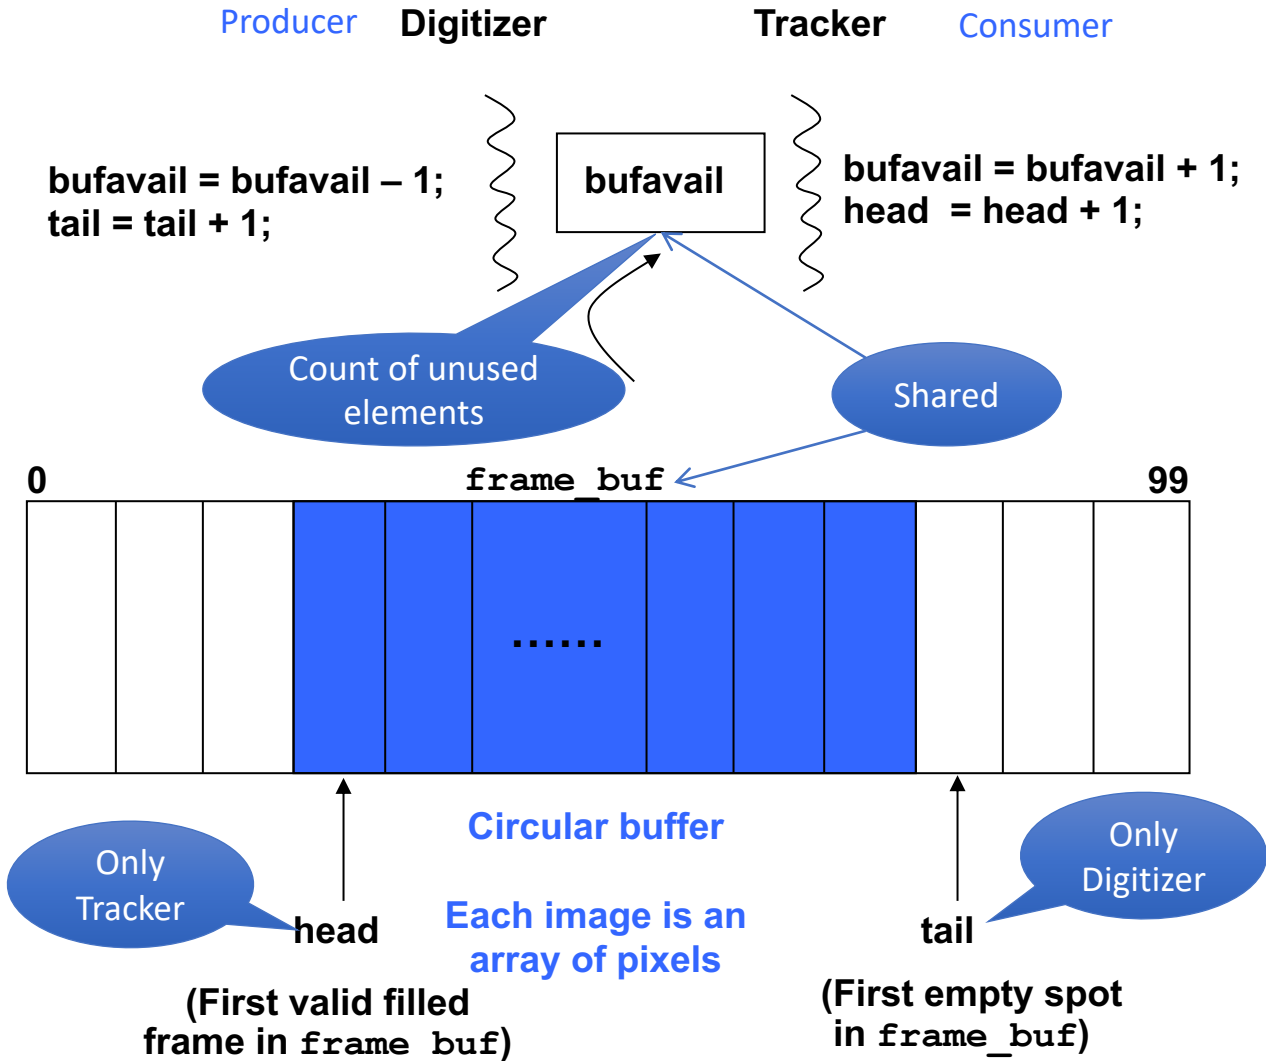
- Hardware software partnership {
- **software**: accomplished by keeping **all** threads in the **same address space** by the OS
 - **hardware**: accomplished by **hardware shared memory** and coherent caches (we will see this later)



Recall: our producer/consumer app



Global variables:
`int bufavail = MAX;`
`image_type frame_buf[MAX];`



Need for Synchronization

```
int bufavail = MAX; // global
image_type frame_buf[MAX]; // global
```

```
digitizer()
{
    image_type dig_image;
    int tail = 0; // private

    loop {
        if (bufavail > 0) {
            grab(dig_image);
            frame_buf[tail] = dig_image;
            tail = (tail + 1) % MAX;
            bufavail = bufavail - 1;
        }
    }
}
```

```
tracker()
{
    image_type track_image;
    int head = 0; // private

    loop {
        if (bufavail < MAX) {
            track_image = frame_buf[head];
            head = (head + 1) % MAX;
            bufavail = bufavail + 1;
            analyze(track_image);
        }
    }
}
```

Problem?

Need for Synchronization

```
int bufavail = MAX; // global
image_type frame_buf[MAX]; // global
```

```
digitizer()
{
    image_type dig_image;
    int tail = 0; // private

    loop {
        if (bufavail > 0) {
            grab(dig_image);
            frame_buf[tail] = dig_image;
            tail = (tail + 1) % MAX;
            bufavail = bufavail - 1;
        }
    }
}
```

```
tracker()
{
    image_type track_image;
    int head = 0; // private

    loop {
        if (bufavail < MAX) {
            track_image = frame_buf[head];
            head = (head + 1) % MAX;
            bufavail = bufavail + 1;
            analyze(track_image);
        }
    }
}
```

Problem?

Manipulating shared variables(!)

What's the issue?

Say that both threads happen to be executing at the blue arrows...

Tracker


$\text{bufavail} = \text{bufavail} - 1$

→ LD R1, bufavail
→ ADDI R1, R1, -1
→ ST R1, bufavail

Digitizer

$\text{bufavail} = \text{bufavail} + 1$

→ LD R1, bufavail
→ ADDI R1, R1, 1
→ ST R1, bufavail



R1 bufavail R1
~~8~~ 7 is it 7 or 9? ~~8~~ 9
But it should be 8 ! – FAIL!

Need for Synchronization

```
int bufavail = MAX; // global
image_type frame_buf[MAX]; // global
```

```
digitizer()
{
    image_type dig_image;
    int tail = 0; // private

    loop {
        if (bufavail > 0) {
            grab(dig_image);
            frame_buf[tail] = dig_image;
            tail = (tail + 1) % MAX;
            bufavail = bufavail - 1;
        }
    }
}
```

```
tracker()
{
    image_type track_image;
    int head = 0; // private

    loop {
        if (bufavail < MAX) {
            track_image = frame_buf[head];
            head = (head + 1) % MAX;
            bufavail = bufavail + 1;
            analyze(track_image);
        }
    }
}
```


Problem?

Manipulating shared variables(!)

Synchronization Primitives

- lock and unlock
 - mutual exclusion among threads
 - busy-waiting vs. blocking
 - `pthread_mutex_trylock`: no blocking
 - `pthread_mutex_lock`: blocking
 - `pthread_mutex_unlock`

Usage: `mutex lock; // data structure`
`mutex_lock(lock); // acquire lock`
`mutex_unlock(lock); //release lock`



- OS has no idea what you do in the **critical section**
- OS guarantees only 1 thread in the critical section (we call this guarantee an **OS invariant**)

Critical Section

- “Code that is executed in a mutually exclusive manner”
- Shared access to data that must be synchronized, so we implement a mutual exclusion lock which is honored by one or more segments of code that access the shared data
- A critical section is not necessarily a single piece of code. Any segment of code that honors the same mutual exclusion lock is called a critical section.

Fix #1 – with locks

```
int bufavail = MAX;
image_type frame_buf[MAX];
mutex buflock;

digitizer()
{
    image_type dig_image;
    int tail = 0;

    loop {
        pthread_mutex_lock(buflock);
        if (bufavail > 0) {
            grab(dig_image);
            frame_buf[tail] = dig_image;
            tail = (tail + 1) % MAX;
            bufavail = bufavail - 1;
        }
        pthread_mutex_unlock(buflock);
    }
}
```

```
tracker()
{
    image_type track_image;
    int head = 0;

    loop {
        pthread_mutex_lock(buflock);
        if (bufavail < MAX) {
            track_image = frame_buf[head];
            head = (head + 1) % MAX;
            bufavail = bufavail + 1;
            analyze(track_image);
        }
        pthread_mutex_unlock(buflock);
    }
}
```



A pthreads mutex lock

30% A. I just want the participation credit

- B. Allows exactly one thread to acquire it at a time
- C. Allows any number of threads to acquire it at a time
- D. Allows a defined number of threads to acquire it at a time
- E. None of the above (B to D)

Fix #1 – with locks

```
int bufavail = MAX;  
image_type frame_buf[MAX];  
mutex buflock;
```

```
digitizer()  
{  
    image_type dig_image;  
    int tail = 0;  
  
    loop {  
        pthread_mutex_lock(buflock);  
        if (bufavail > 0) {  
            grab(dig_image);  
            frame_buf[tail] = dig_image;  
            tail = (tail + 1) % MAX;  
            bufavail = bufavail - 1;  
        }  
        pthread_mutex_unlock(buflock);  
    }  
}
```

Critical section is far
too coarse!

Problem?

```
tracker()  
{  
    image_type track_image;  
    int head = 0;  
  
    loop {  
        pthread_mutex_lock(buflock);  
        if (bufavail < MAX) {  
            track_image = frame_buf[head];  
            head = (head + 1) % MAX;  
            bufavail = bufavail + 1;  
            analyze(track_image);  
        }  
        pthread_mutex_unlock(buflock);  
    }  
}
```

No concurrency!
No performance improvement.

Fix #1 – with locks

```
int bufavail = MAX;
image_type frame_buf[MAX];
mutex buflock;
```

```
digitizer()
```

```
{
    image_type dig_image;
    int tail = 0;
```

```
    loop {
```

```
        pthread_mutex_lock(buflock);
        if (bufavail > 0) {
            grab(dig_image);
            frame_buf[tail] = dig_image;
            tail = (tail + 1) % MAX;
            bufavail = bufavail - 1;
        }
        pthread_mutex_unlock(buflock);
    }
```

```
}
```

No need
for mutex

```
tracker()
```

```
{
    image_type track_image;
    int head = 0;
```

```
    loop {
```

```
        pthread_mutex_lock(buflock);
        if (bufavail < MAX) {
            track_image = frame_buf[head];
            head = (head + 1) % MAX;
            bufavail = bufavail + 1;
            analyze(track_image);
        }
        pthread_mutex_unlock(buflock);
    }
```

```
}
```


Fix #2 – with locks

```
int bufavail = MAX;
image_type frame_buf[MAX];
mutex buflock;
```

```
digitizer()
{
    image_type dig_image;
    int tail = 0;

    loop {
        grab(dig_image);
        thread_mutex_lock(buflock);
        while (bufavail == 0) ; // do nothing
        thread_mutex_unlock(buflock);
        frame_buf[tail] = dig_image;
        tail = (tail + 1) % MAX;
        thread_mutex_lock(buflock);
        bufavail = bufavail - 1;
        thread_mutex_unlock(buflock);
    }
}
```

Problem?

```
tracker()
{
    image_type track_image;
    int head = 0;

    loop {
        thread_mutex_lock(buflock);
        while (bufavail == MAX) ; // do nothing
        thread_mutex_unlock(buflock);
        track_image = frame_buf[head];
        head = (head + 1) % MAX;
        thread_mutex_lock(buflock);
        bufavail = bufavail + 1;
        thread_mutex_unlock(buflock);
        analyze(track_image);
    }
}
```

Deadlock!

Fix #3

```
int bufavail = MAX;
image_type frame_buf[MAX];
mutex buflock;
```

```
digitizer()
```

```
{
    image_type dig_image;
    int tail = 0;

    loop {
        grab(dig_image);
        while (bufavail == 0) ; // do nothing
        frame_buf[tail] = dig_image;
        tail = (tail + 1) % MAX;
        thread_mutex_lock(buflock);
        bufavail = bufavail - 1;
        thread_mutex_unlock(buflock);
    }
}
```

We're only
reading so
no need for
mutex

Problem?

```
tracker()
```

```
{
    image_type track_image;
    int head = 0;

    loop {
        while (bufavail == MAX) ; // do nothing
        track_image = frame_buf[head];
        head = (head + 1) % MAX;
        thread_mutex_lock(buflock);
        bufavail = bufavail + 1;
        thread_mutex_unlock(buflock);
        analyze(track_image);
    }
}
```

Busy waiting → Wastes CPU



We have deadlock when thread A is waiting on thread B and

30% A. I just want the participation credit

B. Thread B then waits on thread A

C. Thread B then waits on thread C which then waits on thread A

D. Thread B then tries to claim a mutex lock which is held by thread A

E. All of the above

What should really happen?


- If frame_buf is full
 - Tracker is slow, so digitizer is waiting for space in frame_buf
 - Tracker should let digitizer know when it makes room in frame_buf
- If frame_buf is empty
 - Digitizer is slow, so tracker is waiting for an image in frame_buf
 - Digitizer should let tracker know when it adds an image to frame_buf
- That would stop the busy-waiting

Add a condition variable

- Condition variable functions
 - `pthread_cond_wait`: block for a signal
 - `pthread_cond_signal`: signal **one** waiting thread
 - `pthread_cond_broadcast`: signal **all** waiting threads
- Semantics (OS invariants)
 - `pthread_cond_wait (cond_var c, mutex m)`
 - Atomically release mutex m
 - Put thread to sleep waiting on a signal to cond_var c
 - Atomically re-lock mutex m on awakening

Say we have 3 threads, T1-T3 that all wait on cond_var c.


This is what the OS does:



Add a condition variable

- Condition variable functions
 - `pthread_cond_wait`: block for a signal
 - `pthread_cond_signal`: signal **one** waiting thread
 - `pthread_cond_broadcast`: signal **all** waiting threads
- Semantics (OS invariants)
 - `pthread_cond_signal (cond_var c)`
 - Wake up one thread waiting on `cond_var c`
 - [The signaled thread will then go on to reclaim its mutex before proceeding.]

We have our 3 threads, T1-T3 all waiting on `cond_var c` when signal is called.



Condition variable

- We use condition variables to avoid busy waiting before entering critical sections
- They are a method of inter-process (or inter-thread) communication
- Condition variables *represent* a particular condition involving shared data, but despite their name, they don't actually test for it
- **We** must actually write the code to test for the condition
- And we must make sure our code doesn't enter the critical section until the condition is true
- Since we can't enter the critical section if the condition is false, we can be certain that **we** can't make the condition true; some other thread must do it
- We depend on a notification from the code in another thread that can **make** the condition true to wake us up when the condition is true!
- We will call this condition an *invariant* for entering the critical section

Back to our surveillance app

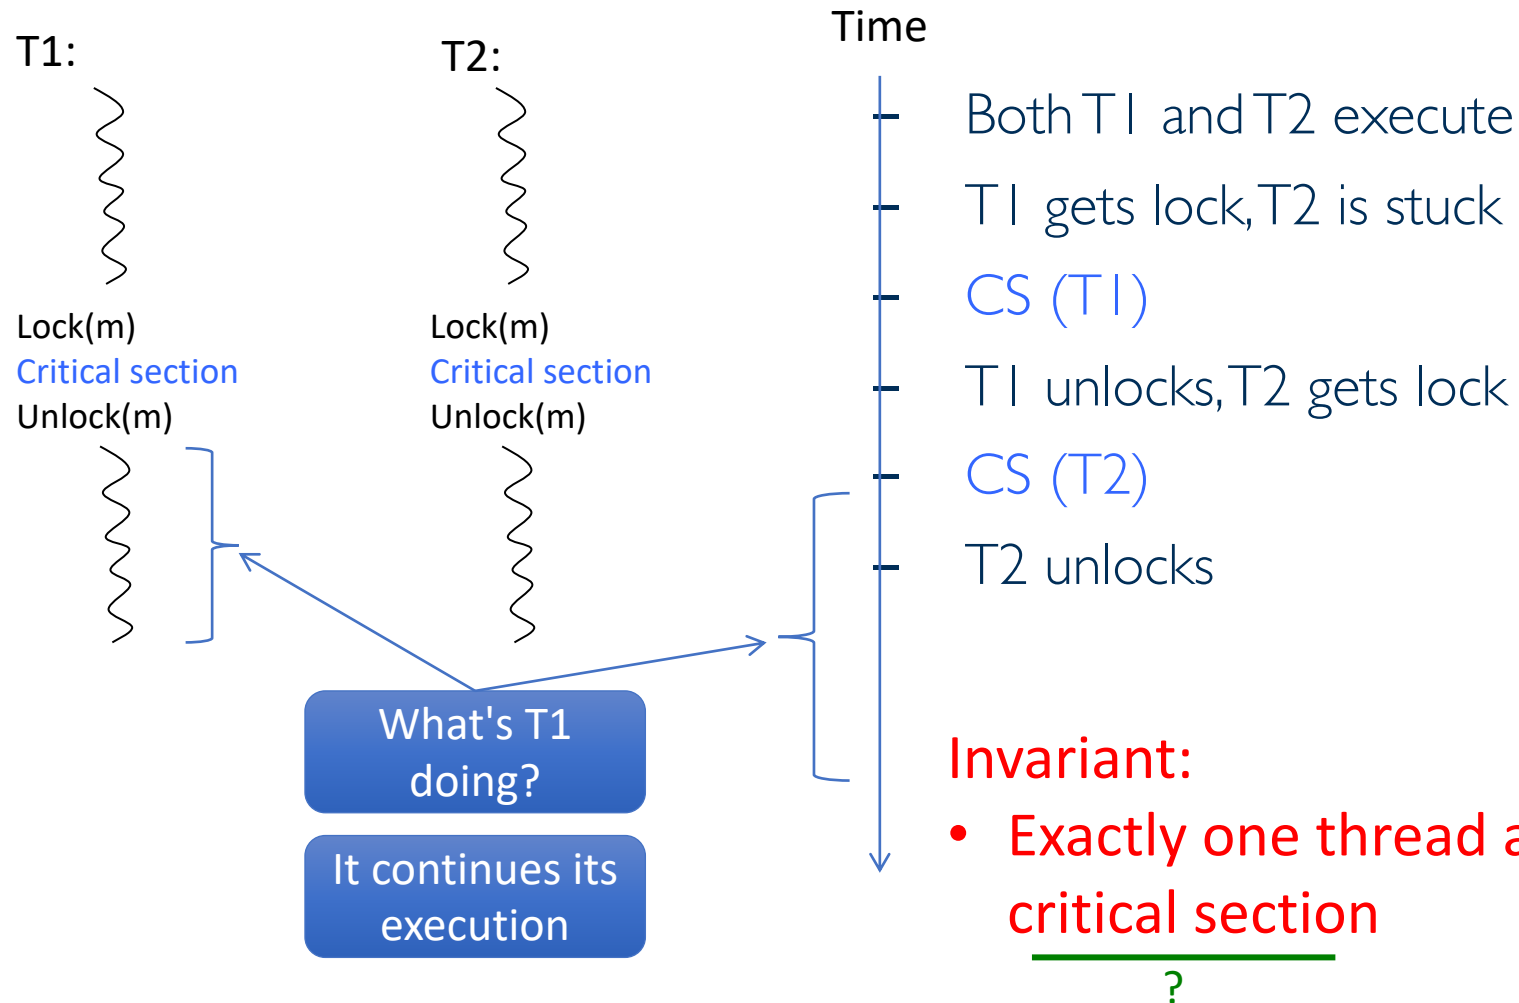
- What were we waiting for?
 - Digitizer waits for frame_buf to be not full
 - Tracker waits for frame_buf to be not empty

Fix #4 – cond_var


```
cond_var buf_not_full, buf_not_empty;
digitizer()
{
    image_type dig_image;
    int tail = 0;
    loop {
        grab(dig_image);
        thread_mutex_lock(buflock);
        if (bufavail == 0)
            thread_cond_wait(buf_not_full, buflock);
        thread_mutex_unlock(buflock);
        frame_buf[tail] = dig_image;
        tail = (tail + 1) % MAX;
        thread_mutex_lock(buflock);
        bufavail = bufavail - 1;
        thread_cond_signal(buf_not_empty);
        thread_mutex_unlock(buflock);
    }
}
```

```
tracker()
{
    image_type track_image;
    int head = 0;
    loop {
        thread_mutex_lock(buflock);
        if (bufavail == MAX)
            thread_cond_wait(buf_not_empty, buflock);
        thread_mutex_unlock(buflock);
        track_image = frame_buf[head];
        head = (head + 1) % MAX;
        thread_mutex_lock(buflock);
        bufavail = bufavail + 1;
        thread_cond_signal(buf_not_full);
        thread_mutex_unlock(buflock);
        analyze(track_image);
    }
}
```

Recall: Mutex locks



More than one critical section?

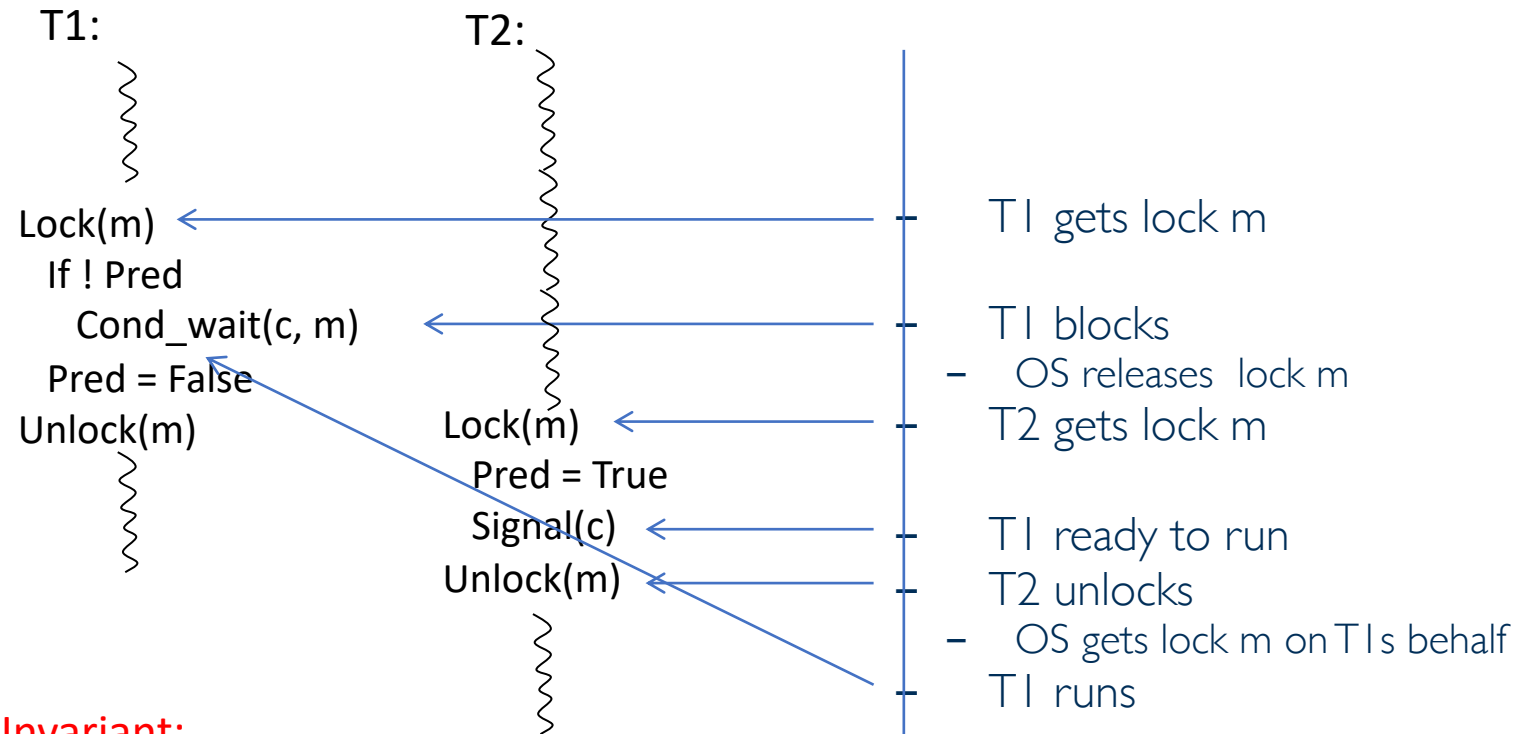


Different locks!
Thus these two critical
sections can run at the
same time!

OS Invariant:

- Exactly one thread at a time in a given critical section
- Different mutex locks create **different** critical sections
- But still at most one thread can run in each!

Recall: Condition variables



Invariant:

- OS ensures T1 gets lock m back
- Anything else?
 - **Pred has to be True**

Who must ensure this?

The programmer!

Back again to our surveillance app

- What were we waiting for?
 - Digitizer waits for frame_buf to be not full
 - Predicate is (bufavail != 0)
 - Tracker waits for frame_buf to be not empty
 - Predicate is (bufavail != MAX)
- So we need two condition variables
 - buf_not_empty and buf_not_full
 - And we know how to test for these conditions using the predicates

Fix #4 – cond_var

```
cond_var buf_not_full, buf_not_empty;
```

```
digitizer()
```

```
{
```

```
    image_type dig_image;
```

```
    int tail = 0;
```

```
    loop {
```

```
        grab(dig_image);
```

```
        thread_mutex_lock(buflock);
```

```
        if (bufavail == 0)
```

```
            thread_cond_wait(buf_not_full, buflock);
```

```
        thread_mutex_unlock(buflock);
```

```
        frame_buf[tail] = dig_image;
```

```
        tail = (tail + 1) % MAX;
```

```
        thread_mutex_lock(buflock);
```

```
        bufavail = bufavail - 1;
```

```
        thread_cond_signal(buf_not_empty);
```

```
        thread_mutex_unlock(buflock);
```

```
}
```

```
} Invariants: Only 1 thread at a time in CS  
              bufavail != 0
```

```
tracker()
```

```
{
```

```
    image_type track_image;
```

```
    int head = 0;
```

```
    loop {
```

```
        thread_mutex_lock(buflock);
```

```
        if (bufavail == MAX)
```

```
            thread_cond_wait(buf_not_empty, buflock);
```

```
        thread_mutex_unlock(buflock);
```

```
        track_image = frame_buf[head];
```

```
        head = (head + 1) % MAX;
```

```
        thread_mutex_lock(buflock);
```

```
        bufavail = bufavail + 1;
```

```
        thread_cond_signal(buf_not_full);
```

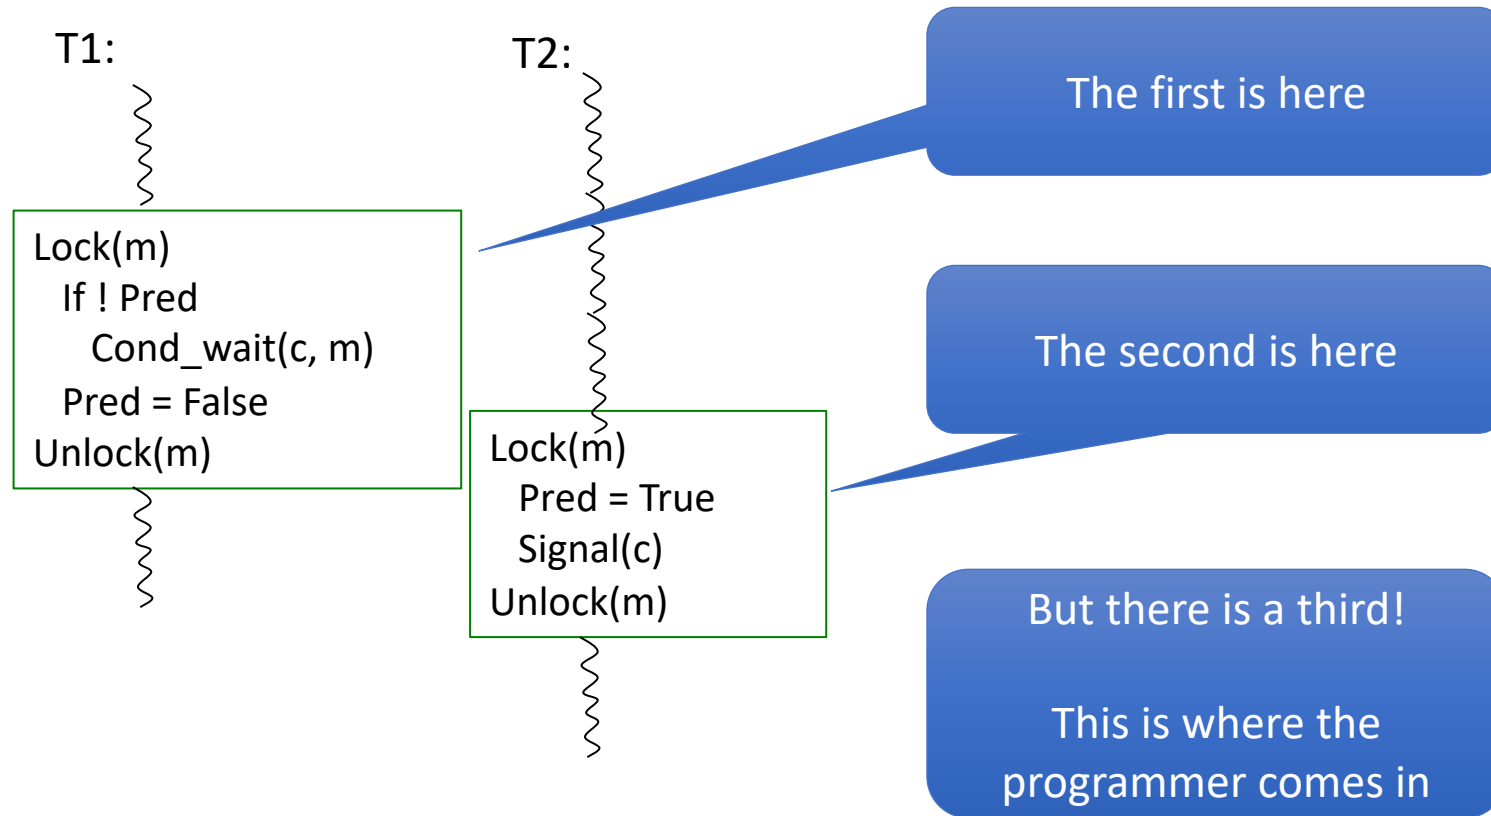
```
        thread_mutex_unlock(buflock);
```

```
        analyze(track_image);
```

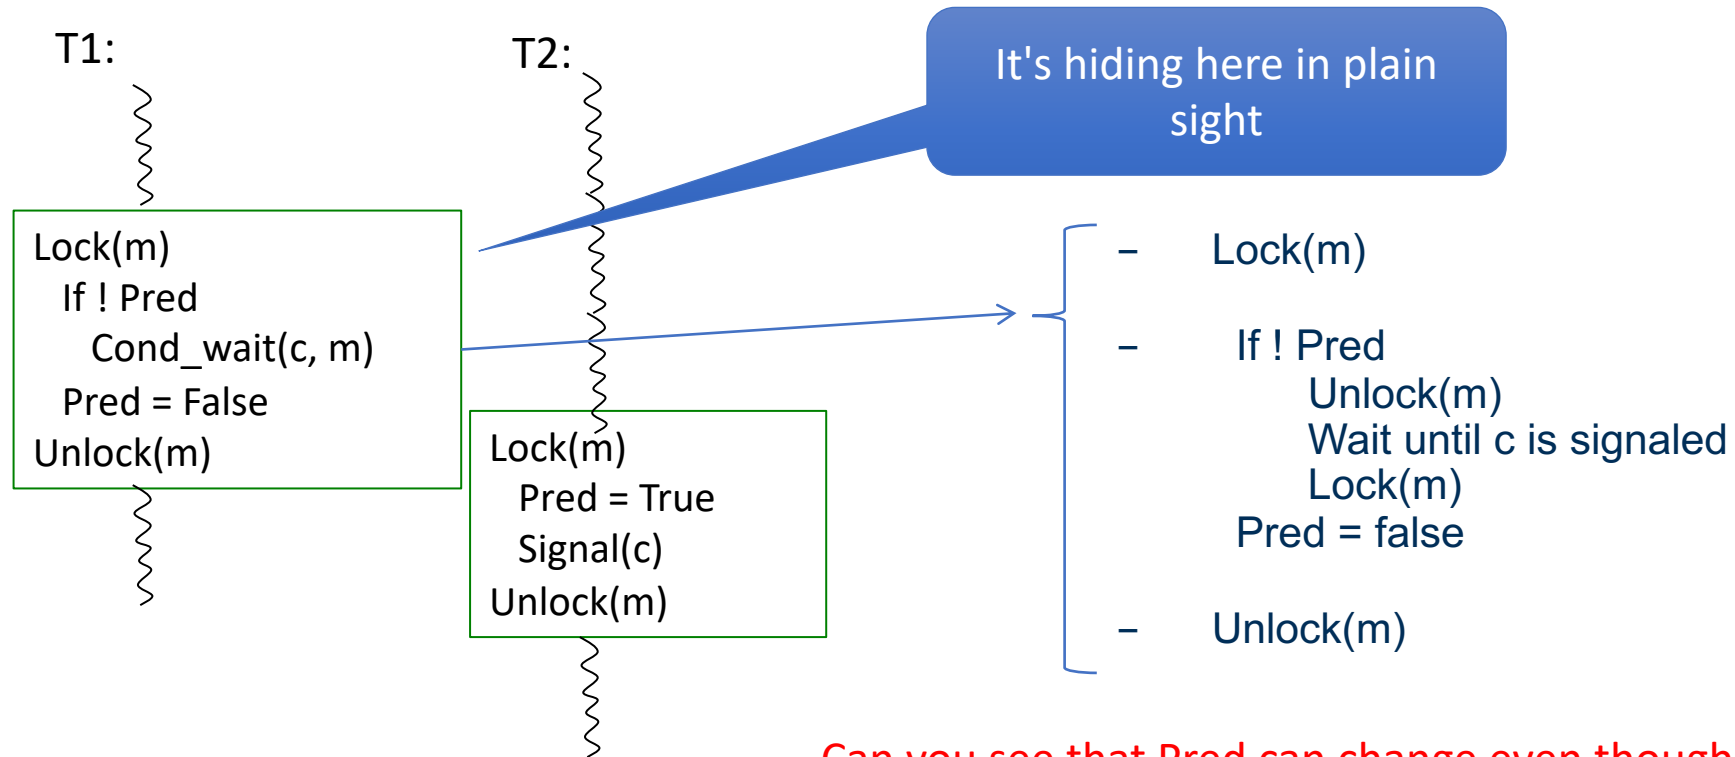
```
}
```

```
} Only 1 thread at a time in CS  
    bufavail != MAX
```

Just how many code blocks are in the critical section here?



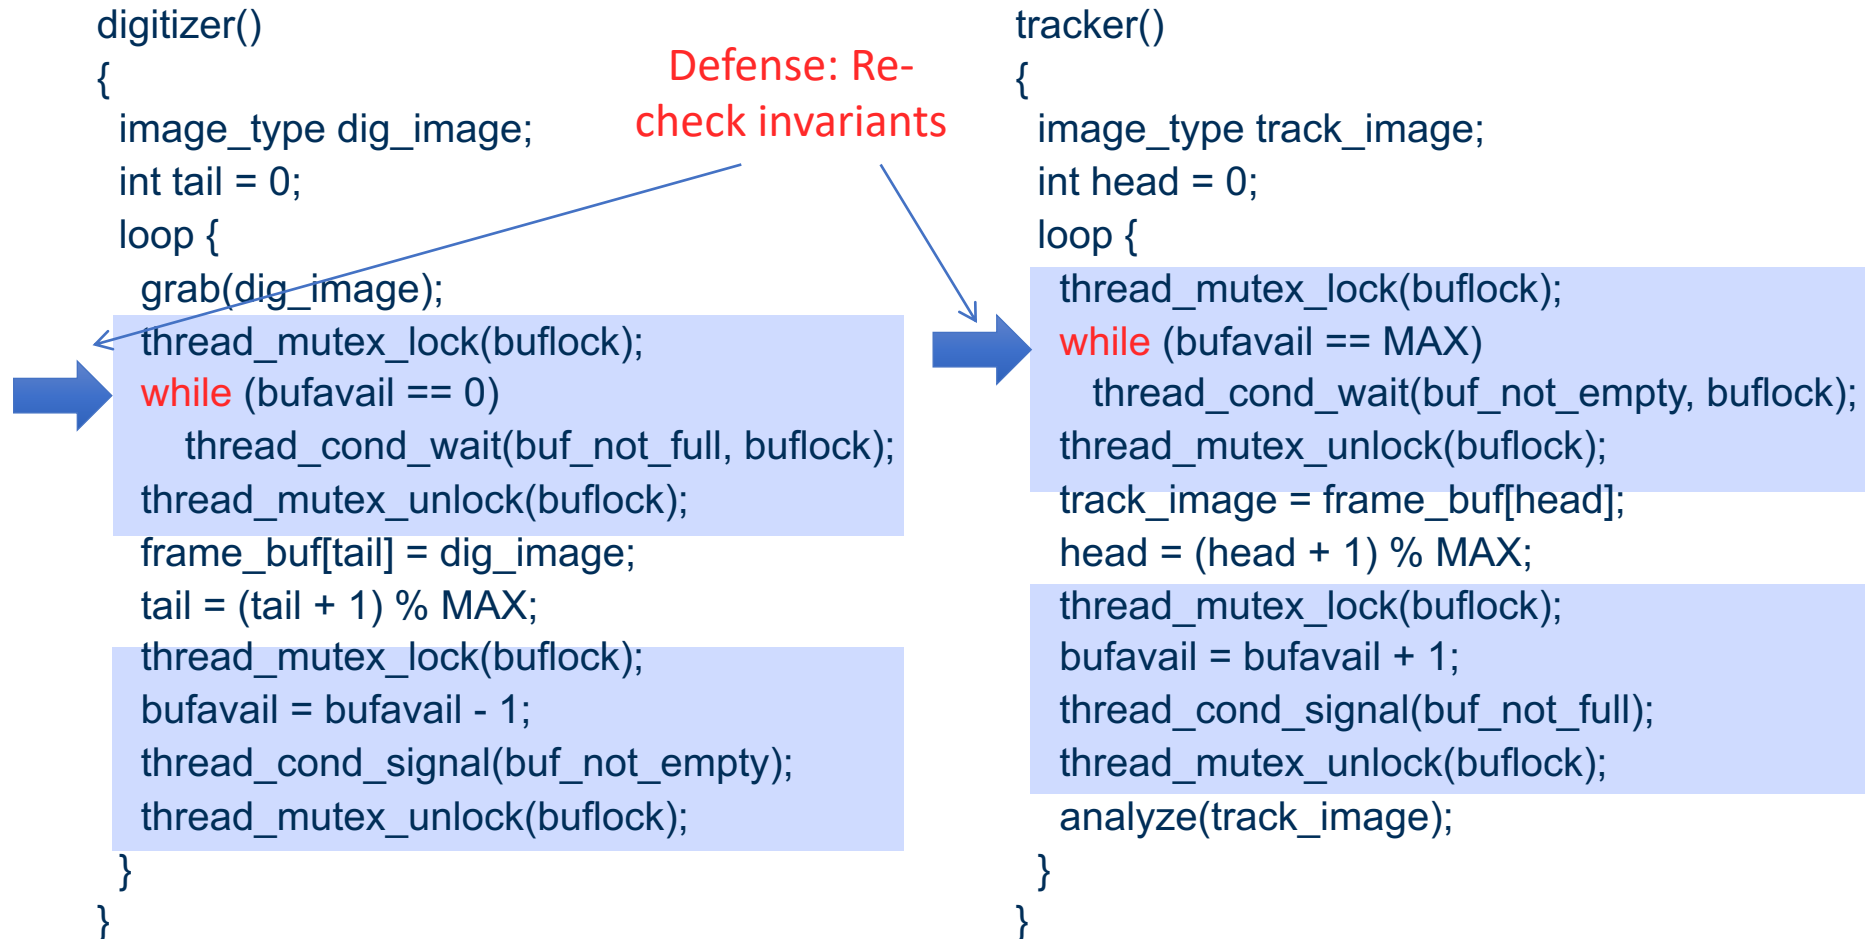
Where is that third block?



Can you see that Pred can change even though this block appears to be a single critical section?

This detail becomes a problem if there are more threads...

Fix # 5 – Defensive programming



The notion of re-checking a flag after waiting is important.




A condition variable...


13%

- A. I just want the participation credit
- B. ... is just another name for a mutex lock
- C. ...enables a thread to wait for a condition to become true without consuming processor cycles
- D. ...enables a thread to enter a critical section
- E. ...none of the above (B to D)

Gotchas in programming with cond vars



(a) Wait before signal



(b) Wait after signal (T1 blocked forever)


Gotchas in programming with cond vars

Say we have three threads that want to share a resource, perhaps a printer...

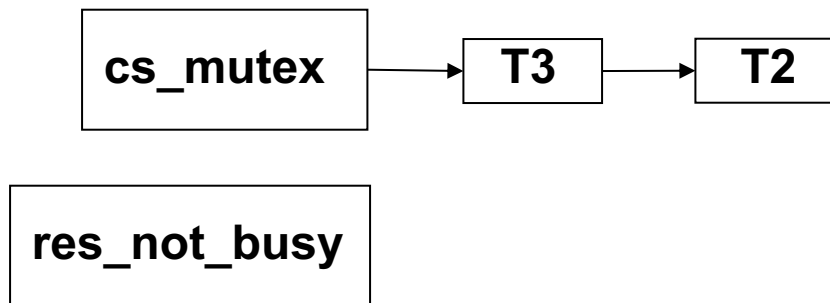
```
acquire_shared_resource()
{
    thread_mutex_lock(cs_mutex); ← T3 is here
    if (res_state == BUSY)
        thread_cond_wait(res_not_busy, cs_mutex); ← T2 is here
    res_state = BUSY;
    thread_mutex_unlock(cs_mutex);
}

release_shared_resource()
{
    thread_mutex_lock(cs_mutex);
    res_state = NOT_BUSY;
    thread_cond_signal(res_not_busy); ← T1 is here
    thread_mutex_unlock(cs_mutex);
}
```

State of waiting queues



(a) Waiting queues before T1 signals



(b) Waiting queues after T1 signals

```
acquire_shared_resource()
{
    thread_mutex_lock(cs_mutex); ← T3 is here
    if (res_state == BUSY)
        thread_cond_wait(res_not_busy, cs_mutex); ← T2 is here
    res_state = BUSY;
    thread_mutex_unlock(cs_mutex);
}

release_shared_resource()
{
    thread_mutex_lock(cs_mutex);
    res_state = NOT_BUSY;
    thread_cond_signal(res_not_busy); ← T1 is here
    thread_mutex_unlock(cs_mutex);
}
```

Gotchas -- what could go wrong?

T1 signals and unlocks
mutex

What if T3 wakes up and
locks the mutex?

T3 sets res_state to BUSY,
unlocks the mutex, and goes
off to use the resource

T2 then locks the mutex

T2 has already tested res_state, so it
unlocks the mutex and goes off to use
the resource(!)

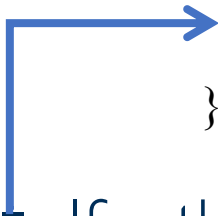
```
acquire_shared_resource()
{
    thread_mutex_lock(cs_mutex); ← T3 is here
    if (res_state == BUSY)
        thread_cond_wait(res_not_busy, cs_mutex); ← T2 is here
    res_state = BUSY;
    thread_mutex_unlock(cs_mutex);
}

release_shared_resource()
{
    thread_mutex_lock(cs_mutex);
    res_state = NOT_BUSY;
    thread_cond_signal(res_not_busy); ← T1 is here
    thread_mutex_unlock(cs_mutex);
}
```

Why did this happen?

- We violated invariants...

```
acquire_shared_resource()
{
    thread_mutex_lock(cs_mutex);
    if (res_state == BUSY)
        thread_cond_wait(res_not_busy, cs_mutex);
    res_state = BUSY;
    thread_mutex_unlock(cs_mutex);
}
```



- If a thread is here, what are the invariants?
 - The thread holds the mutex
 - ➔ the OS ensures that
 - `res_state == NOT_BUSY`
 - ➔ the **programmer** ensures that

Gotchas in programming - 3

- There's yet another surprise...
 - It's possible to have a spurious wake-up of threads by the OS
 - Even without a signal, a thread may be woken up
 - Documented behavior in Linux
 - Turns out to be very hard to avoid this in the kernel
 - Solution: **Defensive programming**

Gotchas— retest predicate

Replace the
"if" with a
"while"

```
acquire_shared_resource()
{
    thread_mutex_lock(cs_mutex);
    while (res_state == BUSY)
        thread_cond_wait(res_not_busy, cs_mutex);
    res_state = BUSY;
    thread_mutex_unlock(cs_mutex);
}

release_shared_resource()
{
    thread_mutex_lock(cs_mutex);
    res_state = NOT_BUSY;
    thread_cond_signal(res_not_busy);
    thread_mutex_unlock(cs_mutex);
}
```

Make T2 recheck predicate


Avoids the "race condition"

Prevents a "timing bug" or non-deterministic result in a parallel program!


Checkpoint

- ~~Pthreads programming~~
- OS issues with threads
- Hardware support for threads


Threads as software structuring abstraction



(a) Dispatcher model



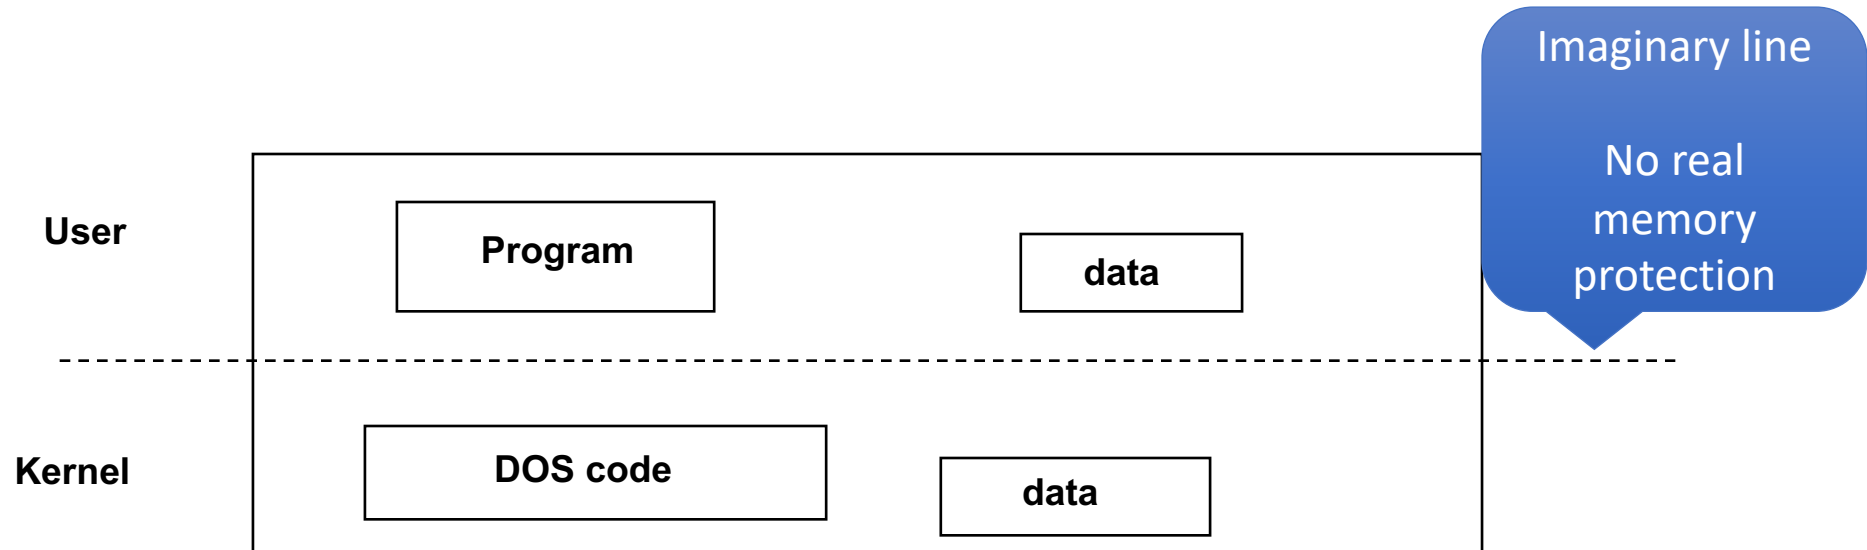
(b) Team model



(c) Pipelined model

Traditional OS: DOS

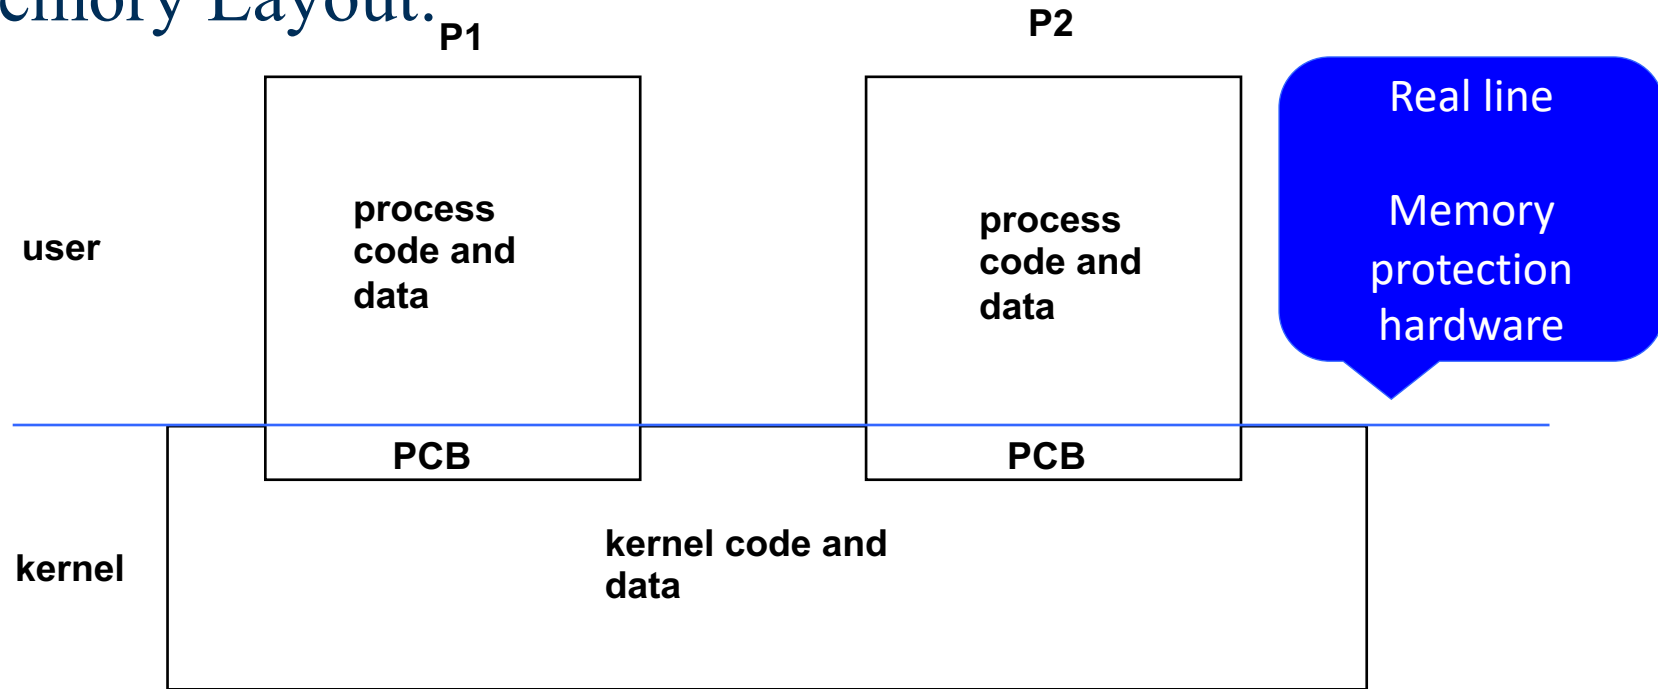
Memory Layout:



- Protection between user and kernel?
- Single process, single thread

Traditional OS: Unix

Memory Layout:



- Protection between user and kernel?
- PCB?
- Multiple processes, one thread each

Tradition

- Programs in these traditional OS are single threaded
 - One PC per program (process), one stack, one set of CPU registers
 - If a process blocks (say disk I/O, network communication, etc.) then no progress for that program as a whole

Multi-Threaded Operating Systems

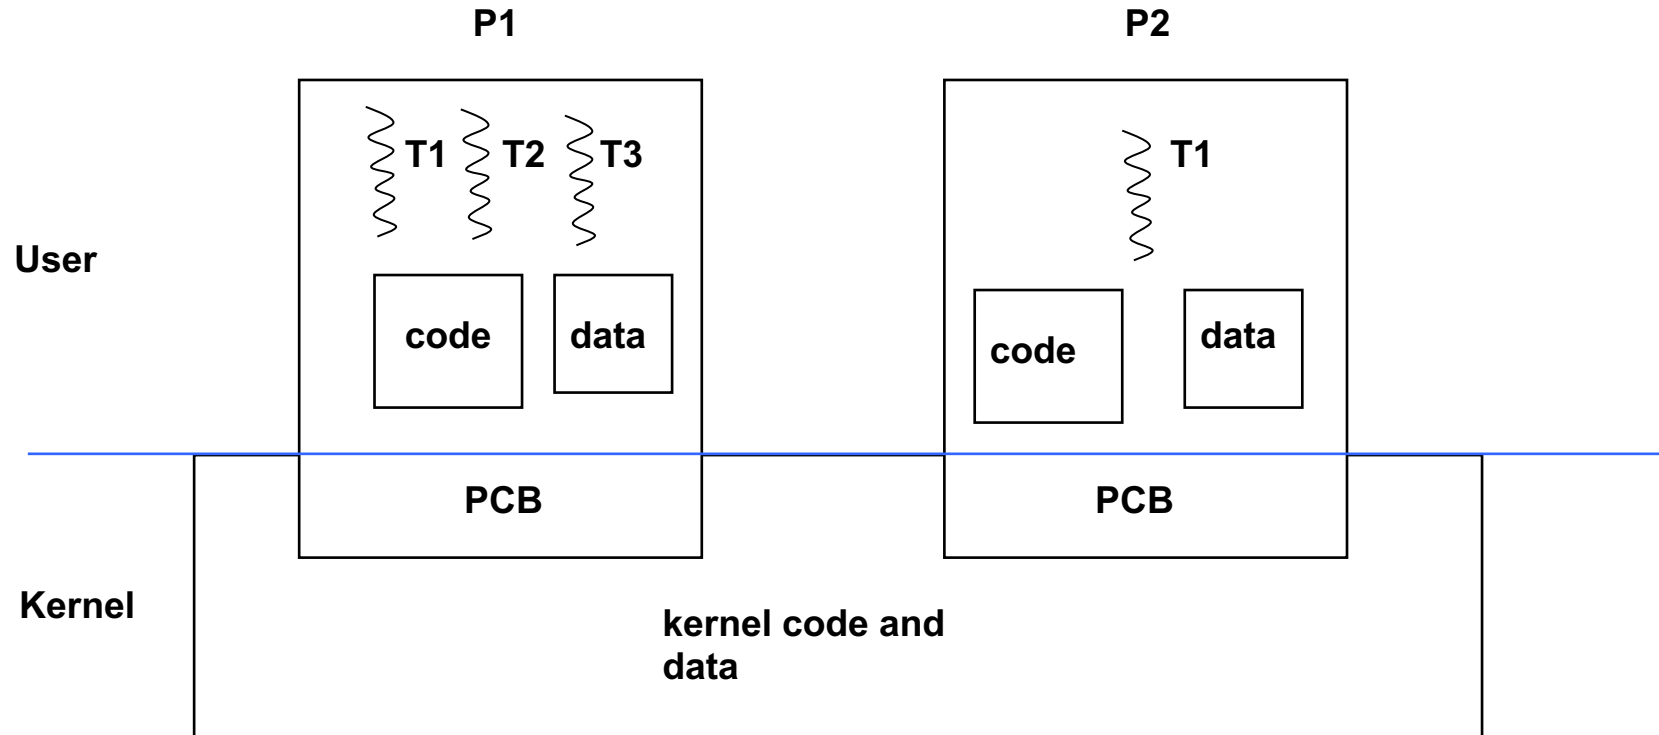
How widespread is support for threads in OS?

- Linux, MacOS, iOS, Android, Windows
- (In other words, every modern operating system)

Process Vs. Thread?

- In a single-threaded program, the state of the executing program is contained in a process
- In a MT program, the state of the executing program is contained in several 'concurrent' threads

Process Vs. Thread




- Computational state (PC, regs, ...) for each thread
- How different from process state?
 - There's a lot of admin info in common

MT Bookkeeping



(a) ST program



(b) MT program

- Can you see why the stack is sometimes called a "cactus stack"?

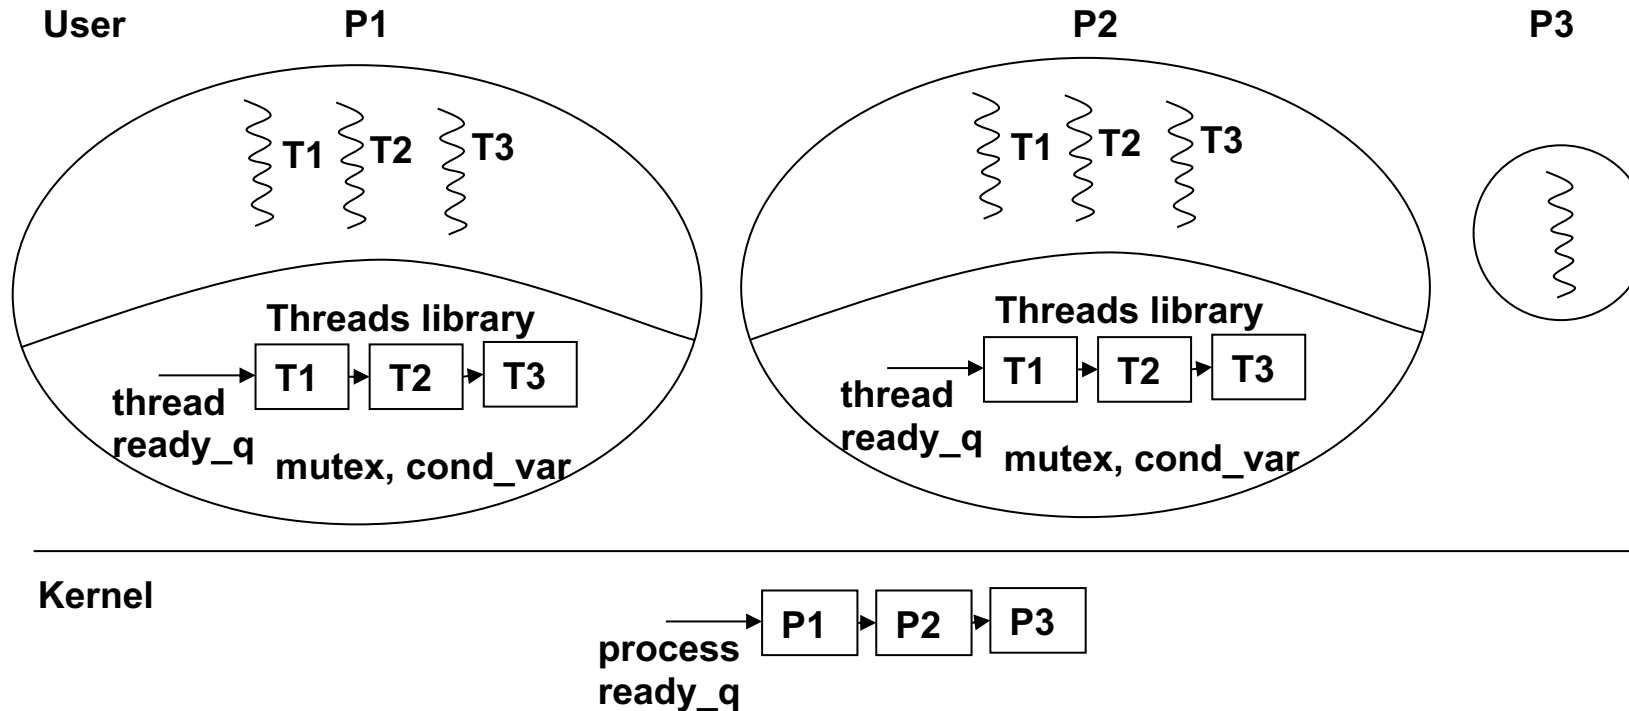
Thread properties

- Threads
 - Share address space of process
 - Cooperate to get job done
- Threads concurrent?
 - Truly parallel if the box is a true multiprocessor
 - Share (time-multiplex) the same CPU on a uniprocessor
- Threaded code different from non-threaded?
 - Protection for data shared among threads
 - Synchronization among threads

User-Level Threads


- OS independent
- Scheduler is part of the user space runtime system
- Thread switching is cheap (just save PC, SP, regs)
- Scheduling is customizable, i.e., more app control
- Blocking call by a thread blocks a process

User-level threads



- OS independent
- Thread library part of application runtime
- Thread switching is cheap
- User-customizable thread scheduling

User-level threads



- Problem?
- Unfortunately, I/O blocks the entire process



User-level threads with process level scheduling...

10%

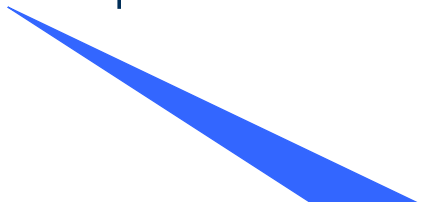
- A. I just want the participation credit
- B. ...serves no purpose since the operating system does not schedule at the thread level
- C. ...is useful for overlapping computation with I/O
- D. ...is useful as a software structuring mechanism at the user level
- E. All of the above

Kernel-level threads

- The norm in most modern operating systems
- Thread switch is more expensive
- Makes sense for blocking calls by threads
- Kernel becomes more complicated dealing with process and thread scheduling
- Thread packages become OS-dependent and non-portable

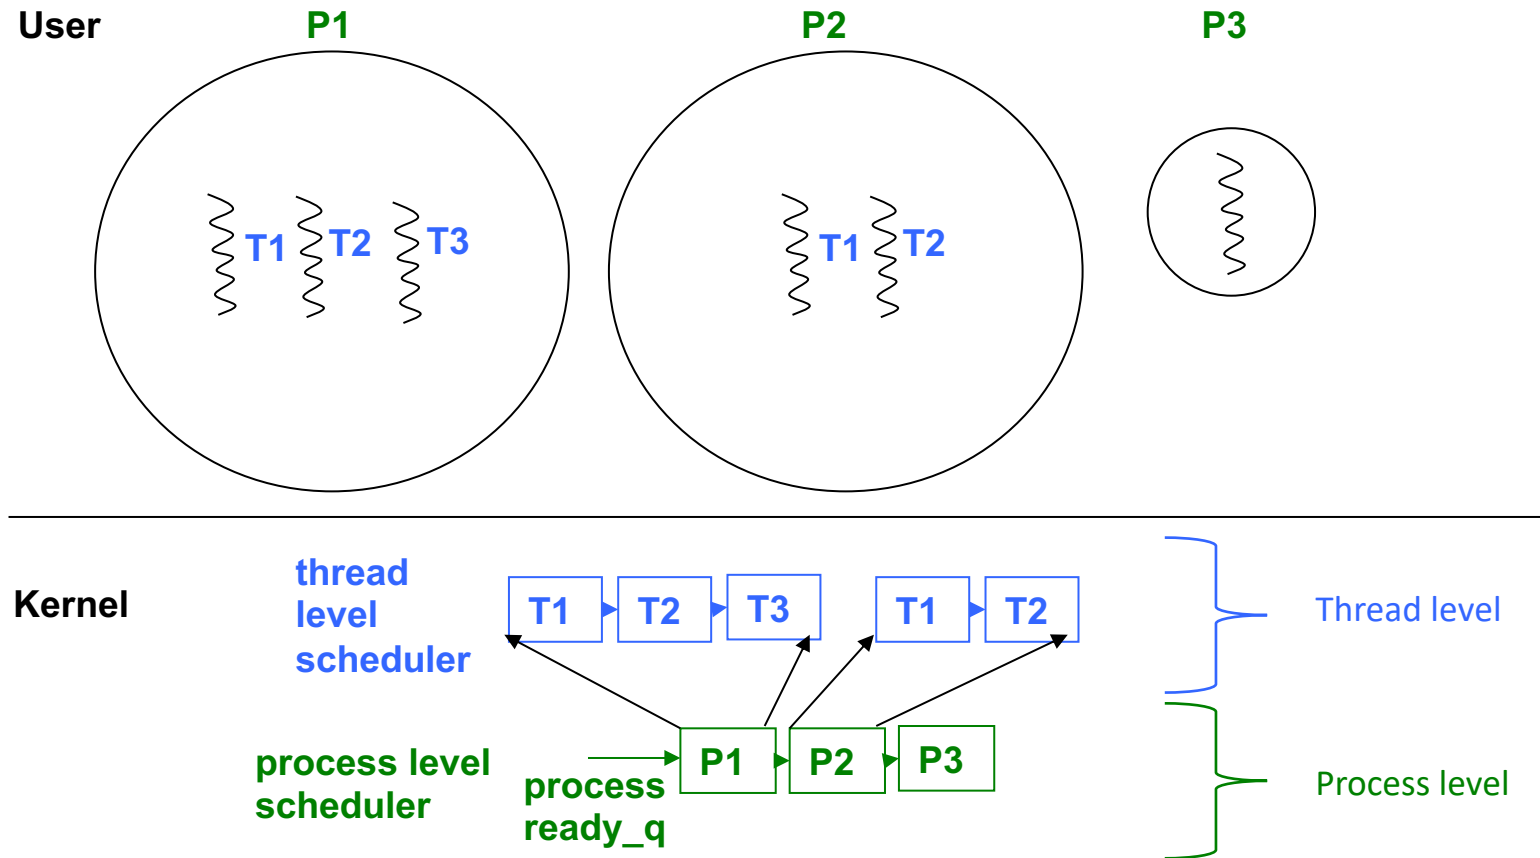


Compelling
reason for
kernel-level
threads



Reason for the
existence of
pthreads (POSIX)
standard

Two-level OS scheduler



- Threading in the application is visible to the OS
- OS provides the thread library

Thread-safe libraries

- Library functions (methods) have concurrency issues when used by user and kernel-level threads
 - All threads in a process share the heap and static data areas
 - Library routines that use static data or the heap are **very likely to implicitly share data** with other threads!
 - Solution is to have thread-safe wrappers to such library calls

Thread safe libraries

<code>/* original version */</code>		<code>/* thread safe version */</code>
		<code>mutex_lock_type cs_mutex;</code>
<code>void *malloc(size_t size)</code>		<code>void *malloc(size_t size)</code>
<code>{</code>		<code>{</code>
		<code> thread_mutex_lock(cs_mutex);</code>
<code> </code>		<code> memory_pointer = malloc(size);</code>
<code> </code>		<code> </code>
		<code> thread_mutex_unlock(cs_mutex);</code>
<code> return(memory_pointer);</code>		<code> return (memory_pointer);</code>
<code>}</code>		<code>}</code>