

CS2200  
Systems and Networks  
Spring 2022

Lecture 26:  
Disk Scheduling

Alexandros (Alex) Daglis  
School of Computer Science  
Georgia Institute of Technology  
[adaglis@gatech.edu](mailto:adaglis@gatech.edu)

# Agenda

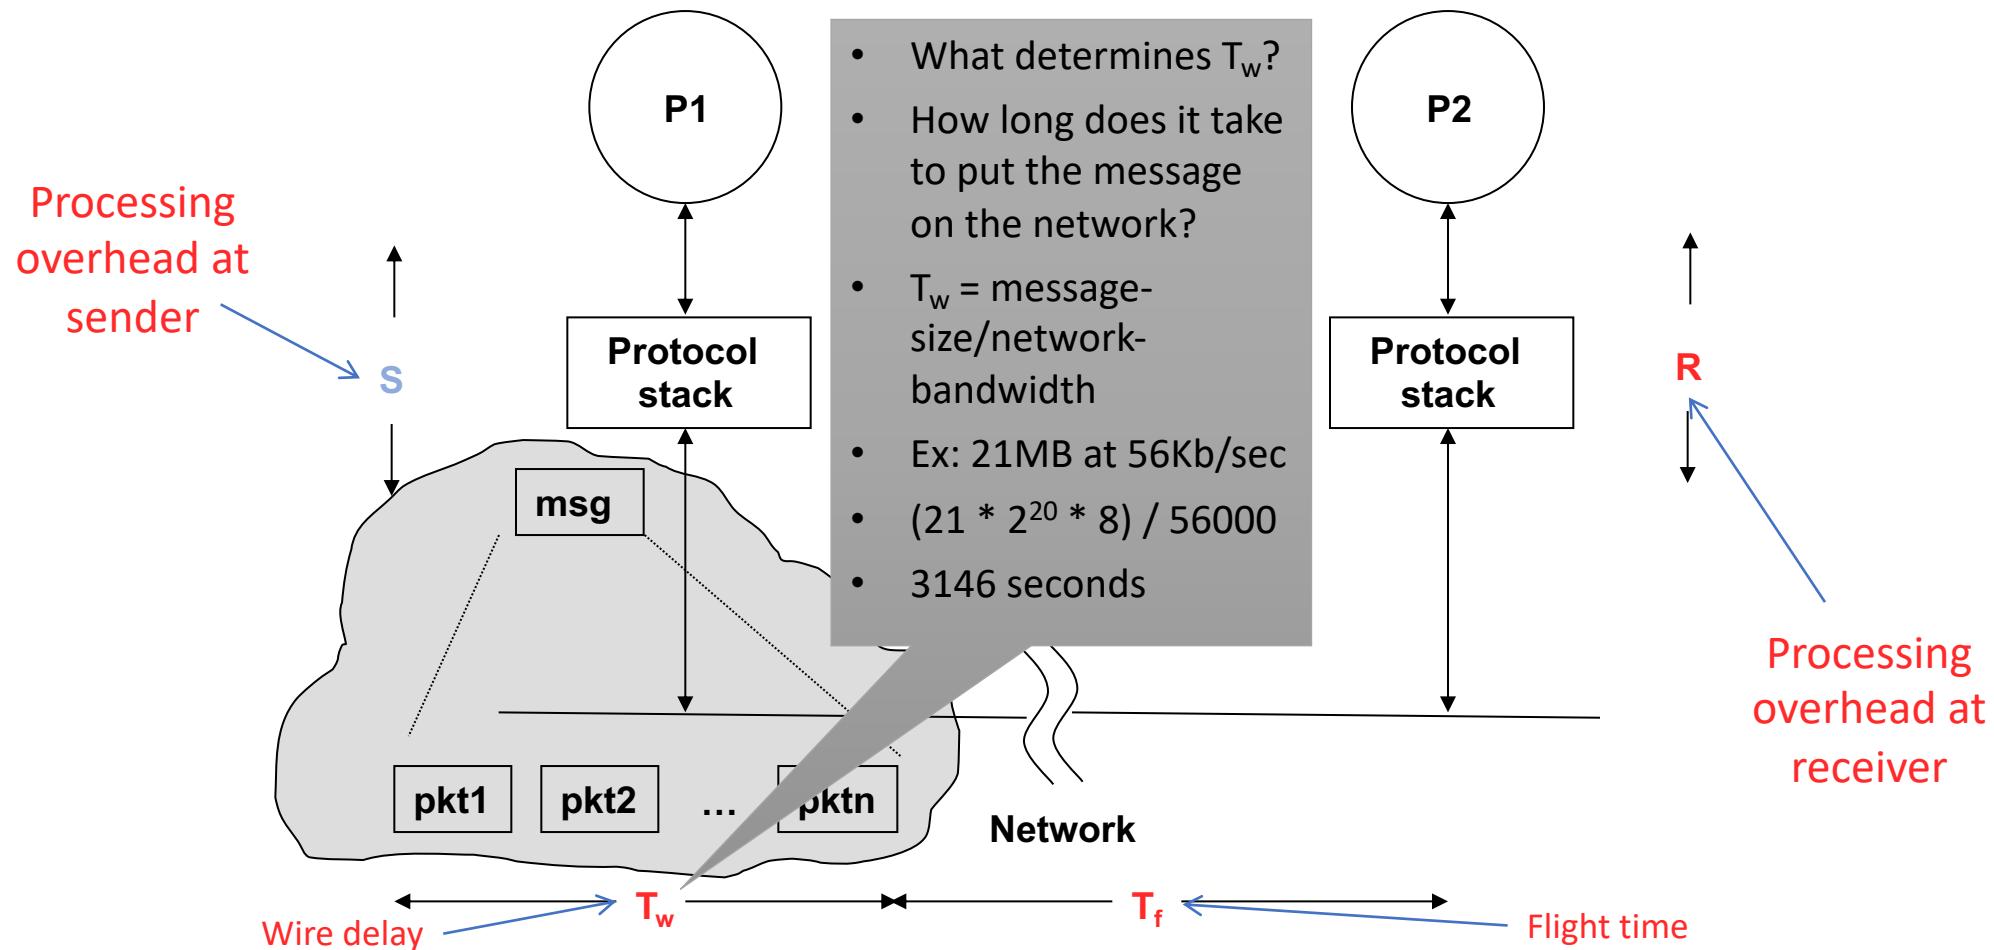
---

- Wrap up Networking
- Disk Scheduling

# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

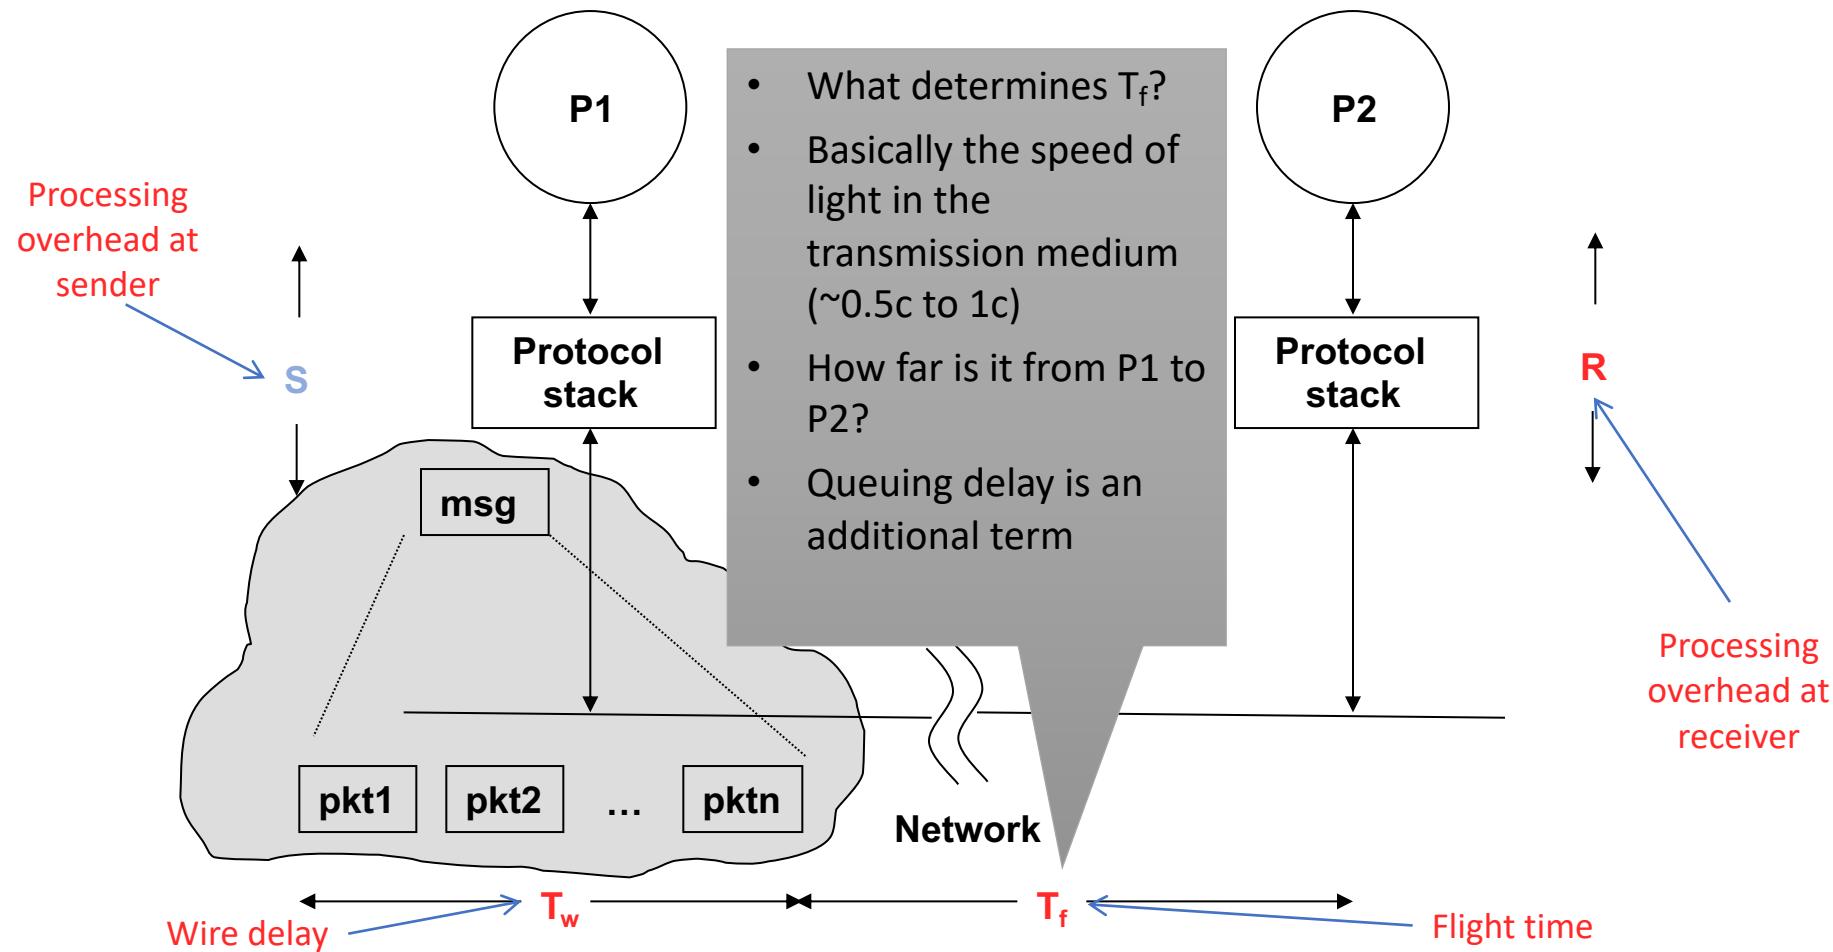
Message throughput = message-size/end-to-end-latency



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

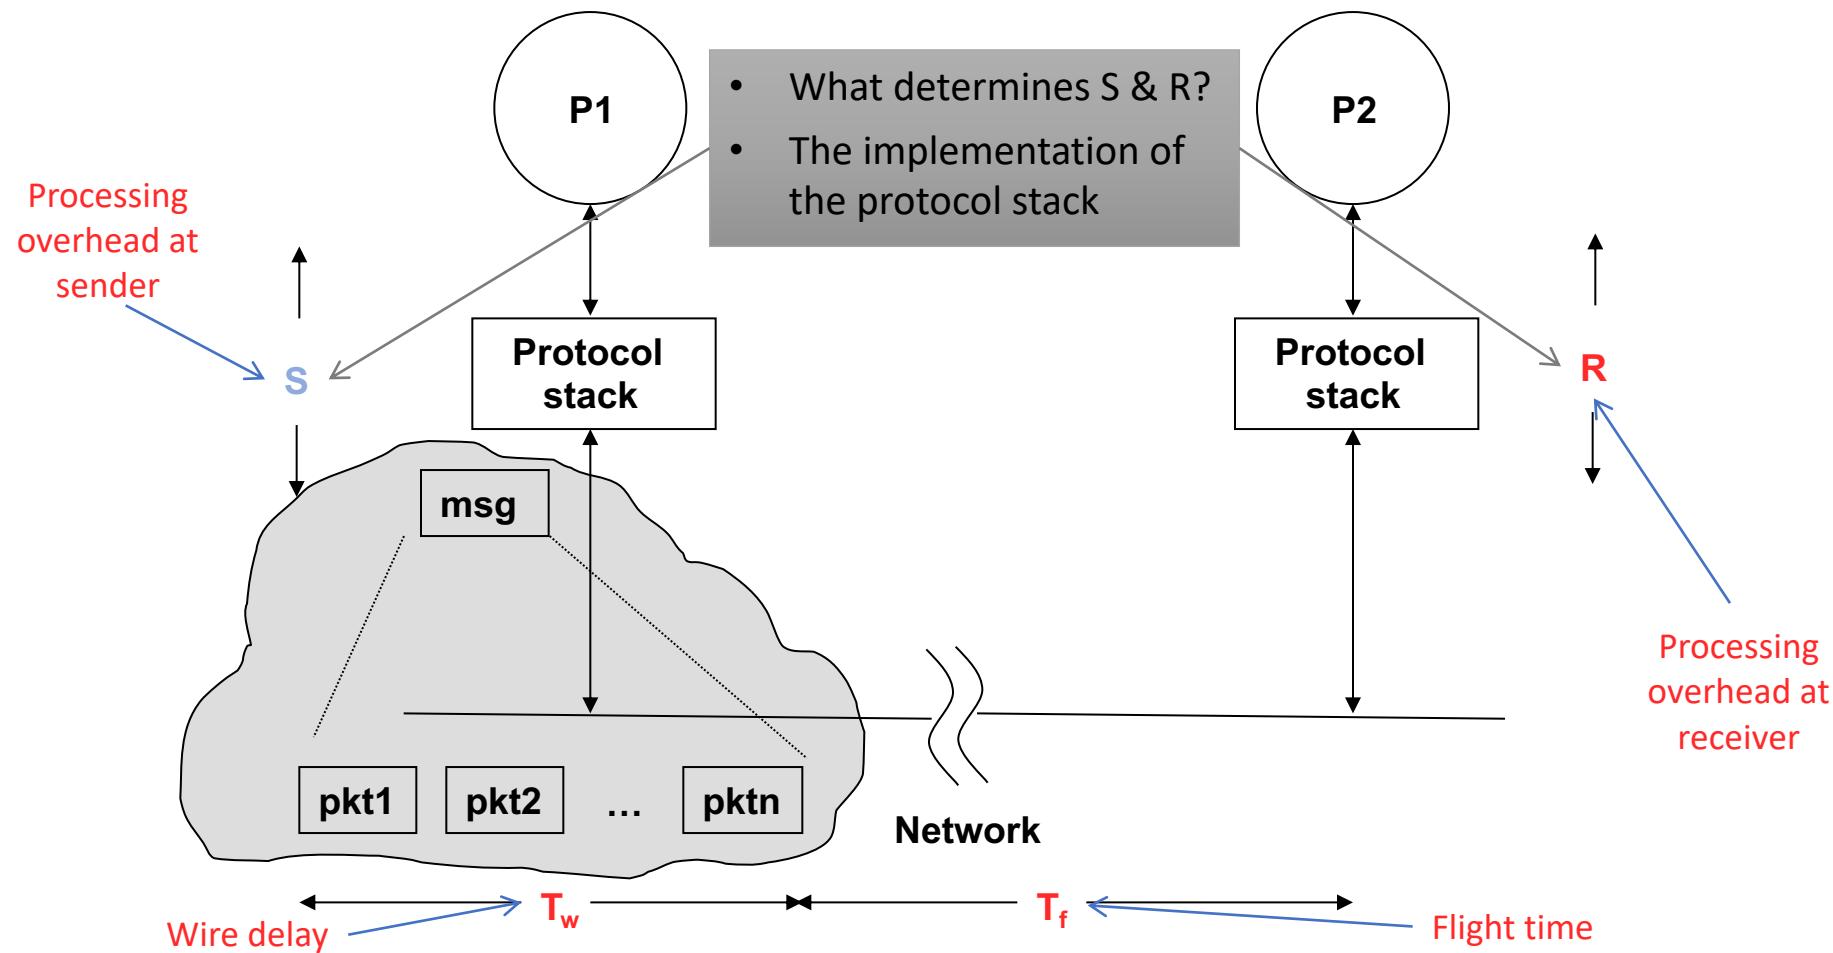
Message throughput = message-size/end-to-end-latency



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

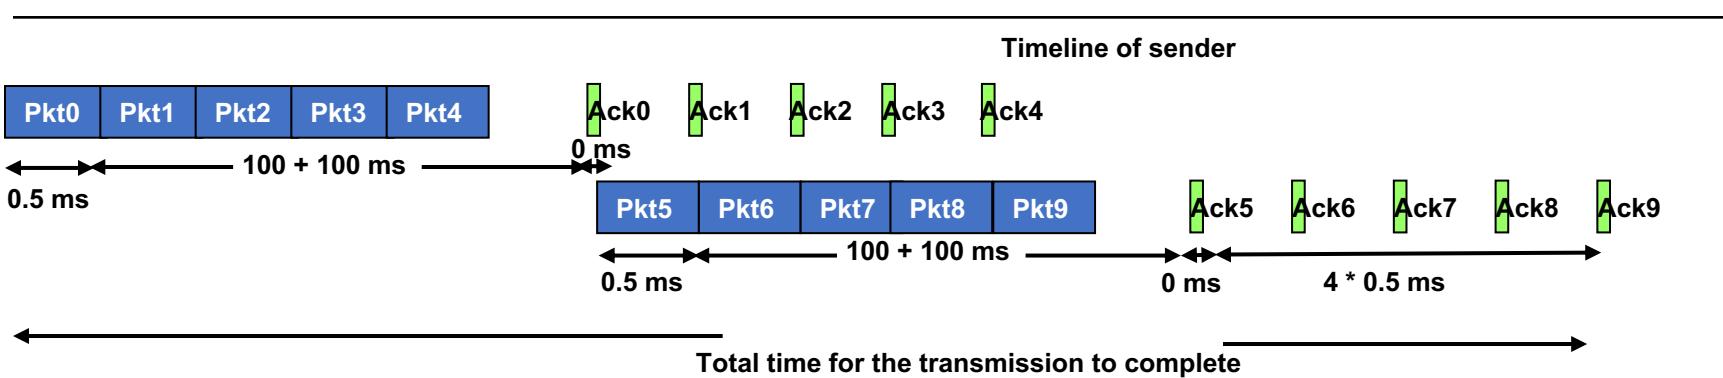
Message throughput = message-size/end-to-end-latency



# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



End-to-end latency for Pkt0

$$\begin{aligned} &= S + T_w + T_f + R \\ &= 0 + 0.5 + 100 + 0 \\ &= \textcolor{green}{100.5 \text{ ms}} \end{aligned}$$

End-to-end latency for Ack0

$$\begin{aligned} &= S + T_w + T_f + R \\ &= 0 + 0.0 + 100 + 0 \\ &= \textcolor{green}{100.0 \text{ ms}} \end{aligned}$$

Total transmission time

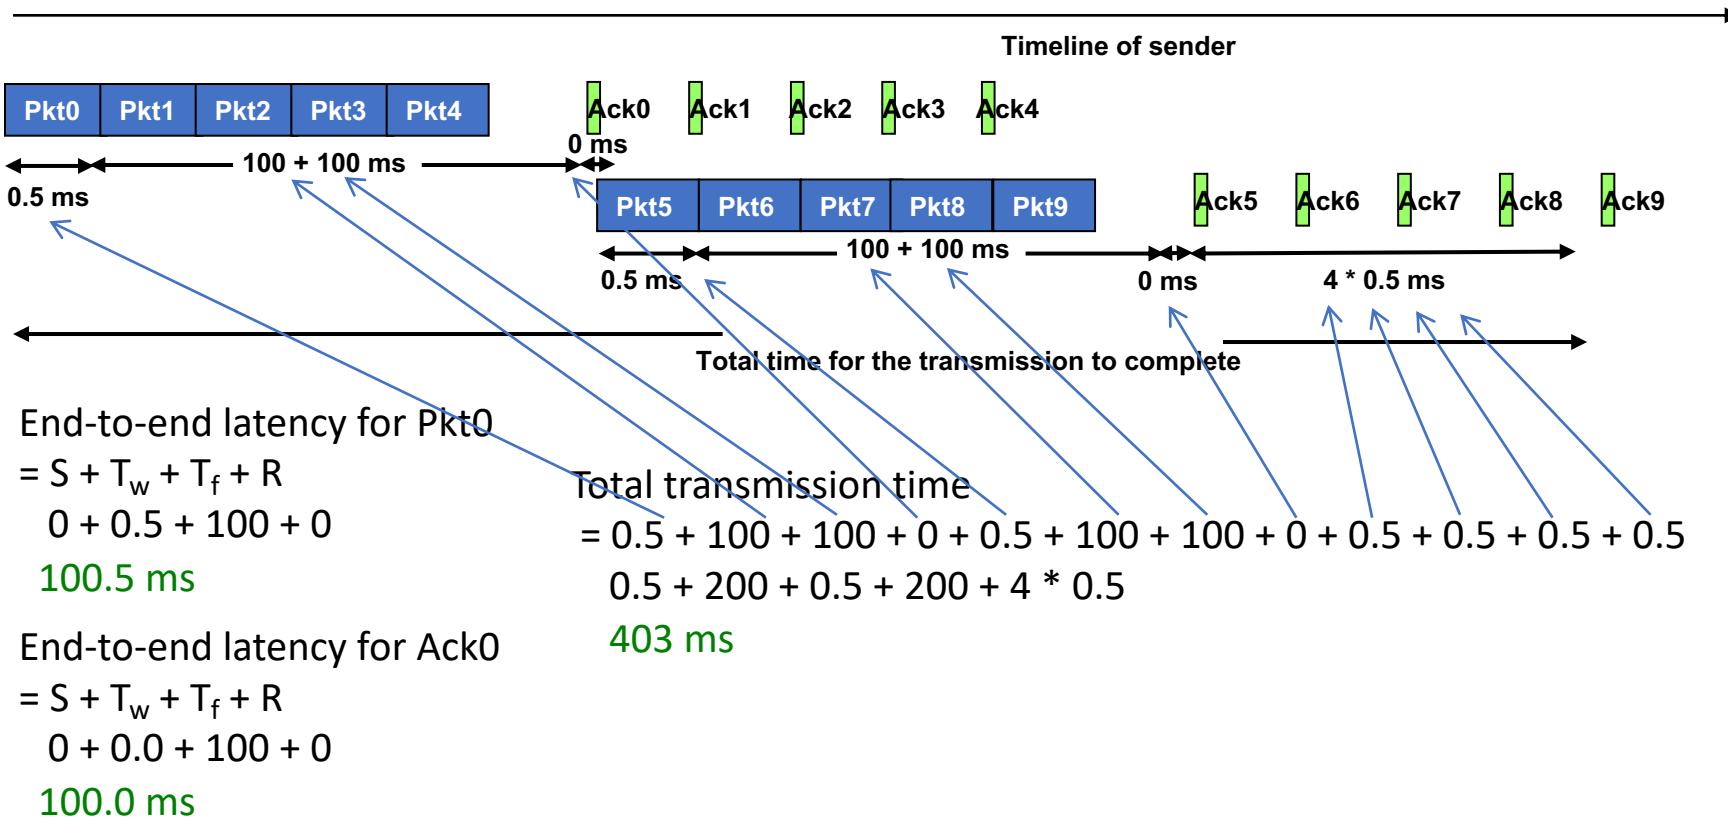
$$\begin{aligned} &= 0.5 + 100 + 100 + 0 + 0.5 + 100 + 100 + 0.5 + 0.5 + 0.5 + 0.5 \\ &= 0.5 + 200 + 0.5 + 200 + 4 * 0.5 \end{aligned}$$

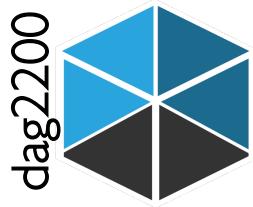
$\textcolor{green}{403 \text{ ms}}$

# Timeline example

100ms latency  
10 packets  
Window size = 5

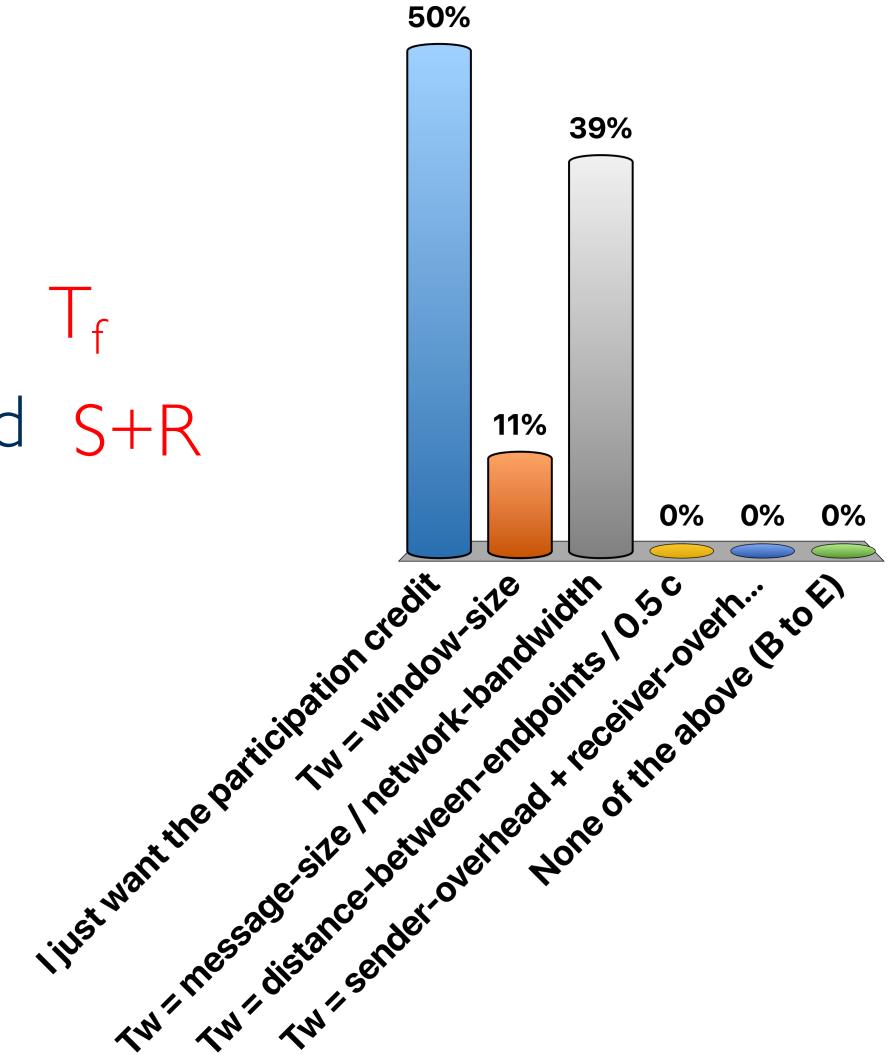
Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0





# In the expression $S + T_w + T_f + R$

- A. I just want the participation credit
- B.  $T_w$  = window-size
- C.  $T_w$  = message-size / network-bandwidth
- D.  $T_w$  = distance-between-endpoints / 0.5 c
- E.  $T_w$  = sender-overhead + receiver-overhead
- F. None of the above (B to E)



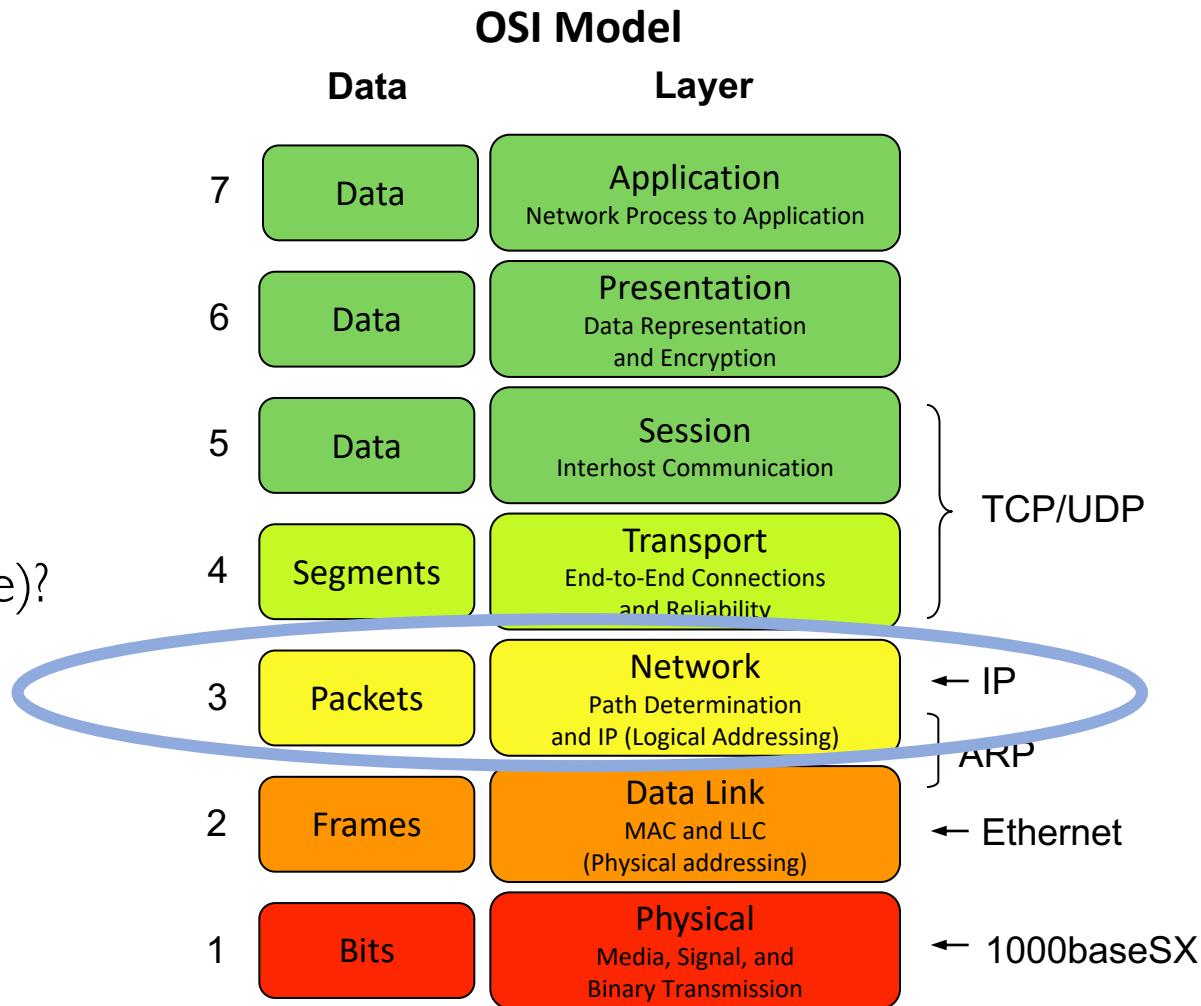
# Acknowledgement example

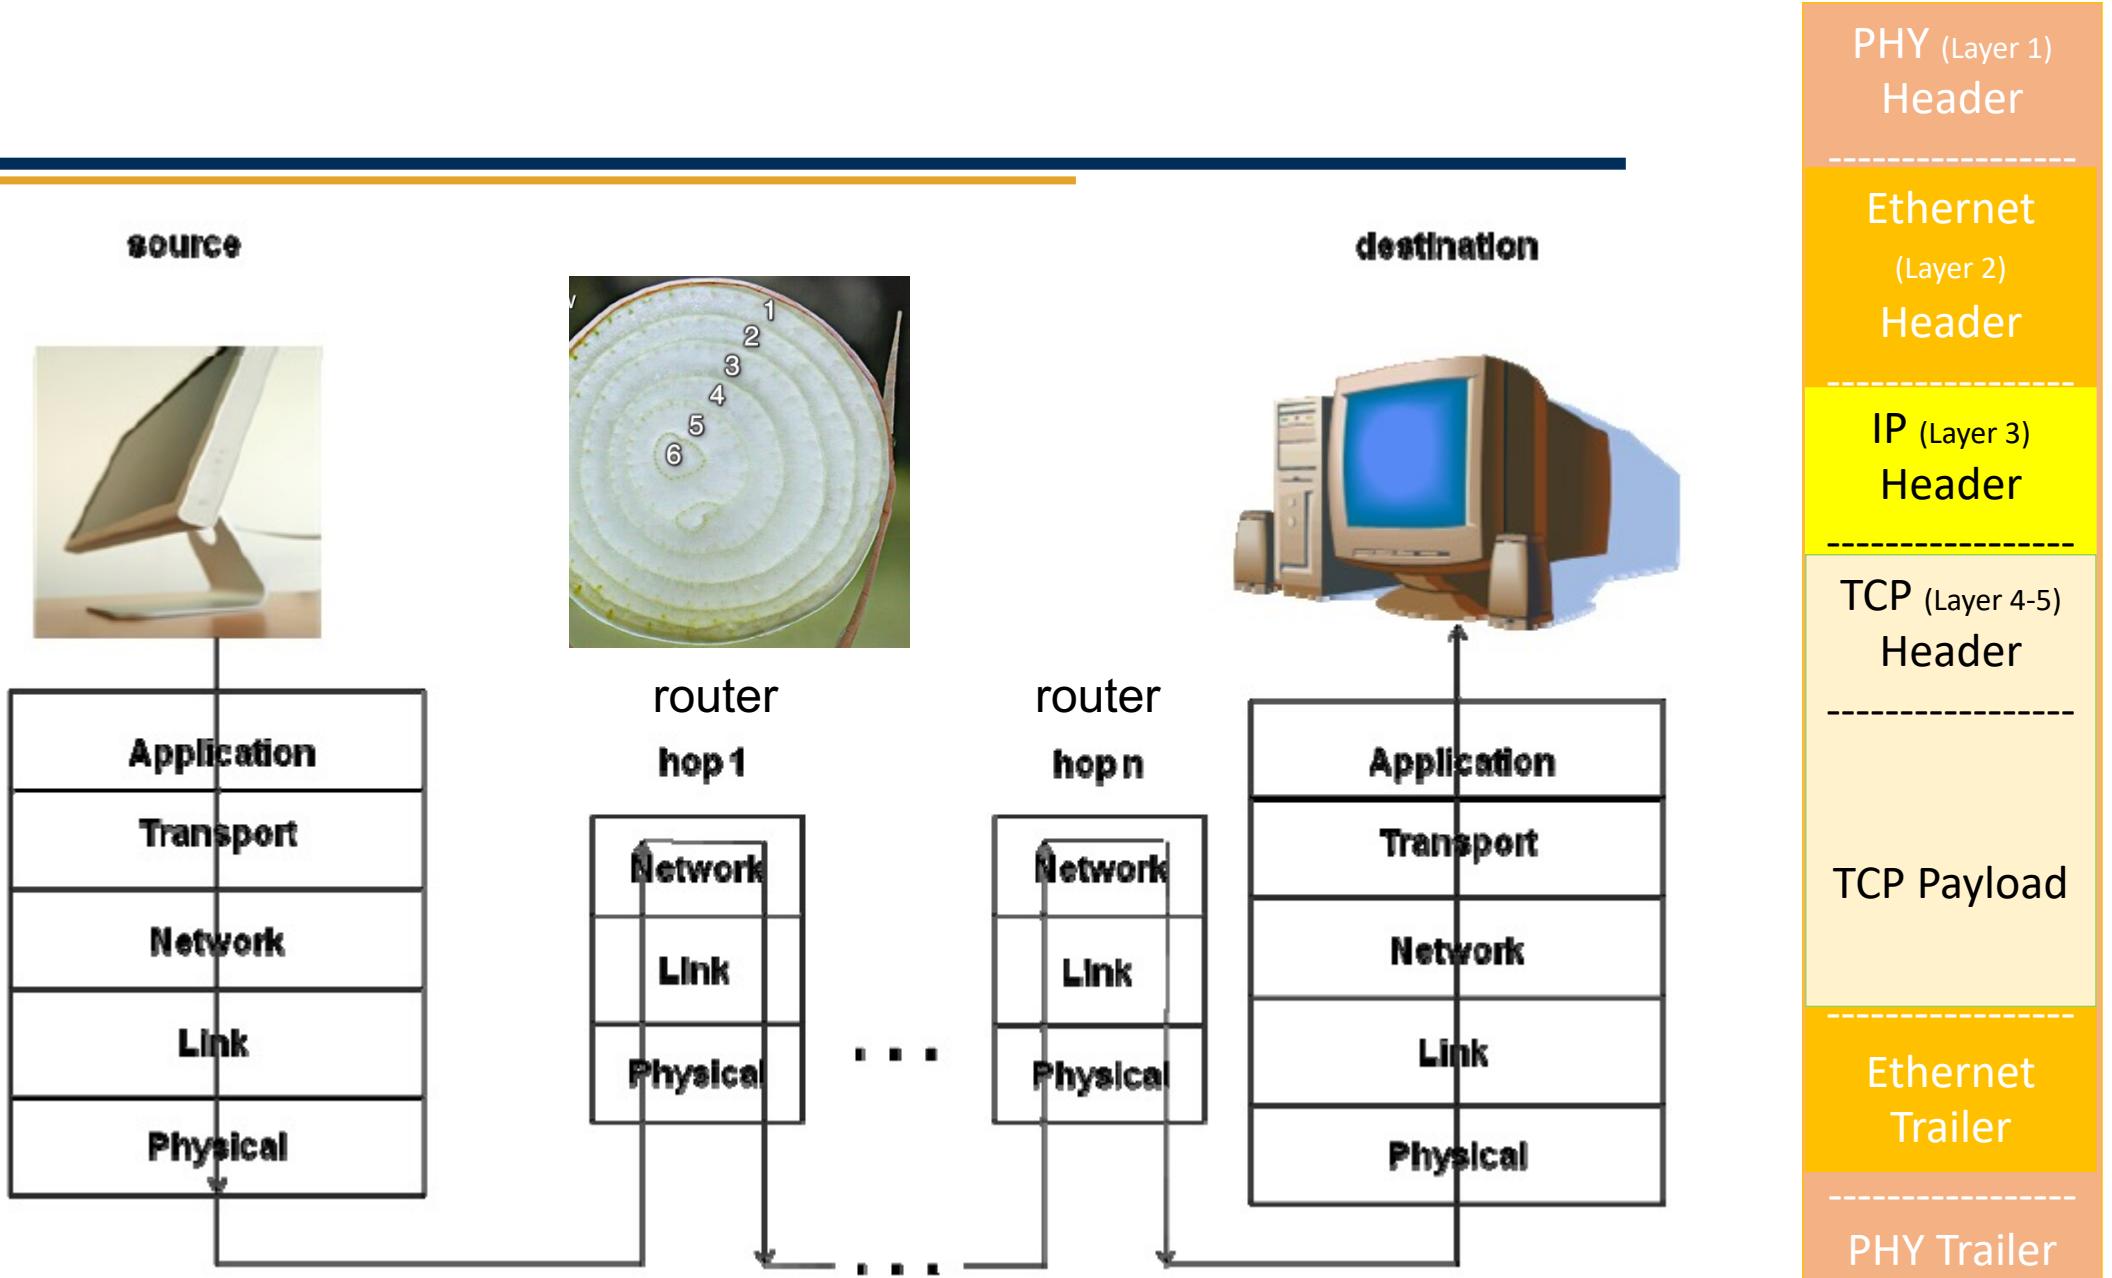
---

- A transport protocol uses a window of 10 packets.
- For a message that consists of 101 packets, how many acknowledgement packets does the destination generate?
- Assume no packets are lost.
- 101

# Back to the network layer

- Why are we back here?
- Recall, IP routing happens here
- We didn't answer a question:  
How does a router know where to route (i.e. who sets up the routing table)?





# Network layer

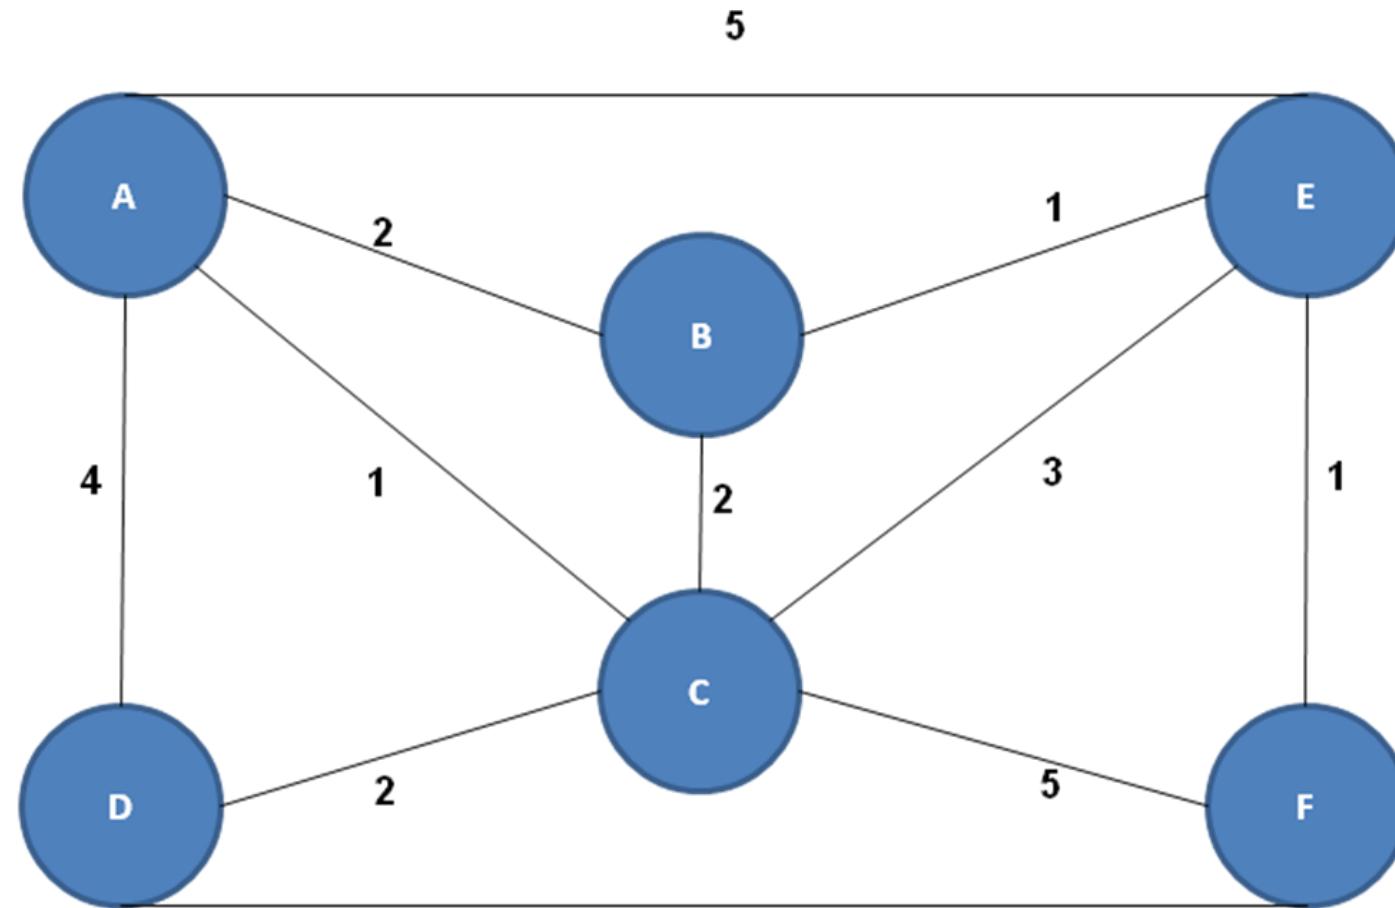
---

- Routing information protocols
  - Link state protocols
  - Distance vector protocols
  - Hierarchical routing
- Addressing

# Intuition behind routing algorithms (link state, distance vector)

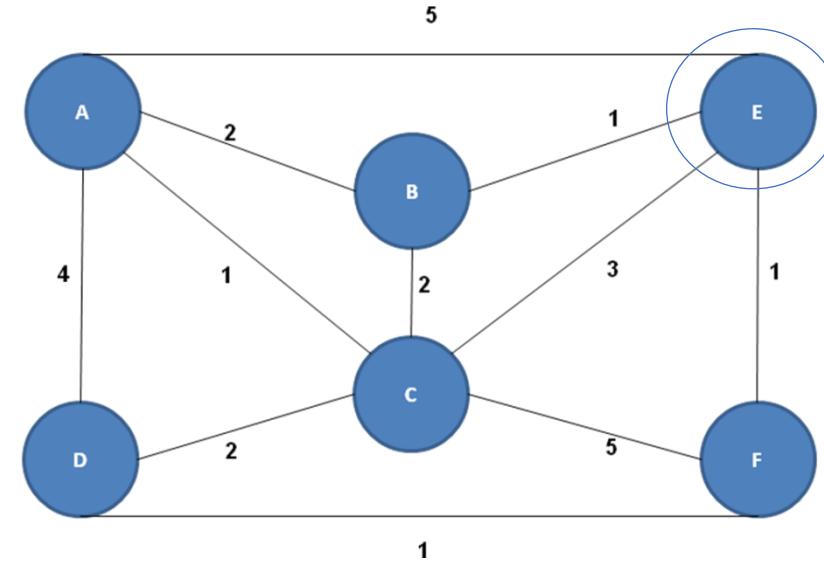
---

What's the latency to get to your neighbor?



# Distance vector in action

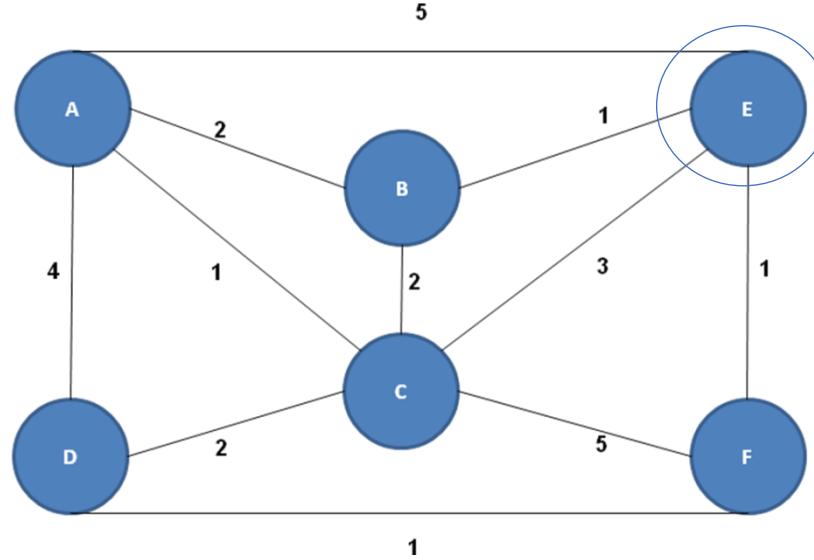
- Uses neighbor info, local algorithm (i.e. every node only sees its neighbors)
- Each node knows the cost to each of its neighbors
- Each node sends its idea of routing to each of its neighbors
- Known colloquially as "routing by rumor"
- For every other node, each node determines the lowest cost next-hop
- Algorithm runs simultaneously on every node
- How does it work?
- Node E knows its costs to its neighbors:



Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	$\infty/$	1/F

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	$\infty$ /	1/F

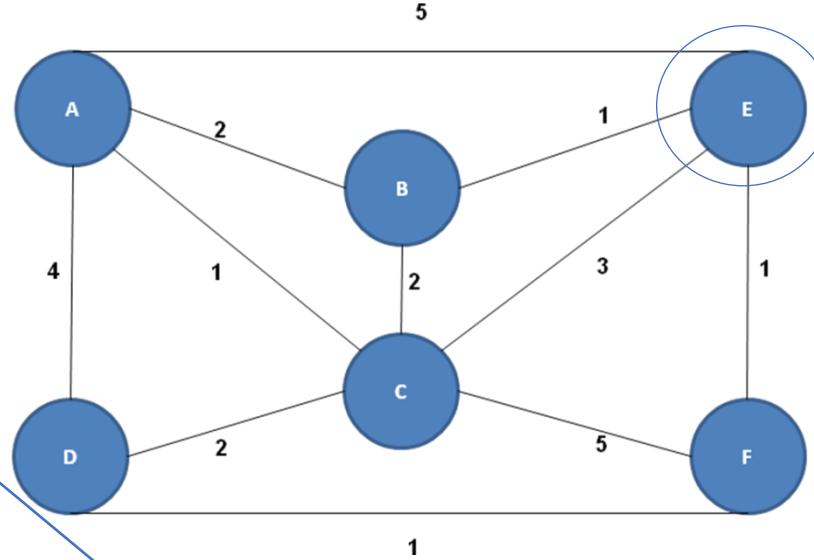


- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

- Reviewing node A's message
  - $5+2/B$ ,  $5+1/C$ ,  $5+4/D$ ,  $5+5/F$
  - Only  $9/D$  is less
  - We'll set the next hop to A

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	9/A	1/F



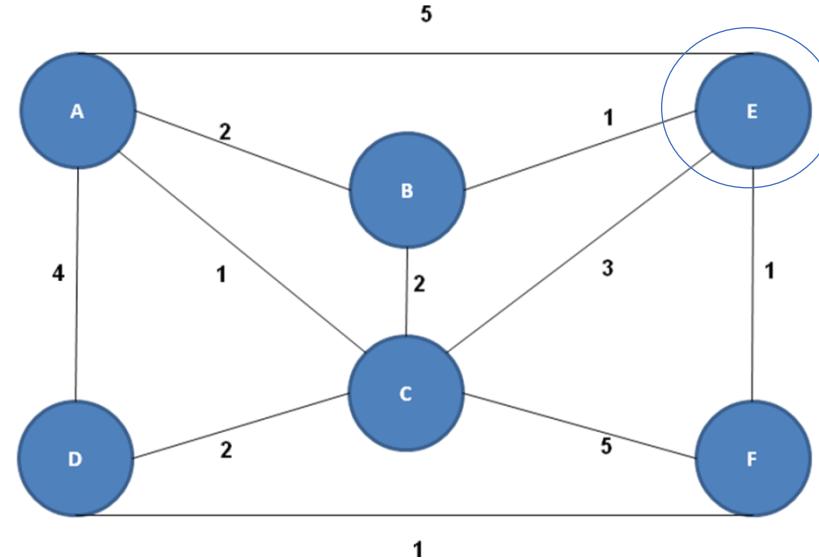
- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

- Reviewing Node B's message
  - 1+2/C, 1+2/A, 1+5/F, 1+4/D
  - 3/A and 5/D are lower cost
  - We set their next hop to B

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

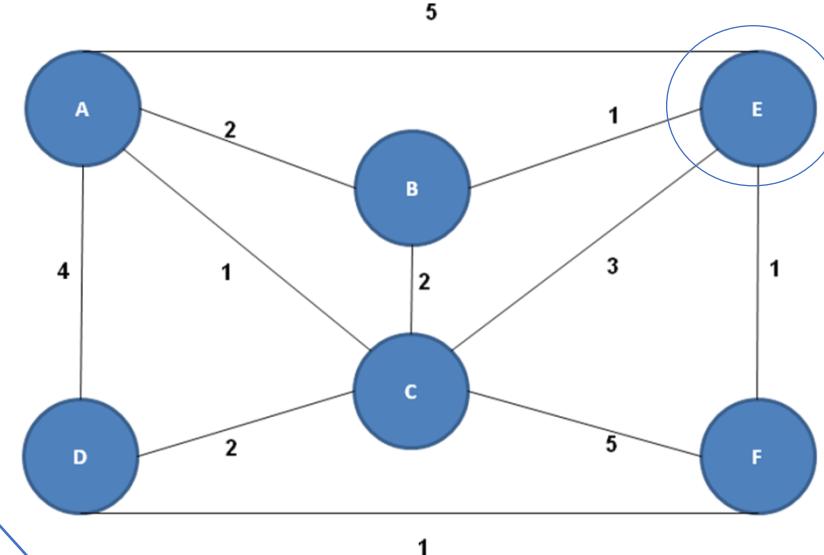


- Reviewing Node C's message
  - $3+1/A$ ,  $3+2/D$ ,  $3+5/F$ ,  $3+2/B$
  - None are lower cost

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

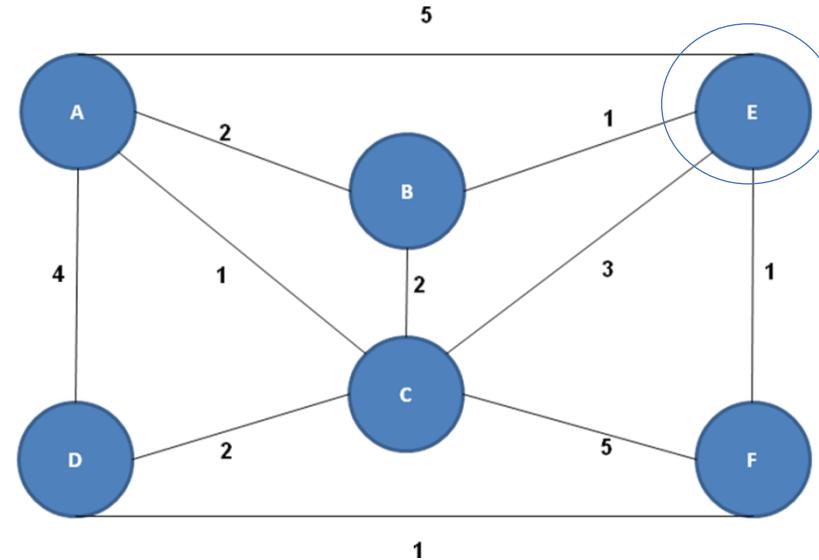


- Reviewing Node F's message
  - $1+5C, 1+1/D, 1+2/B, 1+4/A$
  - $2/D$  is lower cost
  - We'll set the next hop to F

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	2/F	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor



- Do you agree that E now knows the next hop for the lowest cost routes to the other nodes?
- This of course presumed the other nodes knew their lowest cost routes
- When this runs on all nodes simultaneously, they will **converge** on this solution in a relatively short time (convergence depends on diameter of the network)

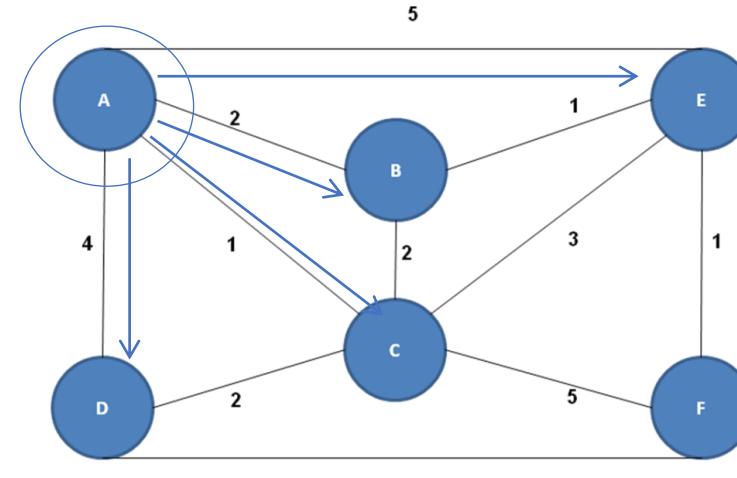
# Distance Vector vs Link State Routing

---

- Distance Vector
  - Each node knows the cost to reach each of its neighbors
  - Each node sends its routing table to only its neighbors
  - Each node revises its own routing table every time it receives a routing table from a neighbor
  - Slower convergence but lower traffic
- Link State
  - Each node advertises its neighbor links and costs to every other node in the network
  - Each node collects the entire topology of the network from these advertisements
  - Using a “shortest path” algorithm, each node calculates its own routing table
  - Faster convergence but more traffic

# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node



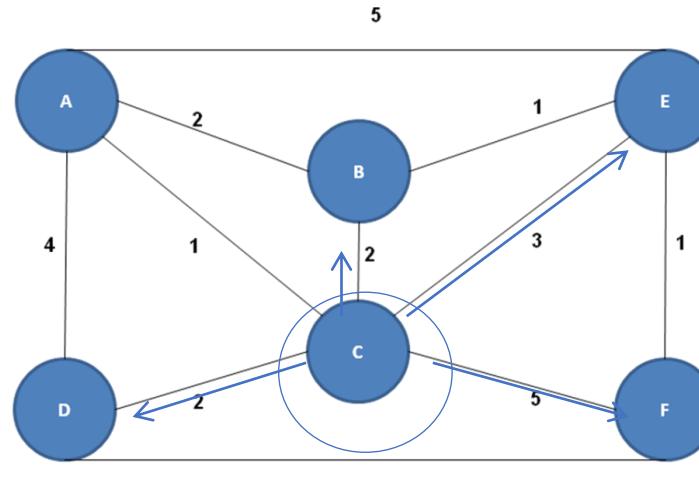
## Calculating on Node A

(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1						
2						
3						
4						
5						

# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node



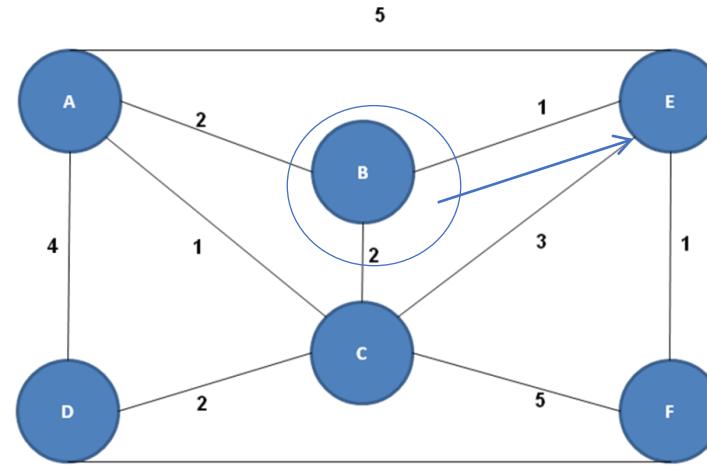
## Calculating on Node A

*(each node does its own calculation)*

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2						
3						
4						
5						

# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node



## Calculating on Node A

(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2	A,C,B	2/AB	-	3/ACD	3/ABE	6/ACF
3						
4						
5						

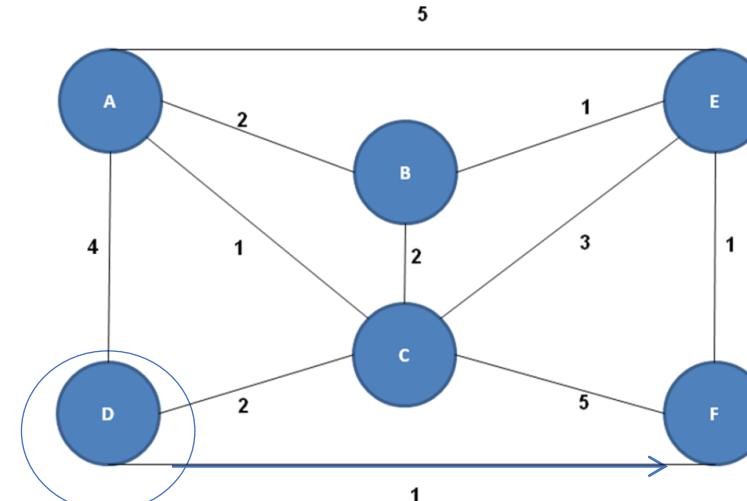
# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node

## Calculating on Node A

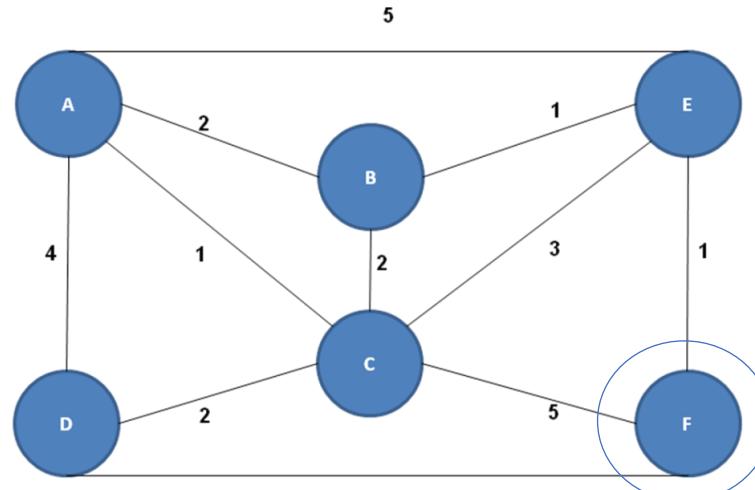
*(each node does its own calculation)*

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2	A,C,B	2/AB	-	3/ACD	3/ABE	6/ACF
3	A,C,B,D	-	-	3/ACD	3/ABE	4/ACDF
4						
5						



# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node



## Calculating on Node A

*(each node does its own calculation)*

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2	A,C,B	2/AB	-	3/ACD	3/ABE	6/ACF
3	A,C,B,D	-	-	3/ACD	3/ABE	4/ACDF
4	A,C,B,D,E	-	-	-	3/ABE	4/ACDF
5	A,C,B,D,E,F	-	-	-	-	4/ACDF

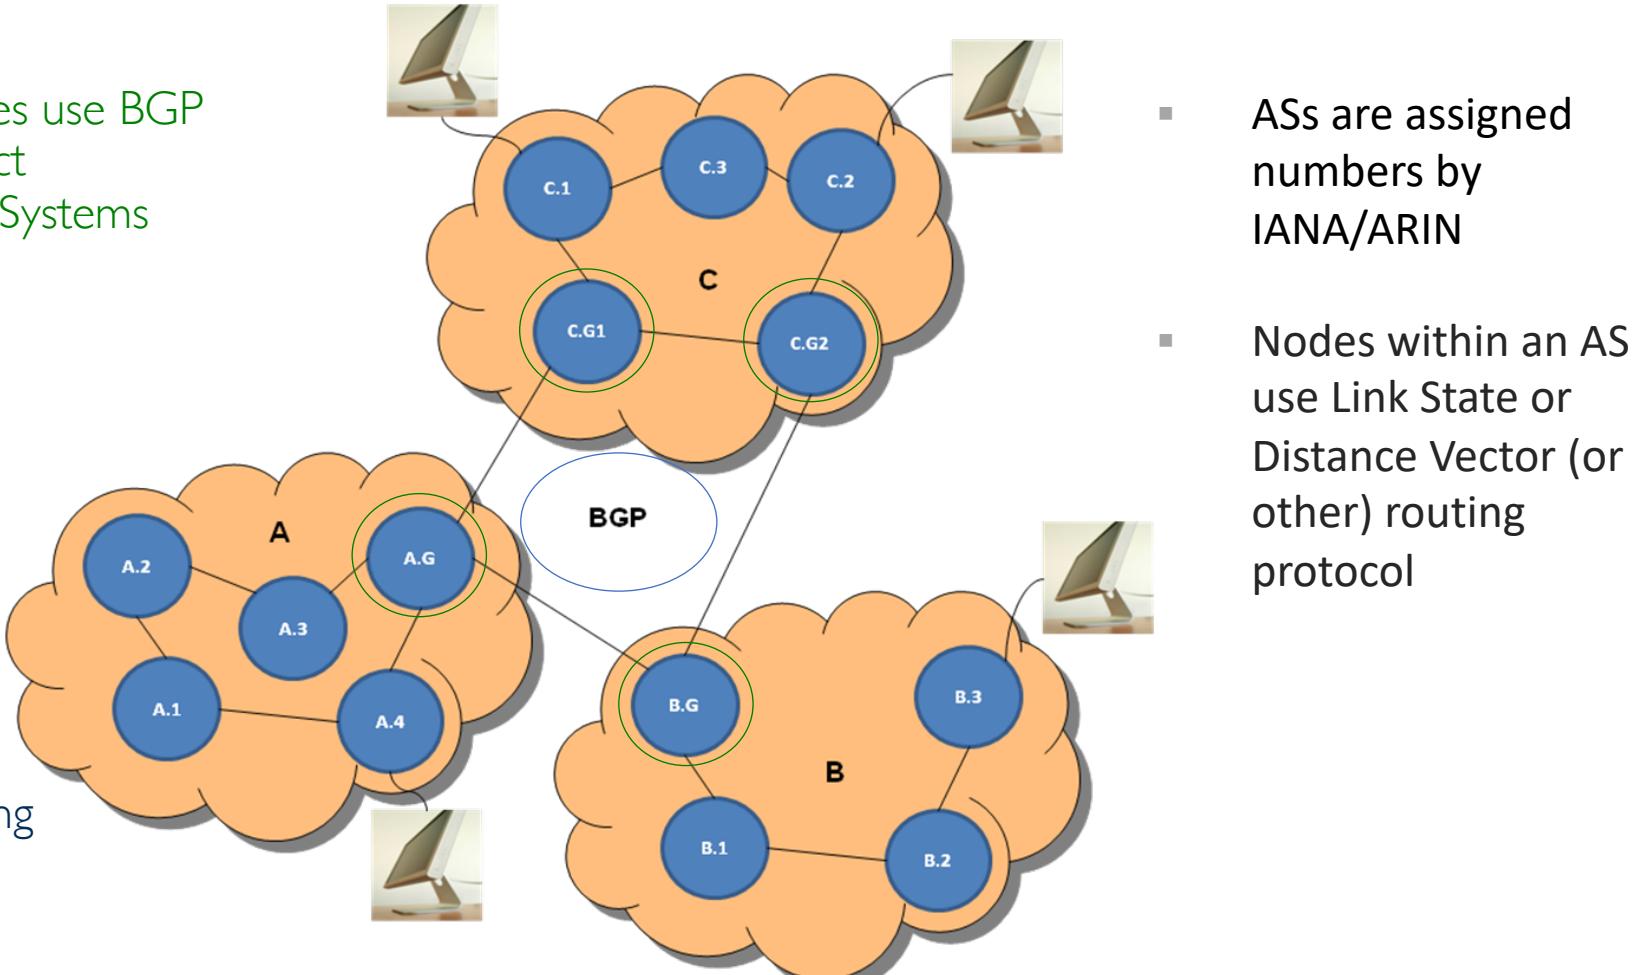
# Hierarchical routing algorithms

---

- Network of networks
- Very large scale, frequent state changes
- Compartmentalize so routing information protocol issues don't spread and cause widespread outages
- What better example than the Internet itself!
- Autonomous Systems (AS)
  - It is a collection of IP routing prefixes under the control of common administrative policy that presents a common, clearly defined routing policy to the internet.
  - Gateway nodes connect to other AS gateway nodes for inter-AS routing

# Hierarchical routing algorithms

- Gateway nodes use BGP to interconnect Autonomous Systems



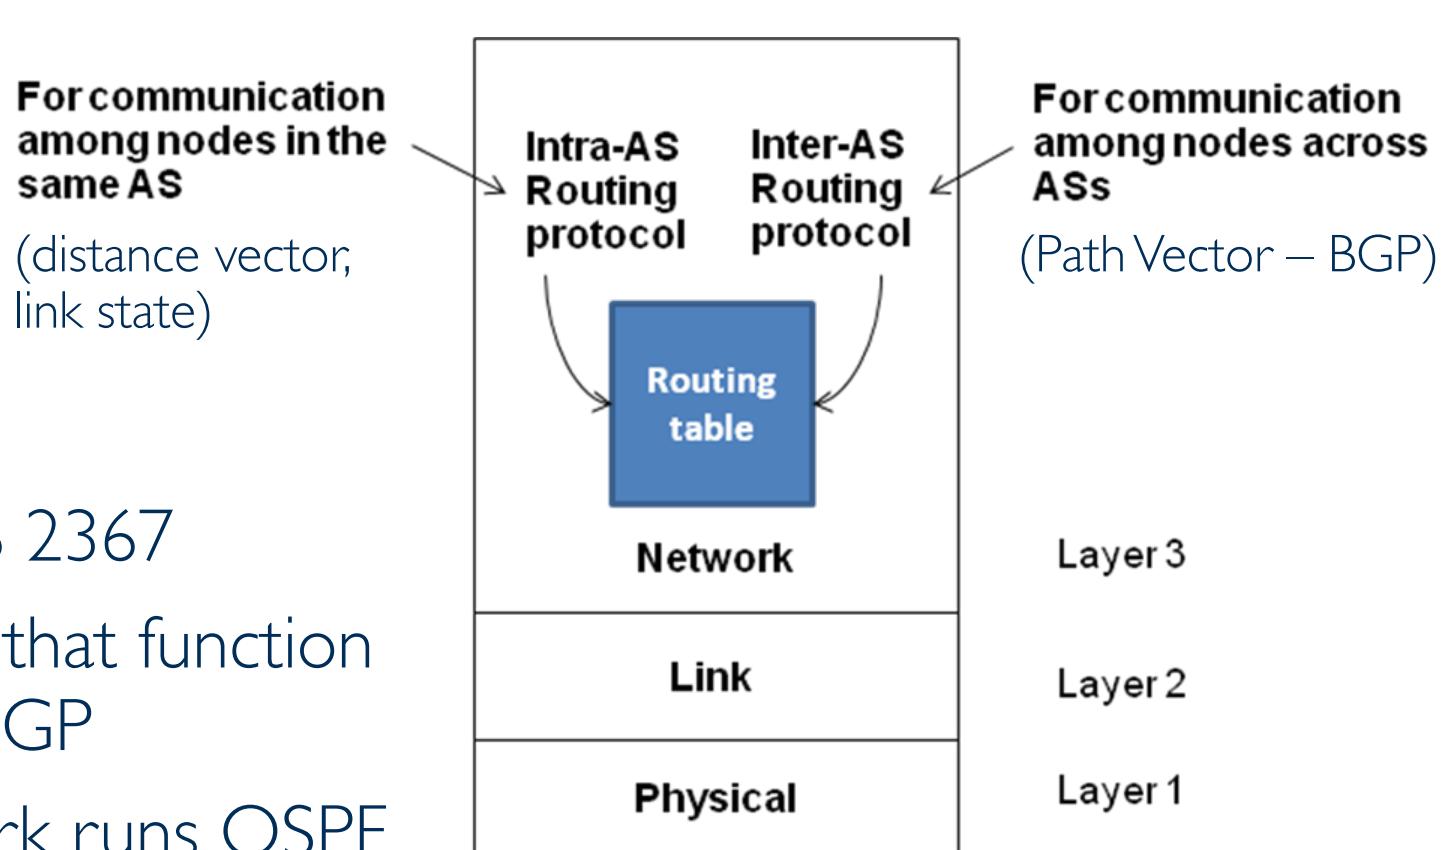
# Border Gateway Protocol (BGP)

---

- BGP is a path-vector protocol. (We're not going into detail about that.)
- The gateway routers send path-vector messages to advertise the reachability of networks.
- Each router that receives a path vector message must verify the advertised path according to its policy.
- If the message is compliant, the router modifies its routing table and the message before sending the message to the next neighbor:
  - It modifies the routing table to maintain the autonomous systems that are traversed in order to reach the destination system.
  - It modifies the message to add its AS number and to replace the next router entry with its identification.

# Routing at GT

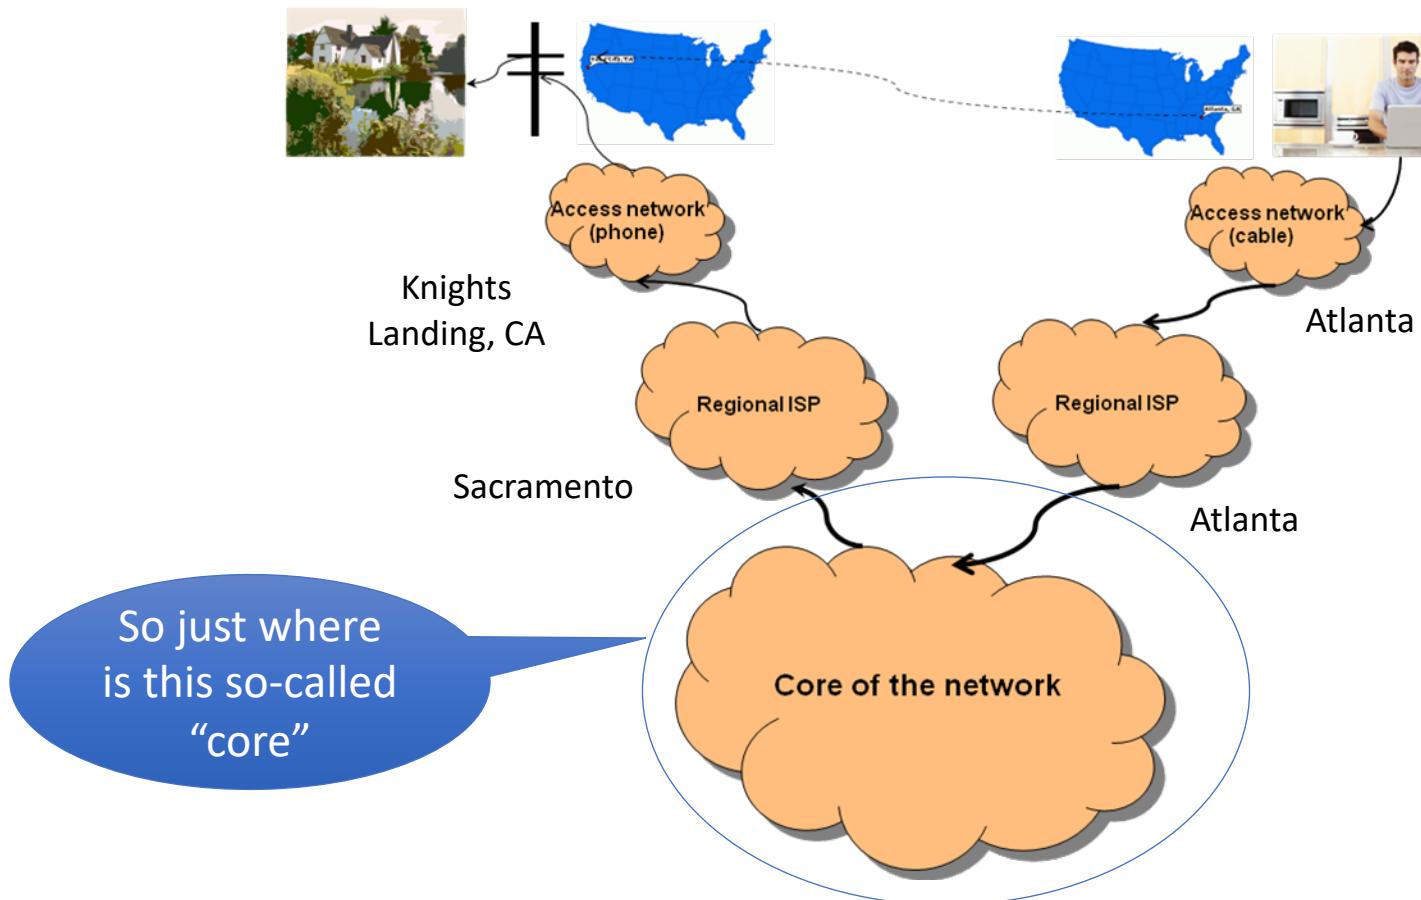
- The GT campus network is AS 2367
- There are two Juniper routers that function as Gateway Routers and run BGP
- The rest of the campus network runs OSPF (a link-state routing protocol)



# Remember this slide?

---

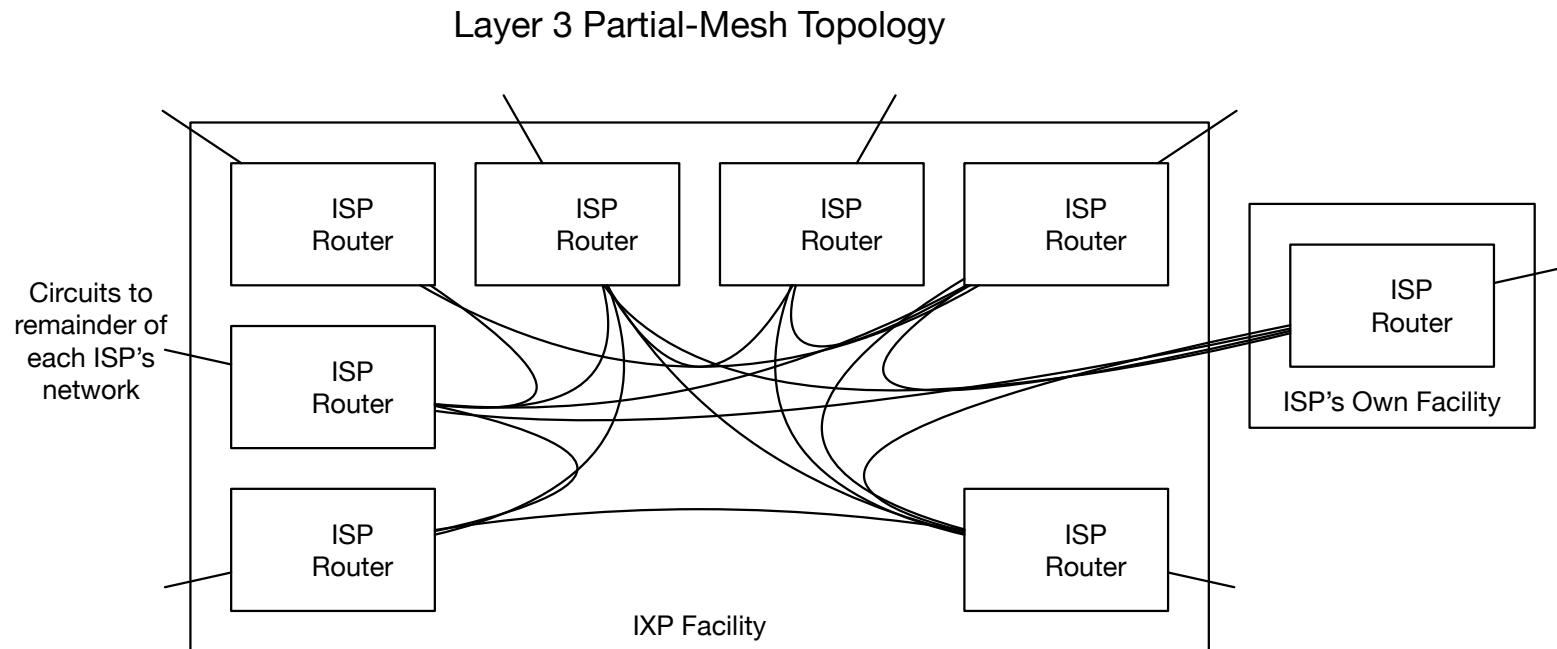
- Now consider an email from Joe to his grandmother





# There really is no internet core ...

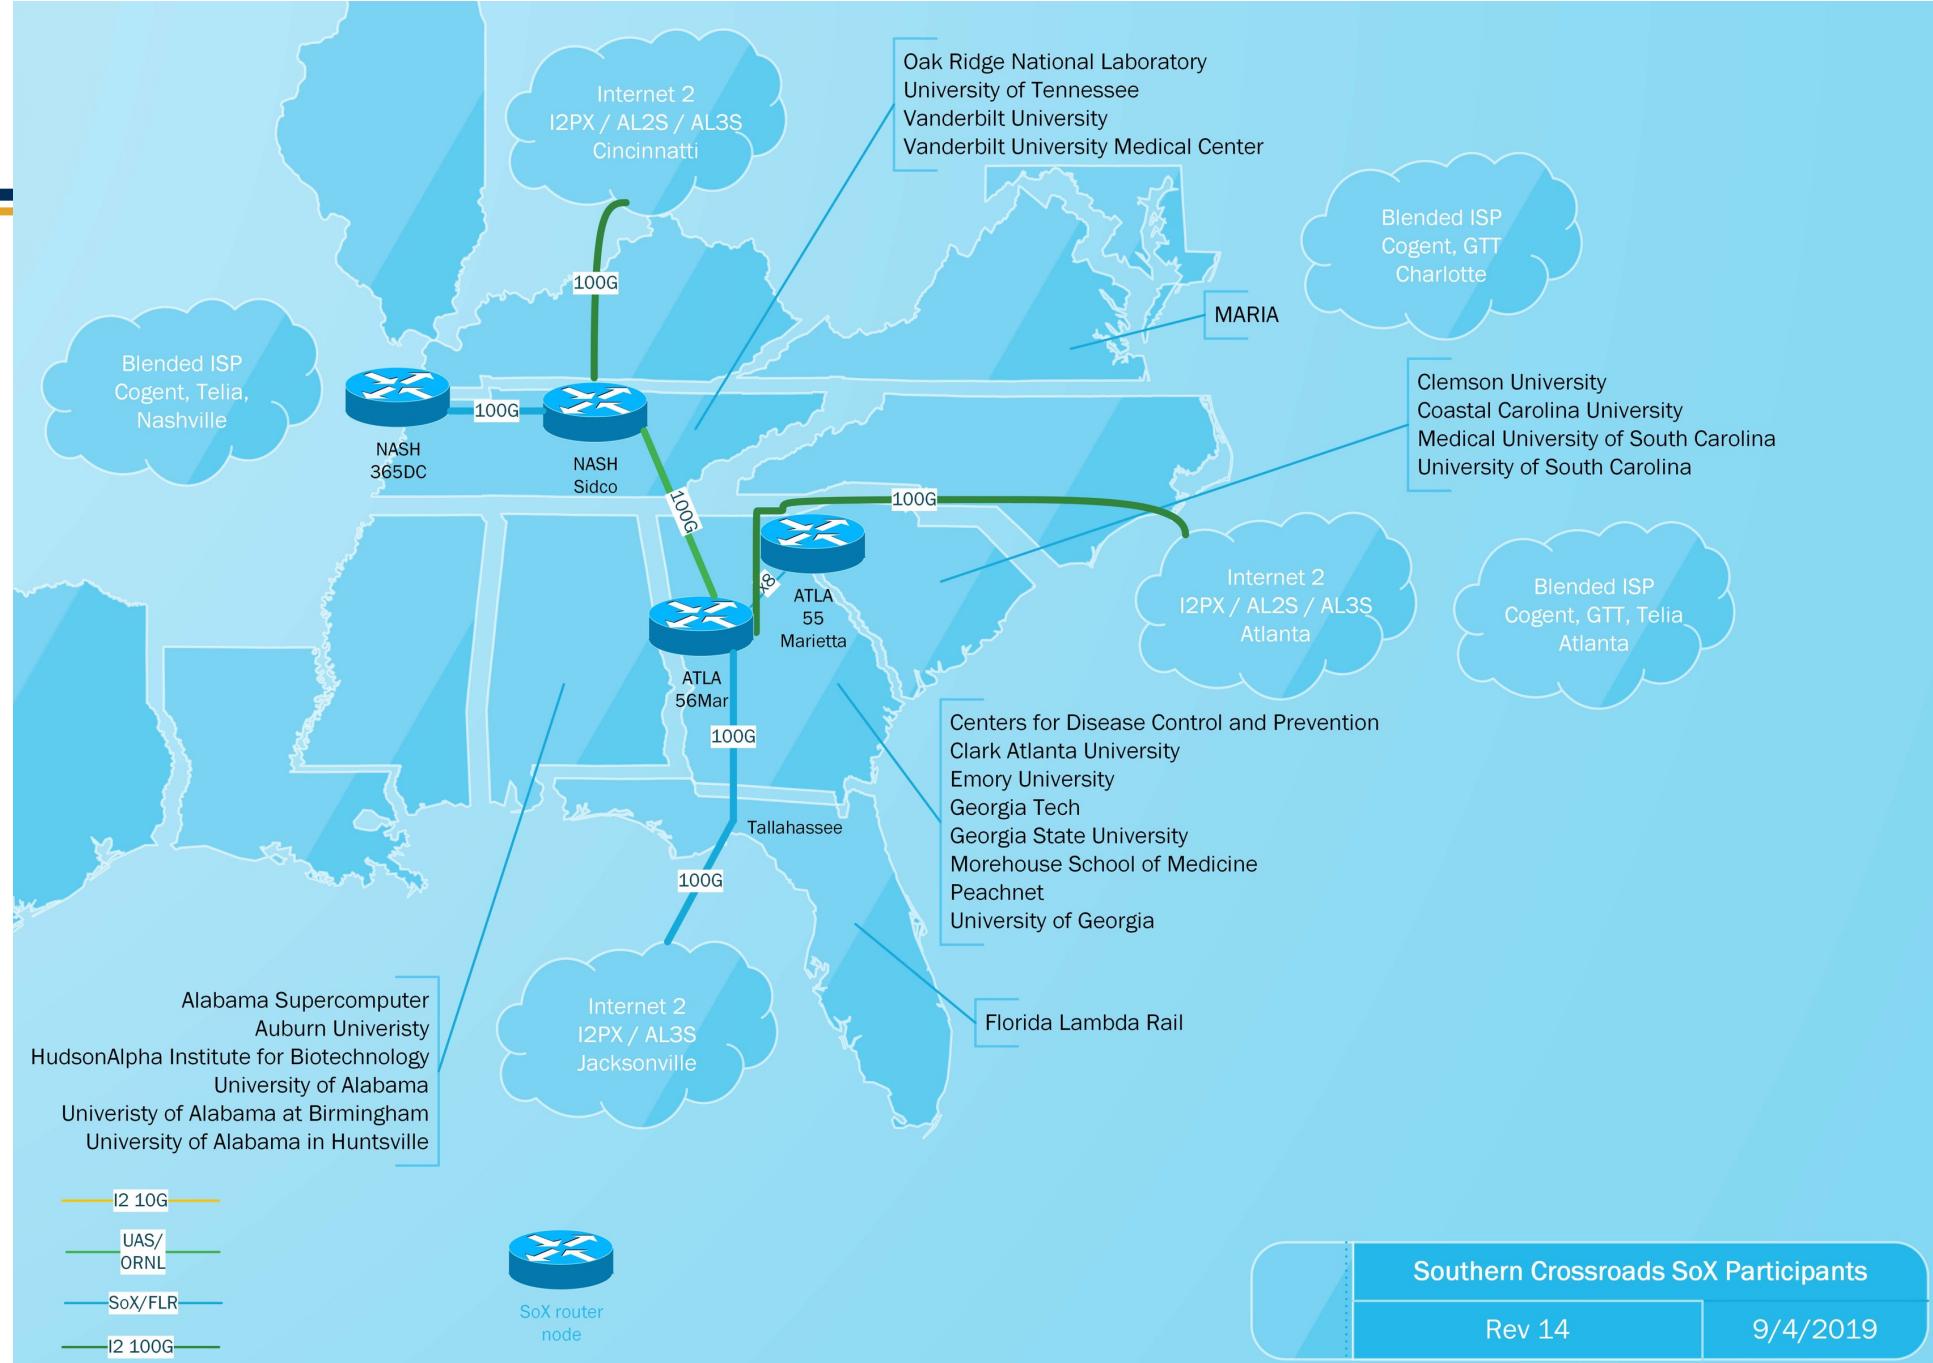
- Instead there are dozens of Internet Exchange Points (IXPs) where groups of Internet Service Providers (ISPs) interconnect their networks



# There really is no internet core ...

---

- The ISPs peer using BGP such that all the gateway routers know multiple routes to every regional ISP
- You can begin to see that the Internet has a highly distributed "core" created by multiple IXPs
- There are at least four IXPs in Atlanta; GT participates in and operates an ISP, SoX, which provides network access for about 3 dozen southeastern universities and research centers
- The diagram on the next slide shows just the IXPs in which SoX participates. Note that traffic is allowed to traverse SoX networks to get from one ISP to another.
- The diagram doesn't show the SoX peering with other entities such as Google, AT&T, NASA, Dept of Energy, and Comcast



# Network Layer Summary

---

Network Terminology	Definition/Use
<b>Circuit switching</b>	A network layer technology used in telephony. Reserves the network resources (link bandwidth in all the links from source to destination) for the duration of the call; no queuing or store-and-forward delays
<b>Packet switching</b>	A network layer technology used in wide area Internet. It supports best effort delivery of packets from source to destination without reserving any network resources en route.
<b>Message switching</b>	Similar to packet switching but at the granularity of the whole message (at the transport level) instead of packets.
<b>Switch/Router</b>	A device that supports the network layer functionality. It may simply be a computer with a number of network interfaces and adequate memory to serve as input and output buffers.
<b>Input buffers</b>	These are buffers associated with each input link to a switch for assembling incoming packets.
<b>Output buffers</b>	These are buffers associated with each outgoing link from a switch if in case the link is busy.
<b>Routing table</b>	This is table that gives the next hop to be used by this switch for an incoming packet based on the destination address. The initial contents of the table as well as periodic updates are a result of routing algorithms in use by the network layer.

# Network Layer Summary

---

Network Terminology	Definition/Use
<b>Delays</b>	The delays experienced by packets in a packet-switched network
<b>Store and forward</b>	This delay is due to the waiting time for the packet to be fully formed in the input buffer before the switch can act on it.
<b>Queuing</b>	This delay accounts for the waiting time experienced by a packet on either the input or the output buffer before it is finally sent out on an outgoing link.
<b>Packet loss</b>	This is due to the switch having to drop a packet due to either the input or the output buffer being full and is indicative of traffic congestion on specific routes of the network.
<b>Service Model</b>	This is the contract between the network layer and the upper layers of the protocol stack. Both the datagram and virtual circuit models used in packet-switched networks provide <b>best effort delivery</b> of packets.
<b>Virtual Circuit (VC)</b>	This model sets up a virtual circuit between the source and destination so that individual packets may simply use this number instead of the destination address. This also helps to simplify the routing decision a switch has to make on an incoming packet. (ATM and X.25)
<b>Datagram</b>	This model does not need any call setup or tear down. Each packet is independent of the others and the switch provides a best effort service model to deliver it to the ultimate destination using information in its routing table. (IP)

# Summary

---

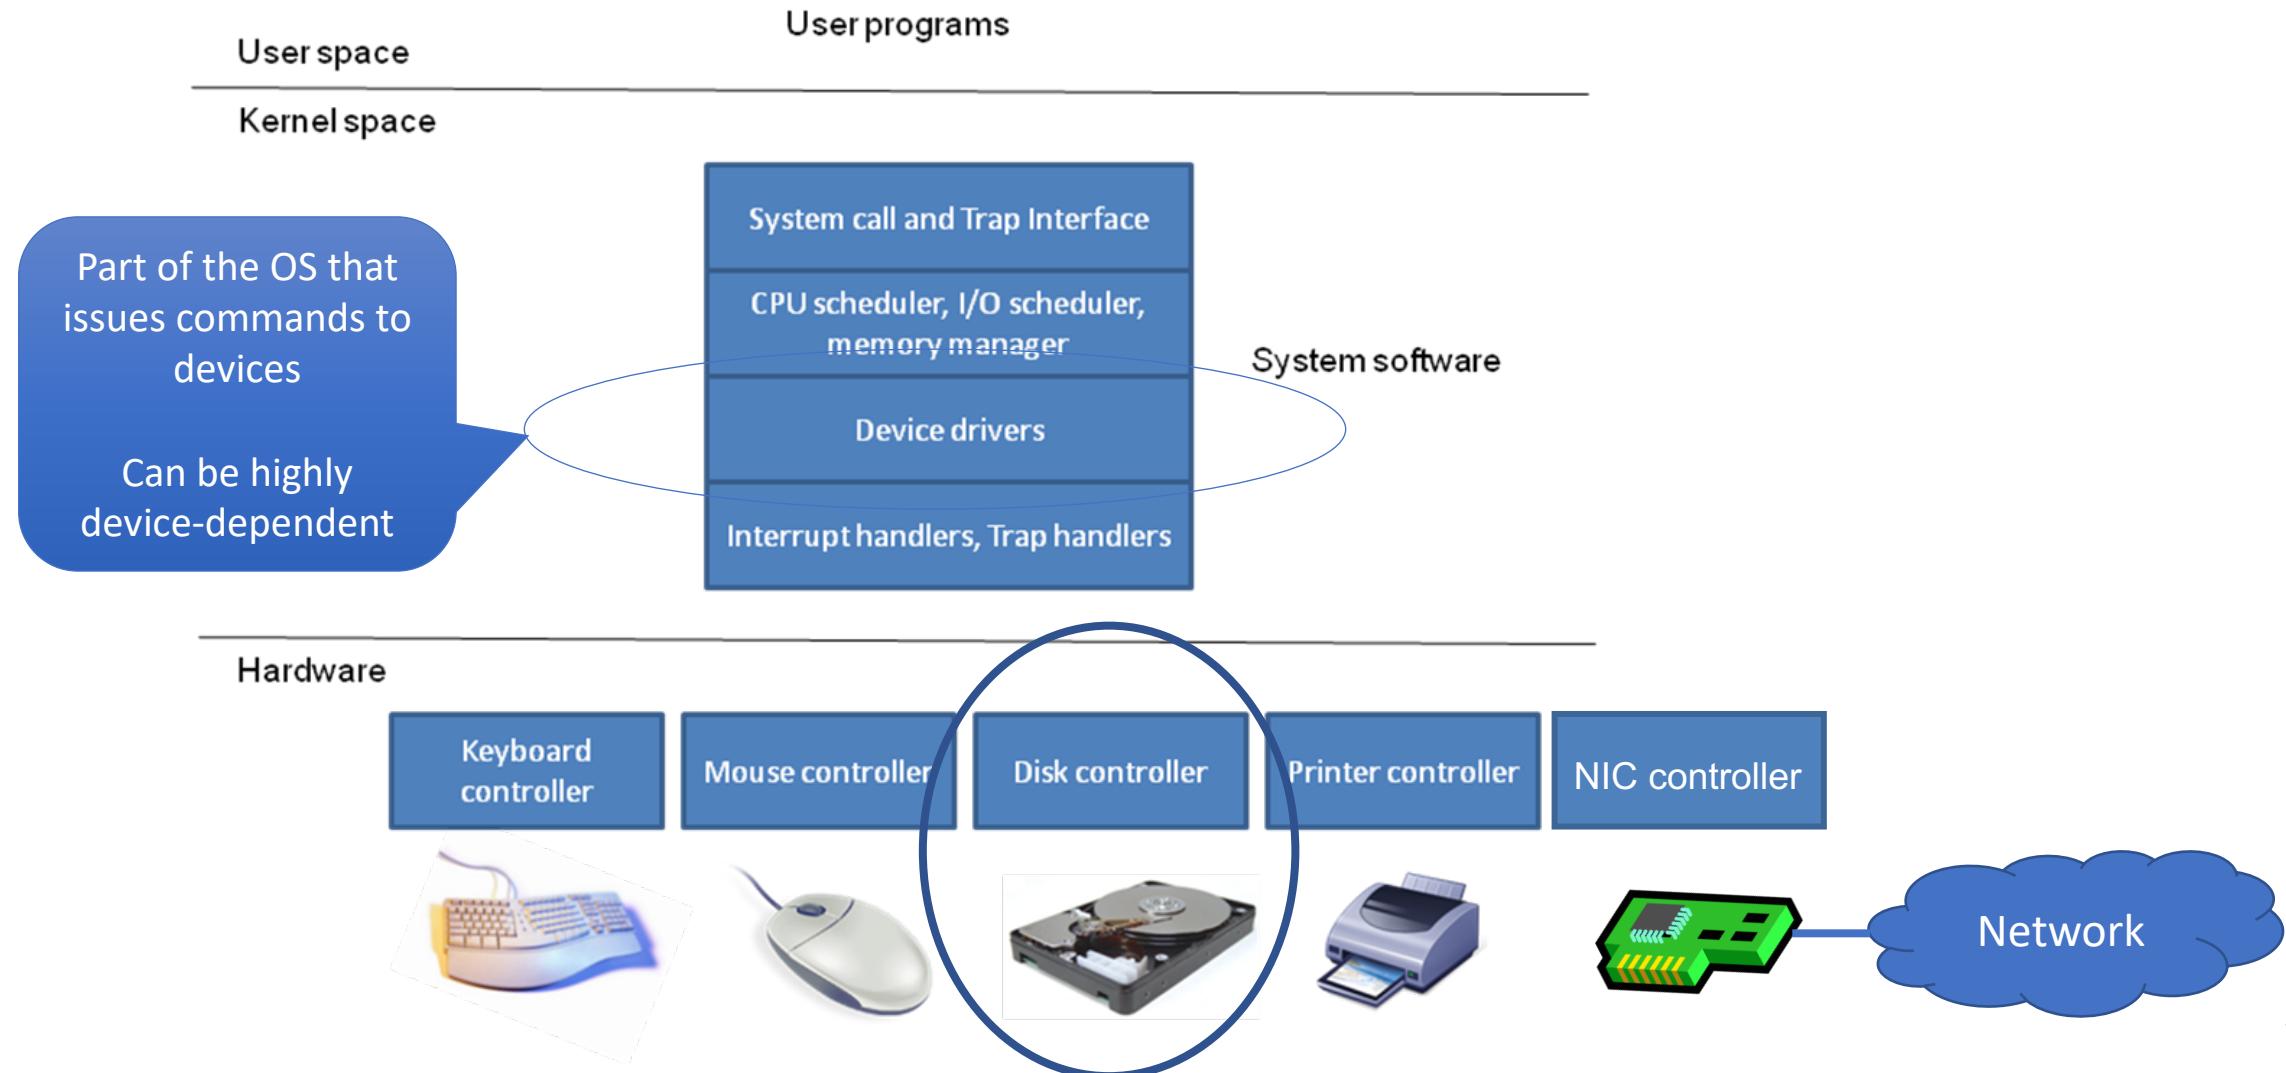
		Data	Layer
Sockets, http	7	Data	Application Network Process to Application
RPC	6	Data	Presentation Data Representation and Encryption
TCP	5	Data	Session Interhost Communication
TCP	4	Segments	Transport End-to-End Connections and Reliability
IP	3	Packets	Network Path Determination and IP (Logical Addressing)
Ethernet	2	Frames	Data Link MAC and LLC (Physical addressing)
Wired, wireless hardware	1	Bits	Physical Media, Signal, and Binary Transmission

# Agenda

---

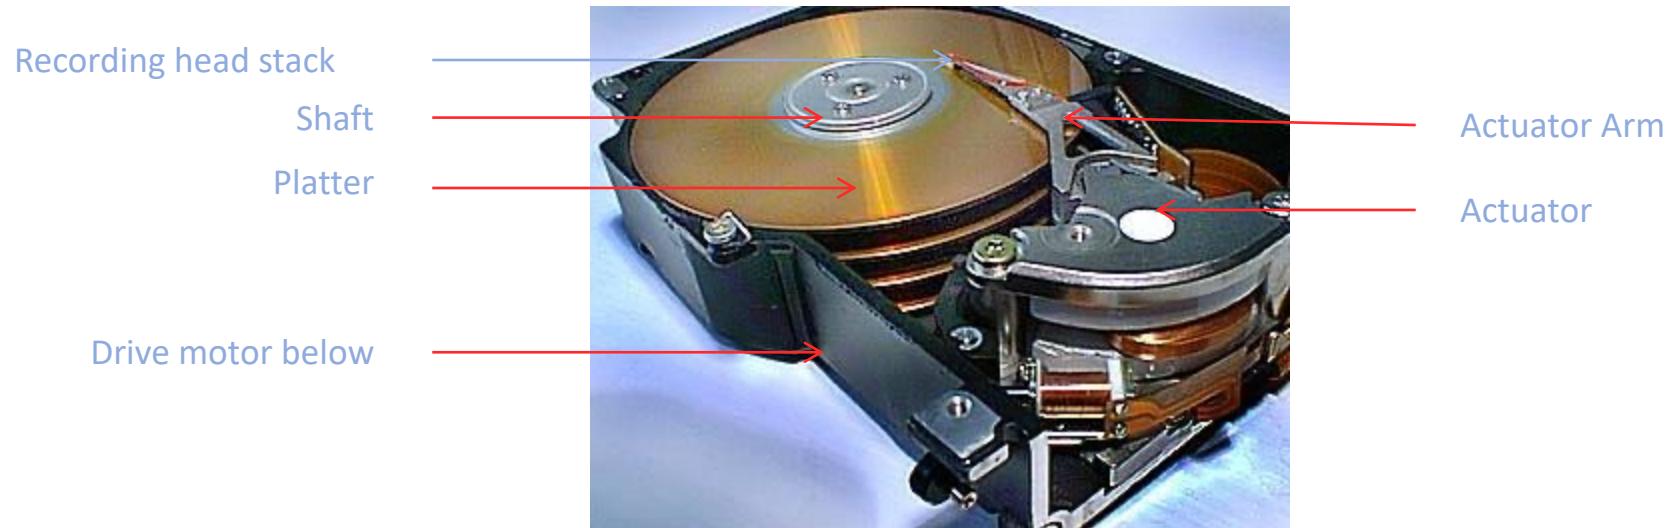
- ~~Wrap up Networking~~
- Disk Scheduling

# Device drivers and OS



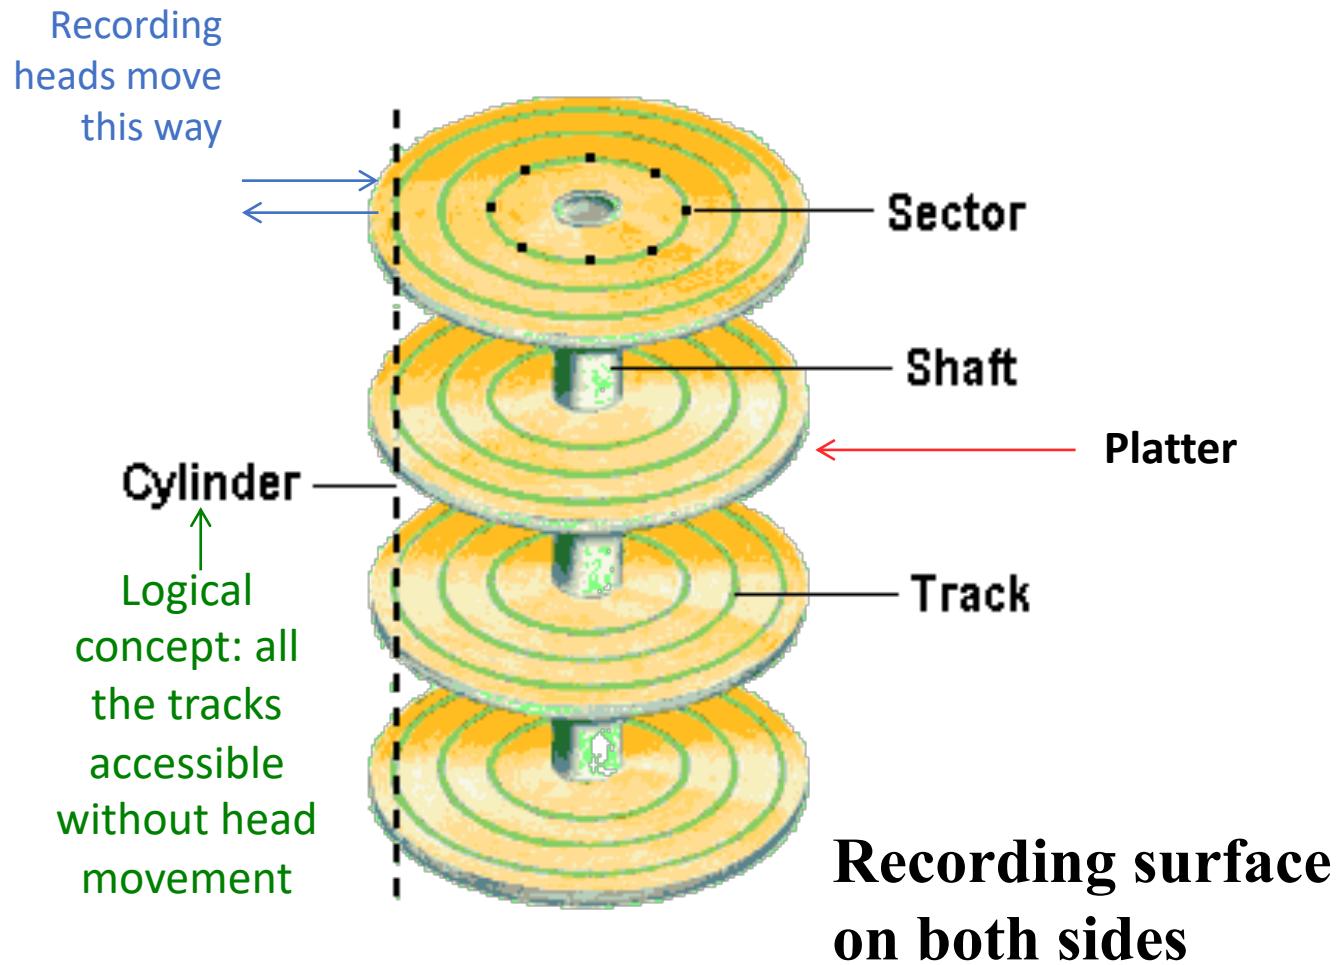
# A modern disk drive

---



# More disk drive terminology

---



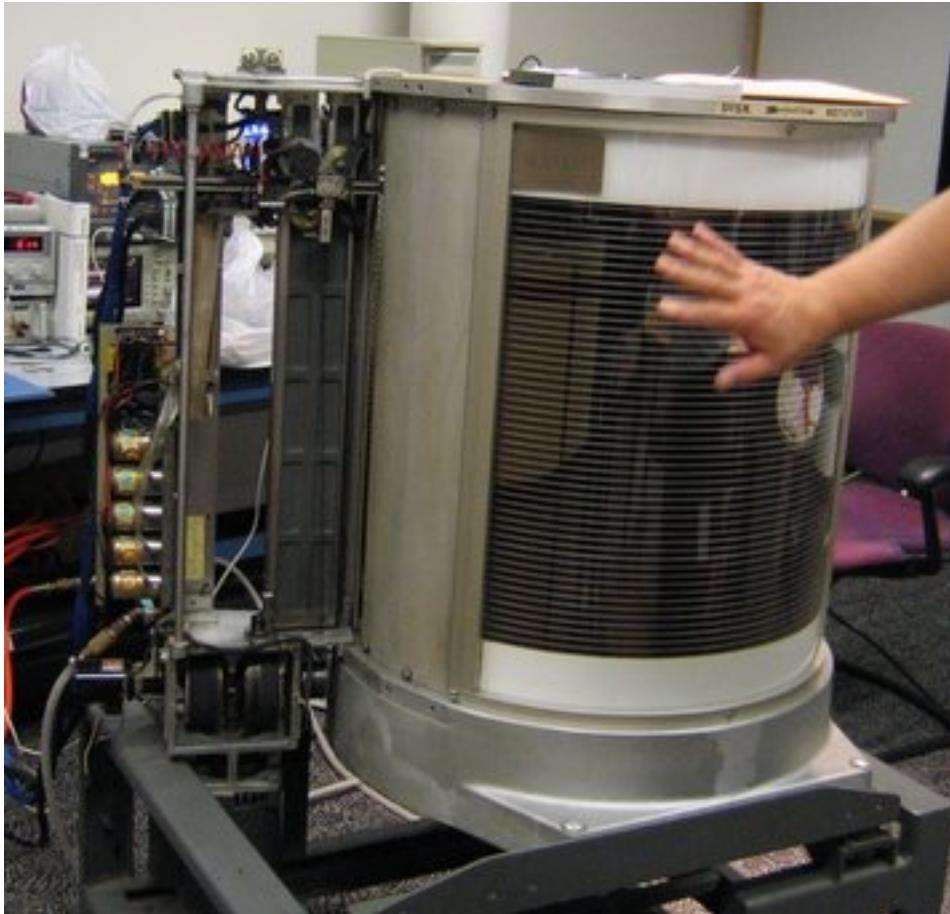
# How a modern disk drive works

---

- Cool 7-min video by Seagate engineer:  
<https://www.youtube.com/watch?v=NtPc0jl2li0>

# The first magnetic disk drive

---

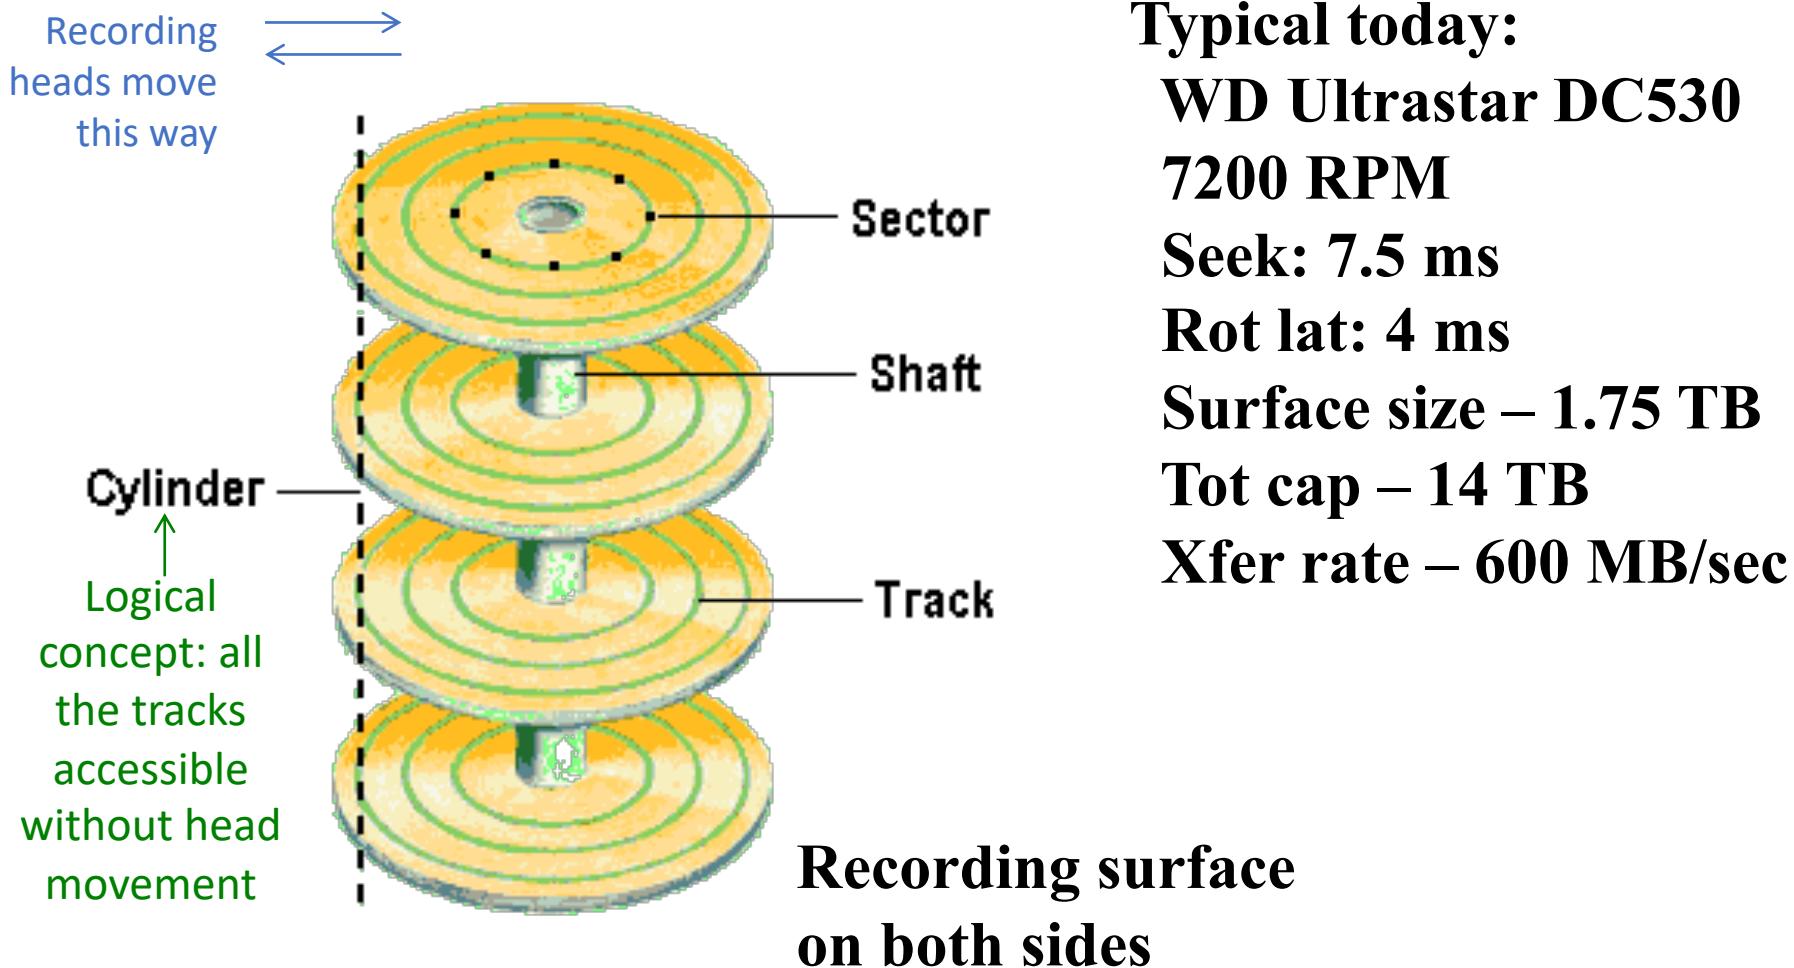


- Announced: 1956
- 50 platters
- 3.75 MB
- 1200 rpm
- 8,800 char/sec
- Weight: 2000 lbs.

IBM 305 RAMAC at the Computer History Museum, Mountain View, CA

<https://www.youtube.com/watch?v=knFRJP3s8WQ>

# More disk drive terminology



# Based on those stats...

---

- How long does it take to find a random block on the disk?
- seek-time + rotational-latency / 2
- In this case,  $7.5\text{ms} + \underline{2\text{ms}} = 9.5\text{ms}$
- Roughly how many random reads can you do per second?
- $1000\text{ms} / 9.5\text{ ms} = \sim 100 \text{ reads(!)}$
- And the data transfer rate if blocks are 4K bytes?
- $100 * 4096 = \underline{409 \text{ KB/sec}}$
- Very different from the maximum transfer rate, right?

**Typical today:**  
**WD Ultrastar DC530**  
**7200 RPM**  
**Seek: 7.5 ms**  
**Rot lat: 4 ms**  
**Surface size – 1.75 TB**  
**Tot cap – 14 TB**  
**Xfer rate – 600 MB/sec**

# Capacity Metrics

---

Let,

$p$  – number of platters,

$n$  – number of surfaces per platter (1 or 2),

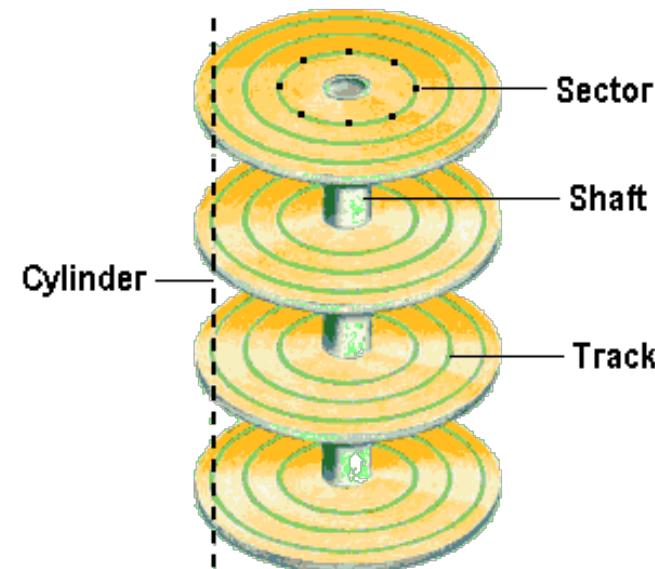
$t$  – number of tracks per surface,

$s$  – number of sectors per track,

$b$  – number of bytes per sector,

The total capacity of the disk:

$$\text{Capacity} = (p * n * t * s * b) \text{ bytes}$$



# Example

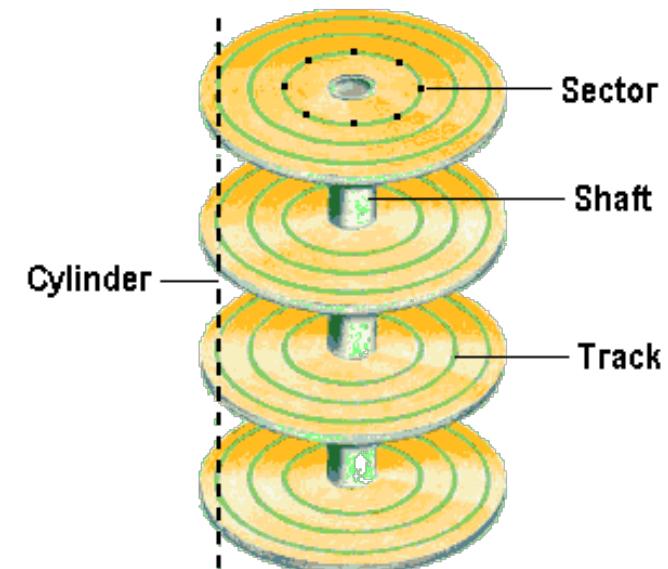
Suppose a disk drive has  $512$  bytes per sector,  $12$  sectors per track,  $100$  tracks per surface,  $2$  surfaces per platter, and  $6$  platters.

What is the total capacity of such a drive in bytes?

How many cylinders does the device have?

$$\begin{aligned}\text{Capacity} &= (p * n * t * s * b) \text{ bytes} \\ &= 6 * 2 * 100 * 12 * 512 \text{ bytes} \\ &= 7200 \text{ KB } (\text{KB} = 2^{10} \text{ bytes})\end{aligned}$$

And it has  $t$  cylinders, or  $100$



# Transfer metrics

---

Let,

$r$  – rotational speed of the disk in RPM,

$s$  – number of sectors per track,

$b$  – number of bytes per sector,

Time for one revolution =  $60/r$  seconds

Amount of data read in one revolution =  $s * b$  bytes

The data transfer rate of the disk:

$$\begin{aligned} & (\text{Amount of data in track}) / (\text{time for one revolution}) \\ &= (s * b) / (60/r) \end{aligned}$$

Data transfer rate

$$= (s * b * r) / 60 \text{ bytes/second}$$

# Example

---

Given the following specifications for a disk drive:

- b 512 bytes per sector
- s 400 sectors per track
- 6000 tracks per surface
- 3 platters
- r Rotational speed 15000 RPM

What is the transfer rate of the disk?

$$\begin{aligned}\text{Data transfer rate} &= (s * b * r) / 60 \text{ bytes/second} \\ &= (400 * 512 * 15000) / 60 \\ &= 51,200,000 \text{ bytes/second}\end{aligned}$$

# Which of these attributes

of a hard disk drive is conceptual rather than physical?

14% A. I just want the participation credit

14% B. p – number of platters

14% C. n – number of surfaces per platter (1 or 2)

14% D. t – number of tracks per surface

14% E. s – number of sectors per track

14% F. b – number of bytes per sector

14% G. c – number of cylinders

# Early 60's disk pack: 5MB

---

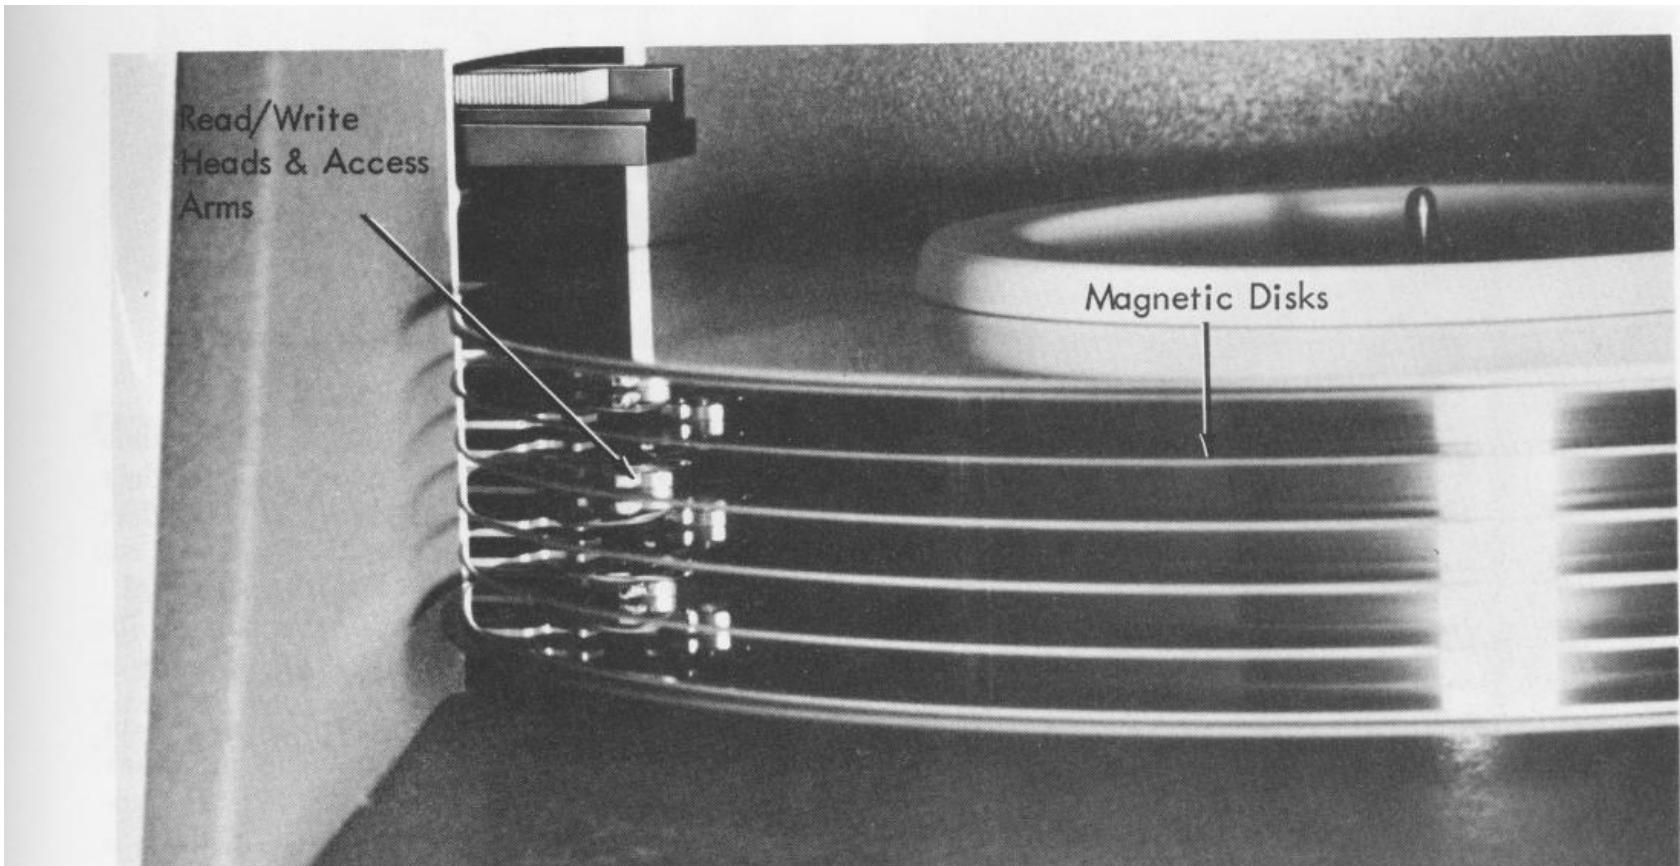


Figure 54. IBM 2311 Disk Storage Drive showing mounted disk pack and access arms

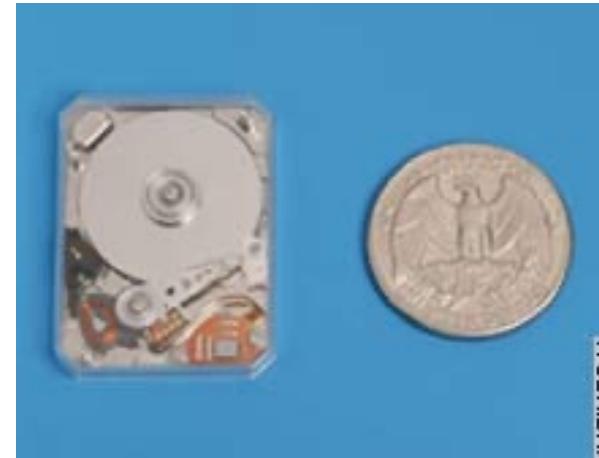
# When computers were big...

---



# Welcome to modern times

---

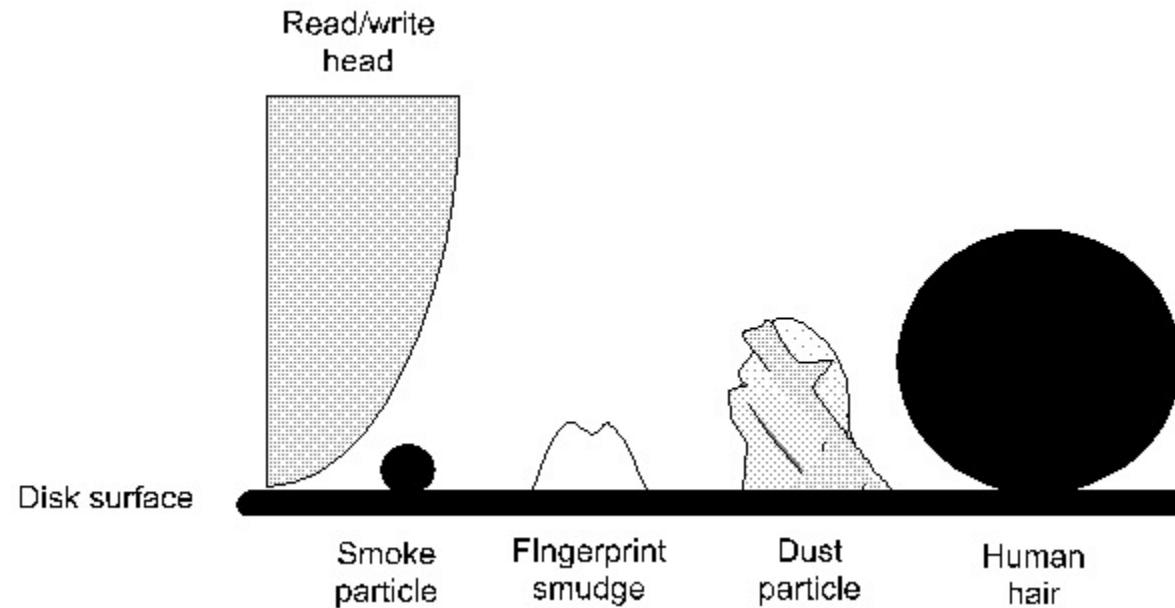


**4 GB drive**

# Why are disk drives sealed?

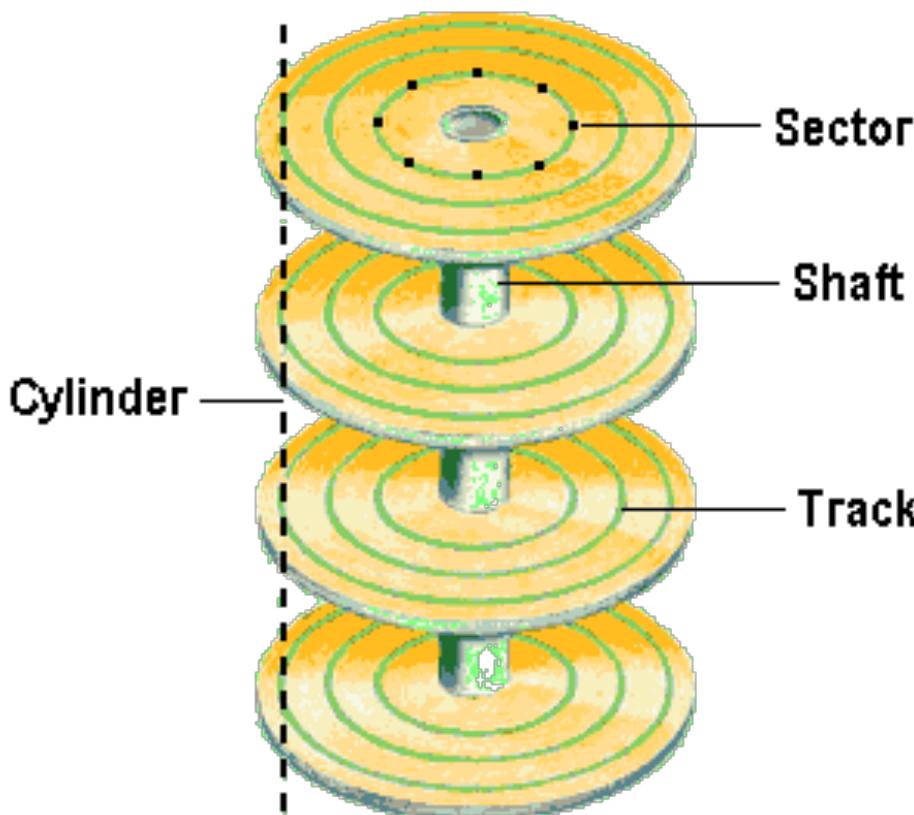
---

Since the early 60s, read/write heads are aerodynamically shaped so they fly on a very thin cushion of air. Obstacles are not well tolerated.

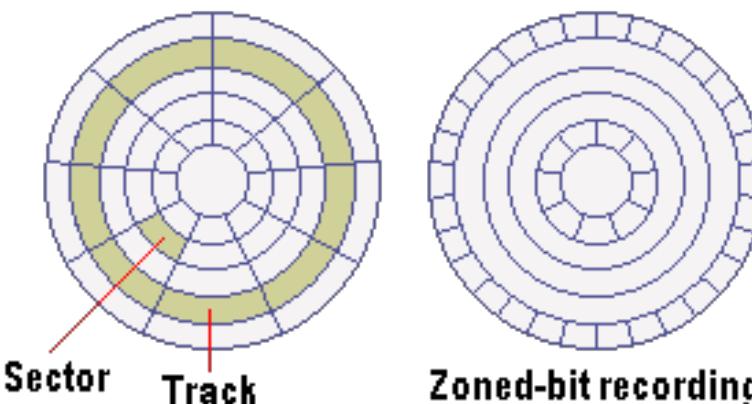


# Zoned recording

---



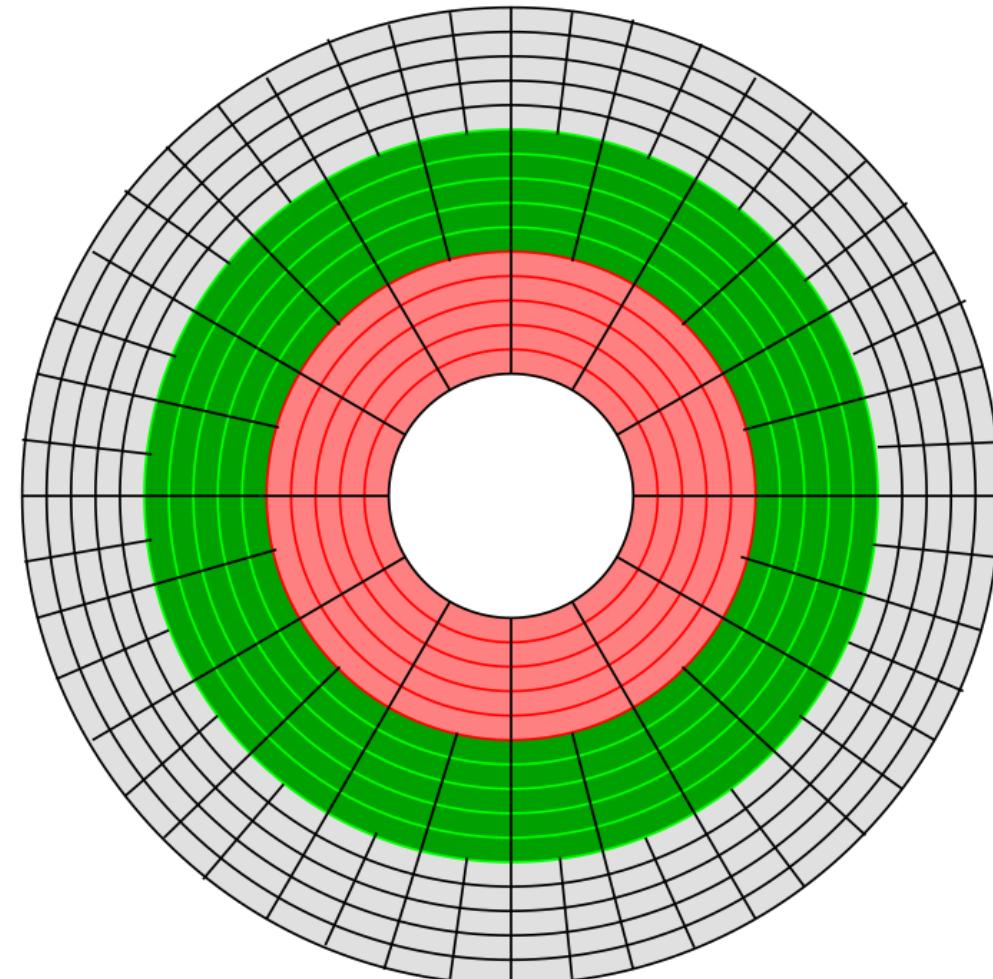
- The outside track of a platter has a larger surface and moves faster than an inside track!
- This allows the outside tracks to hold more data and transfer faster
- Zoned recording (of which there are many schemes) assigns more sectors to the outer tracks



# Example of Zoned Recording

---

- Physical layout of sectors in a zone-bit disc
- As distance from the center increases, the number of sectors in a given angle increases from one (red) to two (green) to four (grey)
- The red zone has 12 sectors/track; green has 24; grey has 48.



---

---

**With zoned recording,**

**$z$  – number of zones,**

**$t_{zi}$  – number of tracks at zone  $z_i$ ,**

**$s_{zi}$  – number of sectors per track at zone  $z_i$ ,**

**The total capacity of the disk with zoned recording:**

**Capacity =  $(p * n * (\sum (t_{zi} * s_{zi}), \text{ for } 1 \leq i \leq z) * b)$  bytes**

# Total access latency

---

Let,

$a$  – average seek time in seconds

$r$  – rotational speed of the disk in RPM

$s$  – number of sectors per track

Rotational latency =  $60/r$  seconds

Average rotational latency =  $(60 / (r * 2))$  seconds

Sector read time = rotational latency / number of sectors per track

Sector read time =  $(60 / (r * s))$  seconds

Time to get to the desired track

= Average seek time

=  $a$  seconds

---

Time to get the head over the desired sector

= Average rotational latency

=  $(60 / (r * 2))$  seconds

Time to read a sector

= Sector read time

=  $(60 / (r * s))$  seconds

Time to read a random sector on the disk

= Time to get to the desired track +

Time to get the head over the desired track +

Time to read a sector

=  $a + (60 / (r * 2)) + (60 / (r * s))$   
seconds

# Example

---

Given the following specifications for a disk drive:

256 bytes per sector

12 sectors per track

20 tracks per surface

3 platters

Average seek time of 20 ms

Rotational speed 3600 RPM

What would be the time to read 6 contiguous sectors from the same track?

What would be the time to read 6 sectors at random?

# Disk scheduling

---

- Ordering of disk requests is quite important for I/O performance just as job scheduling is important for CPU performance
- Seek and rotational delay (on the order of 5ms each) are mechanical and hence huge compared to the speeds of memory (100 ns) and CPU instructions (1 ns). That's  $5 \times 10^{-3}$  compared with  $1 \times 10^{-9}$ , or more than 6 orders of magnitude.
- However, as disk controllers have become more capable, I/O requests are queued for them by the OS as soon as they come in.
- Disk controllers may have tens of pending I/O operations queued and perform their own scheduling on them.

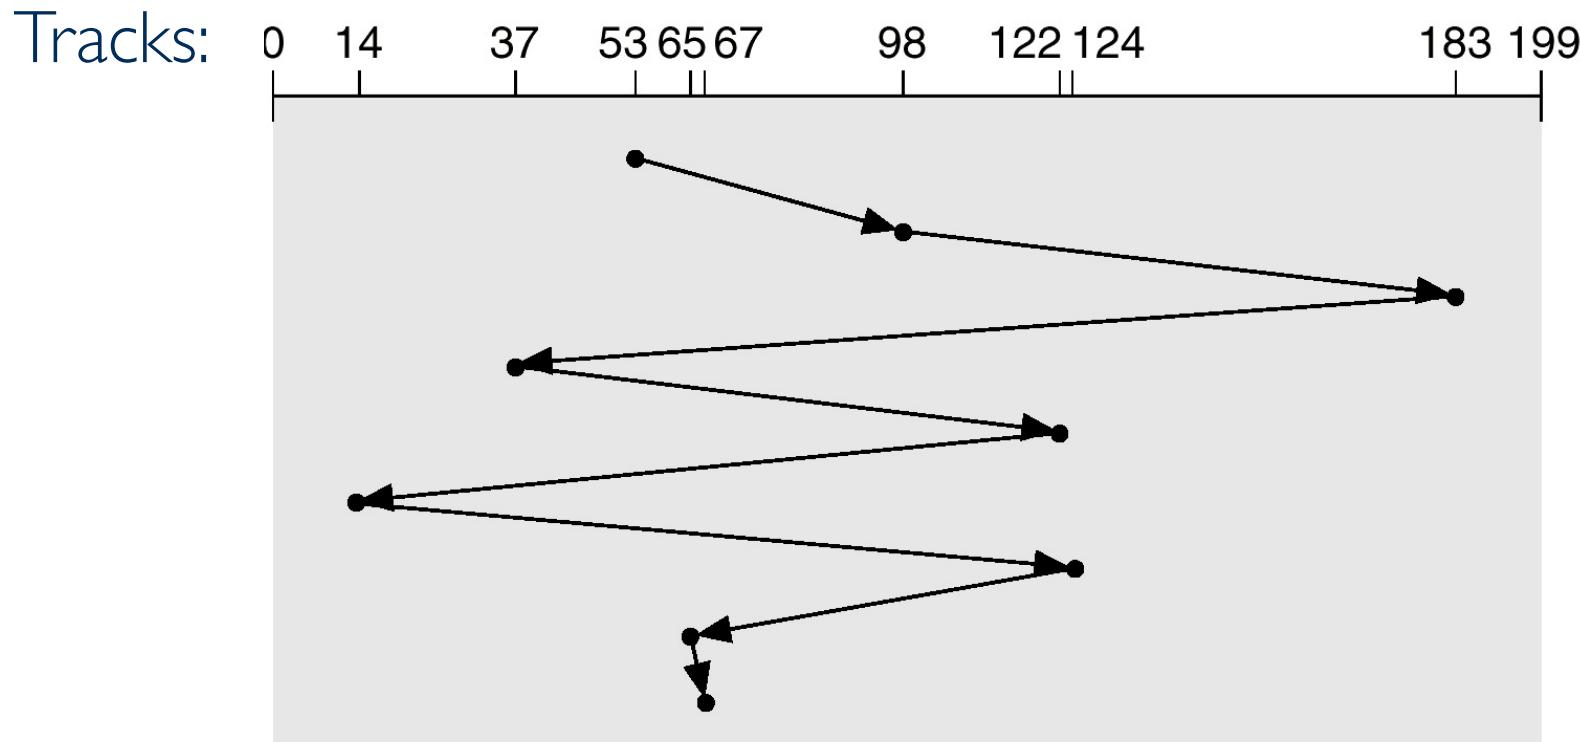
# Disk scheduling

---

- We're going to bypass discussing the metrics for disk scheduling, but know there are many algorithms similar to CPU scheduling
- Some prominent algorithms
  - FCFS – self explanatory
  - SSTF – shortest seek time first
  - LOOK – keep going the same direction until no more requests in that direction (more like a real elevator)
  - SCAN – "elevator algorithm": like LOOK but always go to the top and bottom floors on each trip
  - C-SCAN – always go "up" and don't service any requests on the down cycle
  - C-LOOK – like LOOK but only service requests in one direction

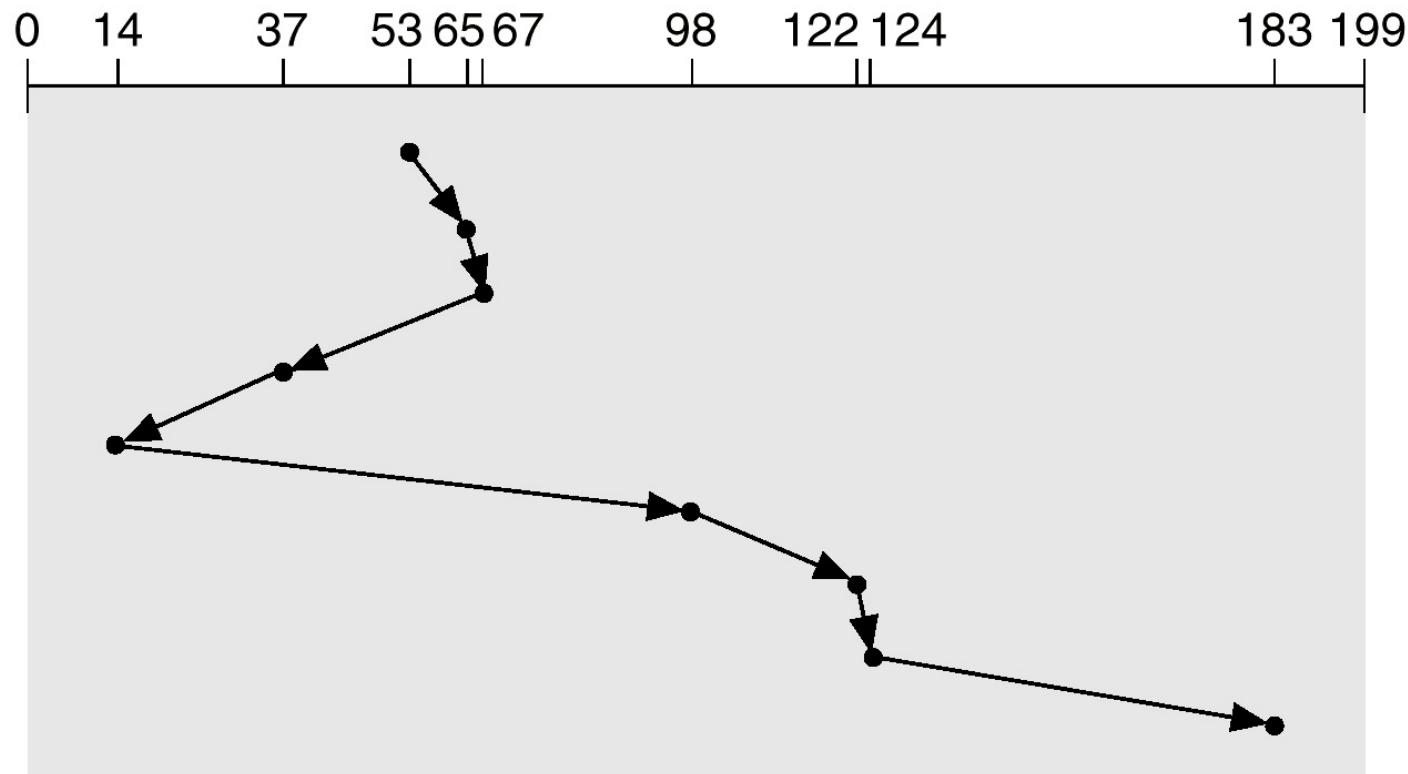
# FCFS

queue = 98, 183, 37, 122, 14, 124, 65, 67  
head starts at 53



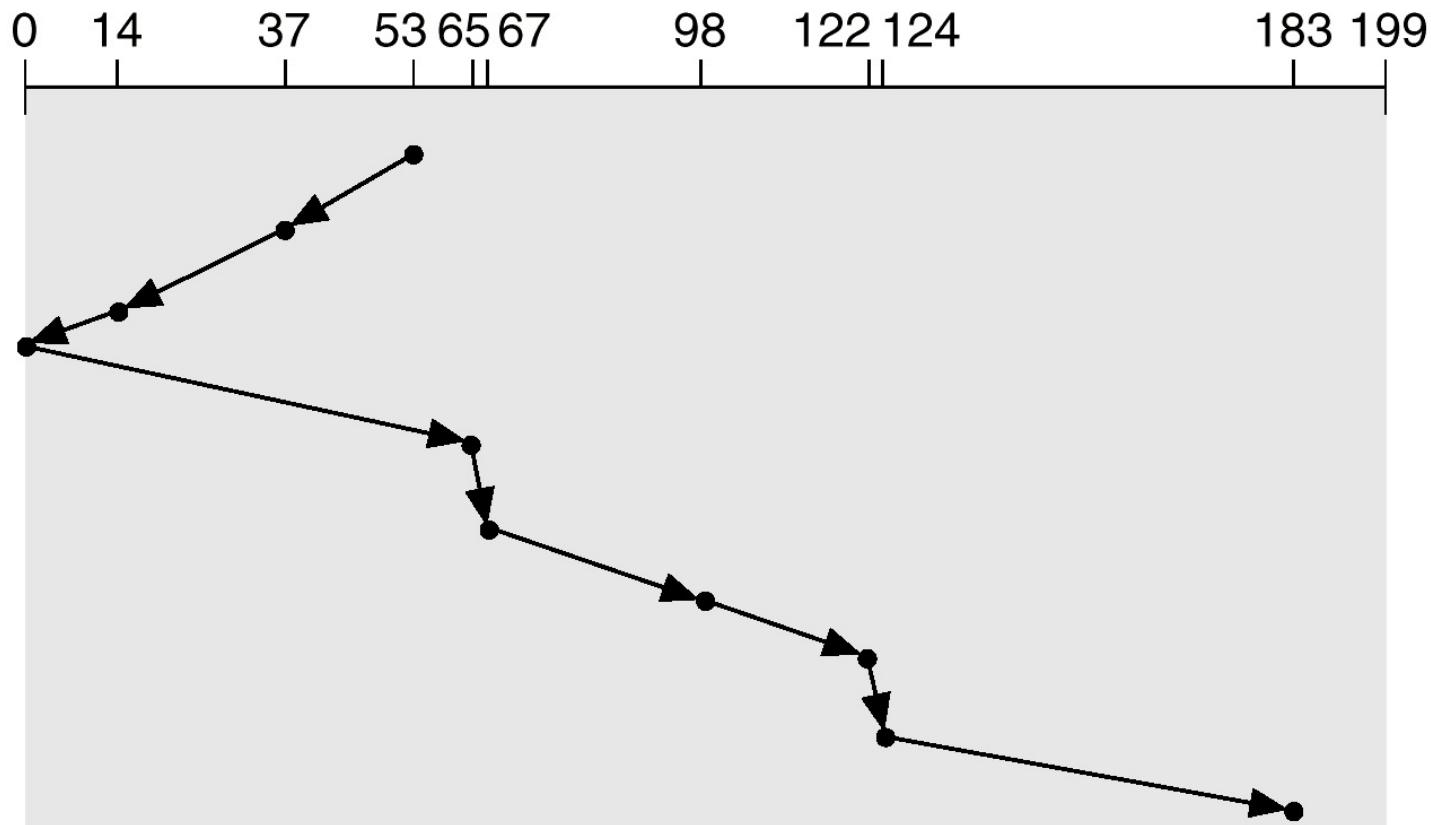
# SSTF

queue = 98, 183, 37, 122, 14, 124, 65, 67  
head starts at 53



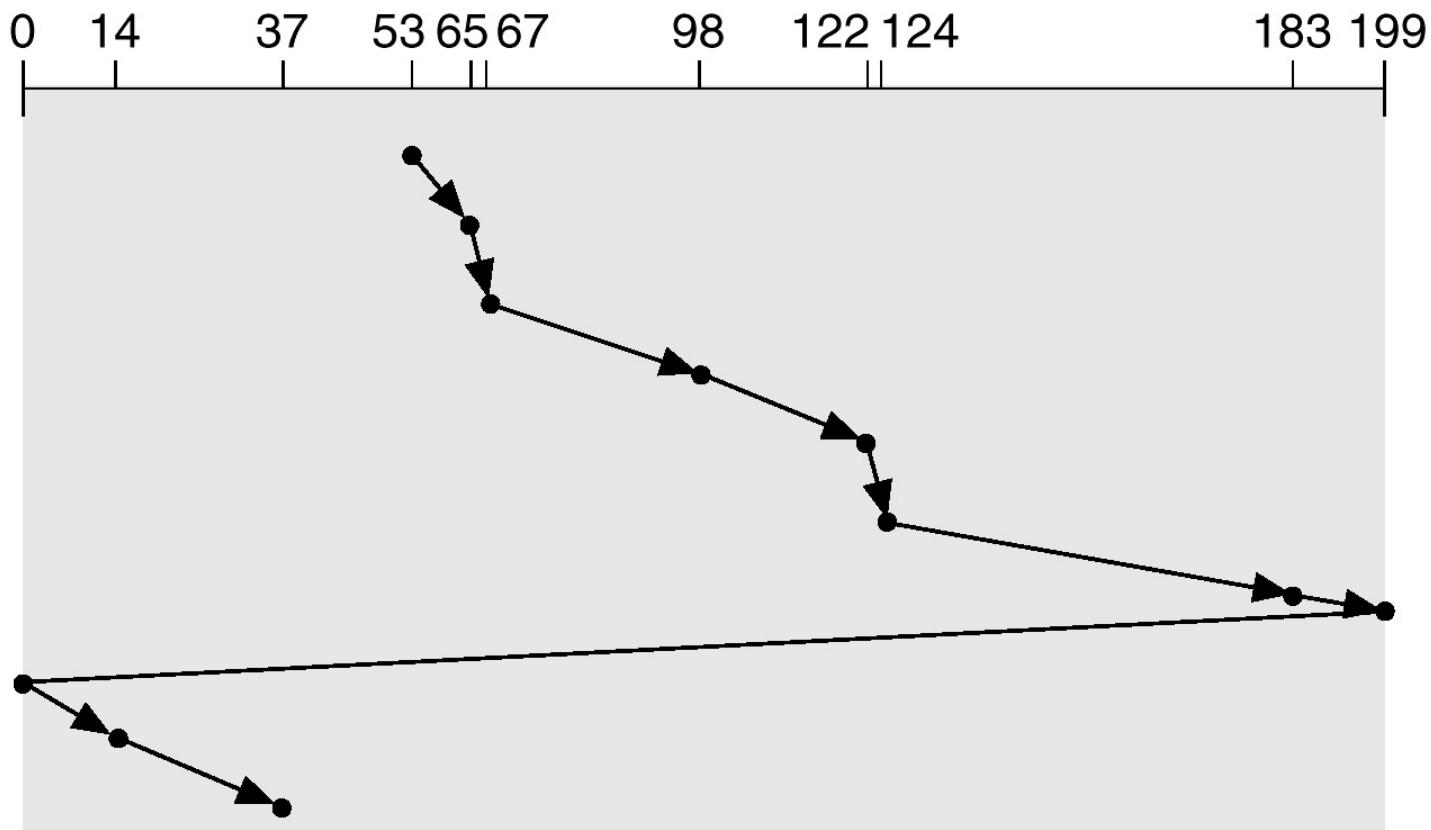
# SCAN

queue = 98, 183, 37, 122, 14, 124, 65, 67  
head starts at 53



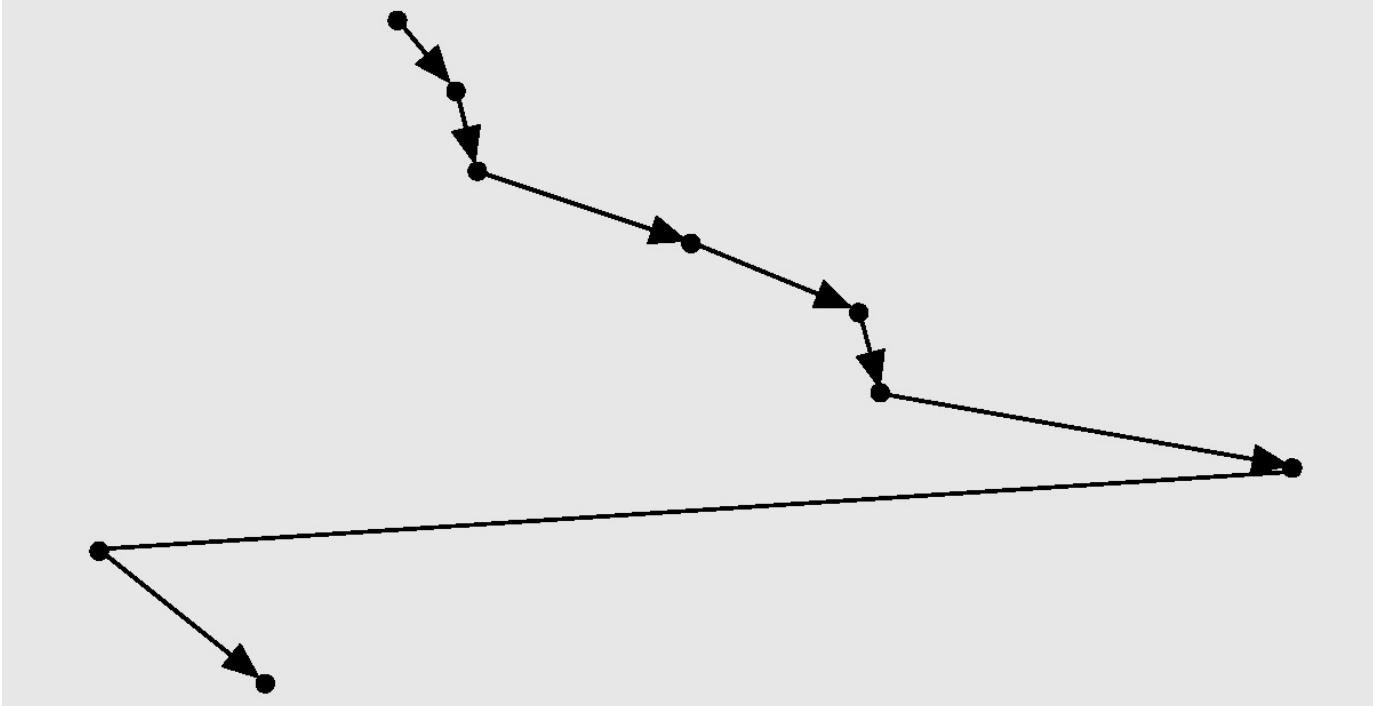
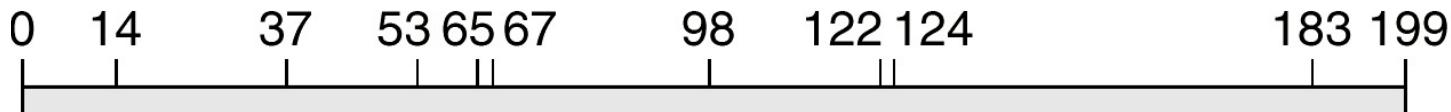
# C-SCAN

queue = 98, 183, 37, 122, 14, 124, 65, 67  
head starts at 53



# C-LOOK

queue = 98, 183, 37, 122, 14, 124, 65, 67  
head starts at 53



# Question

---

Which disk scheduling algorithm isn't called the “elevator algorithm” but works more like a modern elevator?

- A. I just want the participation credit
- B. SCAN
- C. LOOK
- D. C-SCAN
- E. C-LOOK

# Solid-state storage devices (SSDs)

---

- These have rapidly become effective replacements for small & fast storage applications
- They can act like HDDs and are faster but more expensive
- Look for an NVME connected SSD. Direct connection to PCI-E instead of through a SATA/SAS HDD controller.

	HDD	SSD
Capacity	1 - 18TB	0.25 - 15TB
IOPS (I/O per second)	400	50,000-1,000,000
Max transfer rate	Up to 500 MB/sec	Up to 7 GB/sec
Seek time	8 ms	0.1 ms
Energy use	6-16 W	2-5 W
Cost (\$/GB)	0.02	0.10