

# **A Smorgasbord of Embedded and Pervasive Computing Research**

**Kishore Ramachandran**

(part of systems group which includes Ada Gavrilovska, Taesoo Kim, Ling Liu, Calton Pu, and Alexey Tumanov)

## ❑ Current PhD inmates!



Enrique Saurez

Harshit Gupta


Zhuangdi Xu

Adam Hall

Ashish Bijlani



Tyler Landle




Manasvini Sethuraman



Anirudh Sarma



Alan Nussbaum



Difei Cao




Jinsun Yoo

## ❑ Recently escaped!

- Hyojun Kim (IBM Almaden then Startup and now Google); Lateef Yusuf (Amazon then Google); Mungyung Ryu (Google); Kirak Hong (Google and now CTRL-labs); Dave Lillethun (Seattle U.); Dushmanta Mohapatra (Oracle); Wonhee Cho (Microsoft); Beate Ottenwalder (Bosch); Ruben Mayer (TU Munich), Ashish Bijlani (Startup)

## ❑ Plus a number of MS and UGs

# Embedded Pervasive Lab



## Pervasive side of the house



- ❑ Embedded devices treated as black boxes
- ❑ System Support for IoT
  - ❑ Fog/Edge computing

# Current-Generation Applications

*Cloud computing's utility model  
commoditized hardware...*



Amazon  
EC2



Microsoft  
Azure



Google Cloud Platform

*Enabling large-scale apps from  
centralized data centers...*

facebook

NETFLIX



# Next-Generation Applications

Future apps will be *data-driven*, *model-driven*, *machine-in-the-loop*, and far more demanding...



*Autonomous  
Vehicles (AV)*



*Augmented/  
Virtual  
Reality  
(AR/VR)*



*Smart  
Cities*



*Smart  
Camera  
Networks*

*Latency <20 ms*  
*Bandwidth >50 Mbps*  
*Per device*

# Next-Generation Applications

- Sense -> Process -> Actuate
- Common Characteristics
  - Dealing with real-world data streams
  - Real-time interaction among mobile devices
  - Wide-area analytics
- Requirements
  - Dynamic scalability
  - Low-latency communication
  - Efficient in-network processing



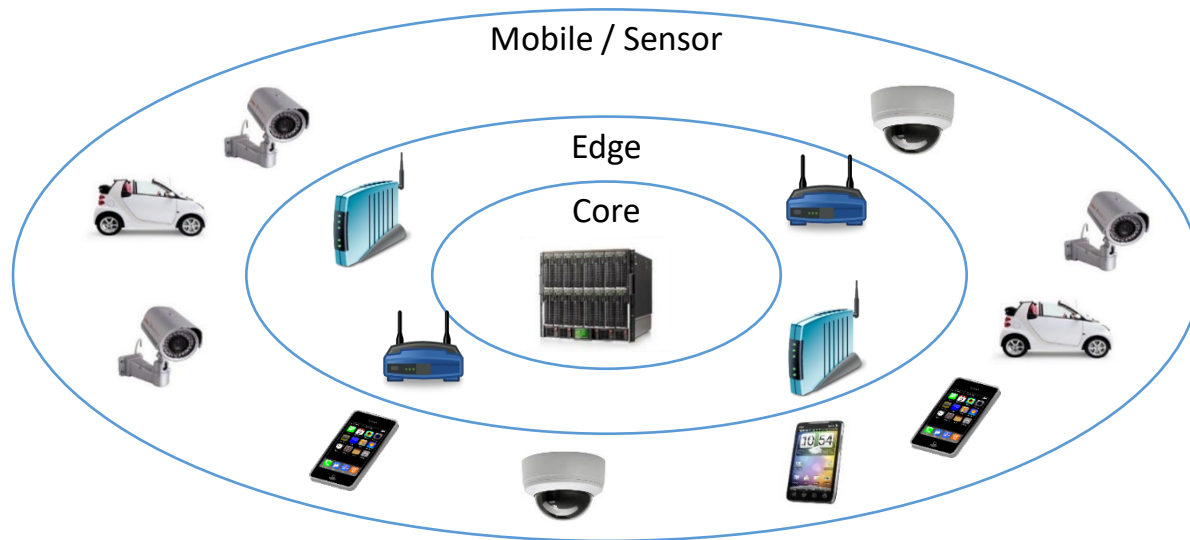
# Cloud Computing

- Good for web apps at human perception speeds
  - Throughput oriented web apps with human in the loop
- Not good for many latency-sensitive IoT apps at computational perception speeds
  - sense -> process -> actuate
- Other considerations
  - Limited by backhaul bandwidth for transporting plethora of 24x7 sensor streams
  - Not all sensor streams meaningful
    - => Quench the streams at the source
  - Privacy and regulatory requirements




# Fog/Edge Computing

- Extending the cloud utility computing to the edge
- Provide utility computing using resources that are
  - Hierarchical
  - Geo-distributed




# Fog/edge computing today

- Edge is slave of the Cloud
  - Platforms: IoT Azure Edge, CISCO Iox, Intel FRD,
- Mobile apps beholden to the Cloud




# Vision for the future



- Elevate Edge to be a peer of the Cloud
  - Prior art: Cloudlets (CMU+Microsoft), MAUI (Microsoft)
- In the limit
  - Make the Edge autonomous even if disconnected from the Cloud


# Why = ?

- Interacting entities (e.g., connected vehicles) connected to different edge nodes
- Horizontal (p2p) interactions among edge nodes essential



# Why ?

- Autonomy of edge (disaster recovery)



# Challenges for making



- Need for powerful frameworks akin to the Cloud at the edge
  - Programming models, storage abstractions, pub/sub systems, ...
- Geo-distributed data replication and consistency models
  - Heterogeneity of network resources
  - Resilience to coordinated power failures
- Rapid deployment of application components, multi-tenancy, and elasticity at the edge
  - Cognizant of limited computational, networking, and storage resources

# Thoughts on Meeting the Challenges

([https://www.cc.gatech.edu/~rama/recent\\_pubs.html](https://www.cc.gatech.edu/~rama/recent_pubs.html))

**Theme:** Elevating the Edge to be a peer of the Cloud

- Geo-distributed programming model for Edge/Cloud continuum
  - OneEdge (ACM SoCC 2021)
  - Foglets (ACM DEBS 2016)
- Geo-distributed data management – replica placement, migration and consistency
  - EPulsar (ACM SEC 2021)
  - FogStore (ACM DEBS 2018)
  - DataFog (HotEdge 2018)
- Efficient Edge runtimes
  - Serverless functions using WebAssembly (ACM IoTDI 2019)
- Applications using autonomous Edge
  - Social Sensing *sans* Cloud (SocialSens 2017)
  - STTR: Space Time Trajectory Registration (ACM DEBS 2018)
  - STVT: Space-Time Vehicle Tracking (HotVideoEdge 2019)
  - Coral-Pie: Space-Time Vehicle Tracking at the Edge (ACM Middleware 2020)
- Vision: “A case for elevating the edge to be a peer of the cloud”, GetMobile, 2020
- Vision: "eCloud: Vision for the Evolution of Edge-Cloud Continuum", IEEE Computer, 2021.

# Ongoing Work

- eCloudSim (with **Daglis, Chatterjee, Dhekne**)
  - eCloud Simulation Infrastructure
- Prescient video prefetching at the edge for AV infotainment (With **Dhekne**)
  - Use route to JIT prefetching and caching for DASH player on vehicle
    - **Foresight (ACM MMSys 2021)**
  - Use mmWAVE (integrated with 5G LTE for edge node selection) to beam to passing vehicle
    - **ClairvoyantEdge (in submission)**
- Edge centric video data management systems for AV (With **Arulraj**)
  - Annotations with video for query processing, multi-tenancy, and sharing
    - **EVA (To appear in ACM SIGMOD)**
- Nimble execution environments for the Edge
  - Analyze cold start times in containers
  - Clean slate exec environment for FaaS
- NFSlicer: dataplane optimizations for processing network functions (With **Daglis**)
  - Selective data movement (e.g., header vs. payload) for NF chaining
    - **NFSlicer (in submission)**
- Performance isolation in the Edge (with **Sameh Elnikety** – Microsoft)
  - Harvest Container: Mixing latency-sensitive and throughput-oriented apps at the edge
- Hardware accelerators in micro-datacenters (with **Tushar Krishna**)
  - Efficiently using accelerators to improve efficiency in edge datacenters.




# Embedded side of the house



- ❑ Infinite storage for mobile devices
- ❑ Optimizing Mobile Video Downloads


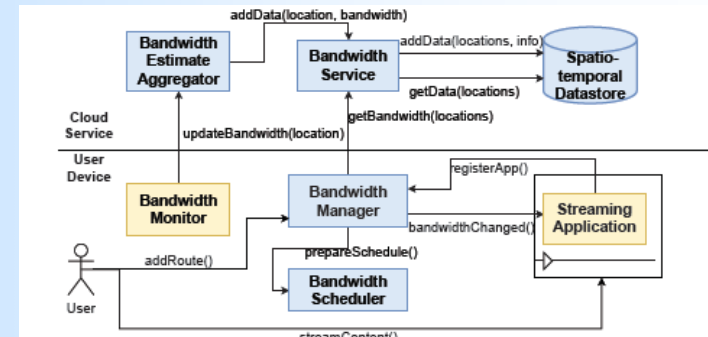
# Infinite Storage for Mobile Devices

- ❑ Seamlessly extend the storage on mobile to the Cloud for any app
  - User space file system
    - APSys 2018, USENIX ATC 2019, Sigmetrics 2021
- ❑ Use machine learning to build user's everyday working set and (off)load (un)wanted data
- ❑ Issues
  - Latency
  - Energy consumption
  - Security and privacy



# Optimizing Mobile Video Downloads

- ❑ Foresight (ACM MMSys 2021)
  - Bandwidth prediction across space and time for mobile users
- ❑ ClairvoyantEdge
  - Short range mmWave augmentation at Edge for high bandwidth video delivery



# Recap



- ❑ Infinite storage for mobile devices
- ❑ Optimizing Mobile Video Downloads
  - ❑ Foresight, ClairvoyantEdge



- ❑ Fog/Edge computing
  - eCloud
  - Foglets, OneEdge, thin virtualization for FaaS
  - Fogstore, DataFog, ePulsar, NFSlicer
  - STTR, Socialsens

# Ongoing Projects

- ❑ **eCloud:** Device-Edge-Cloud continuum
- ❑ **OneEdge:** Device/Edge/Cloud control plane using AV as exemplar
  - Scheduling edge resources, monitoring, migration
- ❑ **Foresight and ClairvoyantEdge:** Prescient video prefetching at the edge for AV infotainment (With Prof. **Dhekne**)
  - Use route to JIT prefetching and caching for DASH player on vehicle
  - Use mmWAVE (integrated with 5G LTE for edge node selection) to beam to passing vehicle
- ❑ **Edge centric video data management systems for AV** (With Prof. **Arulraj**)
  - Annotations with video for query processing, multi-tenancy, and sharing
- ❑ **Nimble execution environments for the Edge**
  - Analyze cold start times in containers
  - Clean slate exec environment for FaaS
- ❑ **NFSlicer:** dataplane for processing network functions (With Prof. **Daglis**)
  - Selective data movement (e.g., header vs. payload) for NF chaining
- ❑ **MicroEdge:** Low-cost edge architecture for camera processing (With Prof. **Krishna**)
- ❑ **Edge computing solution for underserved communities**
  - Smart information services without WAN connectivity

Pubs:

[http://www.cc.gatech.edu/~rama/recent\\_publications.html](http://www.cc.gatech.edu/~rama/recent_publications.html)

E-mail: rama@cc.gatech.edu

Lab:

<http://wiki.cc.gatech.edu/epl>

What does Kishore “**really**” do in his copious spare time when he is not teaching?



Squash anyone?

All virtual!

Soon real?

Table tennis anyone?




wikiHow



# What you should take away?

- ❑ "Kishore" rhymes with "sea shore"
- ❑ Squash/Table-tennis
- ❑ EPL
- ❑ Fog/Edge computing
- ❑ Infinite storage on mobile/EdgeCaching



<https://www.dreamstime.com/namaste-vector-hand-drawn-symbol-yoga-design-image112101574>



# Extra slides on EPL projects underway

Kishore Ramachandran

Two topics

Camera Processing  
@ Edge

Serverless @ Edge

# Camera Network

## Context

- Streaming 24 x 7
- Used by multiple applications
- Processing at ingestion time
- Processing at Edge

## Applications

- Suspicious vehicle tracking
- Smart traffic light
- Camera-assisted navigation for the disability




# Motivating Application


## Space-Time Vehicle Tracking at video ingestion time

- Manual Checking
  - Labor-intensive
  - Error-prone due to lapses in attention
- Intelligent systems
  - Replace manual labor
  - Aid human decision-making

Analyzing



Querying



# Outline

- Recent work on camera applications at Edge
  - Coral-Pie: Space-Time Vehicle Tracking at the Edge (**ACM Middleware 2020**)
- Proposed Work in a nutshell
  - Camera Networks at the Edge
    - MicroEdge: multi-tenancy edge system architecture
    - Accelerator orchestration in MicroEdge
  - Serverless at the Edge
    - Nimble orchestration for FaaS
    - Cross-site load balancing
    - Location-aware auto-scaling

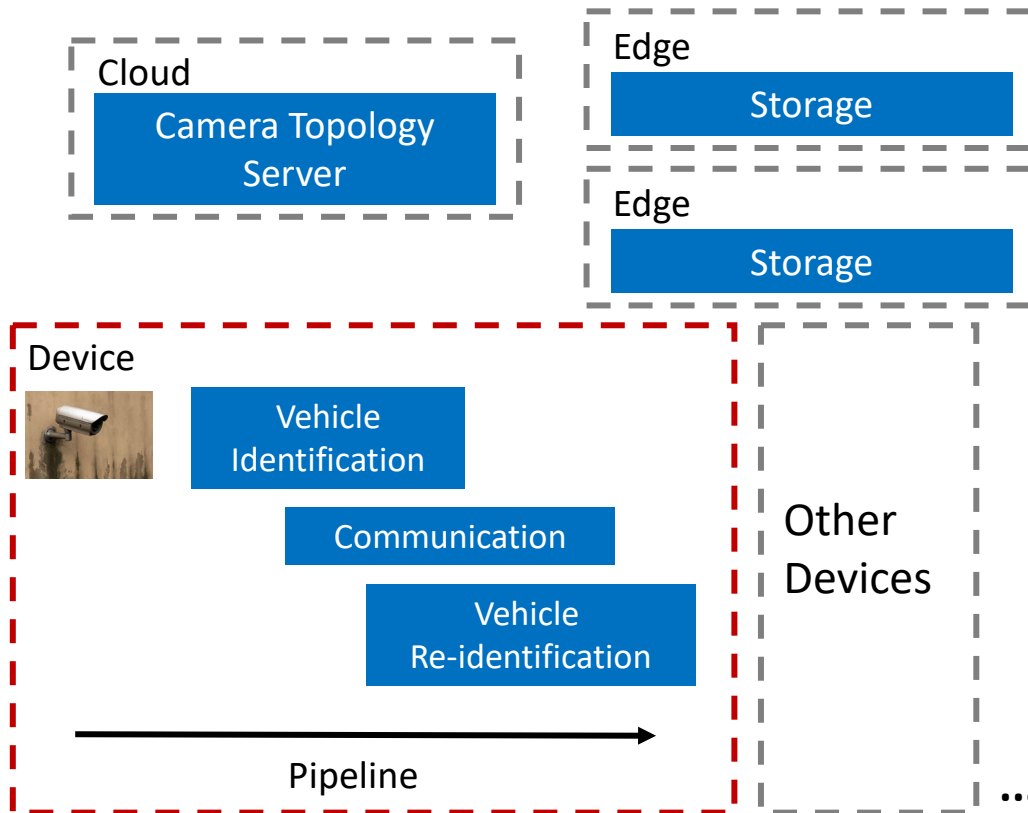
# Recent Work

## Novel Principles

- "Scalable-by-design" camera processing at ingestion time
  - Per-camera latency-based subtask placement
  - Bounded network communication per camera
- Fault tolerance for geo-distributed camera network
  - Automatic camera topology management



# Coral-Pie: Device-Edge-Cloud Continuum



## Device:

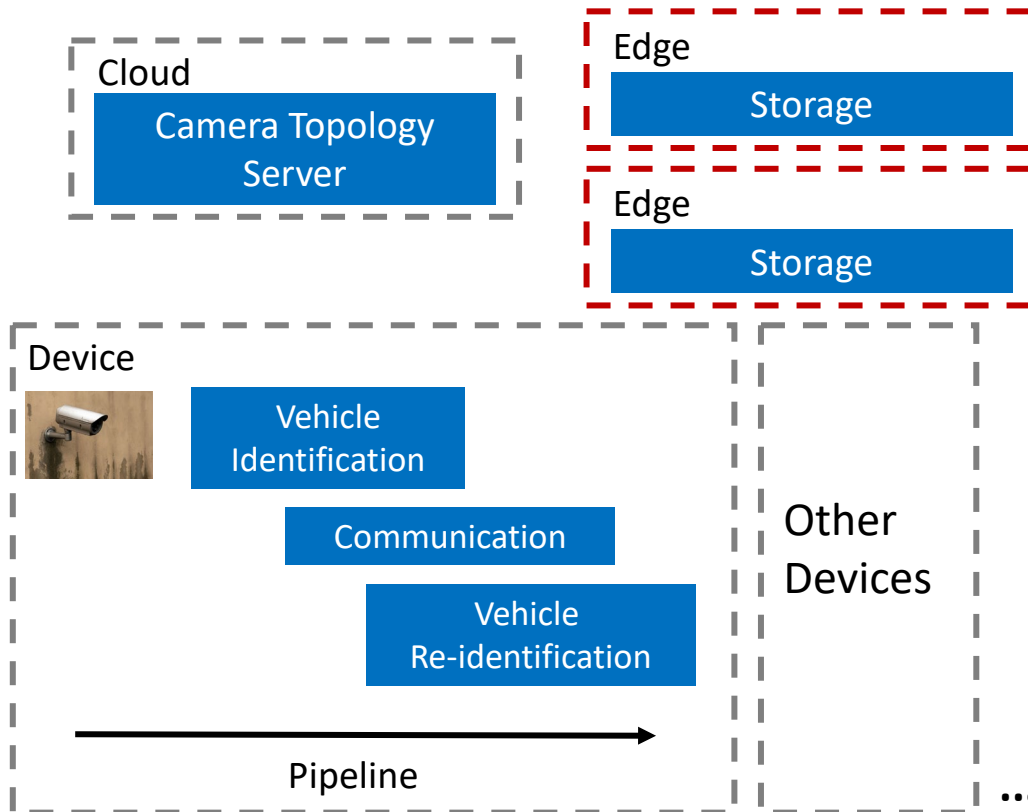
- Camera + 2 RPis + TPU

## Heavy-lifting real-time camera processing

- **Deterministic latency bounds**
- **No pressure on the backhaul network bandwidth**

## Pipelined processing to exploit parallelism

# Coral-Pie: Device-Edge-Cloud Continuum



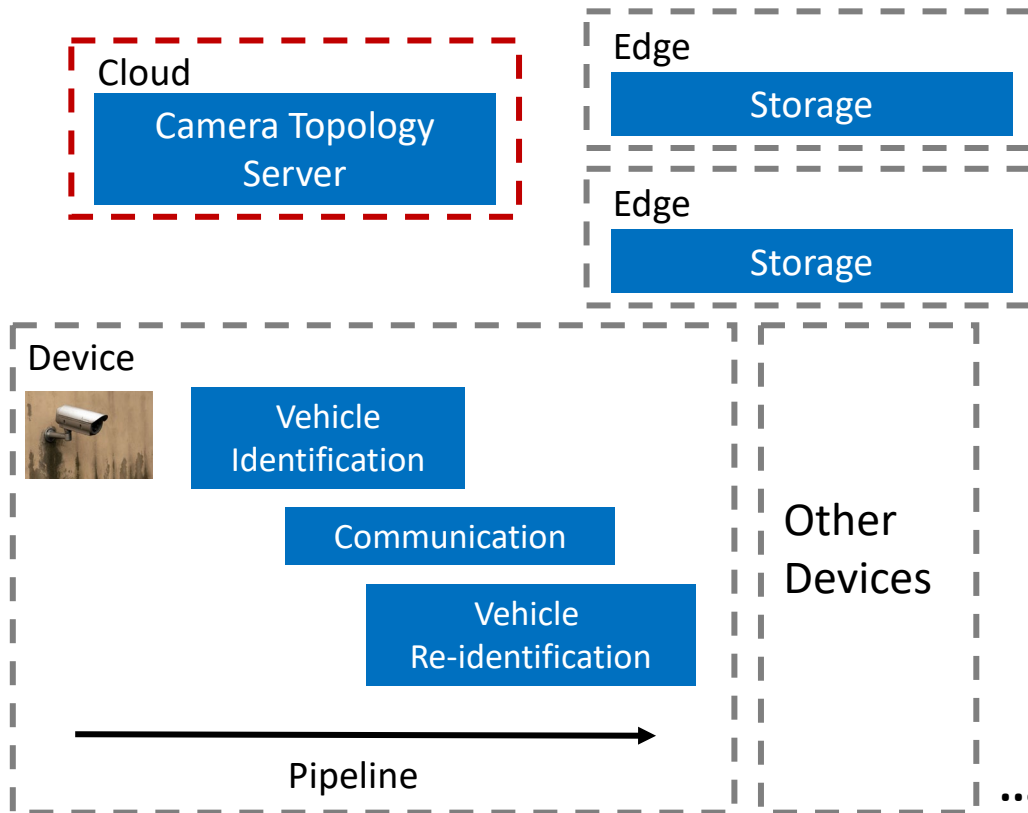
## Edge:

- A multi-tenant micro-datacenter housed in a small footprint location
- **Few network hop(s)** away from Devices

## Offloading storage

- **Asynchronously**
- Helps balance the pipeline of tasks on the Devices
- Backhaul wide-area network (WAN) is not pressured

# Coral-Pie: Device-Edge-Cloud Continuum




## Cloud

- A centralized service provider that all devices can connect to
- **Global knowledge**


Maintains geographical relationship between the cameras

- Infrequently updated (i.e., a new camera is deployed or an old camera is removed)

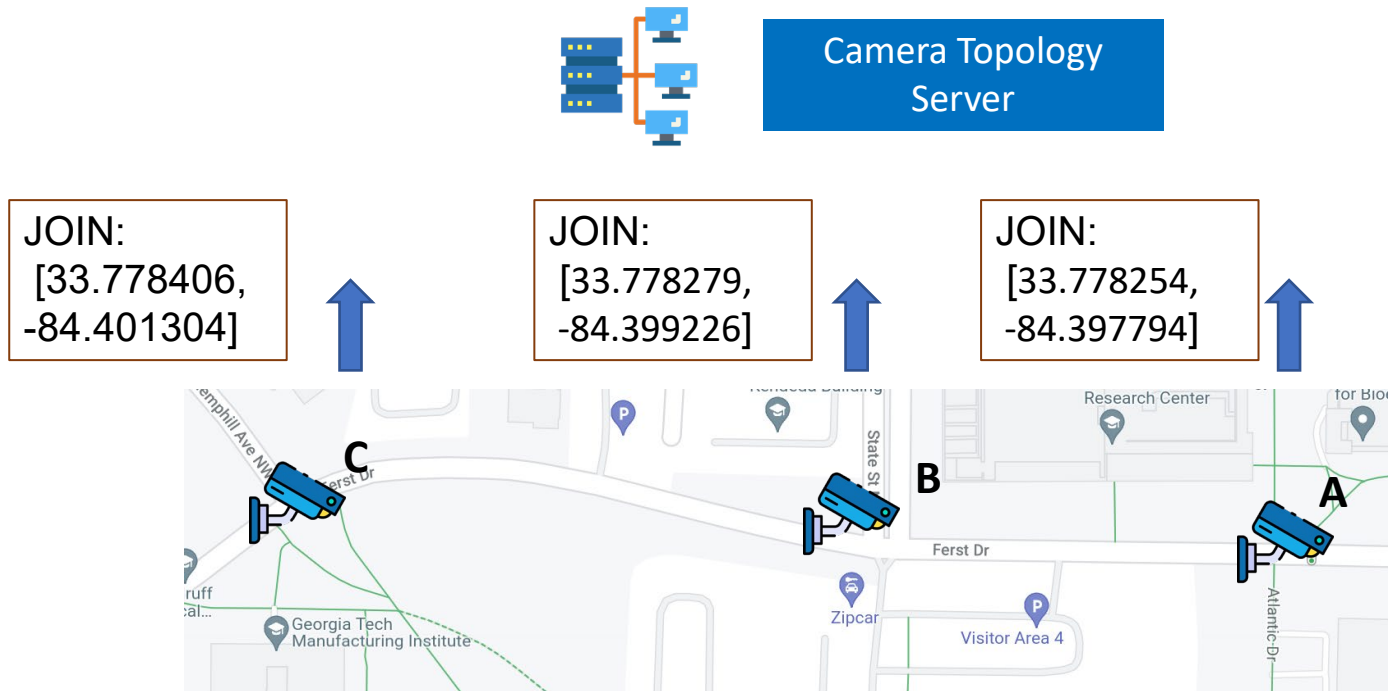
# Coral-Pie in Action




Camera Topology  
Server



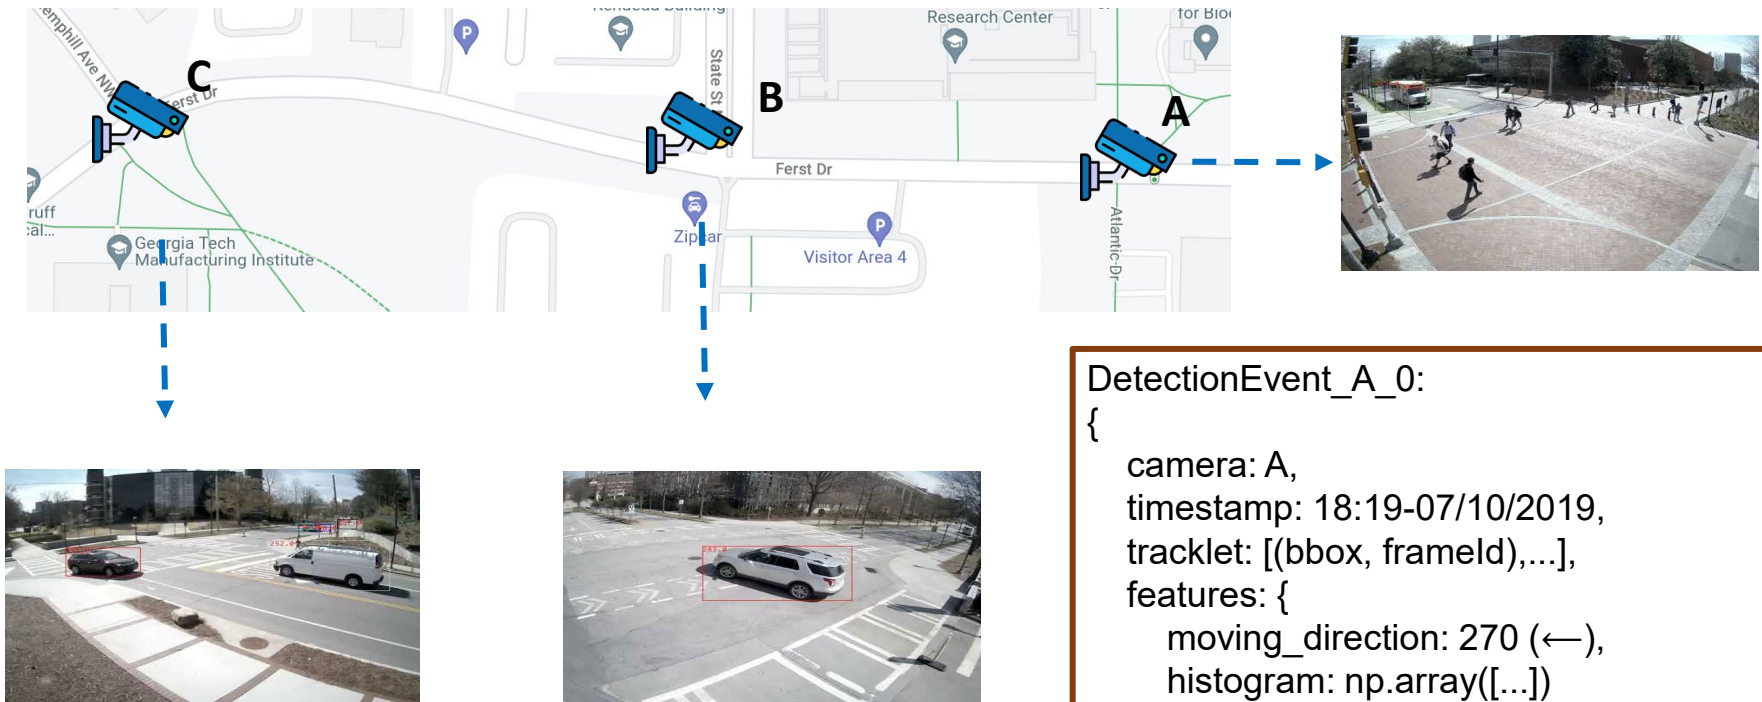
# Cloud: Management of Camera Topology



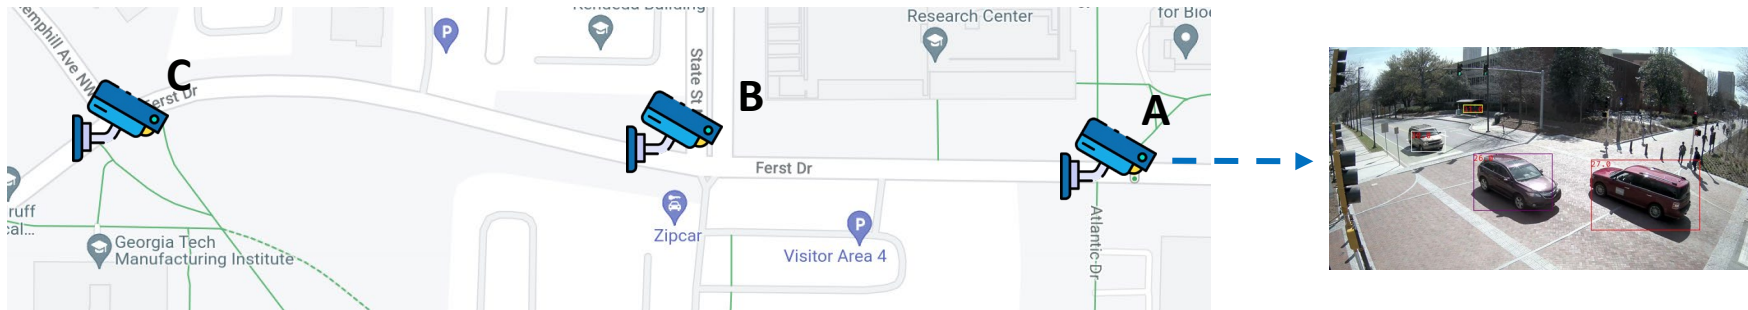
# Cloud: Management of Camera Topology



# Device: Vehicle Identification




# Device: Communication



Camera B's Candidate Pool:

DetectionEvent\_A\_0:

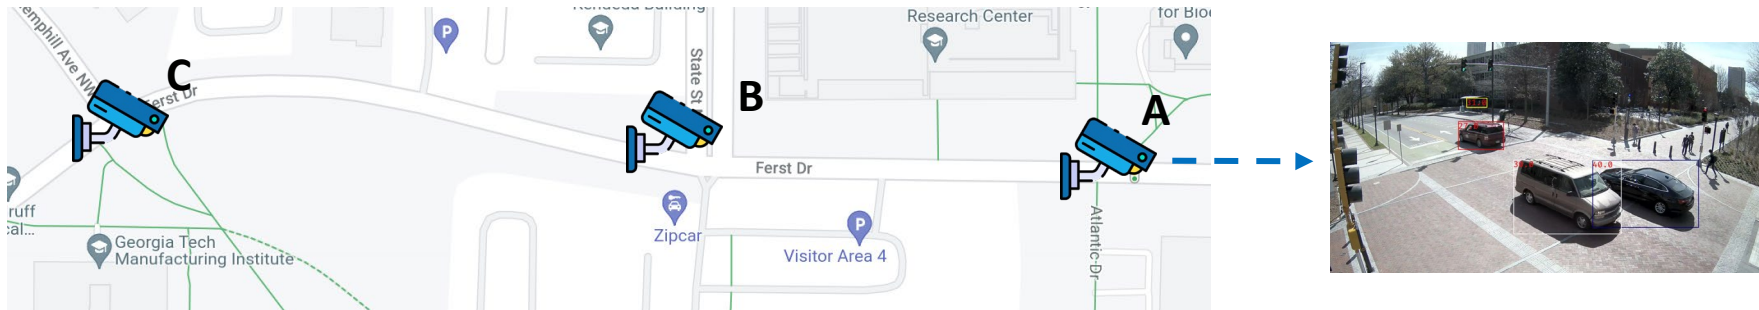


DetectionEvent\_A\_0:

```
{  
  camera: A,  
  timestamp: 18:19-07/10/2019,  
  tracklet: [(bbox, frameId),...],  
  features: {  
    moving_direction: 270 (←),  
    histogram: np.array([...])  
  }  
}
```




# Device: Communication



Camera B's Candidate Pool:

DetectionEvent\_A\_0:



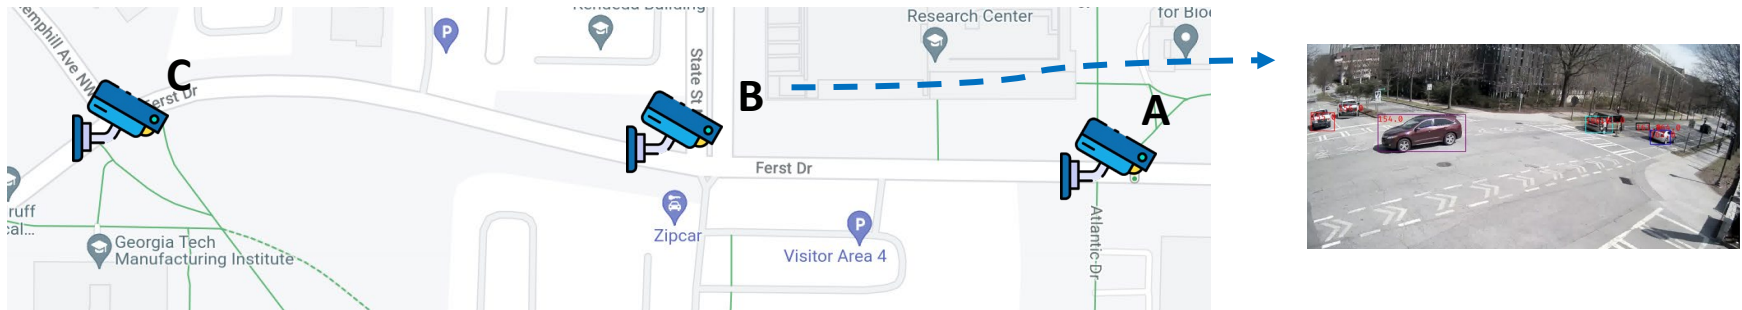
DetectionEvent\_A\_2:



DetectionEvent\_A\_2:

```
{
  camera: A,
  timestamp: 18:19:09-07/10/2019,
  tracklet: [(bbox, frameId),...],
  features: {
    moving_direction: 270 (←),
    histogram: np.array([...])
  }
}
```

# Device: Vehicle Re-Identification at Camera B



Camera B's Candidate Pool:

DetectionEvent\_A\_0:



DetectionEvent\_A\_2:



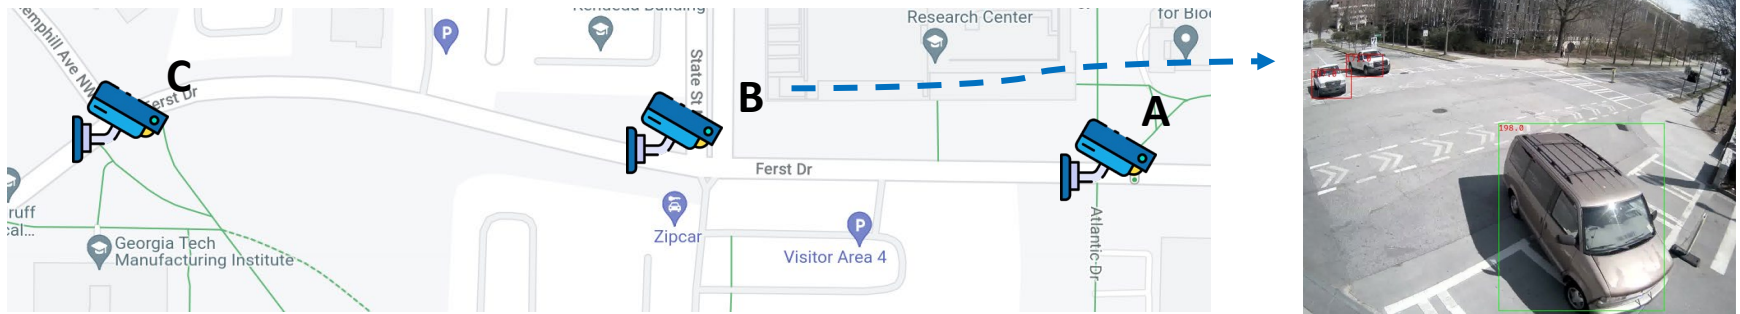
Matching!



DetectionEvent\_B\_7:

```
{  
  camera: B,  
  timestamp: 18:19:26-07/10/2019,  
  tracklet: [(bbox, frameId),...],  
  features: {  
    moving_direction: 270 (←),  
    histogram: np.array([...])  
  }  
}
```

# Device: Vehicle Re-Identification at Camera B



Camera B's Candidate Pool:

DetectionEvent\_A\_0:



DetectionEvent\_A\_2:



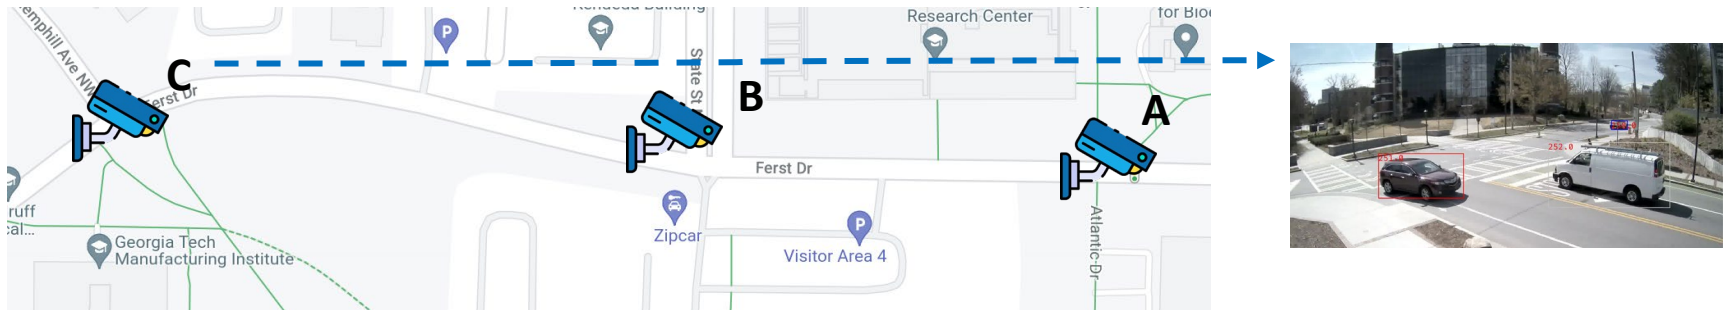
Matching!




DetectionEvent\_B\_9:

```
{
  camera: B,
  timestamp: 18:19:29-07/10/2019,
  tracklet: [(bbox, frameId),...],
  features: {
    moving_direction: 180(↓),
    histogram: np.array([...])
  }
}
```

# Device: Vehicle Re-Identification at Camera C




Camera B's Candidate Pool:



```
DetectionEvent_C_12:
{
  camera: C,
  timestamp: 18:19:46-07/10/2019,
  tracklet: [(bbox, frameId),...],
  features: {
    moving_direction: 315 (↑),
    histogram: np.array([...])
  }
}
```

# Edge: Storage



- Trajectory Storage as Graph
- Frame Storage - traceability

DetectionEvent\_C\_12

DetectionEvent\_B\_7

DetectionEvent\_A\_0

DetectionEvent\_B\_9

DetectionEvent\_A\_2




# From Coral-Pie to MicroEdge

## Limitations of Coral-Pie

- Dedicated hardware resources for each camera processing pipeline
  - Each pipeline may not use the resources all the time
- Camera network intended for multiple applications

## Opportunity

- Share TPU, CPU, and memory resources within and across applications



# Limitations of the State-of-the-art

- Prior art focus on GPU management in resource-rich clusters
  - NVIDIA Docker and NVIDIA MPS
  - Clockwork[1], Heimdall[2], Infaas[3], and Clipper[4]
- Lack of support for multiplexing TPU resources across containers
- Missed opportunity in container-orchestration systems (e.g., Kubernetes, K3s)
  - Under-utilization of TPU resources

[1] Gujarati, Arpan et al. "Serving DNNs like Clockwork: Performance Predictability from the Bottom Up." OSDI (2020).


[2] Yi, Juheon and Youngki Lee. "Heimdall: mobile GPU coordination platform for augmented reality applications." Proceedings of the 26th Annual International Conference on Mobile Computing and Networking (2020).

[3] Romero, Francisco et al. "INFaaS: Automated Model-less Inference Serving." USENIX Annual Technical Conference (2021).

[4] Crankshaw, Daniel et al. "Clipper: A Low-Latency Online Prediction Serving System." NSDI (2017).

# MicroEdge Vision

- A low-cost edge cluster serves computational resources for a set of geo-local cameras.
- A performant **multi-tenancy** architecture.





# Objectives of MicroEdge System Architecture

- Share TPU resources across independent application pipelines
- Performance isolation across independent application pipelines
- Load-balance incoming inference requests across TPU resources
- Co-compile different models on TPUs to reduce model swapping overhead

# MicroEdge System Architecture

## Research Questions

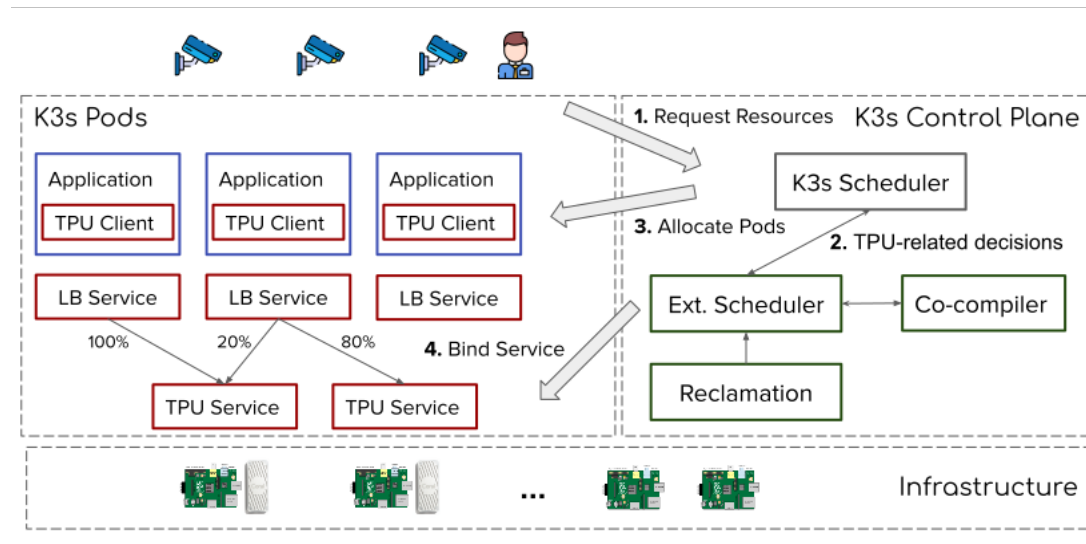
- How do we extend the state of the art?

## Extensions to K3s's control plane

- Orchestrate TPU resources on creation and destruction of application pods
  - Admission Control
- Key components:
  - Extended scheduler, TPU co-compiler, Resource reclamation

## Extensions to K3s's data plane

- Mechanisms to multiplex TPU resources across applications
  - Monitoring TPU usage
- Key components:
  - TPU service, TPU service client library, Load balancing service



# Evaluation Plan

## Application Study

- Space-time vehicle tracking
- Can MicroEdge improve the resource utilization while meeting the SLA of application?

## Trace-based Study

- Azure Function Traces [1]
- Map to ephemeral application pipelines
- Can MicroEdge effectively allocate and reclaim TPU resources?

[1] Shahradd, Mohammad, et al. "Serverless in the wild: Characterizing and optimizing the serverless workload at a large cloud provider." 2020 USENIX Annual Technical Conference (USENIX ATC 20). 2020.

# Serverless at the Edge for camera networks

- Motivation: Workload at the edge is highly variable with time
  - Dedicating resources leads to overprovisioning
  - Scarce resources at the Edge
- Original vision for Coral-Pie was to create space-time track for all vehicles all the time
  - But most of the time, we are interested only in "suspicious" vehicles
  - Increases the chance for resource multiplexing
  - Number of interesting objects detected is sporadic






Image source: Zhang, Miao, et al. "Towards cloud-edge collaborative online video analytics with fine-grained serverless pipelines." *Proceedings of the 12th ACM Multimedia Systems Conference*. 2021.

=> Move toward Serverless at the Edge


# Limitations of State-of-the-Art

## ■ Cloud-based FaaS platforms

- OpenFaaS, KNative (backed by K8s)
- High Latency for OpenFaaS in MicroEdge  
==> High container creation overhead
- Multiple WAN traversals  
==> High WAN overhead
- Location-agnostic scheduling



(b) Latency breakdown with container cold start.




(c) Latency breakdown with pre-warmed containers.

Figure 2: Experimental results with Kubernetes.

WAN latency = 80 ms RTT

# Limitations of State-of-the-Art

- Cloud-based FaaS platforms
  - OpenFaaS, KNative (backed by K8s)
  - High Latency for OpenFaaS in MicroEdge
  - Multiple WAN traversals
  - Location-agnostic scheduling
- Edge-specific FaaS platforms
  - Mu (SoCC'21)
    - Fine-grained queue-size based load-balancing to meet latency requirement
  - CEVAS (MMSys'21)
    - Optimal partitioning of serverless video processing between edge and cloud
  - Both platforms focus only on resources of 1 edge site
  - Load-balancing and Scaling policies are unaware of load on other sites



# Objectives for Serverless @ Edge

- Cross-site load-balancing of function invocation requests
- Location-aware auto-scaling
- Nimble edge execution environment
- Handle heterogeneous network latencies between edge sites




Image source: Xu, Mengwei, et al. "From cloud to edge: a first look at public edge platforms." Proceedings of the 21st ACM Internet Measurement Conference. 2021.

# Initial set of Research Questions

- Cross-site load balancing policy
  - Whether and to which site to offload?
- Global autoscaling policy
  - How to predict spatio-temporal distribution of client requests?
  - Where to provision function containers to minimize resource wastage and meet application requirement?
- Monitoring
  - What metrics are needed to enable above policies?
  - How to efficiently and scalably propagate per-site system state to other sites?
- Agility-oriented optimizations in Kubernetes to reduce cold start overhead
- Network proximity monitoring
  - How to do so accurately, efficiently, and at scale? (Network coordinates?)



# Evaluation Plan

## ■ Load balancing policy

- Proposed network proximity and load-aware approach
- Compare with
  - Coarse-grained monitoring (e.g., uniform load distribution *a la* OpenFaaS)
  - Fine-grained monitoring (e.g., fine-grained queue-sizes *a la* Mu)

## ■ Autoscaling policy

- Proposed spatio-temporal demand distribution based approach
- Compare with
  - Techniques based on CPU utilization/request rate (*a la* OpenFaaS) and latency-aware (*a la* Mu)

## ■ Network proximity

- Is the use of network coordinates accurate, efficient, and scalable ?

# Expected Contributions of the Proposed Research

## **MicroEdge**

- Control and data-plane mechanisms for orchestrating accelerator resources
- Policies for efficient allocation of accelerator resources
- Admission Control and Monitoring policies

## **Serverless @ Edge**

- Cross-site load and latency aware load-balancing policy
- Spatio-temporal demand based cross-site autoscaling policy
- Integrating network coordinates for network proximity monitoring

**END**